

Linux Cross Reference

[Free Electrons](#)

Embedded Linux Experts

• [source navigation](#) • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version: [2.0.40](#) [2.2.26](#) [2.4.37](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#) [3.18](#) [3.19](#) [4.0](#) [4.1](#) [4.2](#)

[Linux/include/linux/netdevice.h](#)

```

1  /*
2  * INET                An implementation of the TCP/IP protocol suite for the LINUX
3  *                    operating system. INET is implemented using the BSD Socket
4  *                    interface as the means of communication with the user level.
5  *
6  *                    Definitions for the Interfaces handler.
7  *
8  * Version:          @(#)dev.h          1.0.10 08/12/93
9  *
10 * Authors:         Ross Biro
11 *                   Fred N. van Kempen, <waltje@uWalt.NL.Mugnet.ORG>
12 *                   Corey Minyard <wf-rch!minyard@relay.EU.net>
13 *                   Donald J. Becker, <becker@cesdis.gsfc.nasa.gov>
14 *                   Alan Cox, <alan@lxorguk.ukuu.org.uk>
15 *                   Bjorn Ekwall. <bjorn@blox.se>
16 *                   Pekka Riikonen <priikone@poseidon.pspt.fi>
17 *
18 *                   This program is free software; you can redistribute it and/or
19 *                   modify it under the terms of the GNU General Public License
20 *                   as published by the Free Software Foundation; either version
21 *                   2 of the License, or (at your option) any later version.
22 *
23 *                   Moved to /usr/include/linux for NET3
24 */
25 #ifndef LINUX\_NETDEVICE\_H
26 #define LINUX\_NETDEVICE\_H
27
28 #include <linux/timer.h>
29 #include <linux/bug.h>
30 #include <linux/delay.h>
31 #include <linux/atomic.h>
32 #include <linux/prefetch.h>
33 #include <asm/cache.h>
34 #include <asm/byteorder.h>
35
36 #include <linux/percpu.h>
37 #include <linux/rculist.h>
38 #include <linux/dmaengine.h>
39 #include <linux/workqueue.h>
40 #include <linux/dynamic_queue_limits.h>
41
42 #include <linux/ethtool.h>
43 #include <net/net_namespace.h>
44 #include <net/dsa.h>
45 #ifdef CONFIG_DCB
46 #include <net/dcbnl.h>
47 #endif
48 #include <net/netprio_cgroup.h>
49
50 #include <linux/netdev_features.h>
51 #include <linux/neighbor.h>
52 #include <uapi/linux/netdevice.h>
53 #include <uapi/linux/if_bonding.h>

```

```

54
55 struct netpoll_info;
56 struct device;
57 struct phy_device;
58 /* 802.11 specific */
59 struct wireless_dev;
60 /* 802.15.4 specific */
61 struct wpan_dev;
62 struct mpls_dev;
63
64 void netdev_set_default_ethtool_ops(struct net_device *dev,
65                                     const struct ethtool_ops *ops);
66
67 /* Backlog congestion levels */
68 #define NET_RX_SUCCESS 0 /* keep 'em coming, baby */
69 #define NET_RX_DROP 1 /* packet dropped */
70
71 /*
72  * Transmit return codes: transmit return codes originate from three different
73  * namespaces:
74  *
75  * - qdisc return codes
76  * - driver transmit return codes
77  * - errno values
78  *
79  * Drivers are allowed to return any one of those in their hard_start_xmit()
80  * function. Real network devices commonly used with qdiscs should only return
81  * the driver transmit return codes though - when qdiscs are used, the actual
82  * transmission happens asynchronously, so the value is not propagated to
83  * higher layers. Virtual network devices transmit synchronously, in this case
84  * the driver transmit return codes are consumed by dev_queue_xmit(), all
85  * others are propagated to higher layers.
86  */
87
88 /* qdisc ->enqueue() return codes. */
89 #define NET_XMIT_SUCCESS 0x00
90 #define NET_XMIT_DROP 0x01 /* skb dropped */
91 #define NET_XMIT_CN 0x02 /* congestion notification */
92 #define NET_XMIT_POLICED 0x03 /* skb is shot by police */
93 #define NET_XMIT_MASK 0x0f /* qdisc flags in net/sch_generic.h */
94
95 /* NET_XMIT_CN is special. It does not guarantee that this packet is lost. It
96  * indicates that the device will soon be dropping packets, or already drops
97  * some packets of the same priority; prompting us to send less aggressively. */
98 #define net_xmit_eval(e) ((e) == NET_XMIT_CN ? 0 : (e))
99 #define net_xmit_errno(e) ((e) != NET_XMIT_CN ? -ENOBUFFS : 0)
100
101 /* Driver transmit return codes */
102 #define NETDEV_TX_MASK 0xf0
103
104 enum netdev_tx {
105     __NETDEV_TX_MIN = INT_MIN, /* make sure enum is signed */
106     NETDEV_TX_OK = 0x00, /* driver took care of packet */
107     NETDEV_TX_BUSY = 0x10, /* driver tx path was busy */
108     NETDEV_TX_LOCKED = 0x20, /* driver tx lock was already taken */
109 };
110 typedef enum netdev_tx netdev_tx_t;
111
112 /*
113  * Current order: NETDEV_TX_MASK > NET_XMIT_MASK >= 0 is significant;
114  * hard_start_xmit() return < NET_XMIT_MASK means skb was consumed.
115  */
116 static inline bool dev_xmit_complete(int rc)
117 {
118     /*
119      * Positive cases with an skb consumed by a driver:
120      * - successful transmission (rc == NETDEV_TX_OK)
121      * - error while transmitting (rc < 0)
122      * - error while queueing to a different device (rc & NET_XMIT_MASK)
123      */
124     if (likely(rc < NET_XMIT_MASK))
125         return true;
126

```

```

127         return false;
128     }
129
130     /*
131     *      Compute the worst case header length according to the protocols
132     *      used.
133     */
134
135     #if defined(CONFIG_WLAN) || IS\_ENABLED(CONFIG_AX25)
136     # if defined(CONFIG_MAC80211_MESH)
137     #   define LL\_MAX\_HEADER 128
138     # else
139     #   define LL\_MAX\_HEADER 96
140     # endif
141     #else
142     # define LL\_MAX\_HEADER 32
143     #endif
144
145     #if !IS\_ENABLED(CONFIG_NET_IPIP) && !IS\_ENABLED(CONFIG_NET_IPGRE) && \
146         !IS\_ENABLED(CONFIG_IPV6_SIT) && !IS\_ENABLED(CONFIG_IPV6_TUNNEL)
147     #define MAX\_HEADER LL\_MAX\_HEADER
148     #else
149     #define MAX\_HEADER (LL\_MAX\_HEADER + 48)
150     #endif
151
152     /*
153     *      Old network device statistics. Fields are native words
154     *      (unsigned long) so they can be read and written atomically.
155     */
156
157     struct net\_device\_stats {
158         unsigned long    rx_packets;
159         unsigned long    tx_packets;
160         unsigned long    rx\_bytes;
161         unsigned long    tx\_bytes;
162         unsigned long    rx_errors;
163         unsigned long    tx_errors;
164         unsigned long    rx_dropped;
165         unsigned long    tx_dropped;
166         unsigned long    multicast;
167         unsigned long    collisions;
168         unsigned long    rx_length_errors;
169         unsigned long    rx_over_errors;
170         unsigned long    rx_crc_errors;
171         unsigned long    rx_frame_errors;
172         unsigned long    rx_fifo_errors;
173         unsigned long    rx_missed_errors;
174         unsigned long    tx_aborted_errors;
175         unsigned long    tx_carrier_errors;
176         unsigned long    tx_fifo_errors;
177         unsigned long    tx_heartbeat_errors;
178         unsigned long    tx_window_errors;
179         unsigned long    rx_compressed;
180         unsigned long    tx_compressed;
181     };
182
183
184     #include <linux/cache.h>
185     #include <linux/skbuff.h>
186
187     #ifdef CONFIG_RPS
188     #include <linux/static_key.h>
189     extern struct static\_key rps\_needed;
190     #endif
191
192     struct neighbour;
193     struct neigh\_parms;
194     struct sk\_buff;
195
196     struct netdev\_hw\_addr {
197         struct list\_head    list;
198         unsigned char        addr[MAX\_ADDR\_LEN];
199         unsigned char        type;

```

```

200 #define NETDEV_HW_ADDR_T_LAN 1
201 #define NETDEV_HW_ADDR_T_SAN 2
202 #define NETDEV_HW_ADDR_T_SLAVE 3
203 #define NETDEV_HW_ADDR_T_UNICAST 4
204 #define NETDEV_HW_ADDR_T_MULTICAST 5
205 bool global_use;
206 int sync_cnt;
207 int refcount;
208 int synced;
209 struct rcu_head rcu_head;
210 };
211
212 struct netdev_hw_addr_list {
213     struct list_head list;
214     int count;
215 };
216
217 #define netdev_hw_addr_list_count(l) ((l)->count)
218 #define netdev_hw_addr_list_empty(l) (netdev_hw_addr_list_count(l) == 0)
219 #define netdev_hw_addr_list_for_each(ha, l) \
220     list_for_each_entry(ha, &(l)->list, list)
221
222 #define netdev_uc_count(dev) netdev_hw_addr_list_count(&(dev)->uc)
223 #define netdev_uc_empty(dev) netdev_hw_addr_list_empty(&(dev)->uc)
224 #define netdev_for_each_uc_addr(ha, dev) \
225     netdev_hw_addr_list_for_each(ha, &(dev)->uc)
226
227 #define netdev_mc_count(dev) netdev_hw_addr_list_count(&(dev)->mc)
228 #define netdev_mc_empty(dev) netdev_hw_addr_list_empty(&(dev)->mc)
229 #define netdev_for_each_mc_addr(ha, dev) \
230     netdev_hw_addr_list_for_each(ha, &(dev)->mc)
231
232 struct hh_cache {
233     u16 hh_len;
234     u16 __pad;
235     seqlock_t hh_lock;
236
237     /* cached hardware header; allow for machine alignment needs. */
238 #define HH_DATA_MOD 16
239 #define HH_DATA_OFF(__len) \
240     (HH_DATA_MOD - (((__len) - 1) & (HH_DATA_MOD - 1)) + 1))
241 #define HH_DATA_ALIGN(__len) \
242     (((__len)+(HH_DATA_MOD-1))&~(HH_DATA_MOD - 1))
243     unsigned long hh_data[HH_DATA_ALIGN(LL_MAX_HEADER) / sizeof(long)];
244 };
245
246 /* Reserve HH_DATA_MOD byte aligned hard_header_len, but at least that much.
247  * Alternative is:
248  * dev->hard_header_len ? (dev->hard_header_len +
249  * (HH_DATA_MOD - 1)) & ~(HH_DATA_MOD - 1) : 0
250  *
251  * We could use other alignment values, but we must maintain the
252  * relationship HH alignment <= LL alignment.
253  */
254 #define LL_RESERVED_SPACE(dev) \
255     (((dev)->hard_header_len+(dev)->needed_headroom)&~(HH_DATA_MOD - 1)) + HH_DATA_MOD
256 #define LL_RESERVED_SPACE_EXTRA(dev,extra) \
257     (((dev)->hard_header_len+(dev)->needed_headroom+(extra))&~(HH_DATA_MOD - 1)) + HH_DATA_MOD
258
259 struct header_ops {
260     int (*create)(struct sk_buff *skb, struct net_device *dev,
261                  unsigned short type, const void *daddr,
262                  const void *saddr, unsigned int len);
263     int (*parse)(const struct sk_buff *skb, unsigned char *haddr);
264     int (*cache)(const struct neighbour *neigh, struct hh_cache *hh, __be16 type);
265     void (*cache_update)(struct hh_cache *hh,
266                          const struct net_device *dev,
267                          const unsigned char *haddr);
268 };
269
270 /* These flag bits are private to the generic network queueing
271  * layer, they may not be explicitly referenced by any other
272  * code.

```

```

273 */
274
275 enum netdev_state_t {
276     __LINK_STATE_START,
277     __LINK_STATE_PRESENT,
278     __LINK_STATE_NOCARRIER,
279     __LINK_STATE_LINKWATCH_PENDING,
280     __LINK_STATE_DORMANT,
281 };
282
283
284 /*
285  * This structure holds at boot time configured netdevice settings. They
286  * are then used in the device probing.
287  */
288 struct netdev_boot_setup {
289     char name[IFNAMSIZ];
290     struct ifmap map;
291 };
292 #define NETDEV_BOOT_SETUP_MAX 8
293
294 int __init netdev_boot_setup(char *str);
295
296 /*
297  * Structure for NAPI scheduling similar to tasklet but with weighting
298  */
299 struct napi_struct {
300     /* The poll_list must only be managed by the entity which
301      * changes the state of the NAPI_STATE_SCHED bit. This means
302      * whoever atomically sets that bit can add this napi_struct
303      * to the per-cpu poll_list, and whoever clears that bit
304      * can remove from the list right before clearing the bit.
305      */
306     struct list_head poll_list;
307
308     unsigned long state;
309     int weight;
310     unsigned int gro_count;
311     int (*poll)(struct napi_struct *, int);
312 #ifdef CONFIG_NETPOLL
313     spinlock_t poll_lock;
314     int poll_owner;
315 #endif
316     struct net_device *dev;
317     struct sk_buff *gro_list;
318     struct sk_buff *skb;
319     struct hrtimer timer;
320     struct list_head dev_list;
321     struct hlist_node napi_hash_node;
322     unsigned int napi_id;
323 };
324
325 enum {
326     NAPI_STATE_SCHED, /* Poll is scheduled */
327     NAPI_STATE_DISABLE, /* Disable pending */
328     NAPI_STATE_NPSVC, /* Netpoll - don't dequeue from poll_list */
329     NAPI_STATE_HASHED, /* In NAPI hash */
330 };
331
332 enum gro_result {
333     GRO_MERGED,
334     GRO_MERGED_FREE,
335     GRO_HELD,
336     GRO_NORMAL,
337     GRO_DROP,
338 };
339 typedef enum gro_result gro_result_t;
340
341 /*
342  * enum rx_handler_result - Possible return values for rx_handlers.
343  * @RX_HANDLER_CONSUMED: skb was consumed by rx_handler, do not process it
344  * further.
345  * @RX_HANDLER_ANOTHER: Do another round in receive path. This is indicated in

```

```

346 * case skb->dev was changed by rx_handler.
347 * @RX_HANDLER_EXACT: Force exact delivery, no wildcard.
348 * @RX_HANDLER_PASS: Do nothing, passe the skb as if no rx_handler was called.
349 *
350 * rx_handlers are functions called from inside __netif_receive_skb(), to do
351 * special processing of the skb, prior to delivery to protocol handlers.
352 *
353 * Currently, a net_device can only have a single rx_handler registered. Trying
354 * to register a second rx_handler will return -EBUSY.
355 *
356 * To register a rx_handler on a net_device, use netdev_rx_handler_register().
357 * To unregister a rx_handler on a net_device, use
358 * netdev_rx_handler_unregister().
359 *
360 * Upon return, rx_handler is expected to tell __netif_receive_skb() what to
361 * do with the skb.
362 *
363 * If the rx_handler consumed to skb in some way, it should return
364 * RX_HANDLER_CONSUMED. This is appropriate when the rx_handler arranged for
365 * the skb to be delivered in some other ways.
366 *
367 * If the rx_handler changed skb->dev, to divert the skb to another
368 * net_device, it should return RX_HANDLER_ANOTHER. The rx_handler for the
369 * new device will be called if it exists.
370 *
371 * If the rx_handler consider the skb should be ignored, it should return
372 * RX_HANDLER_EXACT. The skb will only be delivered to protocol handlers that
373 * are registered on exact device (ptype->dev == skb->dev).
374 *
375 * If the rx_handler didn't changed skb->dev, but want the skb to be normally
376 * delivered, it should return RX_HANDLER_PASS.
377 *
378 * A device without a registered rx_handler will behave as if rx_handler
379 * returned RX_HANDLER_PASS.
380 */
381
382 enum rx_handler_result {
383     RX_HANDLER_CONSUMED,
384     RX_HANDLER_ANOTHER,
385     RX_HANDLER_EXACT,
386     RX_HANDLER_PASS,
387 };
388 typedef enum rx_handler_result rx_handler_result_t;
389 typedef rx_handler_result_t rx_handler_func_t(struct sk_buff **pskb);
390
391 void __napi_schedule(struct napi_struct *n);
392 void __napi_schedule_irqoff(struct napi_struct *n);
393
394 static inline bool napi_disable_pending(struct napi_struct *n)
395 {
396     return test_bit(NAPI_STATE_DISABLE, &n->state);
397 }
398
399 /**
400  * napi_schedule_prep - check if napi can be scheduled
401  * @n: napi context
402  *
403  * Test if NAPI routine is already running, and if not mark
404  * it as running. This is used as a condition variable
405  * insure only one NAPI poll instance runs. We also make
406  * sure there is no pending NAPI disable.
407  */
408 static inline bool napi_schedule_prep(struct napi_struct *n)
409 {
410     return !napi_disable_pending(n) &&
411         !test_and_set_bit(NAPI_STATE_SCHED, &n->state);
412 }
413
414 /**
415  * napi_schedule - schedule NAPI poll
416  * @n: napi context
417  *
418  * Schedule NAPI poll routine to be called if it is not already

```

```

419 * running.
420 */
421 static inline void napi_schedule(struct napi_struct *n)
422 {
423     if (napi_schedule_prep(n))
424         __napi_schedule(n);
425 }
426
427 /**
428  * napi_schedule_irqoff - schedule NAPI poll
429  * @n: napi context
430  *
431  * Variant of napi_schedule(), assuming hard irqs are masked.
432  */
433 static inline void napi_schedule_irqoff(struct napi_struct *n)
434 {
435     if (napi_schedule_prep(n))
436         __napi_schedule_irqoff(n);
437 }
438
439 /* Try to reschedule poll. Called by dev->poll() after napi_complete(). */
440 static inline bool napi_reschedule(struct napi_struct *napi)
441 {
442     if (napi_schedule_prep(napi)) {
443         __napi_schedule(napi);
444         return true;
445     }
446     return false;
447 }
448
449 void __napi_complete(struct napi_struct *n);
450 void napi_complete_done(struct napi_struct *n, int work_done);
451 /**
452  * napi_complete - NAPI processing complete
453  * @n: napi context
454  *
455  * Mark NAPI processing as complete.
456  * Consider using napi_complete_done() instead.
457  */
458 static inline void napi_complete(struct napi_struct *n)
459 {
460     return napi_complete_done(n, 0);
461 }
462
463 /**
464  * napi_by_id - Lookup a NAPI by napi_id
465  * @napi_id: hashed napi_id
466  *
467  * Lookup @napi_id in napi_hash table
468  * must be called under rcu_read_lock()
469  */
470 struct napi_struct *napi_by_id(unsigned int napi_id);
471
472 /**
473  * napi_hash_add - add a NAPI to global hashtable
474  * @napi: napi context
475  *
476  * generate a new napi_id and store a @napi under it in napi_hash
477  */
478 void napi_hash_add(struct napi_struct *napi);
479
480 /**
481  * napi_hash_del - remove a NAPI from global table
482  * @napi: napi context
483  *
484  * Warning: caller must observe rcu grace period
485  * before freeing memory containing @napi
486  */
487 void napi_hash_del(struct napi_struct *napi);
488
489 /**
490  * napi_disable - prevent NAPI from scheduling
491  * @n: napi context

```



```

492 *
493 * Stop NAPI from being scheduled on this context.
494 * Waits till any outstanding processing completes.
495 */
496 void napi_disable(struct napi_struct *n);
497
498 /**
499 *      napi_enable - enable NAPI scheduling
500 *      @n: napi context
501 *
502 * Resume NAPI from being scheduled on this context.
503 * Must be paired with napi_disable.
504 */
505 static inline void napi_enable(struct napi_struct *n)
506 {
507     BUG_ON(!test_bit(NAPI_STATE_SCHED, &n->state));
508     smp_mb_before_atomic();
509     clear_bit(NAPI_STATE_SCHED, &n->state);
510 }
511
512 #ifdef CONFIG_SMP
513 /**
514 *      napi_synchronize - wait until NAPI is not running
515 *      @n: napi context
516 *
517 * Wait until NAPI is done being scheduled on this context.
518 * Waits till any outstanding processing completes but
519 * does not disable future activations.
520 */
521 static inline void napi_synchronize(const struct napi_struct *n)
522 {
523     while (test_bit(NAPI_STATE_SCHED, &n->state))
524         msleep(1);
525 }
526 #else
527 #define napi_synchronize(n)    barrier()
528 #endif
529
530 enum netdev_queue_state_t {
531     __QUEUE_STATE_DRV_XOFF,
532     __QUEUE_STATE_STACK_XOFF,
533     __QUEUE_STATE_FROZEN,
534 };
535
536 #define QUEUE_STATE_DRV_XOFF    (1 << __QUEUE_STATE_DRV_XOFF)
537 #define QUEUE_STATE_STACK_XOFF (1 << __QUEUE_STATE_STACK_XOFF)
538 #define QUEUE_STATE_FROZEN     (1 << __QUEUE_STATE_FROZEN)
539
540 #define QUEUE_STATE_ANY_XOFF    (QUEUE_STATE_DRV_XOFF | QUEUE_STATE_STACK_XOFF)
541 #define QUEUE_STATE_ANY_XOFF_OR_FROZEN (QUEUE_STATE_ANY_XOFF | \
542     QUEUE_STATE_FROZEN)
543 #define QUEUE_STATE_DRV_XOFF_OR_FROZEN (QUEUE_STATE_DRV_XOFF | \
544     QUEUE_STATE_FROZEN)
545
546 /**
547 * __QUEUE_STATE_DRV_XOFF is used by drivers to stop the transmit queue. The
548 * netif_tx_* functions below are used to manipulate this flag. The
549 * __QUEUE_STATE_STACK_XOFF flag is used by the stack to stop the transmit
550 * queue independently. The netif_xmit_*stopped functions below are called
551 * to check if the queue has been stopped by the driver or stack (either
552 * of the XOFF bits are set in the state). Drivers should not need to call
553 * netif_xmit*stopped functions, they should only be using netif_tx_*.
554 */
555
556 struct netdev_queue {
557     /*
558      * read mostly part
559      */
560     struct net_device *dev;
561     struct Qdisc __rcu *qdisc;
562     struct Qdisc *qdisc_sleeping;
563 #ifdef CONFIG_SYSFS
564     struct kobject kobj;

```



```

565 #endif
566 #if defined(CONFIG_XPS) && defined(CONFIG_NUMA)
567     int                numa_node;
568 #endif
569 /*
570  * write mostly part
571  */
572     spinlock_t          _xmit_lock    cacheline_aligned_in_smp;
573     int                xmit_lock_owner;
574     /*
575      * please use this field instead of dev->trans_start
576      */
577     unsigned long      trans_start;
578
579     /*
580      * Number of TX timeouts for this queue
581      * (/sys/class/net/DEV/Q/trans_timeout)
582      */
583     unsigned long      trans_timeout;
584
585     unsigned long      state;
586
587 #ifdef CONFIG_BQL
588     struct dql          dql;
589 #endif
590     unsigned long      tx_maxrate;
591 } cacheline_aligned_in_smp;
592
593 static inline int netdev_queue_numa_node_read(const struct netdev_queue *q)
594 {
595     #if defined(CONFIG_XPS) && defined(CONFIG_NUMA)
596         return q->numa_node;
597     #else
598         return NUMA_NO_NODE;
599     #endif
600 }
601
602 static inline void netdev_queue_numa_node_write(struct netdev_queue *q, int node)
603 {
604     #if defined(CONFIG_XPS) && defined(CONFIG_NUMA)
605         q->numa_node = node;
606     #endif
607 }
608
609 #ifdef CONFIG_RPS
610 /*
611  * This structure holds an RPS map which can be of variable length. The
612  * map is an array of CPUs.
613  */
614 struct rps_map {
615     unsigned int len;
616     struct rcu_head rcu;
617     u16 cpus[0];
618 };
619 #define RPS_MAP_SIZE(_num) (sizeof(struct rps_map) + ((_num) * sizeof(u16)))
620
621 /*
622  * The rps_dev_flow structure contains the mapping of a flow to a CPU, the
623  * tail pointer for that CPU's input queue at the time of last enqueue, and
624  * a hardware filter index.
625  */
626 struct rps_dev_flow {
627     u16 cpu;
628     u16 filter;
629     unsigned int last_qtail;
630 };
631 #define RPS_NO_FILTER 0xffff
632
633 /*
634  * The rps_dev_flow_table structure contains a table of flow mappings.
635  */
636 struct rps_dev_flow_table {
637     unsigned int mask;

```

```

638     struct rcu\_head rcu;
639     struct rps\_dev\_flow flows[0];
640 };
641 #define RPS\_DEV\_FLOW\_TABLE\_SIZE(_num) (sizeof(struct rps\_dev\_flow\_table) + \
642     ((_num) * sizeof(struct rps\_dev\_flow)))
643
644 /*
645  * The rps\_sock\_flow\_table contains mappings of flows to the last CPU
646  * on which they were processed by the application (set in recvmsg).
647  * Each entry is a 32bit value. Upper part is the high order bits
648  * of flow hash, lower part is cpu number.
649  * rps\_cpu\_mask is used to partition the space, depending on number of
650  * possible cpus : rps\_cpu\_mask = roundup\_pow\_of\_two(nr\_cpu\_ids) - 1
651  * For example, if 64 cpus are possible, rps\_cpu\_mask = 0x3f,
652  * meaning we use 32-6=26 bits for the hash.
653  */
654 struct rps\_sock\_flow\_table {
655     u32     mask;
656
657     u32     ents[0] \_\_cacheline\_aligned\_in\_smp;
658 };
659 #define RPS\_SOCKET\_FLOW\_TABLE\_SIZE(_num) (offsetof(struct rps\_sock\_flow\_table, ents[_num]))
660
661 #define RPS\_NO\_CPU 0xffff
662
663 extern u32 rps\_cpu\_mask;
664 extern struct rps\_sock\_flow\_table \_\_rcu *rps\_sock\_flow\_table;
665
666 static inline void rps\_record\_sock\_flow(struct rps\_sock\_flow\_table *table,
667     u32 hash)
668 {
669     if (table && hash) {
670         unsigned int index = hash & table->mask;
671         u32 val = hash & ~rps\_cpu\_mask;
672
673         /* We only give a hint, preemption can change cpu under us */
674         val |= raw\_smp\_processor\_id();
675
676         if (table->ents[index] != val)
677             table->ents[index] = val;
678     }
679 }
680
681 #ifdef CONFIG_RFS_ACCEL
682 bool rps\_may\_expire\_flow(struct net\_device *dev, u16 rxq_index, u32 flow_id,
683     u16 filter_id);
684 #endif
685 #endif /* CONFIG_RPS */
686
687 /* This structure contains an instance of an RX queue. */
688 struct netdev\_rx\_queue {
689     #ifdef CONFIG_RPS
690         struct rps\_map \_\_rcu *rps\_map;
691         struct rps\_dev\_flow\_table \_\_rcu *rps\_flow\_table;
692     #endif
693     struct kobject kobj;
694     struct net\_device *dev;
695 } \_\_cacheline\_aligned\_in\_smp;
696
697 /*
698  * RX queue sysfs structures and functions.
699  */
700 struct rx\_queue\_attribute {
701     struct attribute attr;
702     ssize\_t (*show)(struct netdev\_rx\_queue *queue,
703         struct rx\_queue\_attribute *attr, char *buf);
704     ssize\_t (*store)(struct netdev\_rx\_queue *queue,
705         struct rx\_queue\_attribute *attr, const char *buf, size\_t len);
706 };
707
708 #ifdef CONFIG_XPS
709 /*
710  * This structure holds an XPS map which can be of variable length. The

```

```

711  * map is an array of queues.
712  */
713  struct xps_map {
714      unsigned int len;
715      unsigned int alloc_len;
716      struct rcu_head rcu;
717      u16 queues[0];
718  };
719  #define XPS_MAP_SIZE(_num) (sizeof(struct xps_map) + ((_num) * sizeof(u16)))
720  #define XPS_MIN_MAP_ALLOC ((L1_CACHE_BYTES - sizeof(struct xps_map)) \
721      / sizeof(u16))
722
723  /*
724   * This structure holds all XPS maps for device. Maps are indexed by CPU.
725   */
726  struct xps_dev_maps {
727      struct rcu_head rcu;
728      struct xps_map __rcu *cpu_map[0];
729  };
730  #define XPS_DEV_MAPS_SIZE (sizeof(struct xps_dev_maps) + \
731      (nr_cpu_ids * sizeof(struct xps_map *)))
732  #endif /* CONFIG_XPS */
733
734  #define TC_MAX_QUEUE 16
735  #define TC_BITMASK 15
736  /* HW offloaded queuing disciplines txq count and offset maps */
737  struct netdev_tc_txq {
738      u16 count;
739      u16 offset;
740  };
741
742  #if defined(CONFIG_FCOE) || defined(CONFIG_FCOE_MODULE)
743  /*
744   * This structure is to hold information about the device
745   * configured to run FCoE protocol stack.
746   */
747  struct netdev_fcoe_hbainfo {
748      char manufacturer[64];
749      char serial_number[64];
750      char hardware_version[64];
751      char driver_version[64];
752      char optionrom_version[64];
753      char firmware_version[64];
754      char model[256];
755      char model_description[256];
756  };
757  #endif
758
759  #define MAX_PHYS_ITEM_ID_LEN 32
760
761  /* This structure holds a unique identifier to identify some
762   * physical item (port for example) used by a netdevice.
763   */
764  struct netdev_phys_item_id {
765      unsigned char id[MAX_PHYS_ITEM_ID_LEN];
766      unsigned char id_len;
767  };
768
769  typedef u16 (*select_queue_fallback_t)(struct net_device *dev,
770      struct sk_buff *skb);
771
772  /*
773   * This structure defines the management hooks for network devices.
774   * The following hooks can be defined; unless noted otherwise, they are
775   * optional and can be filled with a null pointer.
776   *
777   * int (*ndo_init)(struct net_device *dev);
778   * This function is called once when network device is registered.
779   * The network device can use this to any late stage initializaton
780   * or semantic validation. It can fail with an error code which will
781   * be propagated back to register_netdev
782   *
783   * void (*ndo_uninit)(struct net_device *dev);

```

```

784 *      This function is called when device is unregistered or when registration
785 *      fails. It is not called if init fails.
786 *
787 * int (*ndo_open)(struct net_device *dev);
788 *      This function is called when network device transistions to the up
789 *      state.
790 *
791 * int (*ndo_stop)(struct net_device *dev);
792 *      This function is called when network device transistions to the down
793 *      state.
794 *
795 * netdev_tx_t (*ndo_start_xmit)(struct sk_buff *skb,
796 *                               struct net_device *dev);
797 *      Called when a packet needs to be transmitted.
798 *      Returns NETDEV_TX_OK. Can return NETDEV_TX_BUSY, but you should stop
799 *      the queue before that can happen; it's for obsolete devices and weird
800 *      corner cases, but the stack really does a non-trivial amount
801 *      of useless work if you return NETDEV_TX_BUSY.
802 *      (can also return NETDEV_TX_LOCKED iff NETIF_F_LLTX)
803 *      Required can not be NULL.
804 *
805 * u16 (*ndo_select_queue)(struct net_device *dev, struct sk_buff *skb,
806 *                        void *accel_priv, select_queue_fallback_t fallback);
807 *      Called to decide which queue to when device supports multiple
808 *      transmit queues.
809 *
810 * void (*ndo_change_rx_flags)(struct net_device *dev, int flags);
811 *      This function is called to allow device receiver to make
812 *      changes to configuration when multicast or promiscuous is enabled.
813 *
814 * void (*ndo_set_rx_mode)(struct net_device *dev);
815 *      This function is called device changes address list filtering.
816 *      If driver handles unicast address filtering, it should set
817 *      IFF_UNICAST_FLT to its priv_flags.
818 *
819 * int (*ndo_set_mac_address)(struct net_device *dev, void *addr);
820 *      This function is called when the Media Access Control address
821 *      needs to be changed. If this interface is not defined, the
822 *      mac address can not be changed.
823 *
824 * int (*ndo_validate_addr)(struct net_device *dev);
825 *      Test if Media Access Control address is valid for the device.
826 *
827 * int (*ndo_do_ioctl)(struct net_device *dev, struct ifreq *ifr, int cmd);
828 *      Called when a user request an ioctl which can't be handled by
829 *      the generic interface code. If not defined ioctl's return
830 *      not supported error code.
831 *
832 * int (*ndo_set_config)(struct net_device *dev, struct ifmap *map);
833 *      Used to set network devices bus interface parameters. This interface
834 *      is retained for legacy reason, new devices should use the bus
835 *      interface (PCI) for Low Level management.
836 *
837 * int (*ndo_change_mtu)(struct net_device *dev, int new_mtu);
838 *      Called when a user wants to change the Maximum Transfer Unit
839 *      of a device. If not defined, any request to change MTU will
840 *      will return an error.
841 *
842 * void (*ndo_tx_timeout)(struct net_device *dev);
843 *      Callback uses when the transmitter has not made any progress
844 *      for dev->watchdog ticks.
845 *
846 * struct rtnl_link_stats64* (*ndo_get_stats64)(struct net_device *dev,
847 *                                              struct rtnl_link_stats64 *storage);
848 * struct net_device_stats* (*ndo_get_stats)(struct net_device *dev);
849 *      Called when a user wants to get the network device usage
850 *      statistics. Drivers must do one of the following:
851 *      1. Define @ndo_get_stats64 to fill in a zero-initialised
852 *         rtnl_link_stats64 structure passed by the caller.
853 *      2. Define @ndo_get_stats to update a net_device_stats structure
854 *         (which should normally be dev->stats) and return a pointer to
855 *         it. The structure may be changed asynchronously only if each
856 *         field is written atomically.

```

```

857 *      3. Update dev->stats asynchronously and atomically, and define
858 *      neither operation.
859 *
860 * int (*ndo_vlan_rx_add_vid)(struct net_device *dev, __be16 proto, u16 vid);
861 *      If device support VLAN filtering this function is called when a
862 *      VLAN id is registered.
863 *
864 * int (*ndo_vlan_rx_kill_vid)(struct net_device *dev, __be16 proto, u16 vid);
865 *      If device support VLAN filtering this function is called when a
866 *      VLAN id is unregistered.
867 *
868 * void (*ndo_poll_controller)(struct net_device *dev);
869 *
870 *      SR-IOV management functions.
871 * int (*ndo_set_vf_mac)(struct net_device *dev, int vf, u8* mac);
872 * int (*ndo_set_vf_vlan)(struct net_device *dev, int vf, u16 vlan, u8 qos);
873 * int (*ndo_set_vf_rate)(struct net_device *dev, int vf, int min_tx_rate,
874 *      int max_tx_rate);
875 * int (*ndo_set_vf_spoofchk)(struct net_device *dev, int vf, bool setting);
876 * int (*ndo_get_vf_config)(struct net_device *dev,
877 *      int vf, struct ifla_vf_info *ivf);
878 * int (*ndo_set_vf_link_state)(struct net_device *dev, int vf, int link_state);
879 * int (*ndo_set_vf_port)(struct net_device *dev, int vf,
880 *      struct nlattr *port[]);
881 *
882 *      Enable or disable the VF ability to query its RSS Redirection Table and
883 *      Hash Key. This is needed since on some devices VF share this information
884 *      with PF and querying it may adduce a theoretical security risk.
885 * int (*ndo_set_vf_rss_query_en)(struct net_device *dev, int vf, bool setting);
886 * int (*ndo_get_vf_port)(struct net_device *dev, int vf, struct sk_buff *skb);
887 * int (*ndo_setup_tc)(struct net_device *dev, u8 tc)
888 *      Called to setup 'tc' number of traffic classes in the net device. This
889 *      is always called from the stack with the rtnl lock held and netif tx
890 *      queues stopped. This allows the netdevice to perform queue management
891 *      safely.
892 *
893 *      Fiber Channel over Ethernet (FCoE) offload functions.
894 * int (*ndo_fcoe_enable)(struct net_device *dev);
895 *      Called when the FCoE protocol stack wants to start using LLD for FCoE
896 *      so the underlying device can perform whatever needed configuration or
897 *      initialization to support acceleration of FCoE traffic.
898 *
899 * int (*ndo_fcoe_disable)(struct net_device *dev);
900 *      Called when the FCoE protocol stack wants to stop using LLD for FCoE
901 *      so the underlying device can perform whatever needed clean-ups to
902 *      stop supporting acceleration of FCoE traffic.
903 *
904 * int (*ndo_fcoe_ddp_setup)(struct net_device *dev, u16 xid,
905 *      struct scatterlist *sgl, unsigned int sgc);
906 *      Called when the FCoE Initiator wants to initialize an I/O that
907 *      is a possible candidate for Direct Data Placement (DDP). The LLD can
908 *      perform necessary setup and returns 1 to indicate the device is set up
909 *      successfully to perform DDP on this I/O, otherwise this returns 0.
910 *
911 * int (*ndo_fcoe_ddp_done)(struct net_device *dev, u16 xid);
912 *      Called when the FCoE Initiator/Target is done with the DDPed I/O as
913 *      indicated by the FC exchange id 'xid', so the underlying device can
914 *      clean up and reuse resources for later DDP requests.
915 *
916 * int (*ndo_fcoe_ddp_target)(struct net_device *dev, u16 xid,
917 *      struct scatterlist *sgl, unsigned int sgc);
918 *      Called when the FCoE Target wants to initialize an I/O that
919 *      is a possible candidate for Direct Data Placement (DDP). The LLD can
920 *      perform necessary setup and returns 1 to indicate the device is set up
921 *      successfully to perform DDP on this I/O, otherwise this returns 0.
922 *
923 * int (*ndo_fcoe_get_hbainfo)(struct net_device *dev,
924 *      struct netdev_fcoe_hbainfo *hbainfo);
925 *      Called when the FCoE Protocol stack wants information on the underlying
926 *      device. This information is utilized by the FCoE protocol stack to
927 *      register attributes with Fiber Channel management service as per the
928 *      FC-GS Fabric Device Management Information(FDMI) specification.
929 *

```



```

930 * int (*ndo_fcoe_get_wwn)(struct net_device *dev, u64 *wwn, int type);
931 *   Called when the underlying device wants to override default World Wide
932 *   Name (WWN) generation mechanism in FCoE protocol stack to pass its own
933 *   World Wide Port Name (WWPN) or World Wide Node Name (WWNN) to the FCoE
934 *   protocol stack to use.
935 *
936 *   RFS acceleration.
937 * int (*ndo_rx_flow_steering)(struct net_device *dev, const struct sk_buff *skb,
938 *   u16 rxq_index, u32 flow_id);
939 *   Set hardware filter for RFS. rxq_index is the target queue index;
940 *   flow_id is a flow ID to be passed to rps_may_expire_flow() later.
941 *   Return the filter ID on success, or a negative error code.
942 *
943 *   Slave management functions (for bridge, bonding, etc).
944 * int (*ndo_add_slave)(struct net_device *dev, struct net_device *slave_dev);
945 *   Called to make another netdev an underling.
946 *
947 * int (*ndo_del_slave)(struct net_device *dev, struct net_device *slave_dev);
948 *   Called to release previously enslaved netdev.
949 *
950 *   Feature/offload setting functions.
951 * netdev_features_t (*ndo_fix_features)(struct net_device *dev,
952 *   netdev_features_t features);
953 *   Adjusts the requested feature flags according to device-specific
954 *   constraints, and returns the resulting flags. Must not modify
955 *   the device state.
956 *
957 * int (*ndo_set_features)(struct net_device *dev, netdev_features_t features);
958 *   Called to update device configuration to new features. Passed
959 *   feature set might be less than what was returned by ndo_fix_features().
960 *   Must return >0 or -errno if it changed dev->features itself.
961 *
962 * int (*ndo_fdb_add)(struct ndmsg *ndm, struct nlattr *tb[],
963 *   struct net_device *dev,
964 *   const unsigned char *addr, u16 vid, u16 flags)
965 *   Adds an FDB entry to dev for addr.
966 * int (*ndo_fdb_del)(struct ndmsg *ndm, struct nlattr *tb[],
967 *   struct net_device *dev,
968 *   const unsigned char *addr, u16 vid)
969 *   Deletes the FDB entry from dev corresponding to addr.
970 * int (*ndo_fdb_dump)(struct sk_buff *skb, struct netlink_callback *cb,
971 *   struct net_device *dev, struct net_device *filter_dev,
972 *   int idx)
973 *   Used to add FDB entries to dump requests. Implementers should add
974 *   entries to skb and update idx with the number of entries.
975 *
976 * int (*ndo_bridge_setlink)(struct net_device *dev, struct nlmsghdr *nlh,
977 *   u16 flags)
978 * int (*ndo_bridge_getlink)(struct sk_buff *skb, u32 pid, u32 seq,
979 *   struct net_device *dev, u32 filter_mask,
980 *   int nlflags)
981 * int (*ndo_bridge_dellink)(struct net_device *dev, struct nlmsghdr *nlh,
982 *   u16 flags);
983 *
984 * int (*ndo_change_carrier)(struct net_device *dev, bool new_carrier);
985 *   Called to change device carrier. Soft-devices (like dummy, team, etc)
986 *   which do not represent real hardware may define this to allow their
987 *   userspace components to manage their virtual carrier state. Devices
988 *   that determine carrier state from physical hardware properties (eg
989 *   network cables) or protocol-dependent mechanisms (eg
990 *   USB_CDC_NOTIFY_NETWORK_CONNECTION) should NOT implement this function.
991 *
992 * int (*ndo_get_phys_port_id)(struct net_device *dev,
993 *   struct netdev_phys_item_id *ppid);
994 *   Called to get ID of physical port of this device. If driver does
995 *   not implement this, it is assumed that the hw is not able to have
996 *   multiple net devices on single physical port.
997 *
998 * void (*ndo_add_vxlan_port)(struct net_device *dev,
999 *   sa_family_t sa_family, __be16 port);
1000 *   Called by vxlan to notify a driver about the UDP port and socket
1001 *   address family that vxlan is listening to. It is called only when
1002 *   a new port starts listening. The operation is protected by the

```

```

1003 *      vxlan_net->sock_lock.
1004 *
1005 * void (*ndo_del_vxlan_port)(struct net_device *dev,
1006 *      sa_family_t sa_family, __be16 port);
1007 *      Called by vxlan to notify the driver about a UDP port and socket
1008 *      address family that vxlan is not listening to anymore. The operation
1009 *      is protected by the vxlan_net->sock_lock.
1010 *
1011 * void* (*ndo_dfwd_add_station)(struct net_device *pdev,
1012 *      struct net_device *dev)
1013 *      Called by upper layer devices to accelerate switching or other
1014 *      station functionality into hardware. 'pdev' is the lowerdev
1015 *      to use for the offload and 'dev' is the net device that will
1016 *      back the offload. Returns a pointer to the private structure
1017 *      the upper layer will maintain.
1018 * void (*ndo_dfwd_del_station)(struct net_device *pdev, void *priv)
1019 *      Called by upper layer device to delete the station created
1020 *      by 'ndo_dfwd_add_station'. 'pdev' is the net device backing
1021 *      the station and priv is the structure returned by the add
1022 *      operation.
1023 * netdev_tx_t (*ndo_dfwd_start_xmit)(struct sk_buff *skb,
1024 *      struct net_device *dev,
1025 *      void *priv);
1026 *      Callback to use for xmit over the accelerated station. This
1027 *      is used in place of ndo_start_xmit on accelerated net
1028 *      devices.
1029 * netdev_features_t (*ndo_features_check) (struct sk_buff *skb,
1030 *      struct net_device *dev
1031 *      netdev_features_t features);
1032 *      Called by core transmit path to determine if device is capable of
1033 *      performing offload operations on a given packet. This is to give
1034 *      the device an opportunity to implement any restrictions that cannot
1035 *      be otherwise expressed by feature flags. The check is called with
1036 *      the set of features that the stack has calculated and it returns
1037 *      those the driver believes to be appropriate.
1038 * int (*ndo_set_tx_maxrate)(struct net_device *dev,
1039 *      int queue_index, u32 maxrate);
1040 *      Called when a user wants to set a max-rate limitation of specific
1041 *      TX queue.
1042 * int (*ndo_get_iflink)(const struct net_device *dev);
1043 *      Called to get the iflink value of this device.
1044 */
1045 struct net_device_ops {
1046     int      (*ndo_init)(struct net_device *dev);
1047     void      (*ndo_uninit)(struct net_device *dev);
1048     int      (*ndo_open)(struct net_device *dev);
1049     int      (*ndo_stop)(struct net_device *dev);
1050     netdev_tx_t (*ndo_start_xmit) (struct sk_buff *skb,
1051     struct net_device *dev);
1052     u16      (*ndo_select_queue)(struct net_device *dev,
1053     struct sk_buff *skb,
1054     void *accel_priv,
1055     select_queue_fallback_t fallback);
1056     void      (*ndo_change_rx_flags)(struct net_device *dev,
1057     int flags);
1058     void      (*ndo_set_rx_mode)(struct net_device *dev);
1059     int      (*ndo_set_mac_address)(struct net_device *dev,
1060     void *addr);
1061     int      (*ndo_validate_addr)(struct net_device *dev);
1062     int      (*ndo_do_ioctl)(struct net_device *dev,
1063     struct ifreq *ifr, int cmd);
1064     int      (*ndo_set_config)(struct net_device *dev,
1065     struct ifmap *map);
1066     int      (*ndo_change_mtu)(struct net_device *dev,
1067     int new_mtu);
1068     int      (*ndo_neigh_setup)(struct net_device *dev,
1069     struct neigh_parms *);
1070     void      (*ndo_tx_timeout) (struct net_device *dev);
1071
1072     struct rtnl_link_stats64* (*ndo_get_stats64)(struct net_device *dev,
1073     struct rtnl_link_stats64 *storage);
1074     struct net_device_stats* (*ndo_get_stats)(struct net_device *dev);
1075

```



```

1076         int                                (*ndo_vlan_rx_add_vid)(struct net\_device *dev,
1077         \_\_be16 proto, u16 vid);
1078         int                                (*ndo_vlan_rx_kill_vid)(struct net\_device *dev,
1079         \_\_be16 proto, u16 vid);
1080 #ifdef CONFIG_NET_POLL_CONTROLLER
1081         void                                (*ndo_poll_controller)(struct net\_device *dev);
1082         int                                (*ndo_netpoll_setup)(struct net\_device *dev,
1083         struct netpoll\_info *info);
1084         void                                (*ndo_netpoll_cleanup)(struct net\_device *dev);
1085 #endif
1086 #ifdef CONFIG_NET_RX_BUSY_POLL
1087         int                                (*ndo_busy_poll)(struct napi\_struct *dev);
1088 #endif
1089         int                                (*ndo_set_vf_mac)(struct net\_device *dev,
1090         int queue, u8 *mac);
1091         int                                (*ndo_set_vf_vlan)(struct net\_device *dev,
1092         int queue, u16 vlan, u8 qos);
1093         int                                (*ndo_set_vf_rate)(struct net\_device *dev,
1094         int vf, int min_tx_rate,
1095         int max_tx_rate);
1096         int                                (*ndo_set_vf_spoofchk)(struct net\_device *dev,
1097         int vf, bool setting);
1098         int                                (*ndo_get_vf_config)(struct net\_device *dev,
1099         int vf,
1100         struct ifla\_vf\_info *ivf);
1101         int                                (*ndo_set_vf_link_state)(struct net\_device *dev,
1102         int vf, int link\_state);
1103         int                                (*ndo_get_vf_stats)(struct net\_device *dev,
1104         int vf,
1105         struct ifla\_vf\_stats
1106         *vf_stats);
1107         int                                (*ndo_set_vf_port)(struct net\_device *dev,
1108         int vf,
1109         struct nlattr *port[]);
1110         int                                (*ndo_get_vf_port)(struct net\_device *dev,
1111         int vf, struct sk\_buff *skb);
1112         int                                (*ndo_set_vf_rss_query_en)(
1113         struct net\_device *dev,
1114         int vf, bool setting);
1115         int                                (*ndo_setup_tc)(struct net\_device *dev, u8 tc);
1116 #if IS\_ENABLED(CONFIG_FCOE)
1117         int                                (*ndo_fcoe_enable)(struct net\_device *dev);
1118         int                                (*ndo_fcoe_disable)(struct net\_device *dev);
1119         int                                (*ndo_fcoe_ddp_setup)(struct net\_device *dev,
1120         u16 xid,
1121         struct scatterlist *sgl,
1122         unsigned int sgc);
1123         int                                (*ndo_fcoe_ddp_done)(struct net\_device *dev,
1124         u16 xid);
1125         int                                (*ndo_fcoe_ddp_target)(struct net\_device *dev,
1126         u16 xid,
1127         struct scatterlist *sgl,
1128         unsigned int sgc);
1129         int                                (*ndo_fcoe_get_hbainfo)(struct net\_device *dev,
1130         struct netdev\_fcoe\_hbainfo *hbainfo);
1131 #endif
1132 #if IS\_ENABLED(CONFIG_LIBFCOE)
1133 #define NETDEV\_FCOE\_WWNN 0
1134 #define NETDEV\_FCOE\_WWPN 1
1135         int                                (*ndo_fcoe_get_wnn)(struct net\_device *dev,
1136         u64 *wnn, int type);
1137 #endif
1138 #ifdef CONFIG_RFS_ACCEL
1139         int                                (*ndo_rx_flow_steer)(struct net\_device *dev,
1140         const struct sk\_buff *skb,
1141         u16 rxq_index,
1142         u32 flow_id);
1143 #endif
1144         int                                (*ndo_add_slave)(struct net\_device *dev,
1145         struct net\_device *slave_dev);
1146         int                                (*ndo_del_slave)(struct net\_device *dev,

```

```

1149         struct net\_device *slave_dev);
1150     netdev\_features\_t (*ndo_fix_features)(struct net\_device *dev,
1151         netdev\_features\_t features);
1152     int (*ndo_set_features)(struct net\_device *dev,
1153         netdev\_features\_t features);
1154     int (*ndo_neigh_construct)(struct neighbour *n);
1155     void (*ndo_neigh_destroy)(struct neighbour *n);
1156
1157     int (*ndo_fdb_add)(struct ndmsg *ndm,
1158         struct nlattr *tb[],
1159         struct net\_device *dev,
1160         const unsigned char *addr,
1161         u16 vid,
1162         u16 flags);
1163     int (*ndo_fdb_del)(struct ndmsg *ndm,
1164         struct nlattr *tb[],
1165         struct net\_device *dev,
1166         const unsigned char *addr,
1167         u16 vid);
1168     int (*ndo_fdb_dump)(struct sk\_buff *skb,
1169         struct netlink\_callback *cb,
1170         struct net\_device *dev,
1171         struct net\_device *filter_dev,
1172         int idx);
1173
1174     int (*ndo_bridge_setlink)(struct net\_device *dev,
1175         struct nlmsg\_hdr *nlh,
1176         u16 flags);
1177     int (*ndo_bridge_getlink)(struct sk\_buff *skb,
1178         u32 pid, u32 seq,
1179         struct net\_device *dev,
1180         u32 filter_mask,
1181         int nlflds);
1182     int (*ndo_bridge_dellink)(struct net\_device *dev,
1183         struct nlmsg\_hdr *nlh,
1184         u16 flags);
1185     int (*ndo_change_carrier)(struct net\_device *dev,
1186         bool new_carrier);
1187     int (*ndo_get_phys_port_id)(struct net\_device *dev,
1188         struct netdev\_phys\_item\_id *ppid);
1189     int (*ndo_get_phys_port_name)(struct net\_device *dev,
1190         char *name, size\_t len);
1191     void (*ndo_add_vxlan_port)(struct net\_device *dev,
1192         sa\_family\_t sa_family,
1193         \_\_be16 port);
1194     void (*ndo_del_vxlan_port)(struct net\_device *dev,
1195         sa\_family\_t sa_family,
1196         \_\_be16 port);
1197
1198     void* (*ndo_dfwd_add_station)(struct net\_device *pdev,
1199         struct net\_device *dev);
1200     void (*ndo_dfwd_del_station)(struct net\_device *pdev,
1201         void *priv);
1202
1203     netdev\_tx\_t (*ndo_dfwd_start_xmit) (struct sk\_buff *skb,
1204         struct net\_device *dev,
1205         void *priv);
1206     int (*ndo_get_lock_subclass)(struct net\_device *dev);
1207     netdev\_features\_t (*ndo_features_check) (struct sk\_buff *skb,
1208         struct net\_device *dev,
1209         netdev\_features\_t features);
1210     int (*ndo_set_tx_maxrate)(struct net\_device *dev,
1211         int queue\_index,
1212         u32 maxrate);
1213     int (*ndo_get_iflink)(const struct net\_device *dev);
1214 };
1215
1216 /**
1217  * enum net_device_priv_flags - &struct net_device priv_flags
1218  *
1219  * These are the &struct net_device, they are only set internally
1220  * by drivers and used in the kernel. These flags are invisible to
1221  * userspace, this means that the order of these flags can change

```

```

1222 * during any kernel release.
1223 *
1224 * You should have a pretty good reason to be extending these flags.
1225 *
1226 * @IFF_802_1Q_VLAN: 802.1Q VLAN device
1227 * @IFF_EBRIDGE: Ethernet bridging device
1228 * @IFF_SLAVE_INACTIVE: bonding slave not the curr. active
1229 * @IFF_MASTER_8023AD: bonding master, 802.3ad
1230 * @IFF_MASTER_ALB: bonding master, balance-alb
1231 * @IFF_BONDING: bonding master or slave
1232 * @IFF_SLAVE_NEEDARP: need ARPs for validation
1233 * @IFF_ISATAP: ISATAP interface (RFC4214)
1234 * @IFF_MASTER_ARPMON: bonding master, ARP mon in use
1235 * @IFF_WAN_HDLC: WAN HDLC device
1236 * @IFF_XMIT_DST_RELEASE: dev_hard_start_xmit() is allowed to
1237 *     release skb->dst
1238 * @IFF_DONT_BRIDGE: disallow bridging this ether dev
1239 * @IFF_DISABLE_NETPOLL: disable netpoll at run-time
1240 * @IFF_MACVLAN_PORT: device used as macvlan port
1241 * @IFF_BRIDGE_PORT: device used as bridge port
1242 * @IFF_OVS_DATAPATH: device used as Open vSwitch datapath port
1243 * @IFF_TX_SKB_SHARING: The interface supports sharing skbs on transmit
1244 * @IFF_UNICAST_FLT: Supports unicast filtering
1245 * @IFF_TEAM_PORT: device used as team port
1246 * @IFF_SUPP_NOFCS: device supports sending custom FCS
1247 * @IFF_LIVE_ADDR_CHANGE: device supports hardware address
1248 *     change when it's running
1249 * @IFF_MACVLAN: Macvlan device
1250 */
1251 enum netdev_priv_flags {
1252     IFF_802_1Q_VLAN           = 1<<0,
1253     IFF_EBRIDGE               = 1<<1,
1254     IFF_SLAVE_INACTIVE        = 1<<2,
1255     IFF_MASTER_8023AD         = 1<<3,
1256     IFF_MASTER_ALB            = 1<<4,
1257     IFF_BONDING                = 1<<5,
1258     IFF_SLAVE_NEEDARP         = 1<<6,
1259     IFF_ISATAP                 = 1<<7,
1260     IFF_MASTER_ARPMON         = 1<<8,
1261     IFF_WAN_HDLC               = 1<<9,
1262     IFF_XMIT_DST_RELEASE       = 1<<10,
1263     IFF_DONT_BRIDGE            = 1<<11,
1264     IFF_DISABLE_NETPOLL        = 1<<12,
1265     IFF_MACVLAN_PORT           = 1<<13,
1266     IFF_BRIDGE_PORT            = 1<<14,
1267     IFF_OVS_DATAPATH           = 1<<15,
1268     IFF_TX_SKB_SHARING         = 1<<16,
1269     IFF_UNICAST_FLT            = 1<<17,
1270     IFF_TEAM_PORT              = 1<<18,
1271     IFF_SUPP_NOFCS              = 1<<19,
1272     IFF_LIVE_ADDR_CHANGE        = 1<<20,
1273     IFF_MACVLAN                 = 1<<21,
1274     IFF_XMIT_DST_RELEASE_PERM   = 1<<22,
1275     IFF_IPVLAN_MASTER           = 1<<23,
1276     IFF_IPVLAN_SLAVE           = 1<<24,
1277 };
1278
1279 #define IFF_802_1Q_VLAN       IFF_802_1Q_VLAN
1280 #define IFF_EBRIDGE           IFF_EBRIDGE
1281 #define IFF_SLAVE_INACTIVE     IFF_SLAVE_INACTIVE
1282 #define IFF_MASTER_8023AD      IFF_MASTER_8023AD
1283 #define IFF_MASTER_ALB         IFF_MASTER_ALB
1284 #define IFF_BONDING             IFF_BONDING
1285 #define IFF_SLAVE_NEEDARP      IFF_SLAVE_NEEDARP
1286 #define IFF_ISATAP              IFF_ISATAP
1287 #define IFF_MASTER_ARPMON       IFF_MASTER_ARPMON
1288 #define IFF_WAN_HDLC            IFF_WAN_HDLC
1289 #define IFF_XMIT_DST_RELEASE    IFF_XMIT_DST_RELEASE
1290 #define IFF_DONT_BRIDGE         IFF_DONT_BRIDGE
1291 #define IFF_DISABLE_NETPOLL     IFF_DISABLE_NETPOLL
1292 #define IFF_MACVLAN_PORT        IFF_MACVLAN_PORT
1293 #define IFF_BRIDGE_PORT         IFF_BRIDGE_PORT
1294 #define IFF_OVS_DATAPATH        IFF_OVS_DATAPATH

```

```

1295 #define IFF_TX_SKB_SHARING IFF_TX_SKB_SHARING
1296 #define IFF_UNICAST_FLT IFF_UNICAST_FLT
1297 #define IFF_TEAM_PORT IFF_TEAM_PORT
1298 #define IFF_SUPP_NOFCS IFF_SUPP_NOFCS
1299 #define IFF_LIVE_ADDR_CHANGE IFF_LIVE_ADDR_CHANGE
1300 #define IFF_MACVLAN IFF_MACVLAN
1301 #define IFF_XMIT_DST_RELEASE_PERM IFF_XMIT_DST_RELEASE_PERM
1302 #define IFF_IPVLAN_MASTER IFF_IPVLAN_MASTER
1303 #define IFF_IPVLAN_SLAVE IFF_IPVLAN_SLAVE
1304
1305 /**
1306  * struct net_device - The DEVICE structure.
1307  *
1308  * Actually, this whole structure is a big mistake. It mixes I/O
1309  * data with strictly "high-level" data, and it has to know about
1310  * almost every data structure used in the INET module.
1311  *
1312  * @name: This is the first field of the "visible" part of this structure
1313  * (i.e. as seen by users in the "Space.c" file). It is the name
1314  * of the interface.
1315  *
1316  * @name_hlist: Device name hash chain, please keep it close to name[]
1317  *
1318  * @ifalias: SNMP alias
1319  *
1320  * @mem_end: Shared memory end
1321  *
1322  * @mem_start: Shared memory start
1323  *
1324  * @base_addr: Device I/O address
1325  *
1326  * @irq: Device IRQ number
1327  *
1328  * @carrier_changes: Stats to monitor carrier on<->off transitions
1329  *
1330  * @state: Generic network queuing layer state, see netdev_state_t
1331  *
1332  * @dev_list: The global list of network devices
1333  *
1334  * @napi_list: List entry, that is used for polling napi devices
1335  *
1336  * @unreg_list: List entry, that is used, when we are unregistering the
1337  * device, see the function unregister_netdev
1338  *
1339  * @close_list: List entry, that is used, when we are closing the device
1340  *
1341  * @adj_list: Directly linked devices, like slaves for bonding
1342  *
1343  * @all_adj_list: ALL linked devices, *including* neighbours
1344  *
1345  * @features: Currently active device features
1346  *
1347  * @hw_features: User-changeable features
1348  *
1349  * @wanted_features: User-requested features
1350  *
1351  * @vlan_features: Mask of features inheritable by VLAN devices
1352  *
1353  * @hw_enc_features: Mask of features inherited by encapsulating devices
1354  * This field indicates what encapsulation
1355  * offloads the hardware is capable of doing,
1356  * and drivers will need to set them appropriately.
1357  *
1358  * @mpls_features: Mask of features inheritable by MPLS
1359  *
1360  * @ifindex: interface index
1361  *
1362  * @group: The group, that the device belongs to
1363  *
1364  * @stats: Statistics struct, which was left as a legacy, use
1365  * rtnl_link_stats64 instead
1366  *
1367  * @rx_dropped: Dropped packets by core network,
1368  * do not use this in drivers
1369  *
1370  * @tx_dropped: Dropped packets by core network,
1371  * do not use this in drivers
1372  *
1373  * @wireless_handlers: List of functions to handle Wireless Extensions,
1374  * instead of ioctl,
1375  * see <net/iw_handler.h> for details.
1376  *
1377  * @wireless_data: Instance data managed by the core of wireless extensions
1378  *
1379  * @netdev_ops: Includes several pointers to callbacks,
1380  * if one wants to override the ndo_*( ) functions
1381  *
1382  * @ethtool_ops: Management operations
1383  *
1384  * @header_ops: Includes callbacks for creating, parsing, caching, etc
1385  * of Layer 2 headers.
1386  *
1387  *

```

```

1368 * @flags:          Interface flags (a la BSD)
1369 * @priv_flags:     Like 'flags' but invisible to userspace,
1370 *                 see if.h for the definitions
1371 * @gflags:         Global flags ( kept as legacy )
1372 * @padded:         How much padding added by alloc_netdev()
1373 * @operstate:      RFC2863 operstate
1374 * @link_mode:      Mapping policy to operstate
1375 * @if_port:        Selectable AUI, TP, ...
1376 * @dma:            DMA channel
1377 * @mtu:            Interface MTU value
1378 * @type:           Interface hardware type
1379 * @hard_header_len: Hardware header length
1380 *
1381 * @needed_headroom: Extra headroom the hardware may need, but not in all
1382 *                   cases can this be guaranteed
1383 * @needed_tailroom: Extra tailroom the hardware may need, but not in all
1384 *                   cases can this be guaranteed. Some cases also use
1385 *                   LL_MAX_HEADER instead to allocate the skb
1386 *
1387 * interface address info:
1388 *
1389 * @perm_addr:       Permanent hw address
1390 * @addr_assign_type: Hw address assignment type
1391 * @addr_len:        Hardware address length
1392 * @neigh_priv_len;  Used in neigh_alloc(),
1393 *                   initialized only in atm/clip.c
1394 * @dev_id:          Used to differentiate devices that share
1395 *                   the same link layer address
1396 * @dev_port:        Used to differentiate devices that share
1397 *                   the same function
1398 * @addr_list_lock:  XXX: need comments on this one
1399 * @uc_promisc:      Counter, that indicates, that promiscuous mode
1400 *                   has been enabled due to the need to listen to
1401 *                   additional unicast addresses in a device that
1402 *                   does not implement ndo_set_rx_mode()
1403 * @uc:              unicast mac addresses
1404 * @mc:              multicast mac addresses
1405 * @dev_addrs:       List of device hw addresses
1406 * @queues_kset:     Group of all Kobjects in the Tx and RX queues
1407 * @promiscuity:     Number of times, the NIC is told to work in
1408 *                   Promiscuous mode, if it becomes 0 the NIC will
1409 *                   exit from working in Promiscuous mode
1410 * @allmulti:        Counter, enables or disables allmulticast mode
1411 *
1412 * @vlan_info:       VLAN info
1413 * @dsa_ptr:         dsa specific data
1414 * @tipc_ptr:        TIPC specific data
1415 * @atalk_ptr:       AppleTalk link
1416 * @ip_ptr:          IPv4 specific data
1417 * @dn_ptr:          DECnet specific data
1418 * @ip6_ptr:         IPv6 specific data
1419 * @ax25_ptr:        AX.25 specific data
1420 * @ieee80211_ptr:   IEEE 802.11 specific data, assign before registering
1421 *
1422 * @last_rx:         Time of Last Rx
1423 * @dev_addr:        Hw address (before bcast,
1424 *                   because most packets are unicast)
1425 *
1426 * @_rx:             Array of RX queues
1427 * @num_rx_queues:   Number of RX queues
1428 *                   allocated at register_netdev() time
1429 * @real_num_rx_queues: Number of RX queues currently active in device
1430 *
1431 * @rx_handler:      handler for received packets
1432 * @rx_handler_data: XXX: need comments on this one
1433 * @ingress_queue:   XXX: need comments on this one
1434 * @broadcast:       hw bcast address
1435 *
1436 * @rx_cpu_rmap:     CPU reverse-mapping for RX completion interrupts,
1437 *                   indexed by RX queue number. Assigned by driver.
1438 *                   This must only be set if the ndo_rx_flow_steer
1439 *                   operation is defined
1440 * @index_hlist:     Device index hash chain

```



```

1441 *
1442 * @_tx: Array of TX queues
1443 * @num_tx_queues: Number of TX queues allocated at alloc_netdev_mq() time
1444 * @real_num_tx_queues: Number of TX queues currently active in device
1445 * @qdisc: Root qdisc from userspace point of view
1446 * @tx_queue_len: Max frames per queue allowed
1447 * @tx_global_lock: XXX: need comments on this one
1448 *
1449 * @xps_maps: XXX: need comments on this one
1450 *
1451 * @trans_start: Time (in jiffies) of last Tx
1452 * @watchdog_timeo: Represents the timeout that is used by
1453 * the watchdog ( see dev_watchdog() )
1454 * @watchdog_timer: List of timers
1455 *
1456 * @pcpu_refcnt: Number of references to this device
1457 * @todo_list: Delayed register/unregister
1458 * @Link_watch_list: XXX: need comments on this one
1459 *
1460 * @reg_state: Register/unregister state machine
1461 * @dismantle: Device is going to be freed
1462 * @rtnl_link_state: This enum represents the phases of creating
1463 * a new link
1464 *
1465 * @destructor: Called from unregister,
1466 * can be used to call free_netdev
1467 * @npinfo: XXX: need comments on this one
1468 * @nd_net: Network namespace this network device is inside
1469 *
1470 * @ml_priv: Mid-layer private
1471 * @lstats: Loopback statistics
1472 * @tstats: Tunnel statistics
1473 * @dstats: Dummy statistics
1474 * @vstats: Virtual ethernet statistics
1475 *
1476 * @garp_port: GARP
1477 * @mrp_port: MRP
1478 *
1479 * @dev: Class/net/name entry
1480 * @sysfs_groups: Space for optional device, statistics and wireless
1481 * sysfs groups
1482 *
1483 * @sysfs_rx_queue_group: Space for optional per-rx queue attributes
1484 * @rtnl_link_ops: Rtnl_Link_ops
1485 *
1486 * @gso_max_size: Maximum size of generic segmentation offload
1487 * @gso_max_segs: Maximum number of segments that can be passed to the
1488 * NIC for GSO
1489 * @gso_min_segs: Minimum number of segments that can be passed to the
1490 * NIC for GSO
1491 *
1492 * @dcbnl_ops: Data Center Bridging netlink ops
1493 * @num_tc: Number of traffic classes in the net device
1494 * @tc_to_txq: XXX: need comments on this one
1495 * @prio_tc_map XXX: need comments on this one
1496 *
1497 * @fcoe_ddp_xid: Max exchange id for FCoE LRO by ddp
1498 *
1499 * @priomap: XXX: need comments on this one
1500 * @phydev: Physical device may attach itself
1501 * for hardware timestamping
1502 *
1503 * @qdisc_tx_busylock: XXX: need comments on this one
1504 *
1505 * FIXME: cleanup struct net_device such that network protocol info
1506 * moves out.
1507 */
1508
1509 struct net_device {
1510     char name[IFNAMSIZ];
1511     struct hlist_node name_hlist;
1512     char *ifalias;
1513     /*

```

```

1514      *      I/O specific fields
1515      *      FIXME: Merge these and struct ifmap into one
1516      */
1517      unsigned long      mem\_end;
1518      unsigned long      mem_start;
1519      unsigned long      base\_addr;
1520      int                 irq;
1521
1522      atomic\_t              carrier_changes;
1523
1524      /*
1525      *      Some hardware also needs these fields (state,dev_list,
1526      *      napi_list,unreg_list,close_list) but they are not
1527      *      part of the usual set specified in Space.c.
1528      */
1529
1530      unsigned long      state;
1531
1532      struct list\_head    dev\_list;
1533      struct list\_head    napi_list;
1534      struct list\_head    unreg_list;
1535      struct list\_head    close_list;
1536      struct list\_head    ptype_all;
1537      struct list\_head    ptype_specific;
1538
1539      struct {
1540          struct list\_head upper;
1541          struct list\_head lower;
1542      } adj\_list;
1543
1544      struct {
1545          struct list\_head upper;
1546          struct list\_head lower;
1547      } all\_adj\_list;
1548
1549      netdev\_features\_t      features;
1550      netdev\_features\_t      hw_features;
1551      netdev\_features\_t      wanted_features;
1552      netdev\_features\_t      vlan_features;
1553      netdev\_features\_t      hw_enc_features;
1554      netdev\_features\_t      mpls_features;
1555
1556      int                    ifindex;
1557      int                    group;
1558
1559      struct net\_device\_stats stats;
1560
1561      atomic\_long\_t          rx_dropped;
1562      atomic\_long\_t          tx_dropped;
1563
1564      #ifdef CONFIG_WIRELESS_EXT
1565          const struct iw\_handler\_def * wireless_handlers;
1566          struct iw\_public\_data * wireless_data;
1567      #endif
1568      const struct net\_device\_ops * netdev\_ops;
1569      const struct ethtool\_ops * ethtool\_ops;
1570      #ifdef CONFIG_NET_SWITCHDEV
1571          const struct switchdev\_ops * switchdev\_ops;
1572      #endif
1573
1574      const struct header\_ops * header\_ops;
1575
1576      unsigned int          flags;
1577      unsigned int          priv_flags;
1578
1579      unsigned short        gflags;
1580      unsigned short        padded;
1581
1582      unsigned char          operstate;
1583      unsigned char          link\_mode;
1584
1585      unsigned char          if_port;
1586      unsigned char          dma;

```



```

1587
1588 unsigned int          mtu;
1589 unsigned short        type;
1590 unsigned short        hard_header_len;
1591
1592 unsigned short        needed_headroom;
1593 unsigned short        needed_tailroom;
1594
1595 /* Interface address info. */
1596 unsigned char          perm_addr[MAX_ADDR_LEN];
1597 unsigned char          addr_assign_type;
1598 unsigned char          addr_len;
1599 unsigned short         neigh_priv_len;
1600 unsigned short         dev_id;
1601 unsigned short         dev_port;
1602 spinlock_t             addr_list_lock;
1603 unsigned char          name_assign_type;
1604 bool                   uc_promisc;
1605 struct netdev_hw_addr_list uc;
1606 struct netdev_hw_addr_list mc;
1607 struct netdev_hw_addr_list dev_addrs;
1608
1609 #ifdef CONFIG_SYSFS
1610 struct kset             *queues_kset;
1611 #endif
1612 unsigned int           promiscuity;
1613 unsigned int           allmulti;
1614
1615
1616 /* Protocol specific pointers */
1617
1618 #if IS_ENABLED(CONFIG_VLAN_8021Q)
1619 struct vlan_info __rcu *vlan_info;
1620 #endif
1621 #if IS_ENABLED(CONFIG_NET_DSA)
1622 struct dsa_switch_tree *dsa_ptr;
1623 #endif
1624 #if IS_ENABLED(CONFIG_TIPC)
1625 struct tipc_bearer __rcu *tipc_ptr;
1626 #endif
1627 void                  *atalk_ptr;
1628 struct in_device __rcu *ip_ptr;
1629 struct dn_dev __rcu   *dn_ptr;
1630 struct inet6_dev __rcu *ip6_ptr;
1631 void                  *ax25_ptr;
1632 struct wireless_dev   *ieee80211_ptr;
1633 struct wpan_dev        *ieee802154_ptr;
1634 #if IS_ENABLED(CONFIG_MPLS_ROUTING)
1635 struct mpls_dev __rcu *mpls_ptr;
1636 #endif
1637
1638 /*
1639  * Cache lines mostly used on receive path (including eth_type_trans())
1640  */
1641 unsigned long          last_rx;
1642
1643 /* Interface address info used in eth_type_trans() */
1644 unsigned char          *dev_addr;
1645
1646
1647 #ifdef CONFIG_SYSFS
1648 struct netdev_rx_queue *_rx;
1649
1650 unsigned int           num_rx_queues;
1651 unsigned int           real_num_rx_queues;
1652
1653 #endif
1654
1655 unsigned long          gro_flush_timeout;
1656 rx_handler_func_t __rcu *rx_handler;
1657 void __rcu             *rx_handler_data;
1658
1659 #ifdef CONFIG_NET_CLS_ACT

```

```

1660 struct tcf\_proto \_\_rcu *ingress_cl_list;
1661 #endif
1662 struct netdev\_queue \_\_rcu *ingress_queue;
1663 #ifdef CONFIG_NETFILTER_INGRESS
1664 struct list\_head nf_hooks_ingress;
1665 #endif
1666
1667 unsigned char broadcast[MAX\_ADDR\_LEN];
1668 #ifdef CONFIG_RFS_ACCEL
1669 struct cpu\_rmap *rx_cpu_rmap;
1670 #endif
1671 struct hlist\_node index_hlist;
1672
1673 /*
1674  * Cache lines mostly used on transmit path
1675  */
1676 struct netdev\_queue *_tx cacheline\_aligned\_in\_smp;
1677 unsigned int num_tx_queues;
1678 unsigned int real_num_tx_queues;
1679 struct Qdisc *qdisc;
1680 unsigned long tx_queue_len;
1681 spinlock\_t tx_global_lock;
1682 int watchdog_timeo;
1683
1684 #ifdef CONFIG_XPS
1685 struct xps\_dev\_maps \_\_rcu *xps_maps;
1686 #endif
1687
1688 /* These may be needed for future network-power-down code. */
1689
1690 /*
1691  * trans_start here is expensive for high speed devices on SMP,
1692  * please use netdev_queue->trans_start instead.
1693  */
1694 unsigned long trans_start;
1695
1696 struct timer\_list watchdog\_timer;
1697
1698 int \_\_percpu *pcpu_refcnt;
1699 struct list\_head todo_list;
1700
1701 struct list\_head link_watch_list;
1702
1703 enum { NETREG_UNINITIALIZED=0,
1704        NETREG_REGISTERED, /* completed register_netdevice */
1705        NETREG_UNREGISTERING, /* called unregister_netdevice */
1706        NETREG_UNREGISTERED, /* completed unregister todo */
1707        NETREG_RELEASED, /* called free_netdev */
1708        NETREG_DUMMY, /* dummy device for NAPI poll */
1709 } reg\_state:8;
1710
1711 bool dismantle;
1712
1713 enum {
1714        RTNL_LINK_INITIALIZED,
1715        RTNL_LINK_INITIALIZING,
1716 } rtnl\_link\_state:16;
1717
1718 void (*destructor)(struct net\_device *dev);
1719
1720 #ifdef CONFIG_NETPOLL
1721 struct netpoll\_info \_\_rcu *npinfo;
1722 #endif
1723
1724 possible\_net\_t nd_net;
1725
1726 /* mid-layer private */
1727 union {
1728        void *ml_priv;
1729        struct pcpu\_lstats \_\_percpu *lstats;
1730        struct pcpu\_sw\_netstats \_\_percpu *tstats;
1731        struct pcpu\_dstats \_\_percpu *dstats;
1732        struct pcpu\_vstats \_\_percpu *vstats;

```

```

1733     };
1734
1735     struct garp\_port __rcu *garp\_port;
1736     struct mrp\_port __rcu *mrp\_port;
1737
1738     struct device dev;
1739     const struct attribute\_group *sysfs_groups[4];
1740     const struct attribute\_group *sysfs_rx_queue_group;
1741
1742     const struct rtnl\_link\_ops *rtnl\_link\_ops;
1743
1744     /* for setting kernel sock attribute on TCP connection setup */
1745 #define GSO\_MAX\_SIZE 65536
1746     unsigned int gso\_max\_size;
1747 #define GSO\_MAX\_SEGS 65535
1748     u16 gso\_max\_segs;
1749     u16 gso\_min\_segs;
1750 #ifdef CONFIG_DCB
1751     const struct dcbnl\_rtnl\_ops *dcbnl\_ops;
1752 #endif
1753     u8 num_tc;
1754     struct netdev\_tc\_txq tc_to_txq[TC\_MAX\_QUEUE];
1755     u8 prio_tc_map[TC\_BITMASK + 1];
1756
1757 #if IS\_ENABLED(CONFIG_FCOE)
1758     unsigned int fcoe\_ddp\_xid;
1759 #endif
1760 #if IS\_ENABLED(CONFIG_CGROUP_NET_PRIO)
1761     struct netprio\_map __rcu *priomap;
1762 #endif
1763     struct phy\_device *phydev;
1764     struct lock\_class\_key *qdisc_tx_busylock;
1765 };
1766 #define to\_net\_dev(d) container\_of(d, struct net\_device, dev)
1767
1768 #define NETDEV\_ALIGN 32
1769
1770 static inline
1771 int netdev\_get\_prio\_tc\_map(const struct net\_device *dev, u32 prio)
1772 {
1773     return dev->prio_tc_map[prio & TC\_BITMASK];
1774 }
1775
1776 static inline
1777 int netdev\_set\_prio\_tc\_map(struct net\_device *dev, u8 prio, u8 tc)
1778 {
1779     if (tc >= dev->num_tc)
1780         return -EINVAL;
1781
1782     dev->prio_tc_map[prio & TC\_BITMASK] = tc & TC\_BITMASK;
1783     return 0;
1784 }
1785
1786 static inline
1787 void netdev\_reset\_tc(struct net\_device *dev)
1788 {
1789     dev->num_tc = 0;
1790     memset(dev->tc_to_txq, 0, sizeof(dev->tc_to_txq));
1791     memset(dev->prio_tc_map, 0, sizeof(dev->prio_tc_map));
1792 }
1793
1794 static inline
1795 int netdev\_set\_tc\_queue(struct net\_device *dev, u8 tc, u16 count, u16 offset)
1796 {
1797     if (tc >= dev->num_tc)
1798         return -EINVAL;
1799
1800     dev->tc_to_txq[tc].count = count;
1801     dev->tc_to_txq[tc].offset = offset;
1802     return 0;
1803 }
1804
1805 static inline

```

```

1806 int netdev_set_num_tc(struct net_device *dev, u8 num_tc)
1807 {
1808     if (num_tc > TC_MAX_QUEUE)
1809         return -EINVAL;
1810
1811     dev->num_tc = num_tc;
1812     return 0;
1813 }
1814
1815 static inline
1816 int netdev_get_num_tc(struct net_device *dev)
1817 {
1818     return dev->num_tc;
1819 }
1820
1821 static inline
1822 struct netdev_queue *netdev_get_tx_queue(const struct net_device *dev,
1823                                         unsigned int index)
1824 {
1825     return &dev->_tx[index];
1826 }
1827
1828 static inline struct netdev_queue *skb_get_tx_queue(const struct net_device *dev,
1829                                                    const struct sk_buff *skb)
1830 {
1831     return netdev_get_tx_queue(dev, skb_get_queue_mapping(skb));
1832 }
1833
1834 static inline void netdev_for_each_tx_queue(struct net_device *dev,
1835                                             void (*f)(struct net_device *,
1836                                                       struct netdev_queue *,
1837                                                       void *),
1838                                             void *arg)
1839 {
1840     unsigned int i;
1841
1842     for (i = 0; i < dev->num_tx_queues; i++)
1843         f(dev, &dev->_tx[i], arg);
1844 }
1845
1846 struct netdev_queue *netdev_pick_tx(struct net_device *dev,
1847                                     struct sk_buff *skb,
1848                                     void *accel_priv);
1849
1850 /*
1851  * Net namespace inlines
1852  */
1853 static inline
1854 struct net *dev_net(const struct net_device *dev)
1855 {
1856     return read_pnet(&dev->nd_net);
1857 }
1858
1859 static inline
1860 void dev_net_set(struct net_device *dev, struct net *net)
1861 {
1862     write_pnet(&dev->nd_net, net);
1863 }
1864
1865 static inline bool netdev_uses_dsa(struct net_device *dev)
1866 {
1867     #if IS_ENABLED(CONFIG_NET_DSA)
1868         if (dev->dsa_ptr != NULL)
1869             return dsa_uses_tagged_protocol(dev->dsa_ptr);
1870     #endif
1871     return false;
1872 }
1873
1874 /**
1875  * netdev_priv - access network device private data
1876  * @dev: network device
1877  *
1878  * Get network device private data

```

```

1879 */
1880 static inline void netdev\_priv(const struct net\_device *dev)
1881 {
1882     return (char *)dev + ALIGN(sizeof(struct net\_device), NETDEV\_ALIGN);
1883 }
1884
1885 /* Set the sysfs physical device reference for the network logical device
1886  * if set prior to registration will cause a symlink during initialization.
1887  */
1888 #define SET\_NETDEV\_DEV(net, pdev) ((net)->dev.parent = (pdev))
1889
1890 /* Set the sysfs device type for the network logical device to allow
1891  * fine-grained identification of different network device types. For
1892  * example Ethernet, Wireless LAN, Bluetooth, WiMAX etc.
1893  */
1894 #define SET\_NETDEV\_DEVTYPE(net, devtype) ((net)->dev.type = (devtype))
1895
1896 /* Default NAPI poll() weight
1897  * Device drivers are strongly advised to not use bigger value
1898  */
1899 #define NAPI\_POLL\_WEIGHT 64
1900
1901 /**
1902  *      netif_napi_add - initialize a napi context
1903  *      @dev: network device
1904  *      @napi: napi context
1905  *      @poll: polling function
1906  *      @weight: default weight
1907  *
1908  *      netif_napi_add() must be used to initialize a napi context prior to calling
1909  *      *any* of the other napi related functions.
1910  */
1911 void netif\_napi\_add(struct net\_device *dev, struct napi\_struct *napi,
1912                    int (*poll)(struct napi\_struct *, int), int weight);
1913
1914 /**
1915  *      netif_napi_del - remove a napi context
1916  *      @napi: napi context
1917  *
1918  *      netif_napi_del() removes a napi context from the network device napi list
1919  */
1920 void netif\_napi\_del(struct napi\_struct *napi);
1921
1922 struct napi\_gro\_cb {
1923     /* Virtual address of skb_shinfo(skb)->frags[0].page + offset. */
1924     void *frag0;
1925
1926     /* Length of frag0. */
1927     unsigned int frag0_len;
1928
1929     /* This indicates where we are processing relative to skb->data. */
1930     int data_offset;
1931
1932     /* This is non-zero if the packet cannot be merged with the new skb. */
1933     u16 flush;
1934
1935     /* Save the IP ID here and check when we get to the transport layer */
1936     u16 flush_id;
1937
1938     /* Number of segments aggregated. */
1939     u16 count;
1940
1941     /* Start offset for remote checksum offload */
1942     u16 gro_remcsun_start;
1943
1944     /* jiffies when first packet was created/queued */
1945     unsigned long age;
1946
1947     /* Used in ipv6_gro_receive() and foo-over-udp */
1948     u16 proto;
1949
1950     /* This is non-zero if the packet may be of the same flow. */
1951     u8 same_flow:1;

```

```

1952
1953 /* Used in udp_gro_receive */
1954 u8      udp_mark:1;
1955
1956 /* GRO checksum is valid */
1957 u8      csum_valid:1;
1958
1959 /* Number of checksums via CHECKSUM_UNNECESSARY */
1960 u8      csum_cnt:3;
1961
1962 /* Free the skb? */
1963 u8      free:2;
1964 #define NAPI_GRO_FREE 1
1965 #define NAPI_GRO_FREE_STOLEN_HEAD 2
1966
1967 /* Used in foo-over-udp, set in udp[46]_gro_receive */
1968 u8      is_ipv6:1;
1969
1970 /* 7 bit hole */
1971
1972 /* used to support CHECKSUM_COMPLETE for tunneling protocols */
1973 __wsum  csum;
1974
1975 /* used in skb_gro_receive() slow path */
1976 struct sk_buff *last;
1977 };
1978
1979 #define NAPI_GRO_CB(skb) ((struct napi_gro_cb *) (skb)->cb)
1980
1981 struct packet_type {
1982     __be16      type; /* This is really htons(ether_type). */
1983     struct net_device *dev; /* NULL is wildcarded here */
1984     int (*func) (struct sk_buff *,
1985                 struct net_device *,
1986                 struct packet_type *,
1987                 struct net_device *);
1988     bool (*id_match) (struct packet_type *ptype,
1989                     struct sock *sk);
1990     void *af_packet_priv;
1991     struct list_head list;
1992 };
1993
1994 struct offload_callbacks {
1995     struct sk_buff *(*gso_segment) (struct sk_buff *skb,
1996                                   netdev_features_t features);
1997     struct sk_buff *(*gro_receive) (struct sk_buff **head,
1998                                   struct sk_buff *skb);
1999     int (*gro_complete) (struct sk_buff *skb, int nhoff);
2000 };
2001
2002 struct packet_offload {
2003     __be16      type; /* This is really htons(ether_type). */
2004     u16      priority;
2005     struct offload_callbacks callbacks;
2006     struct list_head list;
2007 };
2008
2009 struct udp_offload;
2010
2011 struct udp_offload_callbacks {
2012     struct sk_buff *(*gro_receive) (struct sk_buff **head,
2013                                   struct sk_buff *skb,
2014                                   struct udp_offload *uoff);
2015     int (*gro_complete) (struct sk_buff *skb,
2016                        int nhoff,
2017                        struct udp_offload *uoff);
2018 };
2019
2020 struct udp_offload {
2021     __be16      port;
2022     u8      ipproto;
2023     struct udp_offload_callbacks callbacks;
2024 };

```

```

2025
2026 /* often modified stats are per cpu, other are shared (netdev->stats) */
2027 struct pcpu_sw_netstats {
2028     u64 rx_packets;
2029     u64 rx_bytes;
2030     u64 tx_packets;
2031     u64 tx_bytes;
2032     struct u64_stats_sync syncp;
2033 };
2034
2035 #define netdev_alloc_pcpu_stats(type) \
2036 ({ \
2037     typeof(type) __percpu *pcpu_stats = alloc_percpu(type); \
2038     if (pcpu_stats) { \
2039         int __cpu; \
2040         for_each_possible_cpu(__cpu) { \
2041             typeof(type) *stat; \
2042             stat = per_cpu_ptr(pcpu_stats, __cpu); \
2043             u64_stats_init(&stat->syncp); \
2044         } \
2045     } \
2046     pcpu_stats; \
2047 })
2048
2049 #include <linux/notifier.h>
2050
2051 /* netdevice notifier chain. Please remember to update the rtnetlink
2052 * notification exclusion list in rtnetlink_event() when adding new
2053 * types.
2054 */
2055 #define NETDEV_UP 0x0001 /* For now you can't veto a device up/down */
2056 #define NETDEV_DOWN 0x0002
2057 #define NETDEV_REBOOT 0x0003 /* Tell a protocol stack a network interface
2058 detected a hardware crash and restarted
2059 - we can use this eg to kick tcp sessions
2060 once done */
2061 #define NETDEV_CHANGE 0x0004 /* Notify device state change */
2062 #define NETDEV_REGISTER 0x0005
2063 #define NETDEV_UNREGISTER 0x0006
2064 #define NETDEV_CHANGEMTU 0x0007 /* notify after mtu change happened */
2065 #define NETDEV_CHANGEADDR 0x0008
2066 #define NETDEV_GOING_DOWN 0x0009
2067 #define NETDEV_CHANGENAME 0x000A
2068 #define NETDEV_FEAT_CHANGE 0x000B
2069 #define NETDEV_BONDING_FAILOVER 0x000C
2070 #define NETDEV_PRE_UP 0x000D
2071 #define NETDEV_PRE_TYPE_CHANGE 0x000E
2072 #define NETDEV_POST_TYPE_CHANGE 0x000F
2073 #define NETDEV_POST_INIT 0x0010
2074 #define NETDEV_UNREGISTER_FINAL 0x0011
2075 #define NETDEV_RELEASE 0x0012
2076 #define NETDEV_NOTIFY_PEERS 0x0013
2077 #define NETDEV_JOIN 0x0014
2078 #define NETDEV_CHANGEUPPER 0x0015
2079 #define NETDEV_RESEND_IGMP 0x0016
2080 #define NETDEV_PRECHANGEMTU 0x0017 /* notify before mtu change happened */
2081 #define NETDEV_CHANGEINFODATA 0x0018
2082 #define NETDEV_BONDING_INFO 0x0019
2083
2084 int register_netdevice_notifier(struct notifier_block *nb);
2085 int unregister_netdevice_notifier(struct notifier_block *nb);
2086
2087 struct netdev_notifier_info {
2088     struct net_device *dev;
2089 };
2090
2091 struct netdev_notifier_change_info {
2092     struct netdev_notifier_info info; /* must be first */
2093     unsigned int flags_changed;
2094 };
2095
2096 static inline void netdev_notifier_info_init(struct netdev_notifier_info *info,
2097     struct net_device *dev)

```



```

2098 {
2099     info->dev = dev;
2100 }
2101
2102 static inline struct net_device *
2103 netdev_notifier_info_to_dev(const struct netdev_notifier_info *info)
2104 {
2105     return info->dev;
2106 }
2107
2108 int call_netdevice_notifiers(unsigned long val, struct net_device *dev);
2109
2110 extern rwlock_t dev_base_lock; /* Device List Lock */
2111
2112 #define for_each_netdev(net, d) \
2113     list_for_each_entry(d, &(net)->dev_base_head, dev_list)
2114 #define for_each_netdev_reverse(net, d) \
2115     list_for_each_entry_reverse(d, &(net)->dev_base_head, dev_list)
2116 #define for_each_netdev_rcu(net, d) \
2117     list_for_each_entry_rcu(d, &(net)->dev_base_head, dev_list)
2118 #define for_each_netdev_safe(net, d, n) \
2119     list_for_each_entry_safe(d, n, &(net)->dev_base_head, dev_list)
2120 #define for_each_netdev_continue(net, d) \
2121     list_for_each_entry_continue(d, &(net)->dev_base_head, dev_list)
2122 #define for_each_netdev_continue_rcu(net, d) \
2123     list_for_each_entry_continue_rcu(d, &(net)->dev_base_head, dev_list)
2124 #define for_each_netdev_in_bond_rcu(bond, slave) \
2125     for_each_netdev_rcu(&init_net, slave) \
2126     if (netdev_master_upper_dev_get_rcu(slave) == (bond))
2127 #define net_device_entry(lh) list_entry(lh, struct net_device, dev_list)
2128
2129 static inline struct net_device *next_net_device(struct net_device *dev)
2130 {
2131     struct list_head *lh;
2132     struct net *net;
2133
2134     net = dev_net(dev);
2135     lh = dev->dev_list.next;
2136     return lh == &net->dev_base_head ? NULL : net_device_entry(lh);
2137 }
2138
2139 static inline struct net_device *next_net_device_rcu(struct net_device *dev)
2140 {
2141     struct list_head *lh;
2142     struct net *net;
2143
2144     net = dev_net(dev);
2145     lh = rcu_dereference(list_next_rcu(&dev->dev_list));
2146     return lh == &net->dev_base_head ? NULL : net_device_entry(lh);
2147 }
2148
2149 static inline struct net_device *first_net_device(struct net *net)
2150 {
2151     return list_empty(&net->dev_base_head) ? NULL :
2152         net_device_entry(net->dev_base_head.next);
2153 }
2154
2155 static inline struct net_device *first_net_device_rcu(struct net *net)
2156 {
2157     struct list_head *lh = rcu_dereference(list_next_rcu(&net->dev_base_head));
2158
2159     return lh == &net->dev_base_head ? NULL : net_device_entry(lh);
2160 }
2161
2162 int netdev_boot_setup_check(struct net_device *dev);
2163 unsigned long netdev_boot_base(const char *prefix, int unit);
2164 struct net_device *dev_getbyhwaddr_rcu(struct net *net, unsigned short type,
2165                                         const char *hwaddr);
2166 struct net_device *dev_getfirstbyhwtype(struct net *net, unsigned short type);
2167 struct net_device *dev_getfirstbyhwtype(struct net *net, unsigned short type);
2168 void dev_add_pack(struct packet_type *pt);
2169 void dev_remove_pack(struct packet_type *pt);

```

```

2171 void dev\_remove\_pack(struct packet\_type *pt);
2172 void dev\_add\_offload(struct packet\_offload *po);
2173 void dev\_remove\_offload(struct packet\_offload *po);
2174
2175 int dev\_get\_iflink(const struct net\_device *dev);
2176 struct net\_device * dev\_get\_by\_flags(struct net *net, unsigned short flags,
2177                                     unsigned short mask);
2178 struct net\_device * dev\_get\_by\_name(struct net *net, const char *name);
2179 struct net\_device * dev\_get\_by\_name\_rcu(struct net *net, const char *name);
2180 struct net\_device * dev\_get\_by\_name(struct net *net, const char *name);
2181 int dev\_alloc\_name(struct net\_device *dev, const char *name);
2182 int dev\_open(struct net\_device *dev);
2183 int dev\_close(struct net\_device *dev);
2184 int dev\_close\_many(struct list\_head *head, bool unlink);
2185 void dev\_disable\_lro(struct net\_device *dev);
2186 int dev\_loopback\_xmit(struct sock *sk, struct sk\_buff *newskb);
2187 int dev\_queue\_xmit\_sk(struct sock *sk, struct sk\_buff *skb);
2188 static inline int dev\_queue\_xmit(struct sk\_buff *skb)
2189 {
2190     return dev\_queue\_xmit\_sk(skb->sk, skb);
2191 }
2192 int dev\_queue\_xmit\_accel(struct sk\_buff *skb, void *accel_priv);
2193 int register\_netdevice(struct net\_device *dev);
2194 void unregister\_netdevice\_queue(struct net\_device *dev, struct list\_head *head);
2195 void unregister\_netdevice\_many(struct list\_head *head);
2196 static inline void unregister\_netdevice(struct net\_device *dev)
2197 {
2198     unregister\_netdevice\_queue(dev, NULL);
2199 }
2200
2201 int netdev\_refcnt\_read(const struct net\_device *dev);
2202 void free\_netdev(struct net\_device *dev);
2203 void netdev\_freemem(struct net\_device *dev);
2204 void synchronize\_net(void);
2205 int init\_dummy\_netdev(struct net\_device *dev);
2206
2207 DECLARE_PER_CPU(int, xmit_recursion);
2208 static inline int dev\_recursion\_level(void)
2209 {
2210     return this\_cpu\_read(xmit_recursion);
2211 }
2212
2213 struct net\_device * dev\_get\_by\_index(struct net *net, int ifindex);
2214 struct net\_device * dev\_get\_by\_index\_rcu(struct net *net, int ifindex);
2215 struct net\_device * dev\_get\_by\_index\_rcu(struct net *net, int ifindex);
2216 int netdev\_get\_name(struct net *net, char *name, int ifindex);
2217 int dev\_restart(struct net\_device *dev);
2218 int skb\_gro\_receive(struct sk\_buff **head, struct sk\_buff *skb);
2219
2220 static inline unsigned int skb\_gro\_offset(const struct sk\_buff *skb)
2221 {
2222     return NAPI\_GRO\_CB(skb)->data_offset;
2223 }
2224
2225 static inline unsigned int skb\_gro\_len(const struct sk\_buff *skb)
2226 {
2227     return skb->len - NAPI\_GRO\_CB(skb)->data_offset;
2228 }
2229
2230 static inline void skb\_gro\_pull(struct sk\_buff *skb, unsigned int len)
2231 {
2232     NAPI\_GRO\_CB(skb)->data_offset += len;
2233 }
2234
2235 static inline void * skb\_gro\_header\_fast(struct sk\_buff *skb,
2236                                           unsigned int offset)
2237 {
2238     return NAPI\_GRO\_CB(skb)->frag0 + offset;
2239 }
2240
2241 static inline int skb\_gro\_header\_hard(struct sk\_buff *skb, unsigned int hlen)
2242 {
2243     return NAPI\_GRO\_CB(skb)->frag0_len < hlen;

```

```

2244 }
2245
2246 static inline void *skb_gro_header_slow(struct sk_buff *skb, unsigned int hlen,
2247                                         unsigned int offset)
2248 {
2249     if (!pskb_may_pull(skb, hlen))
2250         return NULL;
2251
2252     NAPI_GRO_CB(skb)->frag0 = NULL;
2253     NAPI_GRO_CB(skb)->frag0_len = 0;
2254     return skb->data + offset;
2255 }
2256
2257 static inline void *skb_gro_network_header(struct sk_buff *skb)
2258 {
2259     return (NAPI_GRO_CB(skb)->frag0 ? : skb->data) +
2260         skb_network_offset(skb);
2261 }
2262
2263 static inline void skb_gro_postpull_rcsum(struct sk_buff *skb,
2264                                           const void *start, unsigned int len)
2265 {
2266     if (NAPI_GRO_CB(skb)->csum_valid)
2267         NAPI_GRO_CB(skb)->csum = csum_sub(NAPI_GRO_CB(skb)->csum,
2268                                           csum_partial(start, len, 0));
2269 }
2270
2271 /* GRO checksum functions. These are logical equivalents of the normal
2272 * checksum functions (in skbuff.h) except that they operate on the GRO
2273 * offsets and fields in sk_buff.
2274 */
2275
2276 __sum16 __skb_gro_checksum_complete(struct sk_buff *skb);
2277
2278 static inline bool skb_at_gro_remchecksum_start(struct sk_buff *skb)
2279 {
2280     return (NAPI_GRO_CB(skb)->gro_remchecksum_start - skb_headroom(skb) ==
2281            skb_gro_offset(skb));
2282 }
2283
2284 static inline bool __skb_gro_checksum_validate_needed(struct sk_buff *skb,
2285                                                       bool zero_okay,
2286                                                       __sum16 check)
2287 {
2288     return ((skb->ip_summed != CHECKSUM_PARTIAL ||
2289            skb_checksum_start_offset(skb) <
2290            skb_gro_offset(skb)) &&
2291            !skb_at_gro_remchecksum_start(skb) &&
2292            NAPI_GRO_CB(skb)->csum_cnt == 0 &&
2293            (!zero_okay || check));
2294 }
2295
2296 static inline __sum16 __skb_gro_checksum_validate_complete(struct sk_buff *skb,
2297                                                           __wsum psum)
2298 {
2299     if (NAPI_GRO_CB(skb)->csum_valid &&
2300        !csum_fold(csum_add(psum, NAPI_GRO_CB(skb)->csum)))
2301         return 0;
2302
2303     NAPI_GRO_CB(skb)->csum = psum;
2304
2305     return __skb_gro_checksum_complete(skb);
2306 }
2307
2308 static inline void skb_gro_incr_csum_unnecessary(struct sk_buff *skb)
2309 {
2310     if (NAPI_GRO_CB(skb)->csum_cnt > 0) {
2311         /* Consume a checksum from CHECKSUM_UNNECESSARY */
2312         NAPI_GRO_CB(skb)->csum_cnt--;
2313     } else {
2314         /* Update skb for CHECKSUM_UNNECESSARY and csum_level when we
2315         * verified a new top level checksum or an encapsulated one
2316         * during GRO. This saves work if we fallback to normal path.

```

```

2317         */
2318         __skb_incr_checksum_unnecessary(skb);
2319     }
2320 }
2321
2322 #define __skb_gro_checksum_validate(skb, proto, zero_okay, check, \
2323                                   compute_pseudo) \
2324 ({ \
2325     __sum16 __ret = 0; \
2326     if (__skb_gro_checksum_validate_needed(skb, zero_okay, check)) \
2327         __ret = __skb_gro_checksum_validate_complete(skb, \
2328             compute_pseudo(skb, proto)); \
2329     if (__ret) \
2330         __skb_mark_checksum_bad(skb); \
2331     else \
2332         skb_gro_incr_csum_unnecessary(skb); \
2333     __ret; \
2334 })
2335
2336 #define skb_gro_checksum_validate(skb, proto, compute_pseudo) \
2337     __skb_gro_checksum_validate(skb, proto, false, 0, compute_pseudo)
2338
2339 #define skb_gro_checksum_validate_zero_check(skb, proto, check, \
2340                                             compute_pseudo) \
2341     __skb_gro_checksum_validate(skb, proto, true, check, compute_pseudo)
2342
2343 #define skb_gro_checksum_simple_validate(skb) \
2344     __skb_gro_checksum_validate(skb, 0, false, 0, null_compute_pseudo)
2345
2346 static inline bool __skb_gro_checksum_convert_check(struct sk_buff *skb)
2347 {
2348     return (NAPI_GRO_CB(skb)->csum_cnt == 0 &&
2349         !NAPI_GRO_CB(skb)->csum_valid);
2350 }
2351
2352 static inline void __skb_gro_checksum_convert(struct sk_buff *skb,
2353                                             __sum16 check, __wsum pseudo)
2354 {
2355     NAPI_GRO_CB(skb)->csum = ~pseudo;
2356     NAPI_GRO_CB(skb)->csum_valid = 1;
2357 }
2358
2359 #define skb_gro_checksum_try_convert(skb, proto, check, compute_pseudo) \
2360 do { \
2361     if (__skb_gro_checksum_convert_check(skb)) \
2362         __skb_gro_checksum_convert(skb, check, \
2363             compute_pseudo(skb, proto)); \
2364 } while (0)
2365
2366 struct gro_remcsum {
2367     int offset;
2368     __wsum delta;
2369 };
2370
2371 static inline void skb_gro_remcsum_init(struct gro_remcsum *grc)
2372 {
2373     grc->offset = 0;
2374     grc->delta = 0;
2375 }
2376
2377 static inline void skb_gro_remcsum_process(struct sk_buff *skb, void *ptr,
2378                                           int start, int offset,
2379                                           struct gro_remcsum *grc,
2380                                           bool nopartial)
2381 {
2382     __wsum delta;
2383
2384     BUG_ON(!NAPI_GRO_CB(skb)->csum_valid);
2385
2386     if (!nopartial) {
2387         NAPI_GRO_CB(skb)->gro_remcsum_start =
2388             ((unsigned char *)ptr + start) - skb->head;
2389         return;

```

```

2390     }
2391
2392     delta = remcsum_adjust(ptr, NAPI_GRO_CB(skb)->csum, start, offset);
2393
2394     /* Adjust skb->csum since we changed the packet */
2395     NAPI_GRO_CB(skb)->csum = csum_add(NAPI_GRO_CB(skb)->csum, delta);
2396
2397     grc->offset = (ptr + offset) - (void *)skb->head;
2398     grc->delta = delta;
2399 }
2400
2401 static inline void skb_gro_remcsum_cleanup(struct sk_buff *skb,
2402                                           struct gro_remcsum *grc)
2403 {
2404     if (!grc->delta)
2405         return;
2406
2407     remcsum_unadjust((__sum16 *) (skb->head + grc->offset), grc->delta);
2408 }
2409
2410 static inline int dev_hard_header(struct sk_buff *skb, struct net_device *dev,
2411                                  unsigned short type,
2412                                  const void *daddr, const void *saddr,
2413                                  unsigned int len)
2414 {
2415     if (!dev->header_ops || !dev->header_ops->create)
2416         return 0;
2417
2418     return dev->header_ops->create(skb, dev, type, daddr, saddr, len);
2419 }
2420
2421 static inline int dev_parse_header(const struct sk_buff *skb,
2422                                   unsigned char *haddr)
2423 {
2424     const struct net_device *dev = skb->dev;
2425
2426     if (!dev->header_ops || !dev->header_ops->parse)
2427         return 0;
2428     return dev->header_ops->parse(skb, haddr);
2429 }
2430
2431 typedef int gifconf_func_t(struct net_device *dev, char __user *bufptr, int len);
2432 int register_gifconf(unsigned int family, gifconf_func_t *gifconf);
2433 static inline int unregister_gifconf(unsigned int family)
2434 {
2435     return register_gifconf(family, NULL);
2436 }
2437
2438 #ifdef CONFIG_NET_FLOW_LIMIT
2439 #define FLOW_LIMIT_HISTORY (1 << 7) /* must be ^2 and !overflow buckets */
2440 struct sd_flow_limit {
2441     u64 count;
2442     unsigned int num_buckets;
2443     unsigned int history_head;
2444     u16 history[FLOW_LIMIT_HISTORY];
2445     u8 buckets[];
2446 };
2447
2448 extern int netdev_flow_limit_table_len;
2449 #endif /* CONFIG_NET_FLOW_LIMIT */
2450
2451 /*
2452  * Incoming packets are placed on per-cpu queues
2453  */
2454 struct softnet_data {
2455     struct list_head poll_list;
2456     struct sk_buff_head process_queue;
2457
2458     /* stats */
2459     unsigned int processed;
2460     unsigned int time_squeeze;
2461     unsigned int cpu_collision;
2462     unsigned int received_rps;

```

```

2463 #ifdef CONFIG_RPS
2464     struct softnet\_data      *rps_ipi_list;
2465 #endif
2466 #ifdef CONFIG_NET_FLOW_LIMIT
2467     struct sd\_flow\_limit \_\_rcu *flow_limit;
2468 #endif
2469     struct Qdisc      *output_queue;
2470     struct Qdisc      **output_queue_tailp;
2471     struct sk\_buff     *completion_queue;
2472
2473 #ifdef CONFIG_RPS
2474     /* Elements below can be accessed between CPUs for RPS */
2475     struct call\_single\_data csd \_\_cacheline\_aligned\_in\_smp;
2476     struct softnet\_data      *rps_ipi_next;
2477     unsigned int               cpu;
2478     unsigned int               input_queue_head;
2479     unsigned int               input_queue_tail;
2480 #endif
2481     unsigned int               dropped;
2482     struct sk\_buff\_head       input_pkt_queue;
2483     struct napi\_struct        backlog;
2484
2485 };
2486
2487 static inline void input\_queue\_head\_incr(struct softnet\_data *sd)
2488 {
2489     #ifdef CONFIG_RPS
2490         sd->input_queue_head++;
2491     #endif
2492 }
2493
2494 static inline void input\_queue\_tail\_incr\_save(struct softnet\_data *sd,
2495                                             unsigned int *qtail)
2496 {
2497     #ifdef CONFIG_RPS
2498         *qtail = ++sd->input_queue_tail;
2499     #endif
2500 }
2501
2502 DECLARE_PER_CPU_ALIGNED(struct softnet\_data, softnet\_data);
2503
2504 void \_\_netif\_schedule(struct Qdisc *q);
2505 void netif\_schedule\_queue(struct netdev\_queue *txq);
2506
2507 static inline void netif\_tx\_schedule\_all(struct net\_device *dev)
2508 {
2509     unsigned int i;
2510
2511     for (i = 0; i < dev->num_tx_queues; i++)
2512         netif\_schedule\_queue(netdev\_get\_tx\_queue(dev, i));
2513 }
2514
2515 static inline void netif\_tx\_start\_queue(struct netdev\_queue *dev_queue)
2516 {
2517     clear\_bit(__QUEUE_STATE_DRV_XOFF, &dev_queue->state);
2518 }
2519
2520 /**
2521  *      netif\_start\_queue - allow transmit
2522  *      @dev: network device
2523  *
2524  *      Allow upper layers to call the device hard\_start\_xmit routine.
2525  */
2526 static inline void netif\_start\_queue(struct net\_device *dev)
2527 {
2528     netif\_tx\_start\_queue(netdev\_get\_tx\_queue(dev, 0));
2529 }
2530
2531 static inline void netif\_tx\_start\_all\_queues(struct net\_device *dev)
2532 {
2533     unsigned int i;
2534
2535     for (i = 0; i < dev->num_tx_queues; i++) {

```

```

2536         struct netdev_queue *txq = netdev_get_tx_queue(dev, i);
2537         netif_tx_start_queue(txq);
2538     }
2539 }
2540
2541 void netif_tx_wake_queue(struct netdev_queue *dev_queue);
2542
2543 /**
2544  *   netif_wake_queue - restart transmit
2545  *   @dev: network device
2546  *
2547  *   Allow upper layers to call the device hard_start_xmit routine.
2548  *   Used for flow control when transmit resources are available.
2549  */
2550 static inline void netif_wake_queue(struct net_device *dev)
2551 {
2552     netif_tx_wake_queue(netdev_get_tx_queue(dev, 0));
2553 }
2554
2555 static inline void netif_tx_wake_all_queues(struct net_device *dev)
2556 {
2557     unsigned int i;
2558
2559     for (i = 0; i < dev->num_tx_queues; i++) {
2560         struct netdev_queue *txq = netdev_get_tx_queue(dev, i);
2561         netif_tx_wake_queue(txq);
2562     }
2563 }
2564
2565 static inline void netif_tx_stop_queue(struct netdev_queue *dev_queue)
2566 {
2567     set_bit(__QUEUE_STATE_DRV_XOFF, &dev_queue->state);
2568 }
2569
2570 /**
2571  *   netif_stop_queue - stop transmitted packets
2572  *   @dev: network device
2573  *
2574  *   Stop upper layers calling the device hard_start_xmit routine.
2575  *   Used for flow control when transmit resources are unavailable.
2576  */
2577 static inline void netif_stop_queue(struct net_device *dev)
2578 {
2579     netif_tx_stop_queue(netdev_get_tx_queue(dev, 0));
2580 }
2581
2582 void netif_tx_stop_all_queues(struct net_device *dev);
2583
2584 static inline bool netif_tx_queue_stopped(const struct netdev_queue *dev_queue)
2585 {
2586     return test_bit(__QUEUE_STATE_DRV_XOFF, &dev_queue->state);
2587 }
2588
2589 /**
2590  *   netif_queue_stopped - test if transmit queue is flowblocked
2591  *   @dev: network device
2592  *
2593  *   Test if transmit queue on device is currently unable to send.
2594  */
2595 static inline bool netif_queue_stopped(const struct net_device *dev)
2596 {
2597     return netif_tx_queue_stopped(netdev_get_tx_queue(dev, 0));
2598 }
2599
2600 static inline bool netif_xmit_stopped(const struct netdev_queue *dev_queue)
2601 {
2602     return dev_queue->state & QUEUE_STATE_ANY_XOFF;
2603 }
2604
2605 static inline bool
2606 netif_xmit_frozen_or_stopped(const struct netdev_queue *dev_queue)
2607 {
2608     return dev_queue->state & QUEUE_STATE_ANY_XOFF_OR_FROZEN;

```



```

2609 }
2610
2611 static inline bool
2612 netif_xmit_frozen_or_drv_stopped(const struct netdev_queue *dev_queue)
2613 {
2614     return dev_queue->state & QUEUE_STATE_DRV_XOFF_OR_FROZEN;
2615 }
2616
2617 /**
2618  * netdev_txq_bql_enqueue_prefetchw - prefetch bql data for write
2619  * @dev_queue: pointer to transmit queue
2620  *
2621  * BQL enabled drivers might use this helper in their ndo_start_xmit(),
2622  * to give appropriate hint to the cpu.
2623  */
2624 static inline void netdev_txq_bql_enqueue_prefetchw(struct netdev_queue *dev_queue)
2625 {
2626     #ifdef CONFIG_BQL
2627         prefetchw(&dev_queue->dql.num_queued);
2628     #endif
2629 }
2630
2631 /**
2632  * netdev_txq_bql_complete_prefetchw - prefetch bql data for write
2633  * @dev_queue: pointer to transmit queue
2634  *
2635  * BQL enabled drivers might use this helper in their TX completion path,
2636  * to give appropriate hint to the cpu.
2637  */
2638 static inline void netdev_txq_bql_complete_prefetchw(struct netdev_queue *dev_queue)
2639 {
2640     #ifdef CONFIG_BQL
2641         prefetchw(&dev_queue->dql.limit);
2642     #endif
2643 }
2644
2645 static inline void netdev_tx_sent_queue(struct netdev_queue *dev_queue,
2646                                         unsigned int bytes)
2647 {
2648     #ifdef CONFIG_BQL
2649         dql_queued(&dev_queue->dql, bytes);
2650
2651         if (likely(dql_avail(&dev_queue->dql) >= 0))
2652             return;
2653
2654         set_bit(__QUEUE_STATE_STACK_XOFF, &dev_queue->state);
2655
2656         /*
2657          * The XOFF flag must be set before checking the dql_avail below,
2658          * because in netdev_tx_completed_queue we update the dql_completed
2659          * before checking the XOFF flag.
2660          */
2661         smp_mb();
2662
2663         /* check again in case another CPU has just made room avail */
2664         if (unlikely(dql_avail(&dev_queue->dql) >= 0))
2665             clear_bit(__QUEUE_STATE_STACK_XOFF, &dev_queue->state);
2666     #endif
2667 }
2668
2669 /**
2670  * netdev_sent_queue - report the number of bytes queued to hardware
2671  * @dev: network device
2672  * @bytes: number of bytes queued to the hardware device queue
2673  *
2674  * Report the number of bytes queued for sending/completion to the network
2675  * device hardware queue. @bytes should be a good approximation and should
2676  * exactly match netdev_completed_queue() @bytes
2677  */
2678 static inline void netdev_sent_queue(struct net_device *dev, unsigned int bytes)
2679 {
2680     netdev_tx_sent_queue(netdev_get_tx_queue(dev, 0), bytes);
2681 }

```

```

2682
2683 static inline void netdev\_tx\_completed\_queue(struct netdev\_queue *dev_queue,
2684                                             unsigned int pkts, unsigned int bytes)
2685 {
2686     #ifdef CONFIG_BQL
2687         if (unlikely(!bytes))
2688             return;
2689
2690         dql\_completed(&dev_queue->dql, bytes);
2691
2692         /*
2693          * Without the memory barrier there is a small possibility that
2694          * netdev_tx_sent_queue will miss the update and cause the queue to
2695          * be stopped forever
2696          */
2697         smp\_mb();
2698
2699         if (dql\_avail(&dev_queue->dql) < 0)
2700             return;
2701
2702         if (test\_and\_clear\_bit(__QUEUE_STATE_STACK_XOFF, &dev_queue->state))
2703             netif\_schedule\_queue(dev_queue);
2704     #endif
2705 }
2706
2707 /**
2708  * netdev\_completed\_queue - report bytes and packets completed by device
2709  * @dev: network device
2710  * @pkts: actual number of packets sent over the medium
2711  * @bytes: actual number of bytes sent over the medium
2712  *
2713  * Report the number of bytes and packets transmitted by the network device
2714  * hardware queue over the physical medium, @bytes must exactly match the
2715  * @bytes amount passed to netdev\_sent\_queue()
2716  */
2717 static inline void netdev\_completed\_queue(struct net\_device *dev,
2718                                             unsigned int pkts, unsigned int bytes)
2719 {
2720     netdev\_tx\_completed\_queue(netdev\_get\_tx\_queue(dev, 0), pkts, bytes);
2721 }
2722
2723 static inline void netdev\_tx\_reset\_queue(struct netdev\_queue *q)
2724 {
2725     #ifdef CONFIG_BQL
2726         clear\_bit(__QUEUE_STATE_STACK_XOFF, &q->state);
2727         dql\_reset(&q->dql);
2728     #endif
2729 }
2730
2731 /**
2732  * netdev\_reset\_queue - reset the packets and bytes count of a network device
2733  * @dev_queue: network device
2734  *
2735  * Reset the bytes and packet count of a network device and clear the
2736  * software flow control OFF bit for this network device
2737  */
2738 static inline void netdev\_reset\_queue(struct net\_device *dev_queue)
2739 {
2740     netdev\_tx\_reset\_queue(netdev\_get\_tx\_queue(dev_queue, 0));
2741 }
2742
2743 /**
2744  * netdev\_cap\_txqueue - check if selected tx queue exceeds device queues
2745  * @dev: network device
2746  * @queue_index: given tx queue index
2747  *
2748  * Returns 0 if given tx queue index >= number of device tx queues,
2749  * otherwise returns the originally passed tx queue index.
2750  */
2751 static inline u16 netdev\_cap\_txqueue(struct net\_device *dev, u16 queue_index)
2752 {
2753     if (unlikely(queue_index >= dev->real_num_tx_queues)) {
2754         net\_warn\_ratelimited("%s selects TX queue %d, but real number of TX queues is %d\n",

```

```

2755         dev->name, queue_index,
2756         dev->real_num_tx_queues);
2757     return 0;
2758 }
2759
2760 return queue_index;
2761 }
2762
2763 /**
2764  *      netif_running - test if up
2765  *      @dev: network device
2766  *
2767  *      Test if the device has been brought up.
2768  */
2769 static inline bool netif_running(const struct net_device *dev)
2770 {
2771     return test_bit(__LINK_STATE_START, &dev->state);
2772 }
2773
2774 /**
2775  *      Routines to manage the subqueues on a device. We only need start
2776  *      stop, and a check if it's stopped. All other device management is
2777  *      done at the overall netdevice level.
2778  *      Also test the device if we're multiqueue.
2779  */
2780
2781 /**
2782  *      netif_start_subqueue - allow sending packets on subqueue
2783  *      @dev: network device
2784  *      @queue_index: sub queue index
2785  *
2786  *      Start individual transmit queue of a device with multiple transmit queues.
2787  */
2788 static inline void netif_start_subqueue(struct net_device *dev, u16 queue_index)
2789 {
2790     struct netdev_queue *txq = netdev_get_tx_queue(dev, queue_index);
2791
2792     netif_tx_start_queue(txq);
2793 }
2794
2795 /**
2796  *      netif_stop_subqueue - stop sending packets on subqueue
2797  *      @dev: network device
2798  *      @queue_index: sub queue index
2799  *
2800  *      Stop individual transmit queue of a device with multiple transmit queues.
2801  */
2802 static inline void netif_stop_subqueue(struct net_device *dev, u16 queue_index)
2803 {
2804     struct netdev_queue *txq = netdev_get_tx_queue(dev, queue_index);
2805     netif_tx_stop_queue(txq);
2806 }
2807
2808 /**
2809  *      netif_subqueue_stopped - test status of subqueue
2810  *      @dev: network device
2811  *      @queue_index: sub queue index
2812  *
2813  *      Check individual transmit queue of a device with multiple transmit queues.
2814  */
2815 static inline bool __netif_subqueue_stopped(const struct net_device *dev,
2816                                             u16 queue_index)
2817 {
2818     struct netdev_queue *txq = netdev_get_tx_queue(dev, queue_index);
2819
2820     return netif_tx_queue_stopped(txq);
2821 }
2822
2823 static inline bool netif_subqueue_stopped(const struct net_device *dev,
2824                                           struct sk_buff *skb)
2825 {
2826     return __netif_subqueue_stopped(dev, skb_get_queue_mapping(skb));
2827 }

```

```

2828
2829 void netif_wake_subqueue(struct net_device *dev, u16 queue_index);
2830
2831 #ifdef CONFIG_XPS
2832 int netif_set_xps_queue(struct net_device *dev, const struct cpumask *mask,
2833                        u16 index);
2834 #else
2835 static inline int netif_set_xps_queue(struct net_device *dev,
2836                                     const struct cpumask *mask,
2837                                     u16 index)
2838 {
2839     return 0;
2840 }
2841 #endif
2842
2843 u16 __skb_tx_hash(const struct net_device *dev, struct sk_buff *skb,
2844                  unsigned int num_tx_queues);
2845
2846 /*
2847  * Returns a Tx hash for the given packet when dev->real_num_tx_queues is used
2848  * as a distribution range limit for the returned value.
2849  */
2850 static inline u16 skb_tx_hash(const struct net_device *dev,
2851                              struct sk_buff *skb)
2852 {
2853     return __skb_tx_hash(dev, skb, dev->real_num_tx_queues);
2854 }
2855
2856 /**
2857  * netif_is_multiqueue - test if device has multiple transmit queues
2858  * @dev: network device
2859  *
2860  * Check if device has multiple transmit queues
2861  */
2862 static inline bool netif_is_multiqueue(const struct net_device *dev)
2863 {
2864     return dev->num_tx_queues > 1;
2865 }
2866
2867 int netif_set_real_num_tx_queues(struct net_device *dev, unsigned int txq);
2868
2869 #ifdef CONFIG_SYSFS
2870 int netif_set_real_num_rx_queues(struct net_device *dev, unsigned int rxq);
2871 #else
2872 static inline int netif_set_real_num_rx_queues(struct net_device *dev,
2873                                                unsigned int rxq)
2874 {
2875     return 0;
2876 }
2877 #endif
2878
2879 #ifdef CONFIG_SYSFS
2880 static inline unsigned int get_netdev_rx_queue_index(
2881     struct netdev_rx_queue *queue)
2882 {
2883     struct net_device *dev = queue->dev;
2884     int index = queue - dev->_rx;
2885
2886     BUG_ON(index >= dev->num_rx_queues);
2887     return index;
2888 }
2889 #endif
2890
2891 #define DEFAULT_MAX_NUM_RSS_QUEUES (8)
2892 int netif_get_num_default_rss_queues(void);
2893
2894 enum skb_free_reason {
2895     SKB_REASON_CONSUMED,
2896     SKB_REASON_DROPPED,
2897 };
2898
2899 void __dev_kfree_skb_irq(struct sk_buff *skb, enum skb_free_reason reason);
2900 void __dev_kfree_skb_any(struct sk_buff *skb, enum skb_free_reason reason);

```

```

2901
2902 /*
2903  * It is not allowed to call kfree_skb() or consume_skb() from hardware
2904  * interrupt context or with hardware interrupts being disabled.
2905  * (in_irq() || irqs_disabled())
2906  *
2907  * We provide four helpers that can be used in following contexts :
2908  *
2909  * dev_kfree_skb_irq(skb) when caller drops a packet from irq context,
2910  *   replacing kfree_skb(skb)
2911  *
2912  * dev_consume_skb_irq(skb) when caller consumes a packet from irq context.
2913  *   Typically used in place of consume_skb(skb) in TX completion path
2914  *
2915  * dev_kfree_skb_any(skb) when caller doesn't know its current irq context,
2916  *   replacing kfree_skb(skb)
2917  *
2918  * dev_consume_skb_any(skb) when caller doesn't know its current irq context,
2919  *   and consumed a packet. Used in place of consume_skb(skb)
2920  */
2921 static inline void dev_kfree_skb_irq(struct sk_buff *skb)
2922 {
2923     __dev_kfree_skb_irq(skb, SKB_REASON_DROPPED);
2924 }
2925
2926 static inline void dev_consume_skb_irq(struct sk_buff *skb)
2927 {
2928     __dev_kfree_skb_irq(skb, SKB_REASON_CONSUMED);
2929 }
2930
2931 static inline void dev_kfree_skb_any(struct sk_buff *skb)
2932 {
2933     __dev_kfree_skb_any(skb, SKB_REASON_DROPPED);
2934 }
2935
2936 static inline void dev_consume_skb_any(struct sk_buff *skb)
2937 {
2938     __dev_kfree_skb_any(skb, SKB_REASON_CONSUMED);
2939 }
2940
2941 int netif_rx(struct sk_buff *skb);
2942 int netif_rx_ni(struct sk_buff *skb);
2943 int netif_receive_skb_skb(struct sock *sk, struct sk_buff *skb);
2944 static inline int netif_receive_skb(struct sk_buff *skb)
2945 {
2946     return netif_receive_skb_skb(skb->sk, skb);
2947 }
2948 gro_result_t napi_gro_receive(struct napi_struct *napi, struct sk_buff *skb);
2949 void napi_gro_flush(struct napi_struct *napi, bool flush_old);
2950 struct sk_buff *napi_get_frags(struct napi_struct *napi);
2951 gro_result_t napi_gro_frags(struct napi_struct *napi);
2952 struct packet_offload *gro_find_receive_by_type(__be16 type);
2953 struct packet_offload *gro_find_complete_by_type(__be16 type);
2954
2955 static inline void napi_free_frags(struct napi_struct *napi)
2956 {
2957     kfree_skb(napi->skb);
2958     napi->skb = NULL;
2959 }
2960
2961 int netdev_rx_handler_register(struct net_device *dev,
2962                               rx_handler_func_t *rx_handler,
2963                               void *rx_handler_data);
2964 void netdev_rx_handler_unregister(struct net_device *dev);
2965
2966 bool dev_valid_name(const char *name);
2967 int dev_ioctl(struct net *net, unsigned int cmd, void __user *);
2968 int dev_ethtool(struct net *net, struct ifreq *);
2969 unsigned int dev_get_flags(const struct net_device *);
2970 int __dev_change_flags(struct net_device *, unsigned int flags);
2971 int dev_change_flags(struct net_device *, unsigned int);
2972 void __dev_notify_flags(struct net_device *, unsigned int old_flags,
2973                        unsigned int gchanges);

```

```

2974 int dev_change_name(struct net_device *, const char *);
2975 int dev_set_alias(struct net_device *, const char *, size_t);
2976 int dev_change_net_namespace(struct net_device *, struct net *, const char *);
2977 int dev_set_mtu(struct net_device *, int);
2978 void dev_set_group(struct net_device *, int);
2979 int dev_set_mac_address(struct net_device *, struct sockaddr *);
2980 int dev_change_carrier(struct net_device *, bool new_carrier);
2981 int dev_get_phys_port_id(struct net_device *dev,
2982                          struct netdev_phys_item_id *ppid);
2983 int dev_get_phys_port_name(struct net_device *dev,
2984                            char *name, size_t len);
2985 struct sk_buff *validate_xmit_skb_list(struct sk_buff *skb, struct net_device *dev);
2986 struct sk_buff *dev_hard_start_xmit(struct sk_buff *skb, struct net_device *dev,
2987                                     struct netdev_queue *txq, int *ret);
2988 int __dev_forward_skb(struct net_device *dev, struct sk_buff *skb);
2989 int dev_forward_skb(struct net_device *dev, struct sk_buff *skb);
2990 bool is_skb_forwardable(struct net_device *dev, struct sk_buff *skb);
2991
2992 extern int netdev_budget;
2993
2994 /* Called by rtnetlink.c:rtnl_unlock() */
2995 void netdev_run_todo(void);
2996
2997 /**
2998  * dev_put - release reference to device
2999  * @dev: network device
3000  *
3001  * Release reference to device to allow it to be freed.
3002  */
3003 static inline void dev_put(struct net_device *dev)
3004 {
3005     this_cpu_dec(*dev->pcpu_refcnt);
3006 }
3007
3008 /**
3009  * dev_hold - get reference to device
3010  * @dev: network device
3011  *
3012  * Hold reference to device to keep it from being freed.
3013  */
3014 static inline void dev_hold(struct net_device *dev)
3015 {
3016     this_cpu_inc(*dev->pcpu_refcnt);
3017 }
3018
3019 /* Carrier loss detection, dial on demand. The functions netif_carrier_on
3020  * and _off may be called from IRQ context, but it is caller
3021  * who is responsible for serialization of these calls.
3022  *
3023  * The name carrier is inappropriate, these functions should really be
3024  * called netif_lowerlayer_*() because they represent the state of any
3025  * kind of lower layer not just hardware media.
3026  */
3027
3028 void linkwatch_init_dev(struct net_device *dev);
3029 void linkwatch_fire_event(struct net_device *dev);
3030 void linkwatch_forget_dev(struct net_device *dev);
3031
3032 /**
3033  * netif_carrier_ok - test if carrier present
3034  * @dev: network device
3035  *
3036  * Check if carrier is present on device
3037  */
3038 static inline bool netif_carrier_ok(const struct net_device *dev)
3039 {
3040     return !test_bit(__LINK_STATE_NOCARRIER, &dev->state);
3041 }
3042
3043 unsigned long dev_trans_start(struct net_device *dev);
3044
3045 void __netdev_watchdog_up(struct net_device *dev);
3046

```



```

3047 void netif\_carrier\_on(struct net\_device *dev);
3048
3049 void netif\_carrier\_off(struct net\_device *dev);
3050
3051 /**
3052  *      netif\_dormant\_on - mark device as dormant.
3053  *      @dev: network device
3054  *
3055  *      Mark device as dormant (as per RFC2863).
3056  *
3057  *      The dormant state indicates that the relevant interface is not
3058  *      actually in a condition to pass packets (i.e., it is not 'up') but is
3059  *      in a "pending" state, waiting for some external event. For "on-
3060  *      demand" interfaces, this new state identifies the situation where the
3061  *      interface is waiting for events to place it in the up state.
3062  */
3063
3064 static inline void netif\_dormant\_on(struct net\_device *dev)
3065 {
3066     if (!test\_and\_set\_bit(__LINK_STATE_DORMANT, &dev->state))
3067         linkwatch\_fire\_event(dev);
3068 }
3069
3070 /**
3071  *      netif\_dormant\_off - set device as not dormant.
3072  *      @dev: network device
3073  *
3074  *      Device is not in dormant state.
3075  */
3076
3077 static inline void netif\_dormant\_off(struct net\_device *dev)
3078 {
3079     if (test\_and\_clear\_bit(__LINK_STATE_DORMANT, &dev->state))
3080         linkwatch\_fire\_event(dev);
3081 }
3082
3083 /**
3084  *      netif\_dormant - test if carrier present
3085  *      @dev: network device
3086  *
3087  *      Check if carrier is present on device
3088  */
3089
3090 static inline bool netif\_dormant(const struct net\_device *dev)
3091 {
3092     return test\_bit(__LINK_STATE_DORMANT, &dev->state);
3093 }
3094
3095 /**
3096  *      netif\_oper\_up - test if device is operational
3097  *      @dev: network device
3098  *
3099  *      Check if carrier is operational
3100  */
3101
3102 static inline bool netif\_oper\_up(const struct net\_device *dev)
3103 {
3104     return (dev->operstate == IF_OPER_UP ||
3105             dev->operstate == IF_OPER_UNKNOWN /* backward compat */);
3106 }
3107
3108 /**
3109  *      netif\_device\_present - is device available or removed
3110  *      @dev: network device
3111  *
3112  *      Check if device has not been removed from system.
3113  */
3114
3115 static inline bool netif\_device\_present(struct net\_device *dev)
3116 {
3117     return test\_bit(__LINK_STATE_PRESENT, &dev->state);
3118 }
3119
3120 void netif\_device\_detach(struct net\_device *dev);
3121
3122 void netif\_device\_attach(struct net\_device *dev);

```

```

3120
3121 /*
3122  * Network interface message level settings
3123  */
3124
3125 enum {
3126     NETIF_MSG_DRV           = 0x0001,
3127     NETIF_MSG_PROBE        = 0x0002,
3128     NETIF_MSG_LINK         = 0x0004,
3129     NETIF_MSG_TIMER        = 0x0008,
3130     NETIF_MSG_IFDOWN       = 0x0010,
3131     NETIF_MSG_IFUP         = 0x0020,
3132     NETIF_MSG_RX_ERR       = 0x0040,
3133     NETIF_MSG_TX_ERR       = 0x0080,
3134     NETIF_MSG_TX_QUEUED     = 0x0100,
3135     NETIF_MSG_INTR         = 0x0200,
3136     NETIF_MSG_TX_DONE      = 0x0400,
3137     NETIF_MSG_RX_STATUS    = 0x0800,
3138     NETIF_MSG_PKTDATA      = 0x1000,
3139     NETIF_MSG_HW           = 0x2000,
3140     NETIF_MSG_WOL          = 0x4000,
3141 };
3142
3143 #define netif_msg_drv(p)      ((p)->msg_enable & NETIF_MSG_DRV)
3144 #define netif_msg_probe(p)   ((p)->msg_enable & NETIF_MSG_PROBE)
3145 #define netif_msg_link(p)    ((p)->msg_enable & NETIF_MSG_LINK)
3146 #define netif_msg_timer(p)   ((p)->msg_enable & NETIF_MSG_TIMER)
3147 #define netif_msg_ifdown(p)  ((p)->msg_enable & NETIF_MSG_IFDOWN)
3148 #define netif_msg_ifup(p)    ((p)->msg_enable & NETIF_MSG_IFUP)
3149 #define netif_msg_rx_err(p)  ((p)->msg_enable & NETIF_MSG_RX_ERR)
3150 #define netif_msg_tx_err(p)  ((p)->msg_enable & NETIF_MSG_TX_ERR)
3151 #define netif_msg_tx_queued(p) ((p)->msg_enable & NETIF_MSG_TX_QUEUED)
3152 #define netif_msg_intr(p)    ((p)->msg_enable & NETIF_MSG_INTR)
3153 #define netif_msg_tx_done(p) ((p)->msg_enable & NETIF_MSG_TX_DONE)
3154 #define netif_msg_rx_status(p) ((p)->msg_enable & NETIF_MSG_RX_STATUS)
3155 #define netif_msg_pktdata(p) ((p)->msg_enable & NETIF_MSG_PKTDATA)
3156 #define netif_msg_hw(p)      ((p)->msg_enable & NETIF_MSG_HW)
3157 #define netif_msg_wol(p)     ((p)->msg_enable & NETIF_MSG_WOL)
3158
3159 static inline u32 netif_msg_init(int debug_value, int default_msg_enable_bits)
3160 {
3161     /* use default */
3162     if (debug_value < 0 || debug_value >= (sizeof(u32) * 8))
3163         return default_msg_enable_bits;
3164     if (debug_value == 0) /* no output */
3165         return 0;
3166     /* set low N bits */
3167     return (1 << debug_value) - 1;
3168 }
3169
3170 static inline void __netif_tx_lock(struct netdev_queue *txq, int cpu)
3171 {
3172     spin_lock(&txq->xmit_lock);
3173     txq->xmit_lock_owner = cpu;
3174 }
3175
3176 static inline void __netif_tx_lock_bh(struct netdev_queue *txq)
3177 {
3178     spin_lock_bh(&txq->xmit_lock);
3179     txq->xmit_lock_owner = smp_processor_id();
3180 }
3181
3182 static inline bool __netif_tx_trylock(struct netdev_queue *txq)
3183 {
3184     bool ok = spin_trylock(&txq->xmit_lock);
3185     if (likely(ok))
3186         txq->xmit_lock_owner = smp_processor_id();
3187     return ok;
3188 }
3189
3190 static inline void __netif_tx_unlock(struct netdev_queue *txq)
3191 {
3192     txq->xmit_lock_owner = -1;

```

```

3193         spin_unlock(&txq->_xmit_lock);
3194     }
3195
3196 static inline void __netif_tx_unlock_bh(struct netdev_queue *txq)
3197 {
3198     txq->xmit_lock_owner = -1;
3199     spin_unlock_bh(&txq->_xmit_lock);
3200 }
3201
3202 static inline void txq_trans_update(struct netdev_queue *txq)
3203 {
3204     if (txq->xmit_lock_owner != -1)
3205         txq->trans_start = jiffies;
3206 }
3207
3208 /**
3209  *      netif_tx_lock - grab network device transmit lock
3210  *      @dev: network device
3211  *
3212  *      Get network device transmit lock
3213  */
3214 static inline void netif_tx_lock(struct net_device *dev)
3215 {
3216     unsigned int i;
3217     int cpu;
3218
3219     spin_lock(&dev->tx_global_lock);
3220     cpu = smp_processor_id();
3221     for (i = 0; i < dev->num_tx_queues; i++) {
3222         struct netdev_queue *txq = netdev_get_tx_queue(dev, i);
3223
3224         /* We are the only thread of execution doing a
3225          * freeze, but we have to grab the _xmit_lock in
3226          * order to synchronize with threads which are in
3227          * the ->hard_start_xmit() handler and already
3228          * checked the frozen bit.
3229          */
3230         __netif_tx_lock(txq, cpu);
3231         set_bit(__QUEUE_STATE_FROZEN, &txq->state);
3232         __netif_tx_unlock(txq);
3233     }
3234 }
3235
3236 static inline void netif_tx_lock_bh(struct net_device *dev)
3237 {
3238     local_bh_disable();
3239     netif_tx_lock(dev);
3240 }
3241
3242 static inline void netif_tx_unlock(struct net_device *dev)
3243 {
3244     unsigned int i;
3245
3246     for (i = 0; i < dev->num_tx_queues; i++) {
3247         struct netdev_queue *txq = netdev_get_tx_queue(dev, i);
3248
3249         /* No need to grab the _xmit_lock here. If the
3250          * queue is not stopped for another reason, we
3251          * force a schedule.
3252          */
3253         clear_bit(__QUEUE_STATE_FROZEN, &txq->state);
3254         netif_schedule_queue(txq);
3255     }
3256     spin_unlock(&dev->tx_global_lock);
3257 }
3258
3259 static inline void netif_tx_unlock_bh(struct net_device *dev)
3260 {
3261     netif_tx_unlock(dev);
3262     local_bh_enable();
3263 }
3264
3265 #define HARD_TX_LOCK(dev, txq, cpu) {

```

\

```

3266     if ((dev->features & NETIF_F_LLTX) == 0) {           \
3267         __netif_tx_lock(txq, cpu);                       \
3268     }                                                     \
3269 }
3270
3271 #define HARD_TX_TRYLOCK(dev, txq)                        \
3272     (((dev->features & NETIF_F_LLTX) == 0) ?             \
3273         __netif_tx_trylock(txq) :                       \
3274         true )
3275
3276 #define HARD_TX_UNLOCK(dev, txq) {                      \
3277     if ((dev->features & NETIF_F_LLTX) == 0) {           \
3278         __netif_tx_unlock(txq);                         \
3279     }                                                     \
3280 }
3281
3282 static inline void netif_tx_disable(struct net_device *dev)
3283 {
3284     unsigned int i;
3285     int cpu;
3286
3287     local_bh_disable();
3288     cpu = smp_processor_id();
3289     for (i = 0; i < dev->num_tx_queues; i++) {
3290         struct netdev_queue *txq = netdev_get_tx_queue(dev, i);
3291
3292         __netif_tx_lock(txq, cpu);
3293         netif_tx_stop_queue(txq);
3294         __netif_tx_unlock(txq);
3295     }
3296     local_bh_enable();
3297 }
3298
3299 static inline void netif_addr_lock(struct net_device *dev)
3300 {
3301     spin_lock(&dev->addr_list_lock);
3302 }
3303
3304 static inline void netif_addr_lock_nested(struct net_device *dev)
3305 {
3306     int subclass = SINGLE_DEPTH_NESTING;
3307
3308     if (dev->netdev_ops->ndo_get_lock_subclass)
3309         subclass = dev->netdev_ops->ndo_get_lock_subclass(dev);
3310
3311     spin_lock_nested(&dev->addr_list_lock, subclass);
3312 }
3313
3314 static inline void netif_addr_lock_bh(struct net_device *dev)
3315 {
3316     spin_lock_bh(&dev->addr_list_lock);
3317 }
3318
3319 static inline void netif_addr_unlock(struct net_device *dev)
3320 {
3321     spin_unlock(&dev->addr_list_lock);
3322 }
3323
3324 static inline void netif_addr_unlock_bh(struct net_device *dev)
3325 {
3326     spin_unlock_bh(&dev->addr_list_lock);
3327 }
3328
3329 /*
3330  * dev_addrs walker. Should be used only for read access. Call with
3331  * rcu_read_lock held.
3332  */
3333 #define for_each_dev_addr(dev, ha) \
3334     list_for_each_entry_rcu(ha, &dev->dev_addrs.list, list)
3335
3336 /* These functions live elsewhere (drivers/net/net_init.c, but related) */
3337
3338 void ether_setup(struct net_device *dev);

```

```

3339
3340 /* Support for loadable net-drivers */
3341 struct net_device *alloc_netdev_mqs(int sizeof_priv, const char *name,
3342                                     unsigned char name_assign_type,
3343                                     void (*setup)(struct net_device *),
3344                                     unsigned int txqs, unsigned int rxqs);
3345 #define alloc_netdev(sizeof_priv, name, name_assign_type, setup) \
3346     alloc_netdev_mqs(sizeof_priv, name, name_assign_type, setup, 1, 1)
3347
3348 #define alloc_netdev_mq(sizeof_priv, name, name_assign_type, setup, count) \
3349     alloc_netdev_mqs(sizeof_priv, name, name_assign_type, setup, count, \
3350                      count)
3351
3352 int register_netdev(struct net_device *dev);
3353 void unregister_netdev(struct net_device *dev);
3354
3355 /* General hardware address lists handling functions */
3356 int __hw_addr_sync(struct netdev_hw_addr_list *to_list,
3357                   struct netdev_hw_addr_list *from_list, int addr_len);
3358 void __hw_addr_unsync(struct netdev_hw_addr_list *to_list,
3359                      struct netdev_hw_addr_list *from_list, int addr_len);
3360 int __hw_addr_sync_dev(struct netdev_hw_addr_list *list,
3361                       struct net_device *dev,
3362                       int (*sync)(struct net_device *, const unsigned char *),
3363                       int (*unsync)(struct net_device *,
3364                                     const unsigned char *));
3365 void __hw_addr_unsync_dev(struct netdev_hw_addr_list *list,
3366                          struct net_device *dev,
3367                          int (*unsync)(struct net_device *,
3368                                        const unsigned char *));
3368 void __hw_addr_init(struct netdev_hw_addr_list *list);
3369
3370
3371 /* Functions used for device addresses handling */
3372 int dev_addr_add(struct net_device *dev, const unsigned char *addr,
3373                 unsigned char addr_type);
3374 int dev_addr_del(struct net_device *dev, const unsigned char *addr,
3375                 unsigned char addr_type);
3376 void dev_addr_flush(struct net_device *dev);
3377 int dev_addr_init(struct net_device *dev);
3378
3379 /* Functions used for unicast addresses handling */
3380 int dev_uc_add(struct net_device *dev, const unsigned char *addr);
3381 int dev_uc_add_excl(struct net_device *dev, const unsigned char *addr);
3382 int dev_uc_del(struct net_device *dev, const unsigned char *addr);
3383 int dev_uc_sync(struct net_device *to, struct net_device *from);
3384 int dev_uc_sync_multiple(struct net_device *to, struct net_device *from);
3385 void dev_uc_unsync(struct net_device *to, struct net_device *from);
3386 void dev_uc_flush(struct net_device *dev);
3387 void dev_uc_init(struct net_device *dev);
3388
3389 /**
3390  * __dev_uc_sync - Synchronize device's unicast list
3391  * @dev: device to sync
3392  * @sync: function to call if address should be added
3393  * @unsync: function to call if address should be removed
3394  *
3395  * Add newly added addresses to the interface, and release
3396  * addresses that have been deleted.
3397  */
3398 static inline int __dev_uc_sync(struct net_device *dev,
3399                                int (*sync)(struct net_device *,
3400                                             const unsigned char *),
3401                                int (*unsync)(struct net_device *,
3402                                              const unsigned char *))
3403 {
3404     return __hw_addr_sync_dev(&dev->uc, dev, sync, unsync);
3405 }
3406
3407 /**
3408  * __dev_uc_unsync - Remove synchronized addresses from device
3409  * @dev: device to sync
3410  * @unsync: function to call if address should be removed
3411  *

```

```

3412 * Remove all addresses that were added to the device by dev_uc_sync().
3413 **/
3414 static inline void __dev_uc_unsync(struct net_device *dev,
3415                                   int (*unsync)(struct net_device *,
3416                                                  const unsigned char *))
3417 {
3418     __hw_addr_unsync_dev(&dev->uc, dev, unsync);
3419 }
3420
3421 /* Functions used for multicast addresses handling */
3422 int dev_mc_add(struct net_device *dev, const unsigned char *addr);
3423 int dev_mc_add_global(struct net_device *dev, const unsigned char *addr);
3424 int dev_mc_add_excl(struct net_device *dev, const unsigned char *addr);
3425 int dev_mc_del(struct net_device *dev, const unsigned char *addr);
3426 int dev_mc_del_global(struct net_device *dev, const unsigned char *addr);
3427 int dev_mc_sync(struct net_device *to, struct net_device *from);
3428 int dev_mc_sync_multiple(struct net_device *to, struct net_device *from);
3429 void dev_mc_unsync(struct net_device *to, struct net_device *from);
3430 void dev_mc_flush(struct net_device *dev);
3431 void dev_mc_init(struct net_device *dev);
3432
3433 /**
3434  * __dev_mc_sync - Synchronize device's multicast List
3435  * @dev: device to sync
3436  * @sync: function to call if address should be added
3437  * @unsync: function to call if address should be removed
3438  *
3439  * Add newly added addresses to the interface, and release
3440  * addresses that have been deleted.
3441  */
3442 static inline int __dev_mc_sync(struct net_device *dev,
3443                                 int (*sync)(struct net_device *,
3444                                              const unsigned char *),
3445                                 int (*unsync)(struct net_device *,
3446                                              const unsigned char *))
3447 {
3448     return __hw_addr_sync_dev(&dev->mc, dev, sync, unsync);
3449 }
3450
3451 /**
3452  * __dev_mc_unsync - Remove synchronized addresses from device
3453  * @dev: device to sync
3454  * @unsync: function to call if address should be removed
3455  *
3456  * Remove all addresses that were added to the device by dev_mc_sync().
3457  */
3458 static inline void __dev_mc_unsync(struct net_device *dev,
3459                                    int (*unsync)(struct net_device *,
3460                                                  const unsigned char *))
3461 {
3462     __hw_addr_unsync_dev(&dev->mc, dev, unsync);
3463 }
3464
3465 /* Functions used for secondary unicast and multicast support */
3466 void dev_set_rx_mode(struct net_device *dev);
3467 void __dev_set_rx_mode(struct net_device *dev);
3468 int dev_set_promiscuity(struct net_device *dev, int inc);
3469 int dev_set_allmulti(struct net_device *dev, int inc);
3470 void netdev_state_change(struct net_device *dev);
3471 void netdev_notify_peers(struct net_device *dev);
3472 void netdev_features_change(struct net_device *dev);
3473 /* Load a device via the kmod */
3474 void dev_load(struct net *net, const char *name);
3475 struct rtnl_link_stats64 *dev_get_stats(struct net_device *dev,
3476                                         struct rtnl_link_stats64 *storage);
3477 void netdev_stats_to_stats64(struct rtnl_link_stats64 *stats64,
3478                             const struct net_device_stats *netdev_stats);
3479
3480 extern int netdev_max_backlog;
3481 extern int netdev_tstamp_prequeue;
3482 extern int weight_p;
3483 extern int bpf_jit_enable;
3484

```



```

3485 bool netdev_has_upper_dev(struct net_device *dev, struct net_device *upper_dev);
3486 struct net_device *netdev_upper_get_next_dev_rcu(struct net_device *dev,
3487 struct list_head **iter);
3488 struct net_device *netdev_all_upper_get_next_dev_rcu(struct net_device *dev,
3489 struct list_head **iter);
3490
3491 /* iterate through upper list, must be called under RCU read lock */
3492 #define netdev_for_each_upper_dev_rcu(dev, updev, iter) \
3493     for (iter = &(dev)->adj_list.upper, \
3494          updev = netdev_upper_get_next_dev_rcu(dev, &(iter)); \
3495          updev; \
3496          updev = netdev_upper_get_next_dev_rcu(dev, &(iter)))
3497
3498 /* iterate through upper list, must be called under RCU read lock */
3499 #define netdev_for_each_all_upper_dev_rcu(dev, updev, iter) \
3500     for (iter = &(dev)->all_adj_list.upper, \
3501          updev = netdev_all_upper_get_next_dev_rcu(dev, &(iter)); \
3502          updev; \
3503          updev = netdev_all_upper_get_next_dev_rcu(dev, &(iter)))
3504
3505 void *netdev_lower_get_next_private(struct net_device *dev,
3506 struct list_head **iter);
3507 void *netdev_lower_get_next_private_rcu(struct net_device *dev,
3508 struct list_head **iter);
3509
3510 #define netdev_for_each_lower_private(dev, priv, iter) \
3511     for (iter = (dev)->adj_list.lower.next, \
3512          priv = netdev_lower_get_next_private(dev, &(iter)); \
3513          priv; \
3514          priv = netdev_lower_get_next_private(dev, &(iter)))
3515
3516 #define netdev_for_each_lower_private_rcu(dev, priv, iter) \
3517     for (iter = &(dev)->adj_list.lower, \
3518          priv = netdev_lower_get_next_private_rcu(dev, &(iter)); \
3519          priv; \
3520          priv = netdev_lower_get_next_private_rcu(dev, &(iter)))
3521
3522 void *netdev_lower_get_next(struct net_device *dev,
3523 struct list_head **iter);
3524 #define netdev_for_each_lower_dev(dev, ldev, iter) \
3525     for (iter = &(dev)->adj_list.lower, \
3526          ldev = netdev_lower_get_next(dev, &(iter)); \
3527          ldev; \
3528          ldev = netdev_lower_get_next(dev, &(iter)))
3529
3530 void *netdev_adjacent_get_private(struct list_head *adj_list);
3531 void *netdev_lower_get_first_private_rcu(struct net_device *dev);
3532 struct net_device *netdev_master_upper_dev_get(struct net_device *dev);
3533 struct net_device *netdev_master_upper_dev_get_rcu(struct net_device *dev);
3534 int netdev_upper_dev_link(struct net_device *dev, struct net_device *upper_dev);
3535 int netdev_master_upper_dev_link(struct net_device *dev,
3536 struct net_device *upper_dev);
3537 int netdev_master_upper_dev_link_private(struct net_device *dev,
3538 struct net_device *upper_dev,
3539 void *private);
3540 void netdev_upper_dev_unlink(struct net_device *dev,
3541 struct net_device *upper_dev);
3542 void netdev_adjacent_rename_links(struct net_device *dev, char *oldname);
3543 void *netdev_lower_dev_get_private(struct net_device *dev,
3544 struct net_device *lower_dev);
3545
3546 /* RSS keys are 40 or 52 bytes long */
3547 #define NETDEV_RSS_KEY_LEN 52
3548 extern u8 netdev_rss_key[NETDEV_RSS_KEY_LEN];
3549 void netdev_rss_key_fill(void *buffer, size_t len);
3550
3551 int dev_get_nest_level(struct net_device *dev,
3552 bool (*type_check)(struct net_device *dev));
3553 int skb_checksum_help(struct sk_buff *skb);
3554 struct sk_buff *__skb_gso_segment(struct sk_buff *skb,
3555 netdev_features_t features, bool tx_path);
3556 struct sk_buff *skb_mac_gso_segment(struct sk_buff *skb,
3557 netdev_features_t features);

```

```

3558
3559 struct netdev\_bonding\_info {
3560     ifslave slave;
3561     ifbond master;
3562 };
3563
3564 struct netdev\_notifier\_bonding\_info {
3565     struct netdev\_notifier\_info info; /* must be first */
3566     struct netdev\_bonding\_info bonding\_info;
3567 };
3568
3569 void netdev\_bonding\_info\_change(struct net\_device *dev,
3570                                struct netdev\_bonding\_info *bonding\_info);
3571
3572 static inline
3573 struct sk\_buff *skb\_gso\_segment(struct sk\_buff *skb, netdev\_features\_t features)
3574 {
3575     return \_\_skb\_gso\_segment(skb, features, true);
3576 }
3577
3578 \_\_be16 skb\_network\_protocol(struct sk\_buff *skb, int *depth);
3579
3580 static inline bool can\_checksum\_protocol(netdev\_features\_t features,
3581                                         \_\_be16 protocol)
3582 {
3583     return ((features & NETIF\_F\_GEN\_CSUM) ||
3584            ((features & NETIF\_F\_V4\_CSUM) &&
3585             protocol == htons(ETH\_P\_IP)) ||
3586            ((features & NETIF\_F\_V6\_CSUM) &&
3587             protocol == htons(ETH\_P\_IPV6)) ||
3588            ((features & NETIF\_F\_FCOE\_CRC) &&
3589             protocol == htons(ETH\_P\_FCOE)));
3590 }
3591
3592 #ifdef CONFIG\_BUG
3593 void netdev\_rx\_csum\_fault(struct net\_device *dev);
3594 #else
3595 static inline void netdev\_rx\_csum\_fault(struct net\_device *dev)
3596 {
3597 }
3598 #endif
3599
3600 /* rx skb timestamps */
3601 void net\_enable\_timestamp(void);
3602 void net\_disable\_timestamp(void);
3603
3604 #ifdef CONFIG\_PROC\_FS
3605 int \_\_init dev\_proc\_init(void);
3606 #else
3607 #define dev\_proc\_init() 0
3608 #endif
3609
3610 static inline netdev\_tx\_t \_\_netdev\_start\_xmit(const struct net\_device\_ops *ops,
3611                                              struct sk\_buff *skb, struct net\_device *dev,
3612                                              bool more)
3613 {
3614     skb->xmit\_more = more ? 1 : 0;
3615     return ops->ndo\_start\_xmit(skb, dev);
3616 }
3617
3618 static inline netdev\_tx\_t netdev\_start\_xmit(struct sk\_buff *skb, struct net\_device *dev,
3619                                             struct netdev\_queue *txq, bool more)
3620 {
3621     const struct net\_device\_ops *ops = dev->netdev\_ops;
3622     int rc;
3623
3624     rc = \_\_netdev\_start\_xmit(ops, skb, dev, more);
3625     if (rc == NETDEV\_TX\_OK)
3626         txq\_trans\_update(txq);
3627
3628     return rc;
3629 }
3630
3631 int netdev\_class\_create\_file\_ns(struct class\_attribute *class\_attr,
3632                                const void *ns);

```

```

3631 void netdev\_class\_remove\_file\_ns(struct class\_attribute *class_attr,
3632                                   const void *ns);
3633
3634 static inline int netdev\_class\_create\_file(struct class\_attribute *class_attr)
3635 {
3636     return netdev\_class\_create\_file\_ns(class_attr, NULL);
3637 }
3638
3639 static inline void netdev\_class\_remove\_file(struct class\_attribute *class_attr)
3640 {
3641     netdev\_class\_remove\_file\_ns(class_attr, NULL);
3642 }
3643
3644 extern struct kobj\_ns\_type\_operations net\_ns\_type\_operations;
3645
3646 const char *netdev\_drivername(const struct net\_device *dev);
3647
3648 void linkwatch\_run\_queue(void);
3649
3650 static inline netdev\_features\_t netdev\_intersect\_features(netdev\_features\_t f1,
3651                                                         netdev\_features\_t f2)
3652 {
3653     if (f1 & NETIF\_F\_GEN\_CSUM)
3654         f1 |= (NETIF\_F\_ALL\_CSUM & ~NETIF\_F\_GEN\_CSUM);
3655     if (f2 & NETIF\_F\_GEN\_CSUM)
3656         f2 |= (NETIF\_F\_ALL\_CSUM & ~NETIF\_F\_GEN\_CSUM);
3657     f1 &= f2;
3658     if (f1 & NETIF\_F\_GEN\_CSUM)
3659         f1 &= ~(NETIF\_F\_ALL\_CSUM & ~NETIF\_F\_GEN\_CSUM);
3660
3661     return f1;
3662 }
3663
3664 static inline netdev\_features\_t netdev\_get\_wanted\_features(
3665     struct net\_device *dev)
3666 {
3667     return (dev->features & ~dev->hw_features) | dev->wanted_features;
3668 }
3669 netdev\_features\_t netdev\_increment\_features(netdev\_features\_t all,
3670                                             netdev\_features\_t one, netdev\_features\_t mask);
3671
3672 /* Allow TSO being used on stacked device :
3673  * Performing the GSO segmentation before last device
3674  * is a performance improvement.
3675  */
3676 static inline netdev\_features\_t netdev\_add\_tso\_features(netdev\_features\_t features,
3677                                                         netdev\_features\_t mask)
3678 {
3679     return netdev\_increment\_features(features, NETIF\_F\_ALL\_TSO, mask);
3680 }
3681
3682 int \_\_netdev\_update\_features(struct net\_device *dev);
3683 void netdev\_update\_features(struct net\_device *dev);
3684 void netdev\_change\_features(struct net\_device *dev);
3685
3686 void netif\_stacked\_transfer\_operstate(const struct net\_device *rootdev,
3687                                       struct net\_device *dev);
3688
3689 netdev\_features\_t passthru\_features\_check(struct sk\_buff *skb,
3690                                           struct net\_device *dev,
3691                                           netdev\_features\_t features);
3692 netdev\_features\_t netif\_skb\_features(struct sk\_buff *skb);
3693
3694 static inline bool net\_gso\_ok(netdev\_features\_t features, int gso_type)
3695 {
3696     netdev\_features\_t feature = gso_type << NETIF\_F\_GSO\_SHIFT;
3697
3698     /* check flags correspondence */
3699     BUILD\_BUG\_ON(SKB\_GSO\_TCPV4 != (NETIF\_F\_TSO >> NETIF\_F\_GSO\_SHIFT));
3700     BUILD\_BUG\_ON(SKB\_GSO\_UDP != (NETIF\_F\_UFO >> NETIF\_F\_GSO\_SHIFT));
3701     BUILD\_BUG\_ON(SKB\_GSO\_DODGY != (NETIF\_F\_GSO\_ROBUST >> NETIF\_F\_GSO\_SHIFT));
3702     BUILD\_BUG\_ON(SKB\_GSO\_TCP\_ECN != (NETIF\_F\_TSO\_ECN >> NETIF\_F\_GSO\_SHIFT));
3703     BUILD\_BUG\_ON(SKB\_GSO\_TCPV6 != (NETIF\_F\_TSO6 >> NETIF\_F\_GSO\_SHIFT));

```

```

3704 BUILD_BUG_ON(SKB_GSO_FCOE != (NETIF_F_GSO >> NETIF_F_GSO_SHIFT));
3705 BUILD_BUG_ON(SKB_GSO_GRE != (NETIF_F_GSO_GRE >> NETIF_F_GSO_SHIFT));
3706 BUILD_BUG_ON(SKB_GSO_GRE_CSUM != (NETIF_F_GSO_GRE_CSUM >> NETIF_F_GSO_SHIFT));
3707 BUILD_BUG_ON(SKB_GSO_IPIP != (NETIF_F_GSO_IPIP >> NETIF_F_GSO_SHIFT));
3708 BUILD_BUG_ON(SKB_GSO_SIT != (NETIF_F_GSO_SIT >> NETIF_F_GSO_SHIFT));
3709 BUILD_BUG_ON(SKB_GSO_UDP_TUNNEL != (NETIF_F_GSO_UDP_TUNNEL >> NETIF_F_GSO_SHIFT));
3710 BUILD_BUG_ON(SKB_GSO_UDP_TUNNEL_CSUM != (NETIF_F_GSO_UDP_TUNNEL_CSUM >> NETIF_F_GSO_SHIFT));
3711 BUILD_BUG_ON(SKB_GSO_TUNNEL_REMCSUM != (NETIF_F_GSO_TUNNEL_REMCSUM >> NETIF_F_GSO_SHIFT));
3712
3713 return (features & feature) == feature;
3714 }
3715
3716 static inline bool skb_gso_ok(struct sk_buff *skb, netdev_features_t features)
3717 {
3718     return net_gso_ok(features, skb_shinfo(skb)->gso_type) &&
3719         (!skb_has_frag_list(skb) || (features & NETIF_F_FRAGLIST));
3720 }
3721
3722 static inline bool netif_needs_gso(struct sk_buff *skb,
3723                                   netdev_features_t features)
3724 {
3725     return skb_is_gso(skb) && (!skb_gso_ok(skb, features) ||
3726                                unlikely((skb->ip_summed != CHECKSUM_PARTIAL) &&
3727                                           (skb->ip_summed != CHECKSUM_UNNECESSARY)));
3728 }
3729
3730 static inline void netif_set_gso_max_size(struct net_device *dev,
3731                                           unsigned int size)
3732 {
3733     dev->gso_max_size = size;
3734 }
3735
3736 static inline void skb_gso_error_unwind(struct sk_buff *skb, __be16 protocol,
3737                                         int pulled_hlen, u16 mac_offset,
3738                                         int mac_len)
3739 {
3740     skb->protocol = protocol;
3741     skb->encapsulation = 1;
3742     skb_push(skb, pulled_hlen);
3743     skb_reset_transport_header(skb);
3744     skb->mac_header = mac_offset;
3745     skb->network_header = skb->mac_header + mac_len;
3746     skb->mac_len = mac_len;
3747 }
3748
3749 static inline bool netif_is_macvlan(struct net_device *dev)
3750 {
3751     return dev->priv_flags & IFF_MACVLAN;
3752 }
3753
3754 static inline bool netif_is_macvlan_port(struct net_device *dev)
3755 {
3756     return dev->priv_flags & IFF_MACVLAN_PORT;
3757 }
3758
3759 static inline bool netif_is_ipvlan(struct net_device *dev)
3760 {
3761     return dev->priv_flags & IFF_IPVLAN_SLAVE;
3762 }
3763
3764 static inline bool netif_is_ipvlan_port(struct net_device *dev)
3765 {
3766     return dev->priv_flags & IFF_IPVLAN_MASTER;
3767 }
3768
3769 static inline bool netif_is_bond_master(struct net_device *dev)
3770 {
3771     return dev->flags & IFF_MASTER && dev->priv_flags & IFF_BONDING;
3772 }
3773
3774 static inline bool netif_is_bond_slave(struct net_device *dev)
3775 {
3776     return dev->flags & IFF_SLAVE && dev->priv_flags & IFF_BONDING;

```

```

3777 }
3778
3779 static inline bool netif_supports_nofcs(struct net_device *dev)
3780 {
3781     return dev->priv_flags & IFF_SUPP_NOFCS;
3782 }
3783
3784 /* This device needs to keep skb dst for qdisc enqueue or ndo_start_xmit() */
3785 static inline void netif_keep_dst(struct net_device *dev)
3786 {
3787     dev->priv_flags &= ~(IFF_XMIT_DST_RELEASE | IFF_XMIT_DST_RELEASE_PERM);
3788 }
3789
3790 extern struct pernet_operations __net_initdata loopback_net_ops;
3791
3792 /* Logging, debugging and troubleshooting/diagnostic helpers. */
3793
3794 /* netdev_printk helpers, similar to dev_printk */
3795
3796 static inline const char *netdev_name(const struct net_device *dev)
3797 {
3798     if (!dev->name[0] || strchr(dev->name, '%'))
3799         return "(unnamed net_device)";
3800     return dev->name;
3801 }
3802
3803 static inline const char *netdev_reg_state(const struct net_device *dev)
3804 {
3805     switch (dev->reg_state) {
3806     case NETREG_UNINITIALIZED: return "(uninitialized)";
3807     case NETREG_REGISTERED: return "";
3808     case NETREG_UNREGISTERING: return "(unregistering)";
3809     case NETREG_UNREGISTERED: return "(unregistered)";
3810     case NETREG_RELEASED: return "(released)";
3811     case NETREG_DUMMY: return "(dummy)";
3812     }
3813
3814     WARN_ONCE(1, "%s: unknown reg_state %d\n", dev->name, dev->reg_state);
3815     return "(unknown)";
3816 }
3817
3818 __printf(3, 4)
3819 void netdev_printk(const char *level, const struct net_device *dev,
3820                  const char *format, ...);
3821 __printf(2, 3)
3822 void netdev_emerg(const struct net_device *dev, const char *format, ...);
3823 __printf(2, 3)
3824 void netdev_alert(const struct net_device *dev, const char *format, ...);
3825 __printf(2, 3)
3826 void netdev_crit(const struct net_device *dev, const char *format, ...);
3827 __printf(2, 3)
3828 void netdev_err(const struct net_device *dev, const char *format, ...);
3829 __printf(2, 3)
3830 void netdev_warn(const struct net_device *dev, const char *format, ...);
3831 __printf(2, 3)
3832 void netdev_notice(const struct net_device *dev, const char *format, ...);
3833 __printf(2, 3)
3834 void netdev_info(const struct net_device *dev, const char *format, ...);
3835
3836 #define MODULE_ALIAS_NETDEV(device) \
3837     MODULE_ALIAS("netdev-" device)
3838
3839 #if defined(CONFIG_DYNAMIC_DEBUG)
3840 #define netdev_dbg(__dev, format, args...) \
3841     do { \
3842         dynamic_netdev_dbg(__dev, format, ##args); \
3843     } while (0)
3844 #elif defined(DEBUG)
3845 #define netdev_dbg(__dev, format, args...) \
3846     netdev_printk(KERN_DEBUG, __dev, format, ##args)
3847 #else
3848 #define netdev_dbg(__dev, format, args...) \
3849     ({

```

```

3850         if (0)                                     \
3851             netdev_printk(KERN_DEBUG, __dev, format, ##args); \
3852     })
3853 #endif
3854
3855 #if defined(VERBOSE_DEBUG)
3856 #define netdev_vdbg      netdev_dbg
3857 #else
3858
3859 #define netdev_vdbg(dev, format, args...)          \
3860     ({                                             \
3861         if (0)                                     \
3862             netdev_printk(KERN_DEBUG, dev, format, ##args); \
3863         0;                                         \
3864     })
3865 #endif
3866
3867 /*
3868  * netdev_WARN() acts like dev_printk(), but with the key difference
3869  * of using a WARN/WARN_ON to get the message out, including the
3870  * file/line information and a backtrace.
3871  */
3872 #define netdev_WARN(dev, format, args...)          \
3873     WARN(1, "netdevice: %s%s\n" format, netdev_name(dev), \
3874         netdev_reg_state(dev), ##args)
3875
3876 /* netif printk helpers, similar to netdev_printk */
3877
3878 #define netif_printk(priv, type, level, dev, fmt, args...) \
3879     do {                                                    \
3880         if (netif_msg_##type(priv))                        \
3881             netdev_printk(level, (dev), fmt, ##args);      \
3882     } while (0)
3883
3884 #define netif_level(level, priv, type, dev, fmt, args...) \
3885     do {                                                    \
3886         if (netif_msg_##type(priv))                        \
3887             netdev_##level(dev, fmt, ##args);              \
3888     } while (0)
3889
3890 #define netif_emerg(priv, type, dev, fmt, args...) \
3891     netif_level(emerg, priv, type, dev, fmt, ##args)
3892 #define netif_alert(priv, type, dev, fmt, args...) \
3893     netif_level(alert, priv, type, dev, fmt, ##args)
3894 #define netif_crit(priv, type, dev, fmt, args...) \
3895     netif_level(crit, priv, type, dev, fmt, ##args)
3896 #define netif_err(priv, type, dev, fmt, args...) \
3897     netif_level(err, priv, type, dev, fmt, ##args)
3898 #define netif_warn(priv, type, dev, fmt, args...) \
3899     netif_level(warn, priv, type, dev, fmt, ##args)
3900 #define netif_notice(priv, type, dev, fmt, args...) \
3901     netif_level(notice, priv, type, dev, fmt, ##args)
3902 #define netif_info(priv, type, dev, fmt, args...) \
3903     netif_level(info, priv, type, dev, fmt, ##args)
3904
3905 #if defined(CONFIG_DYNAMIC_DEBUG)
3906 #define netif_dbg(priv, type, netdev, format, args...) \
3907     do {                                                    \
3908         if (netif_msg_##type(priv))                        \
3909             dynamic_netdev_dbg(netdev, format, ##args);    \
3910     } while (0)
3911 #elif defined(DEBUG)
3912 #define netif_dbg(priv, type, dev, format, args...) \
3913     netif_printk(priv, type, KERN_DEBUG, dev, format, ##args)
3914 #else
3915 #define netif_dbg(priv, type, dev, format, args...) \
3916     ({                                             \
3917         if (0)                                     \
3918             netif_printk(priv, type, KERN_DEBUG, dev, format, ##args); \
3919         0;                                         \
3920     })
3921 #endif
3922

```



```

3923 #if defined(VERBOSE_DEBUG)
3924 #define netif_vdbg      netif_dbg
3925 #else
3926 #define netif_vdbg(priv, type, dev, format, args...)      \
3927 ({                                                         \
3928     if (0)                                                 \
3929         netif_printk(priv, type, KERN_DEBUG, dev, format, ##args); \
3930     0;                                                      \
3931 })
3932 #endif
3933
3934 /*
3935  *      The list of packet types we will receive (as opposed to discard)
3936  *      and the routines to invoke.
3937  *
3938  *      Why 16. Because with 16 the only overlap we get on a hash of the
3939  *      low nibble of the protocol value is RARP/SNAP/X.25.
3940  *
3941  *      NOTE: That is no longer true with the addition of VLAN tags. Not
3942  *      sure which should go first, but I bet it won't make much
3943  *      difference if we are running VLANs. The good news is that
3944  *      this protocol won't be in the list unless compiled in, so
3945  *      the average user (w/out VLANs) will not be adversely affected.
3946  *      --BLG
3947  *
3948  *      0800      IP
3949  *      8100      802.1Q VLAN
3950  *      0001      802.3
3951  *      0002      AX.25
3952  *      0004      802.2
3953  *      8035      RARP
3954  *      0005      SNAP
3955  *      0805      X.25
3956  *      0806      ARP
3957  *      8137      IPX
3958  *      0009      Localtalk
3959  *      86DD      IPv6
3960 */
3961 #define PTYPE_HASH_SIZE (16)
3962 #define PTYPE_HASH_MASK (PTYPE_HASH_SIZE - 1)
3963
3964 #endif /* _LINUX_NETDEVICE_H */
3965

```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds

• [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)