

Linux Cross Reference

[Free Electrons](#)

Embedded Linux Experts

• [source navigation](#) • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version: [2.0.40](#) [2.2.26](#) [2.4.37](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#) [3.18](#) [3.19](#) [4.0](#) [4.1](#) [4.2](#)

[Linux](#)/[include](#)/[net](#)/[dst.h](#)

```

1  /*
2  * net/dst.h    Protocol independent destination cache definitions.
3  *
4  * Authors:     Alexey Kuznetsov, <kuznet@ms2.inr.ac.ru>
5  *
6  */
7
8  #ifndef _NET_DST_H
9  #define _NET_DST_H
10
11 #include <net/dst_ops.h>
12 #include <linux/netdevice.h>
13 #include <linux/rtnetlink.h>
14 #include <linux/rcupdate.h>
15 #include <linux/bug.h>
16 #include <linux/jiffies.h>
17 #include <net/neighbour.h>
18 #include <asm/processor.h>
19
20 #define DST_GC_MIN      (HZ/10)
21 #define DST_GC_INC      (HZ/2)
22 #define DST_GC_MAX      (120*HZ)
23
24 /* Each dst_entry has reference count and sits in some parent list(s).
25  * When it is removed from parent list, it is "freed" (dst_free).
26  * After this it enters dead state (dst->obsolete > 0) and if its refcnt
27  * is zero, it can be destroyed immediately, otherwise it is added
28  * to gc list and garbage collector periodically checks the refcnt.
29  */
30
31 struct sk_buff;
32
33 struct dst_entry {
34     struct rcu_head      rcu_head;
35     struct dst_entry     *child;
36     struct net_device     *dev;
37     struct dst_ops       *ops;
38     unsigned long        _metrics;
39     unsigned long        expires;
40     struct dst_entry     *path;
41     struct dst_entry     *from;
42 #ifdef CONFIG_XFRM
43     struct xfrm_state     *xfrm;
44 #else
45     void                 *__pad1;
46 #endif
47     int                  (*input)(struct sk_buff *);
48     int                  (*output)(struct sock *sk, struct sk_buff *skb);
49

```

```

50         unsigned short          flags;
51 #define DST_HOST                  0x0001
52 #define DST_NOXFRM                0x0002
53 #define DST_NOPOLICY              0x0004
54 #define DST_NOHASH                0x0008
55 #define DST_NOCACHE               0x0010
56 #define DST_NOCOUNT               0x0020
57 #define DST_FAKE_RTABLE           0x0040
58 #define DST_XFRM_TUNNEL           0x0080
59 #define DST_XFRM_QUEUE            0x0100
60
61         unsigned short          pending_confirm;
62
63         short                    error;
64
65         /* A non-zero value of dst->obsolete forces by-hand validation
66          * of the route entry. Positive values are set by the generic
67          * dst layer to indicate that the entry has been forcefully
68          * destroyed.
69          *
70          * Negative values are used by the implementation layer code to
71          * force invocation of the dst_ops->check() method.
72          */
73         short                    obsolete;
74 #define DST_OBSOLETE_NONE          0
75 #define DST_OBSOLETE_DEAD          2
76 #define DST_OBSOLETE_FORCE_CHK    -1
77 #define DST_OBSOLETE_KILL         -2
78         unsigned short          header_len;    /* more space at head required */
79         unsigned short          trailer_len;    /* space to reserve at tail */
80 #ifdef CONFIG_IP_ROUTE_CLASSID
81         __u32                   tclassid;
82 #else
83         __u32                   __pad2;
84 #endif
85
86         /*
87          * Align __refcnt to a 64 bytes alignment
88          * (L1_CACHE_SIZE would be too much)
89          */
90 #ifdef CONFIG_64BIT
91         long                    __pad_to_align_refcnt[2];
92 #endif
93         /*
94          * __refcnt wants to be on a different cache line from
95          * input/output/ops or performance tanks badly
96          */
97         atomic_t                __refcnt;      /* client references */
98         int                     __use;
99         unsigned long           lastuse;
100         union {
101             struct dst_entry     *next;
102             struct rtable __rcu   *rt_next;
103             struct rt6_info      *rt6_next;
104             struct dn_route __rcu *dn_next;
105         };
106 };
107
108 __u32 *dst_cow_metrics_generic(struct dst_entry *dst, unsigned long old);
109 extern const __u32 dst_default_metrics[];
110
111 #define DST_METRICS_READ_ONLY      0x1UL
112 #define DST_METRICS_FLAGS          0x3UL
113 #define __DST_METRICS_PTR(Y)      \
114     ((__u32 *)((Y) & ~DST_METRICS_FLAGS))
115 #define DST_METRICS_PTR(X)        __DST_METRICS_PTR((X)->_metrics)
116
117 static inline bool dst_metrics_read_only(const struct dst_entry *dst)
118 {
119     return dst->_metrics & DST_METRICS_READ_ONLY;

```

```

120 }
121
122 void \_\_dst\_destroy\_metrics\_generic(struct dst\_entry *dst, unsigned long old);
123
124 static inline void dst\_destroy\_metrics\_generic(struct dst\_entry *dst)
125 {
126     unsigned long val = dst->_metrics;
127     if (!(val & DST\_METRICS\_READ\_ONLY))
128         \_\_dst\_destroy\_metrics\_generic(dst, val);
129 }
130
131 static inline u32 *dst\_metrics\_write\_ptr(struct dst\_entry *dst)
132 {
133     unsigned long p = dst->_metrics;
134
135     BUG\_ON(!p);
136
137     if (p & DST\_METRICS\_READ\_ONLY)
138         return dst->ops->cow\_metrics(dst, p);
139     return DST\_METRICS\_PTR(p);
140 }
141
142 /* This may only be invoked before the entry has reached global
143 * visibility.
144 */
145 static inline void dst\_init\_metrics(struct dst\_entry *dst,
146                                     const u32 *src\_metrics,
147                                     bool read\_only)
148 {
149     dst->_metrics = ((unsigned long) src\_metrics) |
150                    (read\_only ? DST\_METRICS\_READ\_ONLY : 0);
151 }
152
153 static inline void dst\_copy\_metrics(struct dst\_entry *dest, const struct dst\_entry *src)
154 {
155     u32 *dst\_metrics = dst\_metrics\_write\_ptr(dest);
156
157     if (dst\_metrics) {
158         u32 *src\_metrics = DST\_METRICS\_PTR(src);
159
160         memcpy(dst\_metrics, src\_metrics, RTAX\_MAX * sizeof(u32));
161     }
162 }
163
164 static inline u32 *dst\_metrics\_ptr(struct dst\_entry *dst)
165 {
166     return DST\_METRICS\_PTR(dst);
167 }
168
169 static inline u32
170 dst\_metric\_raw(const struct dst\_entry *dst, const int metric)
171 {
172     u32 *p = DST\_METRICS\_PTR(dst);
173
174     return p[metric-1];
175 }
176
177 static inline u32
178 dst\_metric(const struct dst\_entry *dst, const int metric)
179 {
180     WARN\_ON\_ONCE(metric == RTAX\_HOPLIMIT ||
181                 metric == RTAX\_ADV MSS ||
182                 metric == RTAX\_MTU);
183     return dst\_metric\_raw(dst, metric);
184 }
185
186 static inline u32
187 dst\_metric\_adv mss(const struct dst\_entry *dst)
188 {
189     u32 adv mss = dst\_metric\_raw(dst, RTAX\_ADV MSS);

```

```

190
191     if (!advms)
192         advms = dst->ops->default\_advms\(dst\);
193
194     return advms;
195 }
196
197 static inline void dst\_metric\_set(struct dst\_entry *dst, int metric, u32 val)
198 {
199     u32 *p = dst\_metrics\_write\_ptr\(dst\);
200
201     if (p)
202         p[metric-1] = val;
203 }
204
205 static inline u32
206 dst\_feature(const struct dst\_entry *dst, u32 feature)
207 {
208     return dst\_metric\(dst, RTAX\_FEATURES\) & feature;
209 }
210
211 static inline u32 dst\_mtu(const struct dst\_entry *dst)
212 {
213     return dst->ops->mtu\(dst\);
214 }
215
216 /* RTT metrics are stored in milliseconds for user ABI, but used as jiffies */
217 static inline unsigned long dst\_metric\_rtt(const struct dst\_entry *dst, int metric)
218 {
219     return msecs\_to\_jiffies\(dst\_metric\(dst, metric\)\);
220 }
221
222 static inline u32
223 dst\_allfrag(const struct dst\_entry *dst)
224 {
225     int ret = dst\_feature\(dst, RTAX\_FEATURE\_ALLFRAG\);
226     return ret;
227 }
228
229 static inline int
230 dst\_metric\_locked(const struct dst\_entry *dst, int metric)
231 {
232     return dst\_metric\(dst, RTAX\_LOCK\) & (1<<metric);
233 }
234
235 static inline void dst\_hold(struct dst\_entry *dst)
236 {
237     /*
238     * If your kernel compilation stops here, please check
239     * __pad_to_align_refcnt declaration in struct dst_entry
240     */
241     BUILD\_BUG\_ON(offsetof(struct dst\_entry, __refcnt) & 63);
242     atomic\_inc\(&dst->\_\_refcnt\);
243 }
244
245 static inline void dst\_use(struct dst\_entry *dst, unsigned long time)
246 {
247     dst\_hold\(dst\);
248     dst->\_\_use++;
249     dst->lastuse = time;
250 }
251
252 static inline void dst\_use\_noref(struct dst\_entry *dst, unsigned long time)
253 {
254     dst->\_\_use++;
255     dst->lastuse = time;
256 }
257
258 static inline struct dst\_entry *dst\_clone(struct dst\_entry *dst)
259 {

```

```

260     if (dst)
261         atomic\_inc(&dst->__refcnt);
262     return dst;
263 }
264
265 void dst\_release(struct dst\_entry *dst);
266
267 static inline void refdst\_drop(unsigned long refdst)
268 {
269     if (!(refdst & SKB\_DST\_NOREF))
270         dst\_release((struct dst\_entry *) (refdst & SKB\_DST\_PTRMASK));
271 }
272
273 /**
274  * skb\_dst\_drop - drops skb dst
275  * @skb: buffer
276  *
277  * Drops dst reference count if a reference was taken.
278  */
279 static inline void skb\_dst\_drop(struct sk\_buff *skb)
280 {
281     if (skb->_skb_refdst) {
282         refdst\_drop(skb->_skb_refdst);
283         skb->_skb_refdst = 0UL;
284     }
285 }
286
287 static inline void skb\_dst\_copy(struct sk\_buff *nsk, const struct sk\_buff *osk)
288 {
289     nsk->_skb_refdst = osk->_skb_refdst;
290     if (!(nsk->_skb_refdst & SKB\_DST\_NOREF))
291         dst\_clone(skb\_dst(nsk));
292 }
293
294 /**
295  * skb\_dst\_force - makes sure skb dst is refcounted
296  * @skb: buffer
297  *
298  * If dst is not yet refcounted, let's do it
299  */
300 static inline void skb\_dst\_force(struct sk\_buff *skb)
301 {
302     if (skb\_dst\_is\_noref(skb)) {
303         WARN\_ON(!rcu_read_lock_held());
304         skb->_skb_refdst &= ~SKB\_DST\_NOREF;
305         dst\_clone(skb\_dst(skb));
306     }
307 }
308
309
310 /**
311  * \_\_skb\_tunnel\_rx - prepare skb for rx reinsert
312  * @skb: buffer
313  * @dev: tunnel device
314  * @net: netns for packet i/o
315  *
316  * After decapsulation, packet is going to re-enter (netif_rx()) our stack,
317  * so make some cleanups. (no accounting done)
318  */
319 static inline void \_\_skb\_tunnel\_rx(struct sk\_buff *skb, struct net\_device *dev,
320                                     struct net *net)
321 {
322     skb->dev = dev;
323
324     /*
325      * Clear hash so that we can recalculate the hash for the
326      * encapsulated packet, unless we have already determine the hash
327      * over the L4 4-tuple.
328      */
329     skb\_clear\_hash\_if\_not\_l4(skb);

```

```

330     skb_set_queue_mapping(skb, 0);
331     skb_scrub_packet(skb, !net_eq(net, dev_net(dev)));
332 }
333
334 /**
335  *     skb_tunnel_rx - prepare skb for rx reinsert
336  *     @skb: buffer
337  *     @dev: tunnel device
338  *
339  *     After decapsulation, packet is going to re-enter (netif_rx()) our stack,
340  *     so make some cleanups, and perform accounting.
341  *     Note: this accounting is not SMP safe.
342  */
343 static inline void skb_tunnel_rx(struct sk_buff *skb, struct net_device *dev,
344                                 struct net *net)
345 {
346     /* TODO : stats should be SMP safe */
347     dev->stats.rx_packets++;
348     dev->stats.rx_bytes += skb->len;
349     __skb_tunnel_rx(skb, dev, net);
350 }
351
352 int dst_discard_sk(struct sock *sk, struct sk_buff *skb);
353 static inline int dst_discard(struct sk_buff *skb)
354 {
355     return dst_discard_sk(skb->sk, skb);
356 }
357 void *dst_alloc(struct dst_ops *ops, struct net_device *dev, int initial_ref,
358                int initial_obsolete, unsigned short flags);
359 void __dst_free(struct dst_entry *dst);
360 struct dst_entry *dst_destroy(struct dst_entry *dst);
361
362 static inline void dst_free(struct dst_entry *dst)
363 {
364     if (dst->obsolete > 0)
365         return;
366     if (!atomic_read(&dst->__refcnt)) {
367         dst = dst_destroy(dst);
368         if (!dst)
369             return;
370     }
371     __dst_free(dst);
372 }
373
374 static inline void dst_rcu_free(struct rcu_head *head)
375 {
376     struct dst_entry *dst = container_of(head, struct dst_entry, rcu_head);
377     dst_free(dst);
378 }
379
380 static inline void dst_confirm(struct dst_entry *dst)
381 {
382     dst->pending_confirm = 1;
383 }
384
385 static inline int dst_neigh_output(struct dst_entry *dst, struct neighbour *n,
386                                   struct sk_buff *skb)
387 {
388     const struct hh_cache *hh;
389
390     if (dst->pending_confirm) {
391         unsigned long now = jiffies;
392
393         dst->pending_confirm = 0;
394         /* avoid dirtying neighbour */
395         if (n->confirmed != now)
396             n->confirmed = now;
397     }
398
399     hh = &n->hh;

```

```

400     if ((n->nud_state & NUD_CONNECTED) && hh->hh_len)
401         return neigh_hh_output(hh, skb);
402     else
403         return n->output(n, skb);
404 }
405
406 static inline struct neighbour *dst_neigh_lookup(const struct dst_entry *dst, const void *daddr)
407 {
408     struct neighbour *n = dst->ops->neigh_lookup(dst, NULL, daddr);
409     return IS_ERR(n) ? NULL : n;
410 }
411
412 static inline struct neighbour *dst_neigh_lookup_skb(const struct dst_entry *dst,
413                                                     struct sk_buff *skb)
414 {
415     struct neighbour *n = dst->ops->neigh_lookup(dst, skb, NULL);
416     return IS_ERR(n) ? NULL : n;
417 }
418
419 static inline void dst_link_failure(struct sk_buff *skb)
420 {
421     struct dst_entry *dst = skb_dst(skb);
422     if (dst && dst->ops && dst->ops->link_failure)
423         dst->ops->link_failure(skb);
424 }
425
426 static inline void dst_set_expires(struct dst_entry *dst, int timeout)
427 {
428     unsigned long expires = jiffies + timeout;
429
430     if (expires == 0)
431         expires = 1;
432
433     if (dst->expires == 0 || time_before(expires, dst->expires))
434         dst->expires = expires;
435 }
436
437 /* Output packet to network from transport. */
438 static inline int dst_output_sk(struct sock *sk, struct sk_buff *skb)
439 {
440     return skb_dst(skb)->output(sk, skb);
441 }
442 static inline int dst_output(struct sk_buff *skb)
443 {
444     return dst_output_sk(skb->sk, skb);
445 }
446
447 /* Input packet from network to transport. */
448 static inline int dst_input(struct sk_buff *skb)
449 {
450     return skb_dst(skb)->input(skb);
451 }
452
453 static inline struct dst_entry *dst_check(struct dst_entry *dst, u32 cookie)
454 {
455     if (dst->obsolete)
456         dst = dst->ops->check(dst, cookie);
457     return dst;
458 }
459
460 void dst_init(void);
461
462 /* Flags for xfrm_lookup_flags argument. */
463 enum {
464     XFRM_LOOKUP_ICMP = 1 << 0,
465     XFRM_LOOKUP_QUEUE = 1 << 1,
466     XFRM_LOOKUP_KEEP_DST_REF = 1 << 2,
467 };
468
469 struct flowi;

```

```

470 #ifndef CONFIG_XFRM
471 static inline struct dst\_entry *xfrm\_lookup(struct net *net,
472                                           struct dst\_entry *dst\_orig,
473                                           const struct flowi *fl, struct sock *sk,
474                                           int flags)
475 {
476     return dst\_orig;
477 }
478
479 static inline struct dst\_entry *xfrm\_lookup\_route(struct net *net,
480                                                  struct dst\_entry *dst\_orig,
481                                                  const struct flowi *fl,
482                                                  struct sock *sk,
483                                                  int flags)
484 {
485     return dst\_orig;
486 }
487
488 static inline struct xfrm\_state *dst\_xfrm(const struct dst\_entry *dst)
489 {
490     return NULL;
491 }
492
493 #else
494 struct dst\_entry *xfrm\_lookup(struct net *net, struct dst\_entry *dst\_orig,
495                               const struct flowi *fl, struct sock *sk,
496                               int flags);
497
498 struct dst\_entry *xfrm\_lookup\_route(struct net *net, struct dst\_entry *dst\_orig,
499                                     const struct flowi *fl, struct sock *sk,
500                                     int flags);
501
502 /* skb attached with this dst needs transformation if dst->xfrm is valid */
503 static inline struct xfrm\_state *dst\_xfrm(const struct dst\_entry *dst)
504 {
505     return dst->xfrm;
506 }
507 #endif
508
509 #endif /* _NET_DST_H */
510

```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)