

Linux Cross Reference

[Free Electrons](#)

Embedded Linux Experts

• [source navigation](#) • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version:

[2.0.40](#) [2.2.26](#) [2.4.37](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#) [3.18](#) [3.19](#) [4.0](#) [4.1](#) [4.2](#)

[Linux/include/net/inet_sock.h](#)

```

1  /*
2   * INET           An implementation of the TCP/IP protocol suite for the LINUX
3   *                 operating system.  INET is implemented using the BSD Socket
4   *                 interface as the means of communication with the user level.
5   *
6   *                 Definitions for inet_sock
7   *
8   * Authors:      Many, reorganised here by
9   *                 Arnaldo Carvalho de Melo <acme@mandriva.com>
10  *
11  *                 This program is free software; you can redistribute it and/or
12  *                 modify it under the terms of the GNU General Public License
13  *                 as published by the Free Software Foundation; either version
14  *                 2 of the License, or (at your option) any later version.
15  */
16 #ifndef INET\_SOCK\_H
17 #define INET\_SOCK\_H
18
19 #include <linux/bitops.h>
20 #include <linux/kmemcheck.h>
21 #include <linux/string.h>
22 #include <linux/types.h>
23 #include <linux/jhash.h>
24 #include <linux/netdevice.h>
25
26 #include <net/flow.h>
27 #include <net/sock.h>
28 #include <net/request_sock.h>
29 #include <net/netns/hash.h>
30 #include <net/tcp_states.h>
31
32 /** struct ip_options - IP Options
33  *
34  * @faddr - Saved first hop address
35  * @nexthop - Saved nexthop address in LSRR and SSRR
36  * @is_strictroute - Strict source route
37  * @srr_is_hit - Packet destination addr was our one
38  * @is_changed - IP checksum more not valid
39  * @rr_needaddr - Need to record addr of outgoing dev
40  * @ts_needtime - Need to record timestamp
41  * @ts_needaddr - Need to record addr of outgoing dev
42  */
43 struct ip\_options {

```

```

44     be32          faddr;
45     be32          nexthop;
46     unsigned char optlen;
47     unsigned char srr;
48     unsigned char rr;
49     unsigned char ts;
50     unsigned char is_strictroute:1,
51                  srr_is_hit:1,
52                  is_changed:1,
53                  rr_needaddr:1,
54                  ts_needtime:1,
55                  ts_needaddr:1;
56     unsigned char router_alert;
57     unsigned char cipso;
58     unsigned char __pad2;
59     unsigned char __data[0];
60 };
61
62 struct ip\_options\_rcu {
63     struct rcu\_head rcu;
64     struct ip\_options opt;
65 };
66
67 struct ip\_options\_data {
68     struct ip\_options\_rcu  opt;
69     char data[40];
70 };
71
72 struct inet\_request\_sock {
73     struct request\_sock    req;
74     #define ir\_loc\_addr      req.__req_common.skc_rcv_saddr
75     #define ir\_rmt\_addr      req.__req_common.skc_daddr
76     #define ir\_num           req.__req_common.skc_num
77     #define ir\_rmt\_port      req.__req_common.skc_dport
78     #define ir\_v6\_rmt\_addr    req.__req_common.skc_v6_daddr
79     #define ir\_v6\_loc\_addr    req.__req_common.skc_v6_rcv_saddr
80     #define ir\_iif           req.__req_common.skc_bound_dev_if
81     #define ir\_cookie        req.__req_common.skc_cookie
82     #define ireq\_net         req.__req_common.skc_net
83     #define ireq\_state       req.__req_common.skc_state
84     #define ireq\_family     req.__req_common.skc_family
85
86     kmemcheck bitfield begin(flags);
87     u16                snd_wscale : 4,
88                        rcv_wscale : 4,
89                        tstamp_ok  : 1,
90                        sack_ok    : 1,
91                        wscale_ok  : 1,
92                        ecn_ok     : 1,
93                        acked      : 1,
94                        no_srccheck: 1;
95     kmemcheck bitfield end(flags);
96     u32                ir_mark;
97     union {
98         struct ip\_options\_rcu    *opt;
99         struct sk\_buff          *pktopts;
100    };
101 };
102
103 static inline struct inet\_request\_sock *inet\_rsk(const struct request\_sock *sk)
104 {
105     return (struct inet\_request\_sock *)sk;
106 }
107
108 static inline u32 inet\_request\_mark(const struct sock *sk, struct sk\_buff *skb)

```

```

109 {
110     if (!sk->sk_mark && sock\_net(sk)->ipv4.sysctl_tcp_fwmark_accept)
111         return skb->mark;
112
113     return sk->sk_mark;
114 }
115
116 struct inet\_cork {
117     unsigned int      flags;
118     \_\_be32            addr;
119     struct ip\_options *opt;
120     unsigned int      fragsize;
121     int               length; /* Total Length of all frames */
122     struct dst\_entry  *dst;
123     u8               tx_flags;
124     \_\_u8             ttl;
125     \_\_s16            tos;
126     char             priority;
127 };
128
129 struct inet\_cork\_full {
130     struct inet\_cork    base;
131     struct flowi        fl;
132 };
133
134 struct ip\_mc\_socklist;
135 struct ipv6\_pinfo;
136 struct rtable;
137
138 /** struct inet_sock - representation of INET sockets
139  *
140  * @sk - ancestor class
141  * @pinet6 - pointer to IPv6 control block
142  * @inet_daddr - Foreign IPv4 addr
143  * @inet_rcv_saddr - Bound Local IPv4 addr
144  * @inet_dport - Destination port
145  * @inet_num - Local port
146  * @inet_saddr - Sending source
147  * @uc_ttl - Unicast TTL
148  * @inet_sport - Source port
149  * @inet_id - ID counter for DF pkts
150  * @tos - TOS
151  * @mc_ttl - Multicasting TTL
152  * @is_icsk - is this an inet_connection_sock?
153  * @uc_index - Unicast outgoing device index
154  * @mc_index - Multicast device index
155  * @mc_list - Group array
156  * @cork - info to build ip hdr on each ip frag while socket is corked
157  */
158 struct inet\_sock {
159     /* sk and pinet6 has to be the first two members of inet_sock */
160     struct sock      sk;
161     #if IS\_ENABLED(CONFIG_IPV6)
162     struct ipv6\_pinfo *pinet6;
163     #endif
164     /* Socket demultiplex comparisons on incoming packets. */
165     #define inet\_daddr      sk->__sk_common.skc_daddr
166     #define inet\_rcv\_saddr sk->__sk_common.skc_rcv_saddr
167     #define inet\_dport     sk->__sk_common.skc_dport
168     #define inet\_num       sk->__sk_common.skc_num
169
170     \_\_be32            inet_saddr;
171     \_\_s16            uc_ttl;
172     \_\_u16           cmsg_flags;
173     \_\_be16           inet_sport;

```

```

174         __u16                inet_id;
175
176     struct ip\_options\_rcu __rcu    *inet_opt;
177     int                rx_dst_ifindex;
178     __u8                tos;
179     __u8                min_ttl;
180     __u8                mc_ttl;
181     __u8                pmtudisc;
182     __u8                recverr:1,
183                        is_icsk:1,
184                        freebind:1,
185                        hdrincl:1,
186                        mc_loop:1,
187                        transparent:1,
188                        mc_all:1,
189                        nodefrag:1;
190     __u8                bind_address_no_port:1;
191     __u8                rcv_tos;
192     __u8                convert_csum;
193     int                uc_index;
194     int                mc_index;
195     __be32             mc_addr;
196     struct ip\_mc\_socklist __rcu    *mc_list;
197     struct inet\_cork\_full    cork;
198 };
199
200 #define IPCORK\_OPT        1        /* ip-options has been held in ipcork.opt */
201 #define IPCORK\_ALLFRAG    2        /* always fragment (for ipv6 for now) */
202
203 /* msg flags for inet */
204 #define IP\_CMSG\_PKTINFO    BIT(0)
205 #define IP\_CMSG\_TTL        BIT(1)
206 #define IP\_CMSG\_TOS        BIT(2)
207 #define IP\_CMSG\_RECVOPTS    BIT(3)
208 #define IP\_CMSG\_RETOPTS    BIT(4)
209 #define IP\_CMSG\_PASSESEC    BIT(5)
210 #define IP\_CMSG\_ORIGDSTADDR BIT(6)
211 #define IP\_CMSG\_CHECKSUM    BIT(7)
212
213 static inline struct inet\_sock *inet_sk(const struct sock *sk)
214 {
215     return (struct inet\_sock *)sk;
216 }
217
218 static inline void __inet_sk_copy_descendant(struct sock *sk_to,
219                                             const struct sock *sk_from,
220                                             const int ancestor_size)
221 {
222     memcpy(inet_sk(sk_to) + 1, inet_sk(sk_from) + 1,
223           sk_from->sk_prot->obj_size - ancestor_size);
224 }
225 #if !defined(CONFIG_IPV6)
226 static inline void inet_sk_copy_descendant(struct sock *sk_to,
227                                           const struct sock *sk_from)
228 {
229     __inet_sk_copy_descendant(sk_to, sk_from, sizeof(struct inet\_sock));
230 }
231 #endif
232
233 int inet_sk_rebuild_header(struct sock *sk);
234
235 static inline unsigned int __inet_ehashfn(const __be32 laddr,
236                                           const __u16 lport,
237                                           const __be32 faddr,
238                                           const __be16 fport,

```

```

239                                     u32 initval)
240 {
241     return jhash_3words((__force __u32) laddr,
242                        (__force __u32) faddr,
243                        ((__u32) lport) << 16 | (__force __u32) fport,
244                        initval);
245 }
246
247 struct request_sock *inet_reqsk_alloc(const struct request_sock_ops *ops,
248                                       struct sock *sk_listener);
249
250 static inline __u8 inet_sk_flowi_flags(const struct sock *sk)
251 {
252     __u8 flags = 0;
253
254     if (inet_sk(sk)->transparent || inet_sk(sk)->hdrincl)
255         flags |= FLOWI_FLAG_ANYSRC;
256     return flags;
257 }
258
259 static inline void inet_inc_convert_csum(struct sock *sk)
260 {
261     inet_sk(sk)->convert_csum++;
262 }
263
264 static inline void inet_dec_convert_csum(struct sock *sk)
265 {
266     if (inet_sk(sk)->convert_csum > 0)
267         inet_sk(sk)->convert_csum--;
268 }
269
270 static inline bool inet_get_convert_csum(struct sock *sk)
271 {
272     return !!inet_sk(sk)->convert_csum;
273 }
274
275 #endif /* _INET_SOCK_H */
276

```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)