

Linux Cross Reference

[Free Electrons](#)

Embedded Linux Experts

• [source navigation](#) • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version: [2.0.40](#) [2.2.26](#) [2.4.37](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#) [3.18](#) [3.19](#) [4.0](#) [4.1](#) [4.2](#)

[Linux](#)/[include](#)/[net](#)/[ip.h](#)

```

1  /*
2  * INET          An implementation of the TCP/IP protocol suite for the LINUX
3  *              operating system. INET is implemented using the BSD Socket
4  *              interface as the means of communication with the user level.
5  *
6  *              Definitions for the IP module.
7  *
8  * Version:      @(#)ip.h          1.0.2    05/07/93
9  *
10 * Authors:      Ross Biro
11 *              Fred N. van Kempen, <waltje@uWalt.NL.Mugnet.ORG>
12 *              Alan Cox, <gw4pts@gw4pts.ampr.org>
13 *
14 * Changes:
15 *              Mike McLagan      :      Routing by source
16 *
17 *              This program is free software; you can redistribute it and/or
18 *              modify it under the terms of the GNU General Public License
19 *              as published by the Free Software Foundation; either version
20 *              2 of the License, or (at your option) any later version.
21 */
22 #ifndef IP_H
23 #define IP_H
24
25 #include <linux/types.h>
26 #include <linux/ip.h>
27 #include <linux/in.h>
28 #include <linux/skbuff.h>
29
30 #include <net/inet_sock.h>
31 #include <net/route.h>
32 #include <net/snmp.h>
33 #include <net/flow.h>
34 #include <net/flow_dissector.h>
35
36 struct sock;
37
38 struct inet_skb_parm {
39     struct ip_options opt;          /* Compiled IP options */
40     unsigned char flags;
41
42 #define IPSKB_FORWARDED BIT(0)
43 #define IPSKB_XFRM_TUNNEL_SIZE BIT(1)
44 #define IPSKB_XFRM_TRANSFORMED BIT(2)
45 #define IPSKB_FRAG_COMPLETE BIT(3)
46 #define IPSKB_REROUTED BIT(4)
47 #define IPSKB_DOREDIRECT BIT(5)
48 #define IPSKB_FRAG_PMTU BIT(6)
49
50     u16 frag_max_size;
51 };
52
53 static inline unsigned int ip_hdrlen(const struct sk_buff *skb)
54 {
55     return ip_hdr(skb)->ihl * 4;

```

```

56 }
57
58 struct ipcm_cookie {
59     __be32          addr;
60     int             oif;
61     struct ip_options_rcu *opt;
62     __u8           tx_flags;
63     __u8           ttl;
64     __s16          tos;
65     char           priority;
66 };
67
68 #define IPCB(skb) ((struct inet_skb_parm*)((skb)->cb))
69 #define PKTINFO_SKB_CB(skb) ((struct in_pktinfo *)((skb)->cb))
70
71 struct ip_ra_chain {
72     struct ip_ra_chain __rcu *next;
73     struct sock        *sk;
74     union {
75         void (*destructor)(struct sock *);
76         struct sock *saved_sk;
77     };
78     struct rcu_head      rcu;
79 };
80
81 extern struct ip_ra_chain __rcu *ip_ra_chain;
82
83 /* IP flags. */
84 #define IP_CE          0x8000 /* Flag: "Congestion" */
85 #define IP_DF          0x4000 /* Flag: "Don't Fragment" */
86 #define IP_MF          0x2000 /* Flag: "More Fragments" */
87 #define IP_OFFSET     0x1FFF /* "Fragment Offset" part */
88
89 #define IP_FRAG_TIME   (30 * HZ) /* fragment lifetime */
90
91 struct msghdr;
92 struct net_device;
93 struct packet_type;
94 struct rtable;
95 struct sockaddr;
96
97 int igmp_mc_init(void);
98
99 /*
100  *      Functions provided by ip.c
101  */
102
103 int ip_build_and_send_pkt(struct sk_buff *skb, struct sock *sk,
104     __be32 saddr, __be32 daddr,
105     struct ip_options_rcu *opt);
106 int ip_rcv(struct sk_buff *skb, struct net_device *dev, struct packet_type *pt,
107     struct net_device *orig_dev);
108 int ip_local_deliver(struct sk_buff *skb);
109 int ip_mr_input(struct sk_buff *skb);
110 int ip_output(struct sock *sk, struct sk_buff *skb);
111 int ip_mc_output(struct sock *sk, struct sk_buff *skb);
112 int ip_do_fragment(struct sock *sk, struct sk_buff *skb,
113     int (*output)(struct sock *, struct sk_buff *));
114 void ip_send_check(struct iphdr *ip);
115 int __ip_local_out(struct sk_buff *skb);
116 int ip_local_out_sk(struct sock *sk, struct sk_buff *skb);
117 static inline int ip_local_out(struct sk_buff *skb)
118 {
119     return ip_local_out_sk(skb->sk, skb);
120 }
121
122 int ip_queue_xmit(struct sock *sk, struct sk_buff *skb, struct flowi *fl);
123 void ip_init(void);
124 int ip_append_data(struct sock *sk, struct flowi4 *fl4,
125     int getfrag(void *from, char *to, int offset, int len,
126     int odd, struct sk_buff *skb),
127     void *from, int len, int protolen,
128     struct ipcm_cookie *ipc,
129     struct rtable **rt,
130     unsigned int flags);

```

```

131 int ip_generic_getfrag(void *from, char *to, int offset, int len, int odd,
132                          struct sk_buff *skb);
133 ssize_t ip_append_page(struct sock *sk, struct flowi4 *fl4, struct page *page,
134                        int offset, size_t size, int flags);
135 struct sk_buff * ip_make_skb(struct sock *sk, struct flowi4 *fl4,
136                             struct sk_buff_head *queue,
137                             struct inet_cork *cork);
138 int ip_send_skb(struct net *net, struct sk_buff *skb);
139 int ip_push_pending_frames(struct sock *sk, struct flowi4 *fl4);
140 void ip_flush_pending_frames(struct sock *sk);
141 struct sk_buff *ip_make_skb(struct sock *sk, struct flowi4 *fl4,
142                             int getfrag(void *from, char *to, int offset,
143                                             int len, int odd, struct sk_buff *skb),
144                             void *from, int length, int transhdrlen,
145                             struct ipcm_cookie *ipc, struct rtable **rtp,
146                             unsigned int flags);
147
148 static inline struct sk_buff *ip_finish_skb(struct sock *sk, struct flowi4 *fl4)
149 {
150     return ip_make_skb(sk, fl4, &sk->sk_write_queue, &inet_sk(sk)->cork.base);
151 }
152
153 static inline __u8 get_rttoS(struct ipcm_cookie* ipc, struct inet_sock *inet)
154 {
155     return (ipc->tos != -1) ? RT_TOS(ipc->tos) : RT_TOS(inet->tos);
156 }
157
158 static inline __u8 get_rtconn_flags(struct ipcm_cookie* ipc, struct sock* sk)
159 {
160     return (ipc->tos != -1) ? RT_CONN_FLAGS_TOS(sk, ipc->tos) : RT_CONN_FLAGS(sk);
161 }
162
163 /* datagram.c */
164 int ip4_datagram_connect(struct sock *sk, struct sockaddr *uaddr, int addr_len);
165 int ip4_datagram_connect(struct sock *sk, struct sockaddr *uaddr, int addr_len);
166
167 void ip4_datagram_release_cb(struct sock *sk);
168
169 struct ip_reply_arg {
170     struct kvec iov[1];
171     int flags;
172     __wsum csum;
173     int csumoffset; /* u16 offset of csum in iov[0].iov_base */
174                     /* -1 if not needed */
175     int bound_dev_if;
176     __u8 tos;
177 };
178
179 #define IP_REPLY_ARG_NOSRCHECK 1
180
181 static inline __u8 ip_reply_arg_flowi_flags(const struct ip_reply_arg *arg)
182 {
183     return (arg->flags & IP_REPLY_ARG_NOSRCHECK) ? FLOWI_FLAG_ANYSRC : 0;
184 }
185
186 void ip_send_unicast_reply(struct sock *sk, struct sk_buff *skb,
187                           const struct ip_options *sopt,
188                           __be32 daddr, __be32 saddr,
189                           const struct ip_reply_arg *arg,
190                           unsigned int len);
191
192 #define IP_INC_STATS(net, field) SNMP_INC_STATS64((net)->mib.ip_statistics, field)
193 #define IP_INC_STATS_BH(net, field) SNMP_INC_STATS64_BH((net)->mib.ip_statistics, field)
194 #define IP_ADD_STATS(net, field, val) SNMP_ADD_STATS64((net)->mib.ip_statistics, field, val)
195 #define IP_ADD_STATS_BH(net, field, val) SNMP_ADD_STATS64_BH((net)->mib.ip_statistics, field, val)
196 #define IP_UPD_PO_STATS(net, field, val) SNMP_UPD_PO_STATS64((net)->mib.ip_statistics, field, val)
197 #define IP_UPD_PO_STATS_BH(net, field, val) SNMP_UPD_PO_STATS64_BH((net)->mib.ip_statistics, field, val)
198 #define NET_INC_STATS(net, field) SNMP_INC_STATS((net)->mib.net_statistics, field)
199 #define NET_INC_STATS_BH(net, field) SNMP_INC_STATS_BH((net)->mib.net_statistics, field)
200 #define NET_INC_STATS_USER(net, field) SNMP_INC_STATS_USER((net)->mib.net_statistics, field)
201 #define NET_ADD_STATS(net, field, adnd) SNMP_ADD_STATS((net)->mib.net_statistics, field, adnd)
202 #define NET_ADD_STATS_BH(net, field, adnd) SNMP_ADD_STATS_BH((net)->mib.net_statistics, field, adnd)
203 #define NET_ADD_STATS_USER(net, field, adnd) SNMP_ADD_STATS_USER((net)->mib.net_statistics, field, adnd)
204
205 unsigned long snmp_fold_field(void __percpu *mib, int offt);

```

```

206 #if BITS_PER_LONG==32
207 u64 snmp_fold_field64(void __percpu *mib, int offst, size_t sync_off);
208 #else
209 static inline u64 snmp_fold_field64(void __percpu *mib, int offst, size_t syncp_off)
210 {
211     return snmp_fold_field(mib, offst);
212 }
213 #endif
214
215 void inet_get_local_port_range(struct net *net, int *low, int *high);
216
217 #ifdef CONFIG_SYSCTL
218 static inline int inet_is_local_reserved_port(struct net *net, int port)
219 {
220     if (!net->ipv4.sysctl_local_reserved_ports)
221         return 0;
222     return test_bit(port, net->ipv4.sysctl_local_reserved_ports);
223 }
224
225 static inline bool sysctl_dev_name_is_allowed(const char *name)
226 {
227     return strcmp(name, "default") != 0 && strcmp(name, "all") != 0;
228 }
229
230 #else
231 static inline int inet_is_local_reserved_port(struct net *net, int port)
232 {
233     return 0;
234 }
235 #endif
236
237 /* From inetpeer.c */
238 extern int inet_peer_threshold;
239 extern int inet_peer_minttl;
240 extern int inet_peer_maxttl;
241
242 /* From ip_input.c */
243 extern int sysctl_ip_early_demux;
244
245 /* From ip_output.c */
246 extern int sysctl_ip_dynaddr;
247
248 void ipfrag_init(void);
249
250 void ip_static_sysctl_init(void);
251
252 #define IP4_REPLY_MARK(net, mark) \
253     ((net)->ipv4.sysctl_fwmark_reflect ? (mark) : 0)
254
255 static inline bool ip_is_fragment(const struct iphdr *iph)
256 {
257     return (iph->frag_off & htons(IP_MF | IP_OFFSET)) != 0;
258 }
259
260 #ifdef CONFIG_INET
261 #include <net/dst.h>
262
263 /* The function in 2.2 was invalid, producing wrong result for
264  * check=0xFEFF. It was noticed by Arthur Skawina _year_ ago. --ANK(000625) */
265 static inline
266 int ip_decrease_ttl(struct iphdr *iph)
267 {
268     u32 check = (__force u32)iph->check;
269     check += (__force u32)htons(0x0100);
270     iph->check = (__force __sum16)(check + (check >= 0xFFFF));
271     return --iph->ttl;
272 }
273
274 static inline
275 int ip_dont_fragment(struct sock *sk, struct dst_entry *dst)
276 {
277     return inet_sk(sk)->pmtudisc == IP_PMTUDISC_DO ||
278         (inet_sk(sk)->pmtudisc == IP_PMTUDISC_WANT &&
279          !(dst_metric_locked(dst, RTAX_MTU)));
280 }

```

```

281
282 static inline bool ip_sk_accept_pmtu(const struct sock *sk)
283 {
284     return inet_sk(sk)->pmtudisc != IP_PMTUDISC_INTERFACE &&
285            inet_sk(sk)->pmtudisc != IP_PMTUDISC_OMIT;
286 }
287
288 static inline bool ip_sk_use_pmtu(const struct sock *sk)
289 {
290     return inet_sk(sk)->pmtudisc < IP_PMTUDISC_PROBE;
291 }
292
293 static inline bool ip_sk_ignore_df(const struct sock *sk)
294 {
295     return inet_sk(sk)->pmtudisc < IP_PMTUDISC_DO ||
296            inet_sk(sk)->pmtudisc == IP_PMTUDISC_OMIT;
297 }
298
299 static inline unsigned int ip_dst_mtu_maybe_forward(const struct dst_entry *dst,
300                                                    bool forwarding)
301 {
302     struct net *net = dev_net(dst->dev);
303
304     if (net->ipv4.sysctl_ip_fwd_use_pmtu ||
305         dst_metric_locked(dst, RTAX_MTU) ||
306         !forwarding)
307         return dst_mtu(dst);
308
309     return min(dst->dev->mtu, IP_MAX_MTU);
310 }
311
312 static inline unsigned int ip_skb_dst_mtu(const struct sk_buff *skb)
313 {
314     if (!skb->sk || ip_sk_use_pmtu(skb->sk)) {
315         bool forwarding = IPCB(skb)->flags & IPSKB_FORWARDED;
316         return ip_dst_mtu_maybe_forward(skb_dst(skb), forwarding);
317     } else {
318         return min(skb_dst(skb)->dev->mtu, IP_MAX_MTU);
319     }
320 }
321
322 u32 ip_idsents_reserve(u32 hash, int segs);
323 void __ip_select_ident(struct net *net, struct iphdr *iph, int segs);
324
325 static inline void ip_select_ident_segs(struct net *net, struct sk_buff *skb,
326                                       struct sock *sk, int segs)
327 {
328     struct iphdr *iph = ip_hdr(skb);
329
330     if ((iph->frag_off & htons(IP_DF)) && !skb->ignore_df) {
331         /* This is only to work around buggy Windows95/2000
332          * VJ compression implementations. If the ID field
333          * does not change, they drop every other packet in
334          * a TCP stream using header compression.
335          */
336         if (sk && inet_sk(sk)->inet_daddr) {
337             iph->id = htons(inet_sk(sk)->inet_id);
338             inet_sk(sk)->inet_id += segs;
339         } else {
340             iph->id = 0;
341         }
342     } else {
343         __ip_select_ident(net, iph, segs);
344     }
345 }
346
347 static inline void ip_select_ident(struct net *net, struct sk_buff *skb,
348                                   struct sock *sk)
349 {
350     ip_select_ident_segs(net, skb, sk, 1);
351 }
352
353 static inline __wsum inet_compute_pseudo(struct sk_buff *skb, int proto)
354 {
355     return csum_tcpudp_nofold(ip_hdr(skb)->saddr, ip_hdr(skb)->daddr,

```

```

356         skb->len, proto, 0);
357     }
358
359     /* copy IPv4 saddr & daddr to flow_keys, possibly using 64bit Load/store
360      * Equivalent to :      flow->v4addrs.src = iph->saddr;
361      *                      flow->v4addrs.dst = iph->daddr;
362      */
363     static inline void iph\_to\_flow\_copy\_v4addrs(struct flow\_keys *flow,
364                                               const struct iphdr *iph)
365     {
366         BUILD_BUG_ON(offsetof(typeof(flow)->addrs), v4addrs.dst) !=
367                     offsetof(typeof(flow)->addrs), v4addrs.src) +
368                     sizeof(flow->addrs.v4addrs.src));
369         memcpy(&flow->addrs.v4addrs, &iph->saddr, sizeof(flow->addrs.v4addrs));
370         flow->control.addr_type = FLOW_DISSECTOR_KEY_IPV4_ADDRS;
371     }
372
373     static inline void inet\_set\_txhash(struct sock *sk)
374     {
375         struct inet\_sock *inet = inet\_sk(sk);
376         struct flow\_keys keys;
377
378         memset(&keys, 0, sizeof(keys));
379
380         keys.addrs.v4addrs.src = inet->inet_saddr;
381         keys.addrs.v4addrs.dst = inet->inet_daddr;
382         keys.control.addr\_type = FLOW_DISSECTOR_KEY_IPV4_ADDRS;
383         keys.ports.src = inet->inet_sport;
384         keys.ports.dst = inet->inet_dport;
385
386         sk->sk_txhash = flow\_hash\_from\_keys(&keys);
387     }
388
389     static inline __wsum inet\_gro\_compute\_pseudo(struct sk\_buff *skb, int proto)
390     {
391         const struct iphdr *iph = skb\_gro\_network\_header(skb);
392
393         return csum\_tcpudp\_nofold(iph->saddr, iph->daddr,
394                                   skb\_gro\_len(skb), proto, 0);
395     }
396
397     /*
398      *      Map a multicast IP onto multicast MAC for type ethernet.
399      */
400
401     static inline void ip\_eth\_mc\_map(__be32 naddr, char *buf)
402     {
403         __u32 addr=ntohl(naddr);
404         buf[0]=0x01;
405         buf[1]=0x00;
406         buf[2]=0x5e;
407         buf[5]=addr&0xFF;
408         addr>>=8;
409         buf[4]=addr&0xFF;
410         addr>>=8;
411         buf[3]=addr&0x7F;
412     }
413
414     /*
415      *      Map a multicast IP onto multicast MAC for type IP-over-InfiniBand.
416      *      Leave P_Key as 0 to be filled in by driver.
417      */
418
419     static inline void ip\_ib\_mc\_map(__be32 naddr, const unsigned char *broadcast, char *buf)
420     {
421         __u32 addr;
422         unsigned char scope = broadcast[5] & 0xF;
423
424         buf[0] = 0;           /* Reserved */
425         buf[1] = 0xff;        /* Multicast QPN */
426         buf[2] = 0xff;
427         buf[3] = 0xff;
428         addr = ntohl(naddr);
429         buf[4] = 0xff;
430         buf[5] = 0x10 | scope; /* scope from broadcast address */

```

```

431     buf[6] = 0x40;          /* IPv4 signature */
432     buf[7] = 0x1b;
433     buf[8] = broadcast[8];  /* P_Key */
434     buf[9] = broadcast[9];
435     buf[10] = 0;
436     buf[11] = 0;
437     buf[12] = 0;
438     buf[13] = 0;
439     buf[14] = 0;
440     buf[15] = 0;
441     buf[19] = addr & 0xff;
442     addr >>= 8;
443     buf[18] = addr & 0xff;
444     addr >>= 8;
445     buf[17] = addr & 0xff;
446     addr >>= 8;
447     buf[16] = addr & 0x0f;
448 }
449
450 static inline void ip_ipgre_mc_map(__be32 naddr, const unsigned char *broadcast, char *buf)
451 {
452     if ((broadcast[0] | broadcast[1] | broadcast[2] | broadcast[3]) != 0)
453         memcpy(buf, broadcast, 4);
454     else
455         memcpy(buf, &naddr, sizeof(naddr));
456 }
457
458 #if IS_ENABLED(CONFIG_IPV6)
459 #include <linux/ipv6.h>
460 #endif
461
462 static __inline__ void inet_reset_saddr(struct sock *sk)
463 {
464     inet_sk(sk)->inet_rcv_saddr = inet_sk(sk)->inet_saddr = 0;
465 #if IS_ENABLED(CONFIG_IPV6)
466     if (sk->sk_family == PF_INET6) {
467         struct ipv6_pinfo *np = inet6_sk(sk);
468
469         memset(&np->saddr, 0, sizeof(np->saddr));
470         memset(&sk->sk_v6_rcv_saddr, 0, sizeof(sk->sk_v6_rcv_saddr));
471     }
472 #endif
473 }
474
475 #endif
476
477 bool ip_call_ra_chain(struct sk_buff *skb);
478
479 /*
480  *      Functions provided by ip_fragment.c
481  */
482
483 enum ip_defrag_users {
484     IP_DEFRAG_LOCAL_DELIVER,
485     IP_DEFRAG_CALL_RA_CHAIN,
486     IP_DEFRAG_CONNTRACK_IN,
487     __IP_DEFRAG_CONNTRACK_IN_END = IP_DEFRAG_CONNTRACK_IN + USHRT_MAX,
488     IP_DEFRAG_CONNTRACK_OUT,
489     __IP_DEFRAG_CONNTRACK_OUT_END = IP_DEFRAG_CONNTRACK_OUT + USHRT_MAX,
490     IP_DEFRAG_CONNTRACK_BRIDGE_IN,
491     __IP_DEFRAG_CONNTRACK_BRIDGE_IN = IP_DEFRAG_CONNTRACK_BRIDGE_IN + USHRT_MAX,
492     IP_DEFRAG_VS_IN,
493     IP_DEFRAG_VS_OUT,
494     IP_DEFRAG_VS_FWD,
495     IP_DEFRAG_AF_PACKET,
496     IP_DEFRAG_MACVLAN,
497 };
498
499 /* Return true if the value of 'user' is between 'lower_bond'
500  * and 'upper_bond' inclusively.
501  */
502 static inline bool ip_defrag_user_in_between(u32 user,
503                                             enum ip_defrag_users lower_bond,
504                                             enum ip_defrag_users upper_bond)
505 {

```



```

506     return user >= lower_bond && user <= upper_bond;
507 }
508
509 int ip_defrag(struct sk_buff *skb, u32 user);
510 #ifdef CONFIG_INET
511 struct sk_buff *ip_check_defrag(struct sk_buff *skb, u32 user);
512 #else
513 static inline struct sk_buff *ip_check_defrag(struct sk_buff *skb, u32 user)
514 {
515     return skb;
516 }
517 #endif
518 int ip_frag_mem(struct net *net);
519
520 /*
521  *      Functions provided by ip_forward.c
522  */
523
524 int ip_forward(struct sk_buff *skb);
525
526 /*
527  *      Functions provided by ip_options.c
528  */
529
530 void ip_options_build(struct sk_buff *skb, struct ip_options *opt,
531                      __be32 daddr, struct rtable *rt, int is_frag);
532
533 int ip_options_echo(struct ip_options *dopt, struct sk_buff *skb,
534                     const struct ip_options *sopt);
535 static inline int ip_options_echo(struct ip_options *dopt, struct sk_buff *skb)
536 {
537     return ip_options_echo(dopt, skb, &IPCB(skb)->opt);
538 }
539
540 void ip_options_fragment(struct sk_buff *skb);
541 int ip_options_compile(struct net *net, struct ip_options *opt,
542                        struct sk_buff *skb);
543 int ip_options_get(struct net *net, struct ip_options_rcu **optp,
544                    unsigned char *data, int optlen);
545 int ip_options_get_from_user(struct net *net, struct ip_options_rcu **optp,
546                               unsigned char __user *data, int optlen);
547 void ip_options_undo(struct ip_options *opt);
548 void ip_forward_options(struct sk_buff *skb);
549 int ip_options_rcv_srr(struct sk_buff *skb);
550
551 /*
552  *      Functions provided by ip_sockglue.c
553  */
554
555 void ipv4_pktinfo_prepare(const struct sock *sk, struct sk_buff *skb);
556 void ip_msg_recv_offset(struct msghdr *msg, struct sk_buff *skb, int offset);
557 int ip_msg_send(struct net *net, struct msghdr *msg,
558                 struct ipcm_cookie *ipc, bool allow_ipv6);
559 int ip_setsockopt(struct sock *sk, int level, int optname, char __user *optval,
560                  unsigned int optlen);
561 int ip_getsockopt(struct sock *sk, int level, int optname, char __user *optval,
562                  int __user *optlen);
563 int compat_ip_setsockopt(struct sock *sk, int level, int optname,
564                          char __user *optval, unsigned int optlen);
565 int compat_ip_getsockopt(struct sock *sk, int level, int optname,
566                          char __user *optval, int __user *optlen);
567 int ip_ra_control(struct sock *sk, unsigned char on,
568                  void (*destructor)(struct sock *));
569
570 int ip_recv_error(struct sock *sk, struct msghdr *msg, int len, int *addr_len);
571 void ip_icmp_error(struct sock *sk, struct sk_buff *skb, int err, __be16 port,
572                   u32 info, u8 *payload);
573 void ip_local_error(struct sock *sk, int err, __be32 daddr, __be16 dport,
574                     u32 info);
575
576 static inline void ip_msg_recv(struct msghdr *msg, struct sk_buff *skb)
577 {
578     ip_msg_recv_offset(msg, skb, 0);
579 }
580

```



```
581 bool icmp_global_allow(void);
582 extern int sysctl_icmp_msgs_per_sec;
583 extern int sysctl_icmp_msgs_burst;
584
585 #ifdef CONFIG_PROC_FS
586 int ip_misc_proc_init(void);
587 #endif
588
589 #endif /* _IP_H */
590
```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)