

Linux Cross Reference

[Free Electrons](#)

Embedded Linux Experts

• [source navigation](#) • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version: [2.0.40](#) [2.2.26](#) [2.4.37](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#) [3.18](#) [3.19](#) [4.0](#) [4.1](#) [4.2](#)

[Linux](#)/[include](#)/[net](#)/[sock.h](#)

```

1  /*
2  *  INET          An implementation of the TCP/IP protocol suite for the LINUX
3  *              operating system.  INET is implemented using the BSD Socket
4  *              interface as the means of communication with the user level.
5  *
6  *              Definitions for the AF_INET socket handler.
7  *
8  *  Version:      @(#)sock.h      1.0.4      05/13/93
9  *
10 *  Authors:      Ross Biro
11 *              Fred N. van Kempen, <waltje@uWalt.NL.Mugnet.ORG>
12 *              Corey Minyard <wf-rch!minyard@relay.EU.net>
13 *              Florian La Roche <flla@stud.uni-sb.de>
14 *
15 *  Fixes:
16 *              Alan Cox          :      Volatiles in skbuff pointers. See
17 *                                skbuff comments. May be overdone,
18 *                                better to prove they can be removed
19 *                                than the reverse.
20 *              Alan Cox          :      Added a zapped field for tcp to note
21 *                                a socket is reset and must stay shut up
22 *              Alan Cox          :      New fields for options
23 *              Pauline Middelink :      identd support
24 *              Alan Cox          :      Eliminate Low Level recv/recvfrom
25 *              David S. Miller :      New socket lookup architecture.
26 *              Steve Whitehouse:      Default routines for sock_ops
27 *              Arnaldo C. Melo :      removed net_pinfo, tp_pinfo and made
28 *                                protinfo be just a void pointer, as the
29 *                                protocol specific parts were moved to
30 *                                respective headers and ipv4/v6, etc now
31 *                                use private slabcaches for its socks
32 *              Pedro Hortas      :      New flags field for socket options
33 *
34 *
35 *              This program is free software; you can redistribute it and/or
36 *              modify it under the terms of the GNU General Public License
37 *              as published by the Free Software Foundation; either version
38 *              2 of the License, or (at your option) any later version.
39 */
40 #ifndef _SOCK_H
41 #define _SOCK_H
42
43 #include <linux/hardirq.h>
44 #include <linux/kernel.h>
45 #include <linux/list.h>
46 #include <linux/list_nulls.h>
47 #include <linux/timer.h>

```

```

48 #include <linux/cache.h>
49 #include <linux/bitops.h>
50 #include <linux/lockdep.h>
51 #include <linux/netdevice.h>
52 #include <linux/skbuff.h>          /* struct sk_buff */
53 #include <linux/mm.h>
54 #include <linux/security.h>
55 #include <linux/slab.h>
56 #include <linux/uaccess.h>
57 #include <linux/page_counter.h>
58 #include <linux/memcontrol.h>
59 #include <linux/static_key.h>
60 #include <linux/sched.h>
61
62 #include <linux/filter.h>
63 #include <linux/rculist_nulls.h>
64 #include <linux/poll.h>
65
66 #include <linux/atomic.h>
67 #include <net/dst.h>
68 #include <net/checksum.h>
69 #include <net/tcp_states.h>
70 #include <linux/net_tstamp.h>
71
72 struct cgroup;
73 struct cgroup_subsys;
74 #ifdef CONFIG_NET
75 int mem_cgroup_sockets_init(struct mem_cgroup *memcg, struct cgroup_subsys *ss);
76 void mem_cgroup_sockets_destroy(struct mem_cgroup *memcg);
77 #else
78 static inline
79 int mem_cgroup_sockets_init(struct mem_cgroup *memcg, struct cgroup_subsys *ss)
80 {
81     return 0;
82 }
83 static inline
84 void mem_cgroup_sockets_destroy(struct mem_cgroup *memcg)
85 {
86 }
87 #endif
88 /*
89  * This structure really needs to be cleaned up.
90  * Most of it is for TCP, and not used by any of
91  * the other protocols.
92  */
93
94 /* Define this to get the SOCK_DBG debugging facility. */
95 #define SOCK_DEBUGGING
96 #ifdef SOCK_DEBUGGING
97 #define SOCK_DEBUG(sk, msg...) do { if ((sk) && sock_flag((sk), SOCK_DBG)) \
98                                     printk(KERN_DEBUG msg); } while (0)
99 #else
100 /* Validate arguments and do nothing */
101 static inline __printf(2, 3)
102 void SOCK_DEBUG(const struct sock *sk, const char *msg, ...)
103 {
104 }
105 #endif
106
107 /* This is the per-socket lock. The spinlock provides a synchronization
108  * between user contexts and software interrupt processing, whereas the
109  * mini-semaphore synchronizes multiple users amongst themselves.
110  */
111 typedef struct {
112     spinlock_t          slock;
113     int                 owned;
114     wait_queue_head_t   wq;
115 } /*

```

```

116      * We express the mutex-alike socket_lock semantics
117      * to the lock validator by explicitly managing
118      * the slock as a lock variant (in addition to
119      * the slock itself):
120      */
121 #ifdef CONFIG_DEBUG_LOCK_ALLOC
122     struct lockdep\_map dep_map;
123 #endif
124 } socket\_lock\_t;
125
126 struct sock;
127 struct proto;
128 struct net;
129
130 typedef \_\_u32 \_\_bitwise portpair;
131 typedef \_\_u64 \_\_bitwise addrpair;
132
133 /**
134  * struct sock_common - minimal network layer representation of sockets
135  * @skc_daddr: Foreign IPv4 addr
136  * @skc_rcv_saddr: Bound local IPv4 addr
137  * @skc_hash: hash value used with various protocol lookup tables
138  * @skc_u16hashes: two u16 hash values used by UDP lookup tables
139  * @skc_dport: placeholder for inet_dport/tw_dport
140  * @skc_num: placeholder for inet_num/tw_num
141  * @skc_family: network address family
142  * @skc_state: Connection state
143  * @skc_reuse: %SO_REUSEADDR setting
144  * @skc_reuseport: %SO_REUSEPORT setting
145  * @skc_bound_dev_if: bound device index if != 0
146  * @skc_bind_node: bind hash linkage for various protocol lookup tables
147  * @skc_portaddr_node: second hash linkage for UDP/UDP-Lite protocol
148  * @skc_prot: protocol handlers inside a network family
149  * @skc_net: reference to the network namespace of this socket
150  * @skc_node: main hash linkage for various protocol lookup tables
151  * @skc_nulls_node: main hash linkage for TCP/UDP/UDP-Lite protocol
152  * @skc_tx_queue_mapping: tx queue number for this connection
153  * @skc_refcnt: reference count
154  *
155  * This is the minimal network layer representation of sockets, the header
156  * for struct sock and struct inet_timewait_sock.
157  */
158 struct sock\_common {
159     /* skc_daddr and skc_rcv_saddr must be grouped on a 8 bytes aligned
160      * address on 64bit arches : cf INET_MATCH()
161      */
162     union {
163         \_\_addrpair skc_addrpair;
164         struct {
165             \_\_be32 skc_daddr;
166             \_\_be32 skc_rcv_saddr;
167         };
168     };
169     union {
170         unsigned int skc_hash;
171         \_\_u16 skc_u16hashes[2];
172     };
173     /* skc_dport && skc_num must be grouped as well */
174     union {
175         \_\_portpair skc_portpair;
176         struct {
177             \_\_be16 skc_dport;
178             \_\_u16 skc_num;
179         };
180     };
181
182     unsigned short skc_family;
183     volatile unsigned char skc_state;

```

```

184 unsigned char      skc_reuse:4;
185 unsigned char      skc_reuseport:1;
186 unsigned char      skc_ipv6only:1;
187 unsigned char      skc_net_refcnt:1;
188 int                skc_bound_dev_if;
189 union {
190     struct hlist\_node      skc_bind_node;
191     struct hlist\_nulls\_node skc_portaddr_node;
192 };
193 struct proto            *skc_prot;
194 possible net t      skc_net;
195
196 #if IS\_ENABLED(CONFIG_IPV6)
197     struct in6\_addr      skc_v6_daddr;
198     struct in6\_addr      skc_v6_rcv_saddr;
199 #endif
200
201 atomic64\_t          skc_cookie;
202
203 /*
204  * fields between dontcopy_begin/dontcopy_end
205  * are not copied in sock_copy()
206  */
207 /* private: */
208 int                skc_dontcopy_begin[0];
209 /* public: */
210 union {
211     struct hlist\_node      skc_node;
212     struct hlist\_nulls\_node skc_nulls_node;
213 };
214 int                skc_tx_queue_mapping;
215 atomic t          skc_refcnt;
216 /* private: */
217 int                skc_dontcopy_end[0];
218 /* public: */
219 };
220
221 struct cg\_proto;
222 /**
223  * struct sock - network layer representation of sockets
224  * @__sk_common: shared layout with inet_timewait_sock
225  * @sk_shutdown: mask of %SEND_SHUTDOWN and/or %RCV_SHUTDOWN
226  * @sk_userlocks: %SO_SNDBUF and %SO_RCVBUF settings
227  * @sk_Lock: synchronizer
228  * @sk_rcvbuf: size of receive buffer in bytes
229  * @sk_wq: sock wait queue and async head
230  * @sk_rx_dst: receive input route used by early demux
231  * @sk_dst_cache: destination cache
232  * @sk_dst_lock: destination cache lock
233  * @sk_policy: flow policy
234  * @sk_receive_queue: incoming packets
235  * @sk_wmem_alloc: transmit queue bytes committed
236  * @sk_write_queue: Packet sending queue
237  * @sk_omem_alloc: "o" is "option" or "other"
238  * @sk_wmem_queued: persistent queue size
239  * @sk_forward_alloc: space allocated forward
240  * @sk_napi_id: id of the last napi context to receive data for sk
241  * @sk_ll_usec: usecs to busypoll when there is no data
242  * @sk_allocation: allocation mode
243  * @sk_pacing_rate: Pacing rate (if supported by transport/packet scheduler)
244  * @sk_max_pacing_rate: Maximum pacing rate (%SO_MAX_PACING_RATE)
245  * @sk_sndbuf: size of send buffer in bytes
246  * @sk_flags: %SO_LINGER (l_onoff), %SO_BROADCAST, %SO_KEEPAIVE,
247  *            %SO_OOBINLINE settings, %SO_TIMESTAMPING settings
248  * @sk_no_check_tx: %SO_NO_CHECK setting, set checksum in TX packets
249  * @sk_no_check_rx: allow zero checksum in RX packets
250  * @sk_route_caps: route capabilities (e.g. %NETIF_F_TSO)
251  * @sk_route_nocaps: forbidden route capabilities (e.g. NETIF_F_GSO_MASK)

```

```

252 *      @sk_gso_type: GSO type (e.g. %SKB_GSO_TCPV4)
253 *      @sk_gso_max_size: Maximum GSO segment size to build
254 *      @sk_gso_max_segs: Maximum number of GSO segments
255 *      @sk_lingertime: %SO_LINGER l_linger setting
256 *      @sk_backlog: always used with the per-socket spinlock held
257 *      @sk_callback_lock: used with the callbacks in the end of this struct
258 *      @sk_error_queue: rarely used
259 *      @sk_prot_creator: sk_prot of original sock creator (see ipv6_setsockopt,
260 *                      IPV6_ADDRFORM for instance)
261 *      @sk_err: Last error
262 *      @sk_err_soft: errors that don't cause failure but are the cause of a
263 *                      persistent failure not just 'timed out'
264 *      @sk_drops: raw/udp drops counter
265 *      @sk_ack_backlog: current listen backlog
266 *      @sk_max_ack_backlog: listen backlog set in listen()
267 *      @sk_priority: %SO_PRIORITY setting
268 *      @sk_cgrp_prioidx: socket group's priority map index
269 *      @sk_type: socket type (%SOCK_STREAM, etc)
270 *      @sk_protocol: which protocol this socket belongs in this network family
271 *      @sk_peer_pid: &struct pid for this socket's peer
272 *      @sk_peer_cred: %SO_PEERCRED setting
273 *      @sk_rcvlowat: %SO_RCVLOWAT setting
274 *      @sk_rcvtimeo: %SO_RCVTIMEO setting
275 *      @sk_sndtimeo: %SO_SNDTIMEO setting
276 *      @sk_rxhash: flow hash received from netif layer
277 *      @sk_incoming_cpu: record cpu processing incoming packets
278 *      @sk_txhash: computed flow hash for use on transmit
279 *      @sk_filter: socket filtering instructions
280 *      @sk_timer: sock cleanup timer
281 *      @sk_stamp: time stamp of last packet received
282 *      @sk_tsflags: SO_TIMESTAMPING socket options
283 *      @sk_tskey: counter to disambiguate concurrent tstamp requests
284 *      @sk_socket: Identd and reporting IO signals
285 *      @sk_user_data: RPC Layer private data
286 *      @sk_frag: cached page frag
287 *      @sk_peek_off: current peek_offset value
288 *      @sk_send_head: front of stuff to transmit
289 *      @sk_security: used by security modules
290 *      @sk_mark: generic packet mark
291 *      @sk_classid: this socket's cgroup classid
292 *      @sk_cgrp: this socket's cgroup-specific proto data
293 *      @sk_write_pending: a write to stream socket waits to start
294 *      @sk_state_change: callback to indicate change in the state of the sock
295 *      @sk_data_ready: callback to indicate there is data to be processed
296 *      @sk_write_space: callback to indicate there is bf sending space available
297 *      @sk_error_report: callback to indicate errors (e.g. %MSG_ERRQUEUE)
298 *      @sk_backlog_rcv: callback to process the backlog
299 *      @sk_destruct: called at sock freeing time, i.e. when all refcnt == 0
300 */
301 struct sock {
302     /*
303      * Now struct inet_timewait_sock also uses sock_common, so please just
304      * don't add nothing before this first member (__sk_common) --acme
305      */
306     struct sock_common    __sk_common;
307 #define sk_node            __sk_common.skc_node
308 #define sk_nulls_node     __sk_common.skc_nulls_node
309 #define sk_refcnt         __sk_common.skc_refcnt
310 #define sk_tx_queue_mapping __sk_common.skc_tx_queue_mapping
311 #define sk_dontcopy_begin __sk_common.skc_dontcopy_begin
312 #define sk_dontcopy_end   __sk_common.skc_dontcopy_end
313 #define sk_hash           __sk_common.skc_hash
314 #define sk_portpair       __sk_common.skc_portpair
315 #define sk_num            __sk_common.skc_num
316 #define sk_dport          __sk_common.skc_dport
317 #define sk_addrpair       __sk_common.skc_addrpair
318 #define sk_daddr          __sk_common.skc_daddr

```

```

320 #define sk\_rcv\_saddr          __sk_common.skc_rcv_saddr
321 #define sk\_family              __sk_common.skc_family
322 #define sk\_state                __sk_common.skc_state
323 #define sk\_reuse                __sk_common.skc_reuse
324 #define sk\_reuseport           __sk_common.skc_reuseport
325 #define sk\_ipv6only             __sk_common.skc_ipv6only
326 #define sk\_net\_refcnt           __sk_common.skc_net_refcnt
327 #define sk\_bound\_dev\_if        __sk_common.skc_bound_dev_if
328 #define sk\_bind\_node            __sk_common.skc_bind_node
329 #define sk\_prot                  __sk_common.skc_prot
330 #define sk\_net                    __sk_common.skc_net
331 #define sk\_v6\_daddr              __sk_common.skc_v6_daddr
332 #define sk\_v6\_rcv\_saddr         __sk_common.skc_v6_rcv_saddr
333 #define sk\_cookie               __sk_common.skc_cookie
334
335     socket\_lock\_t                sk_lock;
336     struct sk\_buff\_head         sk_receive_queue;
337     /*
338      * The backlog queue is special, it is always used with
339      * the per-socket spinlock held and requires low latency
340      * access. Therefore we special case it's implementation.
341      * Note : rmem_alloc is in this structure to fill a hole
342      * on 64bit arches, not because its logically part of
343      * backlog.
344      */
345     struct {
346         atomic\_t                    rmem_alloc;
347         int                      len;
348         struct sk\_buff          *head;
349         struct sk\_buff          *tail;
350     } sk_backlog;
351 #define sk\_rmem\_alloc sk_backlog.rmem_alloc
352     int                      sk_forward_alloc;
353 #ifdef CONFIG_RPS
354     u32                      sk_rxhash;
355 #endif
356     u16                      sk_incoming_cpu;
357     /* 16bit hole
358      * Warned : sk_incoming_cpu can be set from softirq,
359      * Do not use this hole without fully understanding possible issues.
360      */
361
362     u32                      sk_txhash;
363 #ifdef CONFIG_NET_RX_BUSY_POLL
364     unsigned int             sk_napi_id;
365     unsigned int             sk_ll_usec;
366 #endif
367     atomic\_t                  sk_drops;
368     int                      sk_rcvbuf;
369
370     struct sk\_filter \_\_rcu    *sk\_filter;
371     struct socket\_wq \_\_rcu   *sk_wq;
372
373 #ifdef CONFIG_XFRM
374     struct xfrm\_policy       *sk_policy[2];
375 #endif
376     unsigned long             sk_flags;
377     struct dst\_entry          *sk_rx_dst;
378     struct dst\_entry \_\_rcu    *sk_dst_cache;
379     spinlock\_t                sk_dst_lock;
380     atomic\_t                  sk_wmem_alloc;
381     atomic\_t                  sk_omem_alloc;
382     int                      sk_sndbuf;
383     struct sk\_buff\_head       sk_write_queue;
384     kmemcheck\_bitfield\_begin(flags);
385     unsigned int              sk_shutdown : 2,
386                             sk_no_check_tx : 1,
387                             sk_no_check_rx : 1,

```



```

388         sk_userlocks : 4,
389         sk_protocol  : 8,
390         sk_type      : 16;
391     kmemcheck bitfield end(flags);
392     int sk_wmem_queued;
393     gfp\_t sk_allocation;
394     u32 sk_pacing_rate; /* bytes per second */
395     u32 sk_max_pacing_rate;
396     netdev\_features\_t sk_route_caps;
397     netdev\_features\_t sk_route_nocaps;
398     int sk_gso_type;
399     unsigned int sk_gso_max_size;
400     u16 sk_gso_max_segs;
401     int sk_rcvlowat;
402     unsigned long sk_lingertime;
403     struct sk\_buff\_head sk_error_queue;
404     struct proto *sk_prot_creator;
405     rwlock\_t sk_callback_lock;
406     int sk_err,
407         sk_err_soft;
408     u32 sk_ack_backlog;
409     u32 sk_max_ack_backlog;
410     u32 sk_priority;
411     #if IS\_ENABLED(CONFIG_CGROUP_NET_PRIO)
412         u32 sk_cgrp_prioidx;
413     #endif
414     struct pid *sk_peer_pid;
415     const struct cred *sk_peer_cred;
416     long sk_rcvtimeo;
417     long sk_sndtimeo;
418     struct timer\_list sk_timer;
419     ktime\_t sk_stamp;
420     u16 sk_tsflags;
421     u32 sk_tskey;
422     struct socket *sk_socket;
423     void *sk_user_data;
424     struct page\_frag sk_frag;
425     struct sk\_buff *sk_send_head;
426     s32 sk_peek_off;
427     int sk_write_pending;
428     #ifdef CONFIG_SECURITY
429     void *sk_security;
430     #endif
431     u32 sk_mark;
432     u32 sk_classid;
433     struct cg\_proto *sk_cgrp;
434     void (*sk_state_change)(struct sock *sk);
435     void (*sk_data_ready)(struct sock *sk);
436     void (*sk_write_space)(struct sock *sk);
437     void (*sk_error_report)(struct sock *sk);
438     int (*sk_backlog_rcv)(struct sock *sk,
439                          struct sk\_buff *skb);
440     void (*sk_destruct)(struct sock *sk);
441 };
442
443 #define \_\_sk\_user\_data(sk) (((void \_\_rcu *)&(sk->sk_user_data)))
444
445 #define rcu\_dereference\_sk\_user\_data(sk) rcu\_dereference(\_\_sk\_user\_data((sk)))
446 #define rcu\_assign\_sk\_user\_data(sk, ptr) rcu\_assign\_pointer(\_\_sk\_user\_data((sk)), ptr)
447
448 /*
449  * SK_CAN_REUSE and SK_NO_REUSE on a socket mean that the socket is OK
450  * or not whether his port will be reused by someone else. SK_FORCE_REUSE
451  * on a socket means that the socket will reuse everybody else's port
452  * without looking at the other's sk_reuse value.
453  */
454
455 #define SK_NO_REUSE 0

```

```

456 #define SK_CAN_REUSE 1
457 #define SK_FORCE_REUSE 2
458
459 static inline int sk_peek_offset(struct sock *sk, int flags)
460 {
461     if ((flags & MSG_PEEK) && (sk->sk_peek_off >= 0))
462         return sk->sk_peek_off;
463     else
464         return 0;
465 }
466
467 static inline void sk_peek_offset_bwd(struct sock *sk, int val)
468 {
469     if (sk->sk_peek_off >= 0) {
470         if (sk->sk_peek_off >= val)
471             sk->sk_peek_off -= val;
472         else
473             sk->sk_peek_off = 0;
474     }
475 }
476
477 static inline void sk_peek_offset_fwd(struct sock *sk, int val)
478 {
479     if (sk->sk_peek_off >= 0)
480         sk->sk_peek_off += val;
481 }
482
483 /*
484  * Hashed Lists helper routines
485  */
486 static inline struct sock *sk_entry(const struct hlist_node *node)
487 {
488     return hlist_entry(node, struct sock, sk_node);
489 }
490
491 static inline struct sock *__sk_head(const struct hlist_head *head)
492 {
493     return hlist_entry(head->first, struct sock, sk_node);
494 }
495
496 static inline struct sock *sk_head(const struct hlist_head *head)
497 {
498     return hlist_empty(head) ? NULL : __sk_head(head);
499 }
500
501 static inline struct sock *__sk_nulls_head(const struct hlist_nulls_head *head)
502 {
503     return hlist_nulls_entry(head->first, struct sock, sk_nulls_node);
504 }
505
506 static inline struct sock *sk_nulls_head(const struct hlist_nulls_head *head)
507 {
508     return hlist_nulls_empty(head) ? NULL : __sk_nulls_head(head);
509 }
510
511 static inline struct sock *sk_next(const struct sock *sk)
512 {
513     return sk->sk_node.next ?
514         hlist_entry(sk->sk_node.next, struct sock, sk_node) : NULL;
515 }
516
517 static inline struct sock *sk_nulls_next(const struct sock *sk)
518 {
519     return (!is_a_nulls(sk->sk_nulls_node.next)) ?
520         hlist_nulls_entry(sk->sk_nulls_node.next,
521             struct sock, sk_nulls_node) :
522         NULL;
523 }

```



```

524
525 static inline bool sk\_unhashed(const struct sock *sk)
526 {
527     return hlist\_unhashed(&sk->sk\_node);
528 }
529
530 static inline bool sk\_hashed(const struct sock *sk)
531 {
532     return !sk\_unhashed(sk);
533 }
534
535 static inline void sk\_node\_init(struct hlist\_node *node)
536 {
537     node->pprev = NULL;
538 }
539
540 static inline void sk\_nulls\_node\_init(struct hlist\_nulls\_node *node)
541 {
542     node->pprev = NULL;
543 }
544
545 static inline void \_\_sk\_del\_node(struct sock *sk)
546 {
547     \_\_hlist\_del(&sk->sk\_node);
548 }
549
550 /* NB: equivalent to hlist_del_init_rcu */
551 static inline bool \_\_sk\_del\_node\_init(struct sock *sk)
552 {
553     if (sk\_hashed(sk)) {
554         \_\_sk\_del\_node(sk);
555         sk\_node\_init(&sk->sk\_node);
556         return true;
557     }
558     return false;
559 }
560
561 /* Grab socket reference count. This operation is valid only
562    when sk is ALREADY grabbed f.e. it is found in hash table
563    or a list and the lookup is made under lock preventing hash table
564    modifications.
565    */
566
567 static inline void sock\_hold(struct sock *sk)
568 {
569     atomic\_inc(&sk->sk\_refcnt);
570 }
571
572 /* Ungrab socket in the context, which assumes that socket refcnt
573    cannot hit zero, f.e. it is true in context of any socketcall.
574    */
575 static inline void \_\_sock\_put(struct sock *sk)
576 {
577     atomic\_dec(&sk->sk\_refcnt);
578 }
579
580 static inline bool sk\_del\_node\_init(struct sock *sk)
581 {
582     bool rc = \_\_sk\_del\_node\_init(sk);
583
584     if (rc) {
585         /* paranoid for a while -acme */
586         WARN\_ON(atomic\_read(&sk->sk\_refcnt) == 1);
587         \_\_sock\_put(sk);
588     }
589     return rc;
590 }
591 #define sk\_del\_node\_init\_rcu(sk)         sk\_del\_node\_init(sk)

```

```

592
593 static inline bool \_\_sk\_nulls\_del\_node\_init\_rcu(struct sock *sk)
594 {
595     if (sk\_hashed(sk)) {
596         hlist\_nulls\_del\_init\_rcu(&sk->sk\_nulls\_node);
597         return true;
598     }
599     return false;
600 }
601
602 static inline bool sk\_nulls\_del\_node\_init\_rcu(struct sock *sk)
603 {
604     bool rc = \_\_sk\_nulls\_del\_node\_init\_rcu(sk);
605
606     if (rc) {
607         /* paranoid for a while -acme */
608         WARN\_ON(atomic\_read(&sk->sk\_refcnt) == 1);
609         \_\_sock\_put(sk);
610     }
611     return rc;
612 }
613
614 static inline void \_\_sk\_add\_node(struct sock *sk, struct hlist\_head *list)
615 {
616     hlist\_add\_head(&sk->sk\_node, list);
617 }
618
619 static inline void sk\_add\_node(struct sock *sk, struct hlist\_head *list)
620 {
621     sock\_hold(sk);
622     \_\_sk\_add\_node(sk, list);
623 }
624
625 static inline void sk\_add\_node\_rcu(struct sock *sk, struct hlist\_head *list)
626 {
627     sock\_hold(sk);
628     hlist\_add\_head\_rcu(&sk->sk\_node, list);
629 }
630
631 static inline void \_\_sk\_nulls\_add\_node\_rcu(struct sock *sk, struct hlist\_nulls\_head *list)
632 {
633     hlist\_nulls\_add\_head\_rcu(&sk->sk\_nulls\_node, list);
634 }
635
636 static inline void sk\_nulls\_add\_node\_rcu(struct sock *sk, struct hlist\_nulls\_head *list)
637 {
638     sock\_hold(sk);
639     \_\_sk\_nulls\_add\_node\_rcu(sk, list);
640 }
641
642 static inline void \_\_sk\_del\_bind\_node(struct sock *sk)
643 {
644     hlist\_del(&sk->sk\_bind\_node);
645 }
646
647 static inline void sk\_add\_bind\_node(struct sock *sk,
648                                     struct hlist\_head *list)
649 {
650     hlist\_add\_head(&sk->sk\_bind\_node, list);
651 }
652
653 #define sk\_for\_each(__sk, list) \
654     hlist\_for\_each\_entry(__sk, list, sk\_node)
655 #define sk\_for\_each\_rcu(__sk, list) \
656     hlist\_for\_each\_entry\_rcu(__sk, list, sk\_node)
657 #define sk\_nulls\_for\_each(__sk, node, list) \
658     hlist\_nulls\_for\_each\_entry(__sk, node, list, sk\_nulls\_node)
659 #define sk\_nulls\_for\_each\_rcu(__sk, node, list) \

```

```

660 hlist\_nulls\_for\_each\_entry\_rcu(__sk, node, list, sk\_nulls\_node)
661 #define sk\_for\_each\_from(__sk) \
662 hlist\_for\_each\_entry\_from(__sk, sk\_node)
663 #define sk\_nulls\_for\_each\_from(__sk, node) \
664     if (__sk && ({ node = &(__sk)->sk\_nulls\_node; 1; })) \
665         hlist\_nulls\_for\_each\_entry\_from(__sk, node, sk\_nulls\_node)
666 #define sk\_for\_each\_safe(__sk, tmp, list) \
667     hlist\_for\_each\_entry\_safe(__sk, tmp, list, sk\_node)
668 #define sk\_for\_each\_bound(__sk, list) \
669     hlist\_for\_each\_entry(__sk, list, sk\_bind\_node)
670
671 /**
672  * sk\_nulls\_for\_each\_entry\_offset - iterate over a list at a given struct offset
673  * @tpos:      the type * to use as a loop cursor.
674  * @pos:       the &struct hlist_node to use as a loop cursor.
675  * @head:      the head for your list.
676  * @offset:    offset of hlist_node within the struct.
677  *
678  */
679 #define sk\_nulls\_for\_each\_entry\_offset(tpos, pos, head, offset) \
680     for (pos = (head)->first; \
681         (!is\_a\_nulls(pos)) && \
682         ({ tpos = (typeof(*tpos) *)((void *)pos - offset); 1;}); \
683         pos = pos->next)
684
685 static inline struct user\_namespace *sk\_user\_ns(struct sock *sk)
686 {
687     /* Careful only use this in a context where these parameters
688      * can not change and must all be valid, such as recvmss from
689      * userspace.
690      */
691     return sk->sk_socket->file->f_cred->user\_ns;
692 }
693
694 /* Sock flags */
695 enum sock\_flags {
696     SOCK_DEAD,
697     SOCK_DONE,
698     SOCK_URGINLINE,
699     SOCK_KEEPOPEN,
700     SOCK_LINGER,
701     SOCK_DESTROY,
702     SOCK_BROADCAST,
703     SOCK_TIMESTAMP,
704     SOCK_ZAPPED,
705     SOCK_USE_WRITE_QUEUE, /* whether to call sk->sk_write_space in sock_wfree */
706     SOCK_DBG, /* %SO_DEBUG setting */
707     SOCK_RCVTSTAMP, /* %SO_TIMESTAMP setting */
708     SOCK_RCVTSTAMPNS, /* %SO_TIMESTAMPNS setting */
709     SOCK_LOCALROUTE, /* route locally only, %SO_DONTRROUTE setting */
710     SOCK_QUEUE_SHRUNK, /* write queue has been shrunk recently */
711     SOCK_MEMALLOC, /* VM depends on this socket for swapping */
712     SOCK_TIMESTAMPING_RX_SOFTWARE, /* %SOF_TIMESTAMPING_RX_SOFTWARE */
713     SOCK_FASYNC, /* fasync() active */
714     SOCK_RXQ_OVFL,
715     SOCK_ZEROCOPY, /* buffers from userspace */
716     SOCK_WIFI_STATUS, /* push wifi status to userspace */
717     SOCK_NOFCS, /* Tell NIC not to do the Ethernet FCS.
718                  * Will use last 4 bytes of packet sent from
719                  * user-space instead.
720                  */
721     SOCK_FILTER_LOCKED, /* Filter cannot be changed anymore */
722     SOCK_SELECT_ERR_QUEUE, /* Wake select on error queue */
723 };
724
725 static inline void sock\_copy\_flags(struct sock *nsk, struct sock *osk)
726 {
727     nsk->sk_flags = osk->sk_flags;

```

```

728 }
729
730 static inline void sock_set_flag(struct sock *sk, enum sock_flags flag)
731 {
732     __set_bit(flag, &sk->sk_flags);
733 }
734
735 static inline void sock_reset_flag(struct sock *sk, enum sock_flags flag)
736 {
737     __clear_bit(flag, &sk->sk_flags);
738 }
739
740 static inline bool sock_flag(const struct sock *sk, enum sock_flags flag)
741 {
742     return test_bit(flag, &sk->sk_flags);
743 }
744
745 #ifdef CONFIG_NET
746 extern struct static_key memalloc_socks;
747 static inline int sk_memalloc_socks(void)
748 {
749     return static_key_false(&memalloc_socks);
750 }
751 #else
752
753 static inline int sk_memalloc_socks(void)
754 {
755     return 0;
756 }
757
758 #endif
759
760 static inline gfp_t sk_gfp_atomic(struct sock *sk, gfp_t gfp_mask)
761 {
762     return GFP_ATOMIC | (sk->sk_allocation & GFP_MEMALLOC);
763 }
764
765 static inline void sk_acceptq_removed(struct sock *sk)
766 {
767     sk->sk_ack_backlog--;
768 }
769
770 static inline void sk_acceptq_added(struct sock *sk)
771 {
772     sk->sk_ack_backlog++;
773 }
774
775 static inline bool sk_acceptq_is_full(const struct sock *sk)
776 {
777     return sk->sk_ack_backlog > sk->sk_max_ack_backlog;
778 }
779
780 /*
781  * Compute minimal free write space needed to queue new packets.
782  */
783 static inline int sk_stream_min_wspace(const struct sock *sk)
784 {
785     return sk->sk_wmem_queued >> 1;
786 }
787
788 static inline int sk_stream_wspace(const struct sock *sk)
789 {
790     return sk->sk_sndbuf - sk->sk_wmem_queued;
791 }
792
793 void sk_stream_write_space(struct sock *sk);
794
795 /* OOB backlog add */

```

```

796 static inline void __sk_add_backlog(struct sock *sk, struct sk_buff *skb)
797 {
798     /* dont let skb dst not refcounted, we are going to leave rcu lock */
799     skb_dst_force(skb);
800
801     if (!sk->sk_backlog.tail)
802         sk->sk_backlog.head = skb;
803     else
804         sk->sk_backlog.tail->next = skb;
805
806     sk->sk_backlog.tail = skb;
807     skb->next = NULL;
808 }
809
810 /*
811  * Take into account size of receive queue and backlog queue
812  * Do not take into account this skb truesize,
813  * to allow even a single big packet to come.
814 */
815 static inline bool sk_rcvqueues_full(const struct sock *sk, unsigned int limit)
816 {
817     unsigned int qsize = sk->sk_backlog.len + atomic_read(&sk->sk_rmem_alloc);
818
819     return qsize > limit;
820 }
821
822 /* The per-socket spinlock must be held here. */
823 static inline __must_check int sk_add_backlog(struct sock *sk, struct sk_buff *skb,
824                                              unsigned int limit)
825 {
826     if (sk_rcvqueues_full(sk, limit))
827         return -ENOBUFS;
828
829     __sk_add_backlog(sk, skb);
830     sk->sk_backlog.len += skb->truesize;
831     return 0;
832 }
833
834 int __sk_backlog_rcv(struct sock *sk, struct sk_buff *skb);
835
836 static inline int sk_backlog_rcv(struct sock *sk, struct sk_buff *skb)
837 {
838     if (sk_memalloc_socks() && skb_pfmemalloc(skb))
839         return __sk_backlog_rcv(sk, skb);
840
841     return sk->sk_backlog_rcv(sk, skb);
842 }
843
844 static inline void sk_incoming_cpu_update(struct sock *sk)
845 {
846     sk->sk_incoming_cpu = raw_smp_processor_id();
847 }
848
849 static inline void sock_rps_record_flow_hash(_u32 hash)
850 {
851     #ifdef CONFIG_RPS
852         struct rps_sock_flow_table *sock_flow_table;
853
854         rcu_read_lock();
855         sock_flow_table = rcu_dereference(rps_sock_flow_table);
856         rps_record_sock_flow(sock_flow_table, hash);
857         rcu_read_unlock();
858     #endif
859 }
860
861 static inline void sock_rps_record_flow(const struct sock *sk)
862 {
863     #ifdef CONFIG_RPS

```

```

864     sock\_rps\_record\_flow\_hash(sk->sk_rxhash);
865 #endif
866 }
867
868 static inline void sock\_rps\_save\_rxhash(struct sock *sk,
869                                         const struct sk\_buff *skb)
870 {
871 #ifdef CONFIG_RPS
872     if (unlikely(sk->sk_rxhash != skb->hash))
873         sk->sk_rxhash = skb->hash;
874 #endif
875 }
876
877 static inline void sock\_rps\_reset\_rxhash(struct sock *sk)
878 {
879 #ifdef CONFIG_RPS
880     sk->sk_rxhash = 0;
881 #endif
882 }
883
884 #define sk\_wait\_event(__sk, __timeo, __condition) \
885     ({ \
886         int __rc; \
887         release\_sock(__sk); \
888         __rc = __condition; \
889         if (!__rc) { \
890             *(__timeo) = schedule\_timeout(*(__timeo)); \
891             sched\_annotate\_sleep(); \
892             lock\_sock(__sk); \
893             __rc = __condition; \
894         } \
895     })
896
897 int sk\_stream\_wait\_connect(struct sock *sk, long *timeo_p);
898 int sk\_stream\_wait\_memory(struct sock *sk, long *timeo_p);
899 void sk\_stream\_wait\_close(struct sock *sk, long timeo_p);
900 int sk\_stream\_error(struct sock *sk, int flags, int err);
901 void sk\_stream\_kill\_queues(struct sock *sk);
902 void sk\_set\_memalloc(struct sock *sk);
903 void sk\_clear\_memalloc(struct sock *sk);
904
905 int sk\_wait\_data(struct sock *sk, long *timeo, const struct sk\_buff *skb);
906
907 struct request\_sock\_ops;
908 struct timewait\_sock\_ops;
909 struct inet\_hashinfo;
910 struct raw\_hashinfo;
911 struct module;
912
913 /*
914  * caches using SLAB_DESTROY_BY_RCU should let .next pointer from nulls nodes
915  * un-modified. Special care is taken when initializing object to zero.
916  */
917 static inline void sk\_prot\_clear\_nulls(struct sock *sk, int size)
918 {
919     if (offsetof(struct sock, sk\_node.next) != 0)
920         memset(sk, 0, offsetof(struct sock, sk\_node.next));
921     memset(&sk->sk\_node.pprev, 0,
922           size - offsetof(struct sock, sk\_node.pprev));
923 }
924
925 /* Networking protocol blocks we attach to sockets.
926  * socket layer -> transport layer interface
927  */
928 struct proto {
929     void (*close)(struct sock *sk,
930                  long timeout);
931     int (*connect)(struct sock *sk,

```



```

932         struct sockaddr *uaddr,
933         int addr_len);
934     int (*disconnect)(struct sock *sk, int flags);
935
936     struct sock * (*accept)(struct sock *sk, int flags, int *err);
937
938     int (*ioctl)(struct sock *sk, int cmd,
939         unsigned long arg);
940     int (*init)(struct sock *sk);
941     void (*destroy)(struct sock *sk);
942     void (*shutdown)(struct sock *sk, int how);
943     int (*setsockopt)(struct sock *sk, int level,
944         int optname, char \_\_user *optval,
945         unsigned int optlen);
946     int (*getsockopt)(struct sock *sk, int level,
947         int optname, char \_\_user *optval,
948         int \_\_user *option);
949 #ifdef CONFIG_COMPAT
950     int (*compat_setsockopt)(struct sock *sk,
951         int level,
952         int optname, char \_\_user *optval,
953         unsigned int optlen);
954     int (*compat_getsockopt)(struct sock *sk,
955         int level,
956         int optname, char \_\_user *optval,
957         int \_\_user *option);
958     int (*compat_ioctl)(struct sock *sk,
959         unsigned int cmd, unsigned long arg);
960 #endif
961     int (*sendmsg)(struct sock *sk, struct msghdr *msg,
962         size\_t len);
963     int (*recvmsg)(struct sock *sk, struct msghdr *msg,
964         size\_t len, int noblock, int flags,
965         int *addr_len);
966     int (*sendpage)(struct sock *sk, struct page *page,
967         int offset, size\_t size, int flags);
968     int (*bind)(struct sock *sk,
969         struct sockaddr *uaddr, int addr_len);
970
971     int (*backlog\_rcv)(struct sock *sk,
972         struct sk\_buff *skb);
973
974     void (*release\_cb)(struct sock *sk);
975
976     /* Keeping track of sk's, looking them up, and port selection methods. */
977     void (*hash)(struct sock *sk);
978     void (*unhash)(struct sock *sk);
979     void (*rehash)(struct sock *sk);
980     int (*get\_port)(struct sock *sk, unsigned short snum);
981     void (*clear\_sk)(struct sock *sk, int size);
982
983     /* Keeping track of sockets in use */
984 #ifdef CONFIG_PROC_FS
985     unsigned int inuse_idx;
986 #endif
987
988     bool (*stream\_memory\_free)(const struct sock *sk);
989     /* Memory pressure */
990     void (*enter\_memory\_pressure)(struct sock *sk);
991     atomic\_long\_t *memory_allocated; /* Current allocated memory. */
992     struct percpu\_counter *sockets_allocated; /* Current number of sockets. */
993     /*
994      * Pressure flag: try to collapse.
995      * Technical note: it is used by multiple contexts non atomically.
996      * All the __sk_mem_schedule() is of this nature: accounting
997      * is strict, actions are advisory and have some latency.
998      */
999     int *memory_pressure;

```

```

1000 long                *sysctl_mem;
1001 int                 *sysctl_wmem;
1002 int                 *sysctl_rmem;
1003 int                 max_header;
1004 bool                no_autobind;
1005
1006 struct kmem\_cache    *slab;
1007 unsigned int         obj_size;
1008 int                  slab_flags;
1009
1010 struct percpu\_counter *orphan_count;
1011
1012 struct request\_sock\_ops *rsk_prot;
1013 struct timewait\_sock\_ops *twsk_prot;
1014
1015 union {
1016     struct inet\_hashinfo    *hashinfo;
1017     struct udp\_table        *udp_table;
1018     struct raw\_hashinfo     *raw_hash;
1019 } h;
1020
1021 struct module          *owner;
1022
1023 char                  name[32];
1024
1025 struct list\_head      node;
1026 #ifdef SOCK_REFCNT_DEBUG
1027     atomic\_t          socks;
1028 #endif
1029 #ifdef CONFIG_MEMCG_KMEM
1030     /*
1031      * cgroup specific init/deinit functions. Called once for all
1032      * protocols that implement it, from cgroups populate function.
1033      * This function has to setup any files the protocol want to
1034      * appear in the kmem cgroup filesystem.
1035      */
1036     int                (*init_cgroup)(struct mem\_cgroup *memcg,
1037                                       struct cgroup\_subsys *ss);
1038     void                (*destroy_cgroup)(struct mem\_cgroup *memcg);
1039     struct cg\_proto     (*proto_cgroup)(struct mem\_cgroup *memcg);
1040 #endif
1041 };
1042
1043 /*
1044  * Bits in struct cg_proto.flags
1045  */
1046 enum cg\_proto\_flags {
1047     /* Currently active and new sockets should be assigned to cgroups */
1048     MEMCG_SOCKET_ACTIVE,
1049     /* It was ever activated; we must disarm static keys on destruction */
1050     MEMCG_SOCKET_ACTIVATED,
1051 };
1052
1053 struct cg\_proto {
1054     struct page\_counter    memory_allocated;    /* Current allocated memory. */
1055     struct percpu\_counter sockets_allocated;    /* Current number of sockets. */
1056     int                memory_pressure;
1057     long                sysctl_mem[3];
1058     unsigned long       flags;
1059     /*
1060      * memcg field is used to find which memcg we belong directly
1061      * Each memcg struct can hold more than one cg_proto, so container_of
1062      * won't really cut.
1063      *
1064      * The elegant solution would be having an inverse function to
1065      * proto_cgroup in struct proto, but that means polluting the structure
1066      * for everybody, instead of just for memcg users.
1067      */

```

```

1068     struct mem\_cgroup      *memcg;
1069 };
1070
1071 int proto\_register(struct proto *prot, int alloc_slab);
1072 void proto\_unregister(struct proto *prot);
1073
1074 static inline bool memcg\_proto\_active(struct cg\_proto *cg_proto)
1075 {
1076     return test\_bit(MEMCG_SOCKET_ACTIVE, &cg_proto->flags);
1077 }
1078
1079 #ifdef SOCK_REFCNT_DEBUG
1080 static inline void sk\_refcnt\_debug\_inc(struct sock *sk)
1081 {
1082     atomic\_inc(&sk->sk_prot->socks);
1083 }
1084
1085 static inline void sk\_refcnt\_debug\_dec(struct sock *sk)
1086 {
1087     atomic\_dec(&sk->sk_prot->socks);
1088     printk(KERN_DEBUG "%s socket %p released, %d are still alive\n",
1089         sk->sk_prot->name, sk, atomic\_read(&sk->sk_prot->socks));
1090 }
1091
1092 static inline void sk\_refcnt\_debug\_release(const struct sock *sk)
1093 {
1094     if (atomic\_read(&sk->sk_refcnt) != 1)
1095         printk(KERN_DEBUG "Destruction of the %s socket %p delayed, refcnt=%d\n",
1096             sk->sk_prot->name, sk, atomic\_read(&sk->sk_refcnt));
1097 }
1098 #else /* SOCK_REFCNT_DEBUG */
1099 #define sk\_refcnt\_debug\_inc(sk) do { } while (0)
1100 #define sk\_refcnt\_debug\_dec(sk) do { } while (0)
1101 #define sk\_refcnt\_debug\_release(sk) do { } while (0)
1102 #endif /* SOCK_REFCNT_DEBUG */
1103
1104 #if defined(CONFIG_MEMCG_KMEM) && defined(CONFIG_NET)
1105 extern struct static\_key memcg\_socket\_limit\_enabled;
1106 static inline struct cg\_proto *parent\_cg\_proto(struct proto *proto,
1107     struct cg\_proto *cg_proto)
1108 {
1109     return proto->proto_cgroup(parent\_mem\_cgroup(cg_proto->memcg));
1110 }
1111 #define mem\_cgroup\_sockets\_enabled static\_key\_false(&memcg\_socket\_limit\_enabled)
1112 #else
1113 #define mem\_cgroup\_sockets\_enabled 0
1114 static inline struct cg\_proto *parent\_cg\_proto(struct proto *proto,
1115     struct cg\_proto *cg_proto)
1116 {
1117     return NULL;
1118 }
1119 #endif
1120
1121 static inline bool sk\_stream\_memory\_free(const struct sock *sk)
1122 {
1123     if (sk->sk_wmem_queued >= sk->sk_sndbuf)
1124         return false;
1125
1126     return sk->sk_prot->stream_memory_free ?
1127         sk->sk_prot->stream_memory_free(sk) : true;
1128 }
1129
1130 static inline bool sk\_stream\_is\_writeable(const struct sock *sk)
1131 {
1132     return sk\_stream\_wspace(sk) >= sk\_stream\_min\_wspace(sk) &&
1133         sk\_stream\_memory\_free(sk);
1134 }
1135

```

```

1136
1137 static inline bool sk\_has\_memory\_pressure(const struct sock *sk)
1138 {
1139     return sk->sk\_prot->memory_pressure != NULL;
1140 }
1141
1142 static inline bool sk\_under\_memory\_pressure(const struct sock *sk)
1143 {
1144     if (!sk->sk\_prot->memory_pressure)
1145         return false;
1146
1147     if (mem\_cgroup\_sockets\_enabled && sk->sk_cgrp)
1148         return !!sk->sk_cgrp->memory_pressure;
1149
1150     return !!*sk->sk\_prot->memory_pressure;
1151 }
1152
1153 static inline void sk\_leave\_memory\_pressure(struct sock *sk)
1154 {
1155     int *memory_pressure = sk->sk\_prot->memory_pressure;
1156
1157     if (!memory_pressure)
1158         return;
1159
1160     if (*memory_pressure)
1161         *memory_pressure = 0;
1162
1163     if (mem\_cgroup\_sockets\_enabled && sk->sk_cgrp) {
1164         struct cg\_proto *cg\_proto = sk->sk_cgrp;
1165         struct proto *prot = sk->sk\_prot;
1166
1167         for (; cg\_proto; cg\_proto = parent\_cg\_proto(prot, cg\_proto))
1168             cg\_proto->memory_pressure = 0;
1169     }
1170 }
1171
1172 static inline void sk\_enter\_memory\_pressure(struct sock *sk)
1173 {
1174     if (!sk->sk\_prot->enter_memory_pressure)
1175         return;
1176
1177     if (mem\_cgroup\_sockets\_enabled && sk->sk_cgrp) {
1178         struct cg\_proto *cg\_proto = sk->sk_cgrp;
1179         struct proto *prot = sk->sk\_prot;
1180
1181         for (; cg\_proto; cg\_proto = parent\_cg\_proto(prot, cg\_proto))
1182             cg\_proto->memory_pressure = 1;
1183     }
1184
1185     sk->sk\_prot->enter_memory_pressure(sk);
1186 }
1187
1188 static inline long sk\_prot\_mem\_limits(const struct sock *sk, int index)
1189 {
1190     long *prot = sk->sk\_prot->sysctl_mem;
1191     if (mem\_cgroup\_sockets\_enabled && sk->sk_cgrp)
1192         prot = sk->sk_cgrp->sysctl_mem;
1193     return prot[index];
1194 }
1195
1196 static inline void memcg\_memory\_allocated\_add(struct cg\_proto *prot,
1197                                             unsigned long amt,
1198                                             int *parent_status)
1199 {
1200     page\_counter\_charge(&prot->memory_allocated, amt);
1201
1202     if (page\_counter\_read(&prot->memory_allocated) >
1203
```

```

1204         prot->memory_allocated.limit)
1205         *parent_status = OVER_LIMIT;
1206     }
1207
1208     static inline void memcg_memory_allocated_sub(struct cg_proto *prot,
1209                                                  unsigned long amt)
1210     {
1211         page_counter_uncharge(&prot->memory_allocated, amt);
1212     }
1213
1214     static inline long
1215     sk_memory_allocated(const struct sock *sk)
1216     {
1217         struct proto *prot = sk->sk_prot;
1218
1219         if (mem_cgroup_sockets_enabled && sk->sk_cgrp)
1220             return page_counter_read(&sk->sk_cgrp->memory_allocated);
1221
1222         return atomic_long_read(prot->memory_allocated);
1223     }
1224
1225     static inline long
1226     sk_memory_allocated_add(struct sock *sk, int amt, int *parent_status)
1227     {
1228         struct proto *prot = sk->sk_prot;
1229
1230         if (mem_cgroup_sockets_enabled && sk->sk_cgrp) {
1231             memcg_memory_allocated_add(sk->sk_cgrp, amt, parent_status);
1232             /* update the root cgroup regardless */
1233             atomic_long_add_return(amt, prot->memory_allocated);
1234             return page_counter_read(&sk->sk_cgrp->memory_allocated);
1235         }
1236
1237         return atomic_long_add_return(amt, prot->memory_allocated);
1238     }
1239
1240     static inline void
1241     sk_memory_allocated_sub(struct sock *sk, int amt)
1242     {
1243         struct proto *prot = sk->sk_prot;
1244
1245         if (mem_cgroup_sockets_enabled && sk->sk_cgrp)
1246             memcg_memory_allocated_sub(sk->sk_cgrp, amt);
1247
1248         atomic_long_sub(amt, prot->memory_allocated);
1249     }
1250
1251     static inline void sk_sockets_allocated_dec(struct sock *sk)
1252     {
1253         struct proto *prot = sk->sk_prot;
1254
1255         if (mem_cgroup_sockets_enabled && sk->sk_cgrp) {
1256             struct cg_proto *cg_proto = sk->sk_cgrp;
1257
1258             for (; cg_proto; cg_proto = parent_cg_proto(prot, cg_proto))
1259                 percpu_counter_dec(&cg_proto->sockets_allocated);
1260         }
1261
1262         percpu_counter_dec(prot->sockets_allocated);
1263     }
1264
1265     static inline void sk_sockets_allocated_inc(struct sock *sk)
1266     {
1267         struct proto *prot = sk->sk_prot;
1268
1269         if (mem_cgroup_sockets_enabled && sk->sk_cgrp) {
1270             struct cg_proto *cg_proto = sk->sk_cgrp;
1271

```

```

1272         for (; cg\_proto; cg\_proto = parent\_cg\_proto(prot, cg\_proto))
1273             percpu\_counter\_inc(&cg\_proto->sockets_allocated);
1274     }
1275
1276     percpu\_counter\_inc(prot->sockets_allocated);
1277 }
1278
1279 static inline int
1280 sk\_sockets\_allocated\_read\_positive(struct sock *sk)
1281 {
1282     struct proto *prot = sk->sk\_prot;
1283
1284     if (mem\_cgroup\_sockets\_enabled && sk->sk\_cgrp)
1285         return percpu\_counter\_read\_positive(&sk->sk\_cgrp->sockets_allocated);
1286
1287     return percpu\_counter\_read\_positive(prot->sockets_allocated);
1288 }
1289
1290 static inline int
1291 proto\_sockets\_allocated\_sum\_positive(struct proto *prot)
1292 {
1293     return percpu\_counter\_sum\_positive(prot->sockets_allocated);
1294 }
1295
1296 static inline long
1297 proto\_memory\_allocated(struct proto *prot)
1298 {
1299     return atomic\_long\_read(prot->memory_allocated);
1300 }
1301
1302 static inline bool
1303 proto\_memory\_pressure(struct proto *prot)
1304 {
1305     if (!prot->memory_pressure)
1306         return false;
1307     return !!*prot->memory_pressure;
1308 }
1309
1310
1311 #ifdef CONFIG_PROC_FS
1312 /* Called with local bh disabled */
1313 void sock\_prot\_inuse\_add(struct net *net, struct proto *prot, int inc);
1314 int sock\_prot\_inuse\_get(struct net *net, struct proto *proto);
1315 #else
1316 static inline void sock\_prot\_inuse\_add(struct net *net, struct proto *prot,
1317                                         int inc)
1318 {
1319 }
1320 #endif
1321
1322
1323 /* With per-bucket locks this operation is not-atomic, so that
1324  * this version is not worse.
1325  */
1326 static inline void \_\_sk\_prot\_rehash(struct sock *sk)
1327 {
1328     sk->sk\_prot->unhash(sk);
1329     sk->sk\_prot->hash(sk);
1330 }
1331
1332 void sk\_prot\_clear\_portaddr\_nulls(struct sock *sk, int size);
1333
1334 /* About 10 seconds */
1335 #define SOCK\_DESTROY\_TIME (10*HZ)
1336
1337 /* Sockets 0-1023 can't be bound to unless you are superuser */
1338 #define PROT\_SOCK 1024
1339

```



```

1340 #define SHUTDOWN_MASK 3
1341 #define RCV_SHUTDOWN 1
1342 #define SEND_SHUTDOWN 2
1343
1344 #define SOCK_SNDBUF_LOCK 1
1345 #define SOCK_RCVBUF_LOCK 2
1346 #define SOCK_BINDADDR_LOCK 4
1347 #define SOCK_BINDPORT_LOCK 8
1348
1349 struct socket_alloc {
1350     struct socket socket;
1351     struct inode vfs_inode;
1352 };
1353
1354 static inline struct socket *SOCKET_I(struct inode *inode)
1355 {
1356     return &container_of(inode, struct socket_alloc, vfs_inode)->socket;
1357 }
1358
1359 static inline struct inode *SOCK_INODE(struct socket *socket)
1360 {
1361     return &container_of(socket, struct socket_alloc, socket)->vfs_inode;
1362 }
1363
1364 /*
1365  * Functions for memory accounting
1366  */
1367 int __sk_mem_schedule(struct sock *sk, int size, int kind);
1368 void __sk_mem_reclaim(struct sock *sk, int amount);
1369
1370 #define SK_MEM_QUANTUM ((int)PAGE_SIZE)
1371 #define SK_MEM_QUANTUM_SHIFT ilog2(SK_MEM_QUANTUM)
1372 #define SK_MEM_SEND 0
1373 #define SK_MEM_RECV 1
1374
1375 static inline int sk_mem_pages(int amt)
1376 {
1377     return (amt + SK_MEM_QUANTUM - 1) >> SK_MEM_QUANTUM_SHIFT;
1378 }
1379
1380 static inline bool sk_has_account(struct sock *sk)
1381 {
1382     /* return true if protocol supports memory accounting */
1383     return !!sk->sk_prot->memory_allocated;
1384 }
1385
1386 static inline bool sk_wmem_schedule(struct sock *sk, int size)
1387 {
1388     if (!sk_has_account(sk))
1389         return true;
1390     return size <= sk->sk_forward_alloc ||
1391         __sk_mem_schedule(sk, size, SK_MEM_SEND);
1392 }
1393
1394 static inline bool
1395 sk_rmem_schedule(struct sock *sk, struct sk_buff *skb, int size)
1396 {
1397     if (!sk_has_account(sk))
1398         return true;
1399     return size <= sk->sk_forward_alloc ||
1400         __sk_mem_schedule(sk, size, SK_MEM_RECV) ||
1401         skb_pfmalloc(skb);
1402 }
1403
1404 static inline void sk_mem_reclaim(struct sock *sk)
1405 {
1406     if (!sk_has_account(sk))
1407         return;

```

```

1408         if (sk->sk_forward_alloc >= SK_MEM_QUANTUM)
1409             __sk_mem_reclaim(sk, sk->sk_forward_alloc);
1410     }
1411
1412     static inline void sk_mem_reclaim_partial(struct sock *sk)
1413     {
1414         if (!sk_has_account(sk))
1415             return;
1416         if (sk->sk_forward_alloc > SK_MEM_QUANTUM)
1417             __sk_mem_reclaim(sk, sk->sk_forward_alloc - 1);
1418     }
1419
1420     static inline void sk_mem_charge(struct sock *sk, int size)
1421     {
1422         if (!sk_has_account(sk))
1423             return;
1424         sk->sk_forward_alloc += size;
1425     }
1426
1427     static inline void sk_mem_uncharge(struct sock *sk, int size)
1428     {
1429         if (!sk_has_account(sk))
1430             return;
1431         sk->sk_forward_alloc -= size;
1432     }
1433
1434     static inline void sk_wmem_free_skb(struct sock *sk, struct sk_buff *skb)
1435     {
1436         sock_set_flag(sk, SOCK_QUEUE_SHRUNK);
1437         sk->sk_wmem_queued -= skb->truesize;
1438         sk_mem_uncharge(sk, skb->truesize);
1439         kfree_skb(skb);
1440     }
1441
1442     /* Used by processes to "lock" a socket state, so that
1443      * interrupts and bottom half handlers won't change it
1444      * from under us. It essentially blocks any incoming
1445      * packets, so that we won't get any new data or any
1446      * packets that change the state of the socket.
1447      *
1448      * While locked, BH processing will add new packets to
1449      * the backlog queue. This queue is processed by the
1450      * owner of the socket lock right before it is released.
1451      *
1452      * Since ~2.3.5 it is also exclusive sleep lock serializing
1453      * accesses from user process context.
1454      */
1455     #define sock_owned_by_user(sk) ((sk)->sk_lock.owned)
1456
1457     static inline void sock_release_ownership(struct sock *sk)
1458     {
1459         sk->sk_lock.owned = 0;
1460     }
1461
1462     /*
1463      * Macro so as to not evaluate some arguments when
1464      * Lockdep is not enabled.
1465      *
1466      * Mark both the sk_lock and the sk_lock.slock as a
1467      * per-address-family lock class.
1468      */
1469     #define sock_lock_init_class_and_name(sk, sname, skey, name, key) \
1470     do { \
1471         sk->sk_lock.owned = 0; \
1472         init_waitqueue_head(&sk->sk_lock.wq); \
1473         spin_lock_init(&(sk)->sk_lock.slock); \
1474         debug_check_no_locks_freed((void *)&(sk)->sk_lock, \
1475                                     sizeof((sk)->sk_lock)); \

```

```

1476     lockdep_set_class_and_name(&(sk)->sk_lock.slock,
1477                               (skey), (sname));
1478     lockdep_init_map(&(sk)->sk_lock.dep_map, (name), (key), 0);
1479 } while (0)
1480
1481 void lock_sock_nested(struct sock *sk, int subclass);
1482
1483 static inline void lock_sock(struct sock *sk)
1484 {
1485     lock_sock_nested(sk, 0);
1486 }
1487
1488 void release_sock(struct sock *sk);
1489
1490 /* BH context may only use the following locking interface. */
1491 #define bh_lock_sock(__sk) spin_lock(&((__sk)->sk_lock.slock))
1492 #define bh_lock_sock_nested(__sk) \
1493     spin_lock_nested(&((__sk)->sk_lock.slock), \
1494                     SINGLE_DEPTH_NESTING)
1495 #define bh_unlock_sock(__sk) spin_unlock(&((__sk)->sk_lock.slock))
1496
1497 bool lock_sock_fast(struct sock *sk);
1498 /**
1499  * unlock_sock_fast - complement of lock_sock_fast
1500  * @sk: socket
1501  * @slow: slow mode
1502  *
1503  * fast unlock socket for user context.
1504  * If slow mode is on, we call regular release_sock()
1505  */
1506 static inline void unlock_sock_fast(struct sock *sk, bool slow)
1507 {
1508     if (slow)
1509         release_sock(sk);
1510     else
1511         spin_unlock_bh(&sk->sk_lock.slock);
1512 }
1513
1514
1515 struct sock *sk_alloc(struct net *net, int family, gfp_t priority,
1516                      struct proto *prot, int kern);
1517 void sk_free(struct sock *sk);
1518 void sk_destruct(struct sock *sk);
1519 struct sock *sk_clone_lock(const struct sock *sk, const gfp_t priority);
1520
1521 struct sk_buff *sock_wmalloc(struct sock *sk, unsigned long size, int force,
1522                             gfp_t priority);
1523 void sock_wfree(struct sk_buff *skb);
1524 void skb_orphan_partial(struct sk_buff *skb);
1525 void sock_rfree(struct sk_buff *skb);
1526 void sock_efree(struct sk_buff *skb);
1527 #ifdef CONFIG_INET
1528 void sock_edemux(struct sk_buff *skb);
1529 #else
1530 #define sock_edemux(skb) sock_efree(skb)
1531 #endif
1532
1533 int sock_setsockopt(struct socket *sock, int level, int op,
1534                    char __user *optval, unsigned int optlen);
1535
1536 int sock_getsockopt(struct socket *sock, int level, int op,
1537                    char __user *optval, int __user *optlen);
1538 struct sk_buff *sock_alloc_send_skb(struct sock *sk, unsigned long size,
1539                                     int noblock, int *errcode);
1540 struct sk_buff *sock_alloc_send_skb(struct sock *sk, unsigned long header_len,
1541                                     unsigned long data_len, int noblock,
1542                                     int *errcode, int max_page_order);
1543 void *sock_kmalloc(struct sock *sk, int size, gfp_t priority);

```

```

1544 void sock\_kfree\_s(struct sock *sk, void *mem, int size);
1545 void sock\_kzfree\_s(struct sock *sk, void *mem, int size);
1546 void sk\_send\_sigurg(struct sock *sk);
1547
1548 /*
1549  * Functions to fill in entries in struct proto_ops when a protocol
1550  * does not implement a particular function.
1551  */
1552 int sock\_no\_bind(struct socket *, struct sockaddr *, int);
1553 int sock\_no\_connect(struct socket *, struct sockaddr *, int, int);
1554 int sock\_no\_socketpair(struct socket *, struct socket *);
1555 int sock\_no\_accept(struct socket *, struct socket *, int);
1556 int sock\_no\_getname(struct socket *, struct sockaddr *, int *, int);
1557 unsigned int sock\_no\_poll(struct file *, struct socket *,
1558                          struct poll\_table\_struct *);
1559 int sock\_no\_ioctl(struct socket *, unsigned int, unsigned long);
1560 int sock\_no\_listen(struct socket *, int);
1561 int sock\_no\_shutdown(struct socket *, int);
1562 int sock\_no\_getsockopt(struct socket *, int, int, char __user *, int __user *);
1563 int sock\_no\_setsockopt(struct socket *, int, int, char __user *, unsigned int);
1564 int sock\_no\_sendmsg(struct socket *, struct msghdr *, size\_t);
1565 int sock\_no\_recvmsg(struct socket *, struct msghdr *, size\_t, int);
1566 int sock\_no\_mmap(struct file *file, struct socket *sock,
1567                 struct vm\_area\_struct *vma);
1568 ssize\_t sock\_no\_sendpage(struct socket *sock, struct page *page, int offset,
1569                          size\_t size, int flags);
1570
1571 /*
1572  * Functions to fill in entries in struct proto_ops when a protocol
1573  * uses the inet style.
1574  */
1575 int sock\_common\_getsockopt(struct socket *sock, int level, int optname,
1576                           char __user *optval, int __user *optlen);
1577 int sock\_common\_recvmsg(struct socket *sock, struct msghdr *msg, size\_t size,
1578                          int flags);
1579 int sock\_common\_setsockopt(struct socket *sock, int level, int optname,
1580                           char __user *optval, unsigned int optlen);
1581 int compat\_sock\_common\_getsockopt(struct socket *sock, int level,
1582                                  int optname, char __user *optval, int __user *optlen);
1583 int compat\_sock\_common\_setsockopt(struct socket *sock, int level,
1584                                  int optname, char __user *optval, unsigned int optlen);
1585
1586 void sk\_common\_release(struct sock *sk);
1587
1588 /*
1589  *      Default socket callbacks and setup code
1590  */
1591
1592 /* Initialise core socket variables */
1593 void sock\_init\_data(struct socket *sock, struct sock *sk);
1594
1595 /*
1596  * Socket reference counting postulates.
1597  *
1598  * * Each user of socket SHOULD hold a reference count.
1599  * * Each access point to socket (an hash table bucket, reference from a list,
1600  *   running timer, skb in flight MUST hold a reference count.
1601  * * When reference count hits 0, it means it will never increase back.
1602  * * When reference count hits 0, it means that no references from
1603  *   outside exist to this socket and current process on current CPU
1604  *   is last user and may/should destroy this socket.
1605  * * sk_free is called from any context: process, BH, IRQ. When
1606  *   it is called, socket has no references from outside -> sk_free
1607  *   may release descendant resources allocated by the socket, but
1608  *   to the time when it is called, socket is NOT referenced by any
1609  *   hash tables, lists etc.
1610  * * Packets, delivered from outside (from network or from another process)
1611  *   and enqueued on receive/error queues SHOULD NOT grab reference count,

```

```

1612 *   when they sit in queue. Otherwise, packets will leak to hole, when
1613 *   socket is looked up by one cpu and unhasing is made by another CPU.
1614 *   It is true for udp/raw, netlink (leak to receive and error queues), tcp
1615 *   (leak to backlog). Packet socket does all the processing inside
1616 *   BR_NETPROTO_LOCK, so that it has not this race condition. UNIX sockets
1617 *   use separate SMP lock, so that they are prone too.
1618 */
1619
1620 /* Ungrab socket and destroy it, if it was the last reference. */
1621 static inline void sock_put(struct sock *sk)
1622 {
1623     if (atomic_dec_and_test(&sk->sk_refcnt))
1624         sk_free(sk);
1625 }
1626 /* Generic version of sock_put(), dealing with all sockets
1627  * (TCP_TIMEWAIT, TCP_NEW_SYN_RECV, ESTABLISHED...)
1628 */
1629 void sock_gen_put(struct sock *sk);
1630
1631 int sk_receive_skb(struct sock *sk, struct sk_buff *skb, const int nested);
1632
1633 static inline void sk_tx_queue_set(struct sock *sk, int tx_queue)
1634 {
1635     sk->sk_tx_queue_mapping = tx_queue;
1636 }
1637
1638 static inline void sk_tx_queue_clear(struct sock *sk)
1639 {
1640     sk->sk_tx_queue_mapping = -1;
1641 }
1642
1643 static inline int sk_tx_queue_get(const struct sock *sk)
1644 {
1645     return sk ? sk->sk_tx_queue_mapping : -1;
1646 }
1647
1648 static inline void sk_set_socket(struct sock *sk, struct socket *sock)
1649 {
1650     sk_tx_queue_clear(sk);
1651     sk->sk_socket = sock;
1652 }
1653
1654 static inline wait_queue_head_t *sk_sleep(struct sock *sk)
1655 {
1656     BUILD_BUG_ON(sizeof(struct socket_wq, wait) != 0);
1657     return &rcu_dereference_raw(sk->sk_wq->wait);
1658 }
1659 /* Detach socket from process context.
1660  * Announce socket dead, detach it from wait queue and inode.
1661  * Note that parent inode held reference count on this struct sock,
1662  * we do not release it in this function, because protocol
1663  * probably wants some additional cleanups or even continuing
1664  * to work with this socket (TCP).
1665 */
1666 static inline void sock_orphan(struct sock *sk)
1667 {
1668     write_lock_bh(&sk->sk_callback_lock);
1669     sock_set_flag(sk, SOCK_DEAD);
1670     sk_set_socket(sk, NULL);
1671     sk->sk_wq = NULL;
1672     write_unlock_bh(&sk->sk_callback_lock);
1673 }
1674
1675 static inline void sock_graft(struct sock *sk, struct socket *parent)
1676 {
1677     write_lock_bh(&sk->sk_callback_lock);
1678     sk->sk_wq = parent->wq;
1679     parent->sk = sk;

```

```

1680     sk\_set\_socket(sk, parent);
1681     security\_sock\_graft(sk, parent);
1682     write\_unlock\_bh(&sk->sk_callback_lock);
1683 }
1684
1685 kuid\_t sock\_i\_uid(struct sock *sk);
1686 unsigned long sock\_i\_ino(struct sock *sk);
1687
1688 static inline struct dst\_entry *
1689 \_\_sk\_dst\_get(struct sock *sk)
1690 {
1691     return rcu\_dereference\_check(sk->sk_dst_cache, sock\_owned\_by\_user(sk) ||
1692                                lockdep\_is\_held(&sk->sk_lock.slock));
1693 }
1694
1695 static inline struct dst\_entry *
1696 sk\_dst\_get(struct sock *sk)
1697 {
1698     struct dst\_entry *dst;
1699
1700     rcu\_read\_lock();
1701     dst = rcu\_dereference(sk->sk_dst_cache);
1702     if (dst && !atomic\_inc\_not\_zero(&dst->__refcnt))
1703         dst = NULL;
1704     rcu\_read\_unlock();
1705     return dst;
1706 }
1707
1708 static inline void dst\_negative\_advice(struct sock *sk)
1709 {
1710     struct dst\_entry *ndst, *dst = \_\_sk\_dst\_get(sk);
1711
1712     if (dst && dst->ops->negative_advice) {
1713         ndst = dst->ops->negative_advice(dst);
1714
1715         if (ndst != dst) {
1716             rcu\_assign\_pointer(sk->sk_dst_cache, ndst);
1717             sk\_tx\_queue\_clear(sk);
1718         }
1719     }
1720 }
1721
1722 static inline void
1723 \_\_sk\_dst\_set(struct sock *sk, struct dst\_entry *dst)
1724 {
1725     struct dst\_entry *old_dst;
1726
1727     sk\_tx\_queue\_clear(sk);
1728     /*
1729      * This can be called while sk is owned by the caller only,
1730      * with no state that can be checked in a rcu_dereference_check() cond
1731      */
1732     old_dst = rcu\_dereference\_raw(sk->sk_dst_cache);
1733     rcu\_assign\_pointer(sk->sk_dst_cache, dst);
1734     dst\_release(old_dst);
1735 }
1736
1737 static inline void
1738 sk\_dst\_set(struct sock *sk, struct dst\_entry *dst)
1739 {
1740     struct dst\_entry *old_dst;
1741
1742     sk\_tx\_queue\_clear(sk);
1743     old_dst = xchg((\_\_force struct dst\_entry **)&sk->sk_dst_cache, dst);
1744     dst\_release(old_dst);
1745 }
1746
1747 static inline void

```



```

1748 \_\_sk\_dst\_reset(struct sock *sk)
1749 {
1750     \_\_sk\_dst\_set(sk, NULL);
1751 }
1752
1753 static inline void
1754 sk\_dst\_reset(struct sock *sk)
1755 {
1756     sk\_dst\_set(sk, NULL);
1757 }
1758
1759 struct dst\_entry * \_\_sk\_dst\_check(struct sock *sk, u32 cookie);
1760
1761 struct dst\_entry * sk\_dst\_check(struct sock *sk, u32 cookie);
1762
1763 bool sk\_mc\_loop(struct sock *sk);
1764
1765 static inline bool sk\_can\_gso(const struct sock *sk)
1766 {
1767     return net\_gso\_ok(sk->sk_route_caps, sk->sk_gso_type);
1768 }
1769
1770 void sk\_setup\_caps(struct sock *sk, struct dst\_entry *dst);
1771
1772 static inline void sk\_nocaps\_add(struct sock *sk, netdev\_features\_t flags)
1773 {
1774     sk->sk_route_nocaps |= flags;
1775     sk->sk_route_caps &= ~flags;
1776 }
1777
1778 static inline int skb\_do\_copy\_data\_nocache(struct sock *sk, struct sk\_buff *skb,
1779                                           struct iov\_iter *from, char *to,
1780                                           int copy, int offset)
1781 {
1782     if (skb->ip_summed == CHECKSUM\_NONE) {
1783         wsum\_csum = 0;
1784         if (csum\_and\_copy\_from\_iter(to, copy, &csum, from) != copy)
1785             return -EFAULT;
1786         skb->csum = csum\_block\_add(skb->csum, csum, offset);
1787     } else if (sk->sk_route_caps & NETIF\_F\_NOCACHE\_COPY) {
1788         if (copy\_from\_iter\_nocache(to, copy, from) != copy)
1789             return -EFAULT;
1790     } else if (copy\_from\_iter(to, copy, from) != copy)
1791         return -EFAULT;
1792
1793     return 0;
1794 }
1795
1796 static inline int skb\_add\_data\_nocache(struct sock *sk, struct sk\_buff *skb,
1797                                       struct iov\_iter *from, int copy)
1798 {
1799     int err, offset = skb->len;
1800
1801     err = skb\_do\_copy\_data\_nocache(sk, skb, from, skb\_put(skb, copy),
1802                                   copy, offset);
1803     if (err)
1804         \_\_skb\_trim(skb, offset);
1805
1806     return err;
1807 }
1808
1809 static inline int skb\_copy\_to\_page\_nocache(struct sock *sk, struct iov\_iter *from,
1810                                           struct sk\_buff *skb,
1811                                           struct page *page,
1812                                           int off, int copy)
1813 {
1814     int err;
1815

```

```

1816     err = skb\_do\_copy\_data\_nocache(sk, skb, from, page\_address(page) + off,
1817                                   copy, skb->len);
1818     if (err)
1819         return err;
1820
1821     skb->len           += copy;
1822     skb->data\_len      += copy;
1823     skb->truesize      += copy;
1824     sk->sk_wmem_queued += copy;
1825     sk\_mem\_charge(sk, copy);
1826     return 0;
1827 }
1828
1829 /**
1830  * sk\_wmem\_alloc\_get - returns write allocations
1831  * @sk: socket
1832  *
1833  * Returns sk\_wmem\_alloc minus initial offset of one
1834  */
1835 static inline int sk\_wmem\_alloc\_get(const struct sock *sk)
1836 {
1837     return atomic\_read(&sk->sk_wmem_alloc) - 1;
1838 }
1839
1840 /**
1841  * sk\_rmem\_alloc\_get - returns read allocations
1842  * @sk: socket
1843  *
1844  * Returns sk\_rmem\_alloc
1845  */
1846 static inline int sk\_rmem\_alloc\_get(const struct sock *sk)
1847 {
1848     return atomic\_read(&sk->sk\_rmem\_alloc);
1849 }
1850
1851 /**
1852  * sk\_has\_allocations - check if allocations are outstanding
1853  * @sk: socket
1854  *
1855  * Returns true if socket has write or read allocations
1856  */
1857 static inline bool sk\_has\_allocations(const struct sock *sk)
1858 {
1859     return sk\_wmem\_alloc\_get(sk) || sk\_rmem\_alloc\_get(sk);
1860 }
1861
1862 /**
1863  * wq\_has\_sleeper - check if there are any waiting processes
1864  * @wq: struct socket_wq
1865  *
1866  * Returns true if socket_wq has waiting processes
1867  *
1868  * The purpose of the wq\_has\_sleeper and sock\_poll\_wait is to wrap the memory
1869  * barrier call. They were added due to the race found within the tcp code.
1870  *
1871  * Consider following tcp code paths:
1872  *
1873  * CPU1                      CPU2
1874  *
1875  * sys_select                 receive packet
1876  * ...                       ...
1877  * __add_wait_queue          update tp->rcv_nxt
1878  * ...                       ...
1879  * tp->rcv_nxt check         sock_def_readable
1880  * ...                       {
1881  * schedule                  rcu_read_lock();
1882  *                           wq = rcu_dereference(sk->sk_wq);
1883  *                           if (wq && waitqueue_active(&wq->wait))

```

```

1884 *                                wake_up_interruptible(&wq->wait)
1885 *                                ...
1886 *                                }
1887 *
1888 * The race for tcp fires when the __add_wait_queue changes done by CPU1 stay
1889 * in its cache, and so does the tp->rcv_nxt update on CPU2 side. The CPU1
1890 * could then endup calling schedule and sleep forever if there are no more
1891 * data on the socket.
1892 *
1893 */
1894 static inline bool wq_has_sleeper(struct socket_wq *wq)
1895 {
1896     /* We need to be sure we are in sync with the
1897      * add_wait_queue modifications to the wait queue.
1898      *
1899      * This memory barrier is paired in the sock_poll_wait.
1900      */
1901     smp_mb();
1902     return wq && waitqueue_active(&wq->wait);
1903 }
1904
1905 /**
1906  * sock_poll_wait - place memory barrier behind the poll_wait call.
1907  * @filp:         file
1908  * @wait_address:  socket wait queue
1909  * @p:            poll_table
1910  *
1911  * See the comments in the wq_has_sleeper function.
1912  */
1913 static inline void sock_poll_wait(struct file *filp,
1914                                   wait_queue_head_t *wait_address, poll_table *p)
1915 {
1916     if (!poll_does_not_wait(p) && wait_address) {
1917         poll_wait(filp, wait_address, p);
1918         /* We need to be sure we are in sync with the
1919          * socket flags modification.
1920          *
1921          * This memory barrier is paired in the wq_has_sleeper.
1922          */
1923         smp_mb();
1924     }
1925 }
1926
1927 static inline void skb_set_hash_from_sk(struct sk_buff *skb, struct sock *sk)
1928 {
1929     if (sk->sk_txhash) {
1930         skb->l4_hash = 1;
1931         skb->hash = sk->sk_txhash;
1932     }
1933 }
1934
1935 /**
1936  * Queue a received datagram if it will fit. Stream and sequenced
1937  * protocols can't normally use this as they need to fit buffers in
1938  * and play with them.
1939  *
1940  * Inlined as it's very short and called for pretty much every
1941  * packet ever received.
1942  */
1943
1944 static inline void skb_set_owner_w(struct sk_buff *skb, struct sock *sk)
1945 {
1946     skb_orphan(skb);
1947     skb->sk = sk;
1948     skb->destructor = sock_wfree;
1949     skb_set_hash_from_sk(skb, sk);
1950     /*
1951      * We used to take a refcount on sk, but following operation

```

```

1952     * is enough to guarantee sk_free() wont free this sock until
1953     * all in-flight packets are completed
1954     */
1955     atomic_add(skb->truesize, &sk->sk_wmem_alloc);
1956 }
1957
1958 static inline void skb_set_owner_r(struct sk_buff *skb, struct sock *sk)
1959 {
1960     skb_orphan(skb);
1961     skb->sk = sk;
1962     skb->destructor = sock_rfree;
1963     atomic_add(skb->truesize, &sk->sk_rmem_alloc);
1964     sk_mem_charge(sk, skb->truesize);
1965 }
1966
1967 void sk_reset_timer(struct sock *sk, struct timer_list *timer,
1968                    unsigned long expires);
1969
1970 void sk_stop_timer(struct sock *sk, struct timer_list *timer);
1971
1972 int sock_queue_rcv_skb(struct sock *sk, struct sk_buff *skb);
1973
1974 int sock_queue_err_skb(struct sock *sk, struct sk_buff *skb);
1975 struct sk_buff *sock_dequeue_err_skb(struct sock *sk);
1976
1977 /*
1978  *      Recover an error report and clear atomically
1979  */
1980
1981 static inline int sock_error(struct sock *sk)
1982 {
1983     int err;
1984     if (likely(!sk->sk_err))
1985         return 0;
1986     err = xchg(&sk->sk_err, 0);
1987     return -err;
1988 }
1989
1990 static inline unsigned long sock_wspace(struct sock *sk)
1991 {
1992     int amt = 0;
1993
1994     if (!(sk->sk_shutdown & SEND_SHUTDOWN)) {
1995         amt = sk->sk_sndbuf - atomic_read(&sk->sk_wmem_alloc);
1996         if (amt < 0)
1997             amt = 0;
1998     }
1999     return amt;
2000 }
2001
2002 static inline void sk_wake_async(struct sock *sk, int how, int band)
2003 {
2004     if (sock_flag(sk, SOCK_FASYNC))
2005         sock_wake_async(sk->sk_socket, how, band);
2006 }
2007
2008 /* Since sk_{r,w}mem_alloc sums skb->truesize, even a small frame might
2009  * need sizeof(sk_buff) + MTU + padding, unless net driver perform copybreak.
2010  * Note: for send buffers, TCP works better if we can build two skbs at
2011  * minimum.
2012  */
2013 #define TCP_SKB_MIN_TRUESIZE    (2048 + SKB_DATA_ALIGN(sizeof(struct sk_buff)))
2014
2015 #define SOCK_MIN_SNDBUF        (TCP_SKB_MIN_TRUESIZE * 2)
2016 #define SOCK_MIN_RCVBUF        TCP_SKB_MIN_TRUESIZE
2017
2018 static inline void sk_stream_moderate_sndbuf(struct sock *sk)
2019 {

```

```

2020     if (!(sk->sk_userlocks & SOCK\_SNDBUF\_LOCK)) {
2021         sk->sk_sndbuf = min(sk->sk_sndbuf, sk->sk_wmem_queued >> 1);
2022         sk->sk_sndbuf = max\_t(u32, sk->sk_sndbuf, SOCK\_MIN\_SNDBUF);
2023     }
2024 }
2025
2026 struct sk\_buff *sk\_stream\_alloc\_skb(struct sock *sk, int size, gfp\_t gfp,
2027                                     bool force_schedule);
2028
2029 /**
2030  * sk\_page\_frag - return an appropriate page_frag
2031  * @sk: socket
2032  *
2033  * If socket allocation mode allows current thread to sleep, it means its
2034  * safe to use the per task page_frag instead of the per socket one.
2035  */
2036 static inline struct page\_frag *sk\_page\_frag(struct sock *sk)
2037 {
2038     if (sk->sk_allocation & GFP\_WAIT)
2039         return &current->task_frag;
2040
2041     return &sk->sk_frag;
2042 }
2043
2044 bool sk\_page\_frag\_refill(struct sock *sk, struct page\_frag *pfrag);
2045
2046 /**
2047  *      Default write policy as shown to user space via poll/select/SIGIO
2048  */
2049 static inline bool sock\_writeable(const struct sock *sk)
2050 {
2051     return atomic\_read(&sk->sk_wmem_alloc) < (sk->sk_sndbuf >> 1);
2052 }
2053
2054 static inline gfp\_t gfp\_any(void)
2055 {
2056     return in\_softirq() ? GFP\_ATOMIC : GFP\_KERNEL;
2057 }
2058
2059 static inline long sock\_rcvtimeo(const struct sock *sk, bool noblock)
2060 {
2061     return noblock ? 0 : sk->sk_rcvtimeo;
2062 }
2063
2064 static inline long sock\_sndtimeo(const struct sock *sk, bool noblock)
2065 {
2066     return noblock ? 0 : sk->sk_sndtimeo;
2067 }
2068
2069 static inline int sock\_rcvlowat(const struct sock *sk, int waitall, int len)
2070 {
2071     return (waitall ? len : min\_t(int, sk->sk_rcvlowat, len)) ? : 1;
2072 }
2073
2074 /* Alas, with timeout socket operations are not restartable.
2075  * Compare this to poll().
2076  */
2077 static inline int sock\_intr\_errno(long timeo)
2078 {
2079     return timeo == MAX\_SCHEDULE\_TIMEOUT ? -ERESTARTSYS : -EINTR;
2080 }
2081
2082 struct sock\_skb\_cb {
2083     u32 dropcount;
2084 };
2085
2086 /* Store sock_skb_cb at the end of skb->cb[] so protocol families
2087  * using skb->cb[] would keep using it directly and utilize its

```

```

2088  * alignement guarantee.
2089  */
2090 #define SOCK_SKB_CB_OFFSET ((FIELD_SIZEOF(struct sk_buff, cb) - \
2091                             sizeof(struct sock_skb_cb)))
2092
2093 #define SOCK_SKB_CB(__skb) ((struct sock_skb_cb *)((__skb)->cb + \
2094                                                     SOCK_SKB_CB_OFFSET))
2095
2096 #define sock_skb_cb_check_size(size) \
2097     BUILD_BUG_ON((size) > SOCK_SKB_CB_OFFSET)
2098
2099 static inline void
2100 sock_skb_set_dropcount(const struct sock *sk, struct sk_buff *skb)
2101 {
2102     SOCK_SKB_CB(skb)->dropcount = atomic_read(&sk->sk_drops);
2103 }
2104
2105 void __sock_rcv_timestamp(struct msghdr *msg, struct sock *sk,
2106                          struct sk_buff *skb);
2107 void __sock_rcv_wifi_status(struct msghdr *msg, struct sock *sk,
2108                             struct sk_buff *skb);
2109
2110 static inline void
2111 sock_rcv_timestamp(struct msghdr *msg, struct sock *sk, struct sk_buff *skb)
2112 {
2113     ktime_t kt = skb->tstamp;
2114     struct sock_shared_hwtstamps *hwtstamps = skb_hwtstamps(skb);
2115
2116     /*
2117      * generate control messages if
2118      * - receive time stamping in software requested
2119      * - software time stamp available and wanted
2120      * - hardware time stamps available and wanted
2121      */
2122     if (sock_flag(sk, SOCK_RCVTSTAMP) ||
2123         (sk->sk_tsflags & SOF_TIMESTAMPING_RX_SOFTWARE) ||
2124         (kt.tv64 && sk->sk_tsflags & SOF_TIMESTAMPING_SOFTWARE) ||
2125         (hwtstamps->hwtstamp.tv64 &&
2126          (sk->sk_tsflags & SOF_TIMESTAMPING_RAW_HARDWARE)))
2127         sock_rcv_timestamp(msg, sk, skb);
2128     else
2129         sk->sk_stamp = kt;
2130
2131     if (sock_flag(sk, SOCK_WIFI_STATUS) && skb->wifi_acked_valid)
2132         sock_rcv_wifi_status(msg, sk, skb);
2133 }
2134
2135 void __sock_rcv_ts_and_drops(struct msghdr *msg, struct sock *sk,
2136                             struct sk_buff *skb);
2137
2138 static inline void sock_rcv_ts_and_drops(struct msghdr *msg, struct sock *sk,
2139                                          struct sk_buff *skb)
2140 {
2141     #define FLAGS_TS_OR_DROPS ((1UL << SOCK_RXQ_OVFL) | \
2142                               (1UL << SOCK_RCVTSTAMP))
2143     #define TSFLAGS_ANY (SOF_TIMESTAMPING_SOFTWARE | \
2144                         SOF_TIMESTAMPING_RAW_HARDWARE)
2145
2146     if (sk->sk_flags & FLAGS_TS_OR_DROPS || sk->sk_tsflags & TSFLAGS_ANY)
2147         sock_rcv_ts_and_drops(msg, sk, skb);
2148     else
2149         sk->sk_stamp = skb->tstamp;
2150 }
2151
2152 void __sock_tx_timestamp(const struct sock *sk, __u8 *tx_flags);
2153
2154 /**
2155  * sock_tx_timestamp - checks whether the outgoing packet is to be time stamped

```



```

2156 * @sk:          socket sending this packet
2157 * @tx_flags:    completed with instructions for time stamping
2158 *
2159 * Note : callers should take care of initial *tx_flags value (usually 0)
2160 */
2161 static inline void sock_tx_timestamp(const struct sock *sk, __u8 *tx_flags)
2162 {
2163     if (unlikely(sk->sk_tsflags))
2164         sock_tx_timestamp(sk, tx_flags);
2165     if (unlikely(sock_flag(sk, SOCK_WIFI_STATUS)))
2166         *tx_flags |= SKBTX_WIFI_STATUS;
2167 }
2168
2169 /**
2170 * sk_eat_skb - Release a skb if it is no longer needed
2171 * @sk: socket to eat this skb from
2172 * @skb: socket buffer to eat
2173 *
2174 * This routine must be called with interrupts disabled or with the socket
2175 * locked so that the sk_buff queue operation is ok.
2176 */
2177 static inline void sk_eat_skb(struct sock *sk, struct sk_buff *skb)
2178 {
2179     skb_unlink(skb, &sk->sk_receive_queue);
2180     kfree_skb(skb);
2181 }
2182
2183 static inline
2184 struct net *sock_net(const struct sock *sk)
2185 {
2186     return read_pnet(&sk->sk_net);
2187 }
2188
2189 static inline
2190 void sock_net_set(struct sock *sk, struct net *net)
2191 {
2192     write_pnet(&sk->sk_net, net);
2193 }
2194
2195 static inline struct sock *skb_steal_sock(struct sk_buff *skb)
2196 {
2197     if (skb->sk) {
2198         struct sock *sk = skb->sk;
2199
2200         skb->destructor = NULL;
2201         skb->sk = NULL;
2202         return sk;
2203     }
2204     return NULL;
2205 }
2206
2207 /* This helper checks if a socket is a full socket,
2208 * ie _not_ a timewait or request socket.
2209 */
2210 static inline bool sk_fullsock(const struct sock *sk)
2211 {
2212     return (1 << sk->sk_state) & ~(TCPF_TIME_WAIT | TCPF_NEW_SYN_RECV);
2213 }
2214
2215 void sock_enable_timestamp(struct sock *sk, int flag);
2216 int sock_get_timestamp(struct sock *, struct timeval __user *);
2217 int sock_get_timestampns(struct sock *, struct timespec __user *);
2218 int sock_rcv_errqueue(struct sock *sk, struct msghdr *msg, int len, int level,
2219 int type);
2220
2221 bool sk_ns_capable(const struct sock *sk,
2222 struct user_namespace *user_ns, int cap);
2223 bool sk_capable(const struct sock *sk, int cap);

```

```
2224 bool sk_net_capable(const struct sock *sk, int cap);
2225
2226 extern __u32 sysctl_wmem_max;
2227 extern __u32 sysctl_rmem_max;
2228
2229 extern int sysctl_timestamp_allow_data;
2230 extern int sysctl_optmem_max;
2231
2232 extern __u32 sysctl_wmem_default;
2233 extern __u32 sysctl_rmem_default;
2234
2235 #endif /* _SOCK_H */
2236
```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)