

Linux Cross Reference

Free Electrons

Embedded Linux Experts

• [source navigation](#) • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version: [2.0.40](#) [2.2.26](#) [2.4.37](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#) [3.18](#) [3.19](#) [4.0](#) [4.1](#) [4.2](#)

Linux/include/uapi/linux/if_packet.h

```

1 #ifndef _LINUX_IF_PACKET_H
2 #define _LINUX_IF_PACKET_H
3
4 #include <linux/types.h>
5
6 struct sockaddr_pkt {
7     unsigned short spkt_family;
8     unsigned char spkt_device[14];
9     __be16 spkt_protocol;
10 };
11
12 struct sockaddr_ll {
13     unsigned short  sll_family;
14     __be16          sll_protocol;
15     int             sll_ifindex;
16     unsigned short  sll_hatype;
17     unsigned char    sll_pkttype;
18     unsigned char    sll_halen;
19     unsigned char    sll_addr[8];
20 };
21
22 /* Packet types */
23
24 #define PACKET_HOST            0          /* To us */
25 #define PACKET_BROADCAST      1          /* To all */
26 #define PACKET_MULTICAST      2          /* To group */
27 #define PACKET_OTHERHOST      3          /* To someone else */
28 #define PACKET_OUTGOING       4          /* Outgoing of any type */
29 #define PACKET_LOOPBACK       5          /* MC/BRD frame looped back */
30 #define PACKET_USER           6          /* To user space */
31 #define PACKET_KERNEL         7          /* To kernel space */
32 /* Unused, PACKET_FASTROUTE and PACKET_LOOPBACK are invisible to user space */
33 #define PACKET_FASTROUTE      6          /* Fastrouted frame */
34
35 /* Packet socket options */
36
37 #define PACKET_ADD_MEMBERSHIP  1
38 #define PACKET_DROP_MEMBERSHIP 2
39 #define PACKET_RECV_OUTPUT    3
40 /* Value 4 is still used by obsolete turbo-packet. */
41 #define PACKET_RX_RING        5
42 #define PACKET_STATISTICS     6
43 #define PACKET_COPY_THRESH    7
44 #define PACKET_AUXDATA        8
45 #define PACKET_ORIGDEV        9
46 #define PACKET_VERSION        10
47 #define PACKET_HDRLEN         11
48 #define PACKET_RESERVE        12
49 #define PACKET_TX_RING        13
50 #define PACKET_LOSS           14
51 #define PACKET_VNET_HDR       15
52 #define PACKET_TX_TIMESTAMP    16
53 #define PACKET_TIMESTAMP      17
54 #define PACKET_FANOUT         18
55 #define PACKET_TX_HAS_OFF     19
56 #define PACKET_QDISC_BYPASS    20

```

```

57 #define PACKET_ROLLOVER_STATS 21
58
59 #define PACKET_FANOUT_HASH 0
60 #define PACKET_FANOUT_LB 1
61 #define PACKET_FANOUT_CPU 2
62 #define PACKET_FANOUT_ROLLOVER 3
63 #define PACKET_FANOUT_RND 4
64 #define PACKET_FANOUT_QM 5
65 #define PACKET_FANOUT_FLAG_ROLLOVER 0x1000
66 #define PACKET_FANOUT_FLAG_DEFRAG 0x8000
67
68 struct tpacket_stats {
69     unsigned int    tp_packets;
70     unsigned int    tp_drops;
71 };
72
73 struct tpacket_stats_v3 {
74     unsigned int    tp_packets;
75     unsigned int    tp_drops;
76     unsigned int    tp_freeze_q_cnt;
77 };
78
79 struct tpacket_rollover_stats {
80     __aligned_u64    tp_all;
81     __aligned_u64    tp_huge;
82     __aligned_u64    tp_failed;
83 };
84
85 union tpacket_stats_u {
86     struct tpacket_stats stats1;
87     struct tpacket_stats_v3 stats3;
88 };
89
90 struct tpacket_auxdata {
91     __u32            tp_status;
92     __u32            tp_len;
93     __u32            tp_snaplen;
94     __u16            tp_mac;
95     __u16            tp_net;
96     __u16            tp_vlan_tci;
97     __u16            tp_vlan_tpid;
98 };
99
100 /* Rx ring - header status */
101 #define TP_STATUS_KERNEL 0
102 #define TP_STATUS_USER (1 << 0)
103 #define TP_STATUS_COPY (1 << 1)
104 #define TP_STATUS_LOSING (1 << 2)
105 #define TP_STATUS_CSUMNOTREADY (1 << 3)
106 #define TP_STATUS_VLAN_VALID (1 << 4) /* auxdata has valid tp_vlan_tci */
107 #define TP_STATUS_BLK_TMO (1 << 5)
108 #define TP_STATUS_VLAN_TPID_VALID (1 << 6) /* auxdata has valid tp_vlan_tpid */
109 #define TP_STATUS_CSUM_VALID (1 << 7)
110
111 /* Tx ring - header status */
112 #define TP_STATUS_AVAILABLE 0
113 #define TP_STATUS_SEND_REQUEST (1 << 0)
114 #define TP_STATUS_SENDING (1 << 1)
115 #define TP_STATUS_WRONG_FORMAT (1 << 2)
116
117 /* Rx and Tx ring - header status */
118 #define TP_STATUS_TS_SOFTWARE (1 << 29)
119 #define TP_STATUS_TS_SYS_HARDWARE (1 << 30) /* deprecated, never set */
120 #define TP_STATUS_TS_RAW_HARDWARE (1 << 31)
121
122 /* Rx ring - feature request bits */
123 #define TP_FT_REQ_FILL_RXHASH 0x1
124
125 struct tpacket_hdr {
126     unsigned long    tp_status;
127     unsigned int     tp_len;
128     unsigned int     tp_snaplen;
129     unsigned short    tp_mac;
130     unsigned short    tp_net;
131     unsigned int     tp_sec;
132     unsigned int     tp_usec;

```

```

133 };
134
135 #define TPACKET_ALIGNMENT 16
136 #define TPACKET_ALIGN(x) (((x)+TPACKET_ALIGNMENT-1)&~(TPACKET_ALIGNMENT-1))
137 #define TPACKET_HDRLEN (TPACKET_ALIGN(sizeof(struct tpacket_hdr)) + sizeof(struct sockaddr_ll))
138
139 struct tpacket2_hdr {
140     __u32 tp_status;
141     __u32 tp_len;
142     __u32 tp_snaplen;
143     __u16 tp_mac;
144     __u16 tp_net;
145     __u32 tp_sec;
146     __u32 tp_nsec;
147     __u16 tp_vlan_tci;
148     __u16 tp_vlan_tpid;
149     __u8 tp_padding[4];
150 };
151
152 struct tpacket_hdr_variant1 {
153     __u32 tp_rhash;
154     __u32 tp_vlan_tci;
155     __u16 tp_vlan_tpid;
156     __u16 tp_padding;
157 };
158
159 struct tpacket3_hdr {
160     __u32 tp_next_offset;
161     __u32 tp_sec;
162     __u32 tp_nsec;
163     __u32 tp_snaplen;
164     __u32 tp_len;
165     __u32 tp_status;
166     __u16 tp_mac;
167     __u16 tp_net;
168     /* pkt_hdr variants */
169     union {
170         struct tpacket_hdr_variant1 hv1;
171     };
172     __u8 tp_padding[8];
173 };
174
175 struct tpacket_bd_ts {
176     unsigned int ts_sec;
177     union {
178         unsigned int ts_usec;
179         unsigned int ts_nsec;
180     };
181 };
182
183 struct tpacket_hdr_v1 {
184     __u32 block_status;
185     __u32 num_pkts;
186     __u32 offset_to_first_pkt;
187
188     /* Number of valid bytes (including padding)
189      * blk_len <= tp_block_size
190      */
191     __u32 blk_len;
192
193     /*
194      * Quite a few uses of sequence number:
195      * 1. Make sure cache flush etc worked.
196      *    Well, one can argue - why not use the increasing ts below?
197      *    But look at 2. below first.
198      * 2. When you pass around blocks to other user space decoders,
199      *    you can see which blk[s] is[are] outstanding etc.
200      * 3. Validate kernel code.
201      */
202     __aligned_u64 seq_num;
203
204     /*
205      * ts_last_pkt:
206      *
207      * Case 1. Block has 'N'(N >=1) packets and TMO'd(timed out)
208      *          ts_last_pkt == 'time-stamp of last packet' and NOT the

```

```

209      *           time when the timer fired and the block was closed.
210      *           By providing the ts of the last packet we can absolutely
211      *           guarantee that time-stamp wise, the first packet in the
212      *           next block will never precede the last packet of the
213      *           previous block.
214      * Case 2.   Block has zero packets and TMO'd
215      *           ts_last_pkt = time when the timer fired and the block
216      *           was closed.
217      * Case 3.   Block has 'N' packets and NO TMO.
218      *           ts_last_pkt = time-stamp of the last pkt in the block.
219      *
220      * ts_first_pkt:
221      *           Is always the time-stamp when the block was opened.
222      *           Case a) ZERO packets
223      *                   No packets to deal with but atleast you know the
224      *                   time-interval of this block.
225      *           Case b) Non-zero packets
226      *                   Use the ts of the first packet in the block.
227      *
228      */
229      struct tpacket_bd_ts    ts_first_pkt, ts_last_pkt;
230 };
231
232 union tpacket_bd_header_u {
233     struct tpacket_hdr_v1 bh1;
234 };
235
236 struct tpacket_block_desc {
237     __u32 version;
238     __u32 offset_to_priv;
239     union tpacket_bd_header_u hdr;
240 };
241
242 #define TPACKET2_HDRLEN    (TPACKET_ALIGN(sizeof(struct tpacket2_hdr)) + sizeof(struct sockaddr_ll))
243 #define TPACKET3_HDRLEN    (TPACKET_ALIGN(sizeof(struct tpacket3_hdr)) + sizeof(struct sockaddr_ll))
244
245 enum tpacket_versions {
246     TPACKET_V1,
247     TPACKET_V2,
248     TPACKET_V3
249 };
250
251 /*
252  * Frame structure:
253  *
254  * - Start. Frame must be aligned to TPACKET_ALIGNMENT=16
255  * - struct tpacket_hdr
256  * - pad to TPACKET_ALIGNMENT=16
257  * - struct sockaddr_ll
258  * - Gap, chosen so that packet data (Start+tp_net) alignes to TPACKET_ALIGNMENT=16
259  * - Start+tp_mac: [ Optional MAC header ]
260  * - Start+tp_net: Packet data, aligned to TPACKET_ALIGNMENT=16.
261  * - Pad to align to TPACKET_ALIGNMENT=16
262  */
263
264 struct tpacket_req {
265     unsigned int    tp_block_size;    /* Minimal size of contiguous block */
266     unsigned int    tp_block_nr;     /* Number of blocks */
267     unsigned int    tp_frame_size;   /* Size of frame */
268     unsigned int    tp_frame_nr;     /* Total number of frames */
269 };
270
271 struct tpacket_req3 {
272     unsigned int    tp_block_size;    /* Minimal size of contiguous block */
273     unsigned int    tp_block_nr;     /* Number of blocks */
274     unsigned int    tp_frame_size;   /* Size of frame */
275     unsigned int    tp_frame_nr;     /* Total number of frames */
276     unsigned int    tp_retire_blk_tov; /* timeout in msecs */
277     unsigned int    tp_sizeof_priv;   /* offset to private data area */
278     unsigned int    tp_feature_req_word;
279 };
280
281 union tpacket_req_u {
282     struct tpacket_req    req;
283     struct tpacket_req3   req3;
284 };

```

```
285
286 struct packet_mreq {
287     int      mr_ifindex;
288     unsigned short mr_type;
289     unsigned short mr_alen;
290     unsigned char mr_address[8];
291 };
292
293 #define PACKET_MR_MULTICAST    0
294 #define PACKET_MR_PROMISC     1
295 #define PACKET_MR_ALLMULTI    2
296 #define PACKET_MR_UNICAST     3
297
298 #endif
299
```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)