

Linux Cross Reference

[Free Electrons](#)

Embedded Linux Experts

• [source navigation](#) • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version: [2.0.40](#) [2.2.26](#) [2.4.37](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#) [3.18](#) [3.19](#) [4.0](#) [4.1](#) [4.2](#)

[Linux/net/ipv4/ip_input.c](#)

```

1  /*
2  * INET          An implementation of the TCP/IP protocol suite for the LINUX
3  *              operating system. INET is implemented using the BSD Socket
4  *              interface as the means of communication with the user level.
5  *
6  *              The Internet Protocol (IP) module.
7  *
8  * Authors:      Ross Biro
9  *              Fred N. van Kempen, <waltje@uWalt.NL.Mugnet.ORG>
10 *              Donald Becker, <becker@super.org>
11 *              Alan Cox, <alan@lxorguk.ukuu.org.uk>
12 *              Richard Underwood
13 *              Stefan Becker, <stefanb@yello.ping.de>
14 *              Jorge Cwik, <jorge@laser.satlink.net>
15 *              Arnt Gulbrandsen, <agulbra@nvg.unit.no>
16 *
17 *
18 * Fixes:
19 *              Alan Cox      :      Commented a couple of minor bits of surplus code
20 *              Alan Cox      :      Undefined IP_FORWARD doesn't include the code
21 *                                (just stops a compiler warning).
22 *              Alan Cox      :      Frames with >=MAX_ROUTE record routes, strict routes or loose routes
23 *                                are junked rather than corrupting things.
24 *              Alan Cox      :      Frames to bad broadcast subnets are dumped
25 *                                We used to process them non broadcast and
26 *                                boy could that cause havoc.
27 *              Alan Cox      :      ip_forward sets the free flag on the
28 *                                new frame it queues. Still crap because
29 *                                it copies the frame but at least it
30 *                                doesn't eat memory too.
31 *              Alan Cox      :      Generic queue code and memory fixes.
32 *              Fred Van Kempen :      IP fragment support (borrowed from NET2E)
33 *              Gerhard Koerting:      Forward fragmented frames correctly.
34 *              Gerhard Koerting:      Fixes to my fix of the above 8-).
35 *              Gerhard Koerting:      IP interface addressing fix.
36 *              Linus Torvalds :      More robustness checks
37 *              Alan Cox      :      Even more checks: Still not as robust as it ought to be
38 *              Alan Cox      :      Save IP header pointer for later
39 *              Alan Cox      :      ip option setting
40 *              Alan Cox      :      Use ip_tos/ip_ttl settings
41 *              Alan Cox      :      Fragmentation bogosity removed
42 *                                (Thanks to Mark.Bush@prg.ox.ac.uk)
43 *              Dmitry Gorodchanin :      Send of a raw packet crash fix.
44 *              Alan Cox      :      Silly ip bug when an overlength
45 *                                fragment turns up. Now frees the
46 *                                queue.
47 *              Linus Torvalds/ :      Memory Leakage on fragmentation
48 *              Alan Cox      :      handling.
49 *              Gerhard Koerting:      Forwarding uses IP priority hints
50 *              Teemu Rantanen :      Fragment problems.
51 *              Alan Cox      :      General cleanup, comments and reformat
52 *              Alan Cox      :      SNMP statistics
53 *              Alan Cox      :      BSD address rule semantics. Also see
54 *                                UDP as there is a nasty checksum issue
55 *                                if you do things the wrong way.
56 *              Alan Cox      :      Always defrag, moved IP_FORWARD to the config.in file
57 *              Alan Cox      :      IP options adjust sk->priority.
58 *              Pedro Roque   :      Fix mtu/length error in ip_forward.

```

```

59 *           Alan Cox           :   Avoid ip_chk_addr when possible.
60 *   Richard Underwood         :   IP multicasting.
61 *           Alan Cox           :   Cleaned up multicast handlers.
62 *           Alan Cox           :   RAW sockets demultiplex in the BSD style.
63 *           Gunther Mayer      :   Fix the SNMP reporting typo
64 *           Alan Cox           :   Always in group 224.0.0.1
65 *   Pauline Middelink         :   Fast ip_checksum update when forwarding
66 *                               :   Masquerading support.
67 *           Alan Cox           :   Multicast loopback error for 224.0.0.1
68 *           Alan Cox           :   IP_MULTICAST_LOOP option.
69 *           Alan Cox           :   Use notifiers.
70 *           Bjorn Ekwall        :   Removed ip_csum (from slhc.c too)
71 *           Bjorn Ekwall        :   Moved ip_fast_csum to ip.h (inline!)
72 *           Stefan Becker       :   Send out ICMP_HOST_REDIRECT
73 *   Arnt Gulbrandsen          :   ip_build_xmit
74 *           Alan Cox           :   Per socket routing cache
75 *           Alan Cox           :   Fixed routing cache, added header cache.
76 *           Alan Cox           :   Loopback didn't work right in original ip_build_xmit - fixed it.
77 *           Alan Cox           :   Only send ICMP_REDIRECT if src/dest are the same net.
78 *           Alan Cox           :   Incoming IP option handling.
79 *           Alan Cox           :   Set saddr on raw output frames as per BSD.
80 *           Alan Cox           :   Stopped broadcast source route explosions.
81 *           Alan Cox           :   Can disable source routing
82 *           Takeshi Sone        :   Masquerading didn't work.
83 *   Dave Bonn,Alan Cox        :   Faster IP forwarding whenever possible.
84 *           Alan Cox           :   Memory Leaks, tramples, misc debugging.
85 *           Alan Cox           :   Fixed multicast (by popular demand 8))
86 *           Alan Cox           :   Fixed forwarding (by even more popular demand 8))
87 *           Alan Cox           :   Fixed SNMP statistics [I think]
88 *   Gerhard Koerting          :   IP fragmentation forwarding fix
89 *           Alan Cox           :   Device Lock against page fault.
90 *           Alan Cox           :   IP_HDRINCL facility.
91 *   Werner Almesberger        :   Zero fragment bug
92 *           Alan Cox           :   RAW IP frame length bug
93 *           Alan Cox           :   Outgoing firewall on build_xmit
94 *           A.N.Kuznetsov       :   IP_OPTIONS support throughout the kernel
95 *           Alan Cox           :   Multicast routing hooks
96 *           Jos Vos             :   Do accounting *before* call_in_firewall
97 *   Willy Konynenberg         :   Transparent proxying support
98 *
99 *
100 *
101 * To Fix:
102 *           IP fragmentation wants rewriting cleanly. The RFC815 algorithm is much more efficient
103 *           and could be made very efficient with the addition of some virtual memory hacks to permit
104 *           the allocation of a buffer that can then be 'grown' by twiddling page tables.
105 *           Output fragmentation wants updating along with the buffer management to use a single
106 *           interleaved copy algorithm so that fragmenting has a one copy overhead. Actual packet
107 *           output should probably do its own fragmentation at the UDP/RAW layer. TCP shouldn't cause
108 *           fragmentation anyway.
109 *
110 *           This program is free software; you can redistribute it and/or
111 *           modify it under the terms of the GNU General Public License
112 *           as published by the Free Software Foundation; either version
113 *           2 of the License, or (at your option) any later version.
114 */
115
116 #define pr_fmt(fmt) "IPv4: " fmt
117
118 #include <linux/module.h>
119 #include <linux/types.h>
120 #include <linux/kernel.h>
121 #include <linux/string.h>
122 #include <linux/errno.h>
123 #include <linux/slab.h>
124
125 #include <linux/net.h>
126 #include <linux/socket.h>
127 #include <linux/sockios.h>
128 #include <linux/in.h>
129 #include <linux/inet.h>
130 #include <linux/inetdevice.h>
131 #include <linux/netdevice.h>
132 #include <linux/etherdevice.h>
133
134 #include <net/snmp.h>
135 #include <net/ip.h>
136 #include <net/protocol.h>

```

```

137 #include <net/route.h>
138 #include <linux/skbuff.h>
139 #include <net/sock.h>
140 #include <net/arp.h>
141 #include <net/icmp.h>
142 #include <net/raw.h>
143 #include <net/checksum.h>
144 #include <net/inet_ecn.h>
145 #include <linux/netfilter_ipv4.h>
146 #include <net/xfrm.h>
147 #include <linux/mroute.h>
148 #include <linux/netlink.h>
149
150 /*
151  *      Process Router Attention IP option (RFC 2113)
152  */
153 bool ip_call_ra_chain(struct sk_buff *skb)
154 {
155     struct ip_ra_chain *ra;
156     u8 protocol = ip_hdr(skb)->protocol;
157     struct sock *last = NULL;
158     struct net_device *dev = skb->dev;
159
160     for (ra = rcu_dereference(ip_ra_chain); ra; ra = rcu_dereference(ra->next)) {
161         struct sock *sk = ra->sk;
162
163         /* If socket is bound to an interface, only report
164          * the packet if it came from that interface.
165          */
166         if (sk && inet_sk(sk)->inet_num == protocol &&
167             (!sk->sk_bound_dev_if ||
168              sk->sk_bound_dev_if == dev->ifindex) &&
169             net_eq(sock_net(sk), dev_net(dev))) {
170             if (ip_is_fragment(ip_hdr(skb))) {
171                 if (ip_defrag(skb, IP_DEFRAG_CALL_RA_CHAIN))
172                     return true;
173             }
174             if (last) {
175                 struct sk_buff *skb2 = skb_clone(skb, GFP_ATOMIC);
176                 if (skb2)
177                     raw_rcv(last, skb2);
178             }
179             last = sk;
180         }
181     }
182
183     if (last) {
184         raw_rcv(last, skb);
185         return true;
186     }
187     return false;
188 }
189
190 static int ip_local_deliver_finish(struct sock *sk, struct sk_buff *skb)
191 {
192     struct net *net = dev_net(skb->dev);
193
194     __skb_pull(skb, skb_network_header_len(skb));
195
196     rcu_read_lock();
197     {
198         int protocol = ip_hdr(skb)->protocol;
199         const struct net_protocol *ipprot;
200         int raw;
201
202         resubmit:
203         raw = raw_local_deliver(skb, protocol);
204
205         ipprot = rcu_dereference(inet_protos[protocol]);
206         if (ipprot) {
207             int ret;
208
209             if (!ipprot->no_policy) {
210                 if (!xfrm4_policy_check(NULL, XFRM_POLICY_IN, skb)) {
211                     kfree_skb(skb);
212                     goto out;
213                 }
214             }
215             nf_reset(skb);
216         }
217     }
218
219     out:
220     rcu_read_unlock();
221
222     if (raw)
223         raw_rcv(sk, skb);
224
225     return 0;
226 }

```

```

215     }
216     ret = ipprot->handler\(skb\);
217     if (ret < 0) {
218         protocol = -ret;
219         goto resubmit;
220     }
221     IP\_INC\_STATS\_BH(net, IPSTATS_MIB_INDELIVERS);
222 } else {
223     if (!raw) {
224         if (xfrm4\_policy\_check(NULL, XFRM_POLICY_IN, skb)) {
225             IP\_INC\_STATS\_BH(net, IPSTATS_MIB_INUNKNOWNPROTOS);
226             icmp\_send(skb, ICMP_DEST_UNREACH,
227                     ICMP_PROT_UNREACH, 0);
228         }
229         kfree\_skb(skb);
230     } else {
231         IP\_INC\_STATS\_BH(net, IPSTATS_MIB_INDELIVERS);
232         consume\_skb(skb);
233     }
234 }
235 }
236 out:
237     rcu\_read\_unlock();
238
239     return 0;
240 }
241
242 /*
243  * Deliver IP Packets to the higher protocol layers.
244  */
245 int ip\_local\_deliver(struct sk\_buff *skb)
246 {
247     /*
248      * Reassemble IP fragments.
249      */
250
251     if (ip\_is\_fragment(ip\_hdr(skb))) {
252         if (ip\_defrag(skb, IP_DEFRAG_LOCAL_DELIVER))
253             return 0;
254     }
255
256     return NF\_HOOK(NFPROTO_IPV4, NF_INET_LOCAL_IN, NULL, skb,
257                   skb->dev, NULL,
258                   ip\_local\_deliver\_finish);
259 }
260
261 static inline bool ip\_rcv\_options(struct sk\_buff *skb)
262 {
263     struct ip\_options *opt;
264     const struct iphdr *iph;
265     struct net\_device *dev = skb->dev;
266
267     /* It looks as overkill, because not all
268      * IP options require packet mangling.
269      * But it is the easiest for now, especially taking
270      * into account that combination of IP options
271      * and running sniffer is extremely rare condition.
272      * --ANK (980813)
273     */
274     if (skb\_cow(skb, skb\_headroom(skb))) {
275         IP\_INC\_STATS\_BH(dev\_net(dev), IPSTATS_MIB_INDISCARDS);
276         goto drop;
277     }
278
279     iph = ip\_hdr(skb);
280     opt = &(IPCB(skb)->opt);
281     opt->optlen = iph->ihl*4 - sizeof(struct iphdr);
282
283     if (ip\_options\_compile(dev\_net(dev), opt, skb)) {
284         IP\_INC\_STATS\_BH(dev\_net(dev), IPSTATS_MIB_INHDRERRORS);
285         goto drop;
286     }
287
288     if (unlikely(opt->srr)) {
289         struct in\_device *in_dev = \_\_in\_dev\_get\_rcu(dev);
290
291         if (in_dev) {
292             if (!IN\_DEV\_SOURCE\_ROUTE(in_dev)) {

```

```

293         if (IN_DEV_LOG_MARTIANS(in_dev))
294             net_info_ratelimited("source route option %pI4 -> %pI4\n",
295                                 &iph->saddr,
296                                 &iph->daddr);
297         goto drop;
298     }
299 }
300
301     if (ip_options_rcv_srr(skb))
302         goto drop;
303 }
304
305 return false;
306 drop:
307 return true;
308 }
309
310 int sysctl_ip_early_demux_read_mostly = 1;
311 EXPORT_SYMBOL(sysctl_ip_early_demux);
312
313 static int ip_rcv_finish(struct sock *sk, struct sk_buff *skb)
314 {
315     const struct iphdr *iph = ip_hdr(skb);
316     struct rtable *rt;
317
318     if (sysctl_ip_early_demux && !skb_dst(skb) && !skb->sk) {
319         const struct net_protocol *ipprot;
320         int protocol = iph->protocol;
321
322         ipprot = rcu_dereference(inet_protos[protocol]);
323         if (ipprot && ipprot->early_demux) {
324             ipprot->early_demux(skb);
325             /* must reload iph, skb->head might have changed */
326             iph = ip_hdr(skb);
327         }
328     }
329
330     /*
331     *   Initialise the virtual path cache for the packet. It describes
332     *   how the packet travels inside Linux networking.
333     */
334     if (!skb_dst(skb)) {
335         int err = ip_route_input_noref(skb, iph->daddr, iph->saddr,
336                                       iph->tos, skb->dev);
337         if (unlikely(err)) {
338             if (err == -EXDEV)
339                 NET_INC_STATS_BH(dev_net(skb->dev),
340                                LINUX_MIB_IPRPFILTER);
341             goto drop;
342         }
343     }
344
345     #ifdef CONFIG_IP_ROUTE_CLASSID
346     if (unlikely(skb_dst(skb)->tclassid)) {
347         struct ip_rt_acct *st = this_cpu_ptr(ip_rt_acct);
348         u32 idx = skb_dst(skb)->tclassid;
349         st[idx&0xFF].o_packets++;
350         st[idx&0xFF].o_bytes += skb->len;
351         st[(idx>>16)&0xFF].i_packets++;
352         st[(idx>>16)&0xFF].i_bytes += skb->len;
353     }
354     #endif
355
356     if (iph->ihl > 5 && ip_rcv_options(skb))
357         goto drop;
358
359     rt = skb_rtable(skb);
360     if (rt->rt_type == RTN_MULTICAST) {
361         IP_UPD_PO_STATS_BH(dev_net(rt->dst.dev), IPSTATS_MIB_INMCAST,
362                           skb->len);
363     } else if (rt->rt_type == RTN_BROADCAST)
364         IP_UPD_PO_STATS_BH(dev_net(rt->dst.dev), IPSTATS_MIB_INBCAST,
365                           skb->len);
366
367     return dst_input(skb);
368
369 drop:
370     kfree_skb(skb);

```

```

371     return NET\_RX\_DROP;
372 }
373
374 /*
375  *      Main IP Receive routine.
376  */
377 int ip\_rcv(struct sk\_buff *skb, struct net\_device *dev, struct packet\_type *pt, struct net\_device *orig_dev)
378 {
379     const struct iphdr *iph;
380     u32 len;
381
382     /* When the interface is in promisc. mode, drop all the crap
383      * that it receives, do not try to analyse it.
384      */
385     if (skb->pkt_type == PACKET\_OTHERHOST)
386         goto drop;
387
388     IP\_UPD\_PO\_STATS\_BH(dev\_net(dev), IPSTATS\_MIB\_IN, skb->len);
389
390     skb = skb\_share\_check(skb, GFP\_ATOMIC);
391     if (!skb) {
392         IP\_INC\_STATS\_BH(dev\_net(dev), IPSTATS\_MIB\_INDISCARDS);
393         goto out;
394     }
395
396     if (!pskb\_may\_pull(skb, sizeof(struct iphdr)))
397         goto inhdr\_error;
398
399     iph = ip\_hdr(skb);
400
401     /*
402      *      RFC1122: 3.2.1.2 MUST silently discard any IP frame that fails the checksum.
403      *
404      *      Is the datagram acceptable?
405      *
406      *      1.      Length at Least the size of an ip header
407      *      2.      Version of 4
408      *      3.      Checksums correctly. [Speed optimisation for later, skip Loopback checksums]
409      *      4.      Doesn't have a bogus length
410      */
411
412     if (iph->ihl < 5 || iph->version != 4)
413         goto inhdr\_error;
414
415     BUILD\_BUG\_ON(IPSTATS\_MIB\_ECT1PKTS != IPSTATS\_MIB\_NOECTPKTS + INET\_ECN\_ECT\_1);
416     BUILD\_BUG\_ON(IPSTATS\_MIB\_ECT0PKTS != IPSTATS\_MIB\_NOECTPKTS + INET\_ECN\_ECT\_0);
417     BUILD\_BUG\_ON(IPSTATS\_MIB\_CEPKTS != IPSTATS\_MIB\_NOECTPKTS + INET\_ECN\_CE);
418     IP\_ADD\_STATS\_BH(dev\_net(dev),
419                    IPSTATS\_MIB\_NOECTPKTS + (iph->tos & INET\_ECN\_MASK),
420                    max\_t(unsigned short, 1, skb\_shinfo(skb)->gso_segs));
421
422     if (!pskb\_may\_pull(skb, iph->ihl*4))
423         goto inhdr\_error;
424
425     iph = ip\_hdr(skb);
426
427     if (unlikely(ip\_fast\_csum((u8 *)iph, iph->ihl)))
428         goto csum\_error;
429
430     len = ntohs(iph->tot_len);
431     if (skb->len < len) {
432         IP\_INC\_STATS\_BH(dev\_net(dev), IPSTATS\_MIB\_INTRUNCATEDPKTS);
433         goto drop;
434     } else if (len < (iph->ihl*4))
435         goto inhdr\_error;
436
437     /* Our transport medium may have padded the buffer out. Now we know it
438      * is IP we can trim to the true length of the frame.
439      * Note this now means skb->len holds ntohs(iph->tot_len).
440      */
441     if (pskb\_trim\_rsum(skb, len)) {
442         IP\_INC\_STATS\_BH(dev\_net(dev), IPSTATS\_MIB\_INDISCARDS);
443         goto drop;
444     }
445
446     skb->transport_header = skb->network_header + iph->ihl*4;
447
448

```

```
449      /* Remove any debris in the socket control block */
450      memset(IPCB(skb), 0, sizeof(struct inet_skb_parm));
451
452      /* Must drop socket now because of tproxy. */
453      skb_orphan(skb);
454
455      return NF_HOOK(NFPROTO_IPV4, NF_INET_PRE_ROUTING, NULL, skb,
456                    dev, NULL,
457                    ip_rcv_finish);
458
459  csum_error:
460      IP_INC_STATS_BH(dev_net(dev), IPSTATS_MIB_CSUMERRORS);
461  inhdr_error:
462      IP_INC_STATS_BH(dev_net(dev), IPSTATS_MIB_INHDRERRORS);
463  drop:
464      kfree_skb(skb);
465  out:
466      return NET_RX_DROP;
467  }
468
```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)