

Linux Cross Reference

[Free Electrons](#)

Embedded Linux Experts

• [source navigation](#) • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version: [2.0.40](#) [2.2.26](#) [2.4.37](#) [3.1](#) [3.2](#) [3.3](#) [3.4](#) [3.5](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#)

[Linux/net/ipv4/tcp.c](#)

```

1  /*
2  * INET          An implementation of the TCP/IP protocol suite for the LINUX
3  *              operating system.  INET is implemented using the BSD Socket
4  *              interface as the means of communication with the user level.
5  *
6  *              Implementation of the Transmission Control Protocol(TCP).
7  *
8  * Authors:      Ross Biro
9  *              Fred N. van Kempen, <waltje@uWalt.NL.Mugnet.ORG>
10 *              Mark Evans, <evansmp@uhura.aston.ac.uk>
11 *              Corey Minyard <wf-rch!minyard@relay.EU.net>
12 *              Florian La Roche, <flla@stud.uni-sb.de>
13 *              Charles Hedrick, <hedrick@klinzhai.rutgers.edu>
14 *              Linus Torvalds, <torvalds@cs.helsinki.fi>
15 *              Alan Cox, <gw4pts@gw4pts.ampr.org>
16 *              Matthew Dillon, <dillon@apollo.west.oic.com>
17 *              Arnt Gulbrandsen, <agulbra@nvg.unit.no>
18 *              Jorge Cwik, <jorge@Laser.satlink.net>
19 *
20 * Fixes:
21 *              Alan Cox          :      Numerous verify_area() calls
22 *              Alan Cox          :      Set the ACK bit on a reset
23 *              Alan Cox          :      Stopped it crashing if it closed while
24 *              sk->inuse=1 and was trying to connect
25 *              (tcp_err()).
26 *              Alan Cox          :      All icmp error handling was broken
27 *              pointers passed where wrong and the
28 *              socket was looked up backwards. Nobody
29 *              tested any icmp error code obviously.
30 *              Alan Cox          :      tcp_err() now handled properly. It
31 *              wakes people on errors. poll
32 *              behaves and the icmp error race
33 *              has gone by moving it into sock.c
34 *              Alan Cox          :      tcp_send_reset() fixed to work for
35 *              everything not just packets for
36 *              unknown sockets.
37 *              Alan Cox          :      tcp option processing.
38 *              Alan Cox          :      Reset tweaked (still not 100%) [Had
39 *              syn rule wrong]
40 *              Herp Rosmanith    :      More reset fixes
41 *              Alan Cox          :      No longer acks invalid rst frames.
42 *              Acking any kind of RST is right out.
43 *              Alan Cox          :      Sets an ignore me flag on an rst
44 *              receive otherwise odd bits of prattle
45 *              escape still
46 *              Alan Cox          :      Fixed another acking RST frame bug.
47 *              Should stop LAN workplace lockups.
48 *              Alan Cox          :      Some tidyups using the new skb list
49 *              facilities
50 *              Alan Cox          :      sk->keepopen now seems to work
51 *              Alan Cox          :      Pulls options out correctly on accepts
52 *              Alan Cox          :      Fixed assorted sk->rqueue->next errors
53 *              Alan Cox          :      PSH doesn't end a TCP read. Switched a
54 *              bit to skb ops.
55 *              Alan Cox          :      Tidied tcp_data to avoid a potential
56 *              nasty.
57 *              Alan Cox          :      Added some better commenting, as the
58 *              tcp is hard to follow
59 *              Alan Cox          :      Removed incorrect check for 20 * psh
60 *              Michael O'Reilly   :      ack < copied bug fix.
61 *              Johannes Stille    :      Misc tcp fixes (not all in yet).
62 *              Alan Cox          :      FIN with no memory -> CRASH
63 *              Alan Cox          :      Added socket option proto entries.
64 *              Also added awareness of them to accept.
65 *              Alan Cox          :      Added TCP options (SOL_TCP)

```

```

66 *      Alan Cox      :      Switched wakeup calls to callbacks,
67 *                      so the kernel can layer network
68 *                      sockets.
69 *      Alan Cox      :      Use ip_tos/ip_ttl settings.
70 *      Alan Cox      :      Handle FIN (more) properly (we hope).
71 *      Alan Cox      :      RST frames sent on unsynchronised
72 *                      state ack error.
73 *      Alan Cox      :      Put in missing check for SYN bit.
74 *      Alan Cox      :      Added tcp_select_window() aka NET2E
75 *                      window non shrink trick.
76 *      Alan Cox      :      Added a couple of small NET2E timer
77 *                      fixes
78 *      Charles Hedrick :      TCP fixes
79 *      Toomas Tamm    :      TCP window fixes
80 *      Alan Cox      :      Small URG fix to rlogin ^C ack fight
81 *      Charles Hedrick :      Rewrote most of it to actually work
82 *      Linus          :      Rewrote tcp_read() and URG handling
83 *                      completely
84 *      Gerhard Koerting:      Fixed some missing timer handling
85 *      Matthew Dillon :      Reworked TCP machine states as per RFC
86 *      Gerhard Koerting:      PC/TCP workarounds
87 *      Adam Caldwell  :      Assorted timer/timing errors
88 *      Matthew Dillon :      Fixed another RST bug
89 *      Alan Cox      :      Move to kernel side addressing changes.
90 *      Alan Cox      :      Beginning work on TCP fastpathing
91 *                      (not yet usable)
92 *      Arnt Gulbrandsen:      Turbocharged tcp_check() routine.
93 *      Alan Cox      :      TCP fast path debugging
94 *      Alan Cox      :      Window clamping
95 *      Michael Riepe  :      Bug in tcp_check()
96 *      Matt Dillon   :      More TCP improvements and RST bug fixes
97 *      Matt Dillon   :      Yet more small nasties remove from the
98 *                      TCP code (Be very nice to this man if
99 *                      tcp finally works 100%) 8)
100 *      Alan Cox      :      BSD accept semantics.
101 *      Alan Cox      :      Reset on closedown bug.
102 *      Peter De Schrijver :      ENOTCONN check missing in tcp_sendto().
103 *      Michael Pall   :      Handle poll() after URG properly in
104 *                      all cases.
105 *      Michael Pall   :      Undo the last fix in tcp_read_urg()
106 *                      (multi URG PUSH broke rlogin).
107 *      Michael Pall   :      Fix the multi URG PUSH problem in
108 *                      tcp_readable(), poll() after URG
109 *                      works now.
110 *      Michael Pall   :      recv(...,MSG_OOB) never blocks in the
111 *                      BSD api.
112 *      Alan Cox      :      Changed the semantics of sk->socket to
113 *                      fix a race and a signal problem with
114 *                      accept() and async I/O.
115 *      Alan Cox      :      Relaxed the rules on tcp_sendto().
116 *      Yury Shevchuk  :      Really fixed accept() blocking problem.
117 *      Craig I. Hagan :      Allow for BSD compatible TIME_WAIT for
118 *                      clients/servers which listen in on
119 *                      fixed ports.
120 *      Alan Cox      :      Cleaned the above up and shrank it to
121 *                      a sensible code size.
122 *      Alan Cox      :      Self connect lockup fix.
123 *      Alan Cox      :      No connect to multicast.
124 *      Ross Biro      :      Close unaccepted children on master
125 *                      socket close.
126 *      Alan Cox      :      Reset tracing code.
127 *      Alan Cox      :      Spurious resets on shutdown.
128 *      Alan Cox      :      Giant 15 minute/60 second timer error
129 *      Alan Cox      :      Small whoops in polling before an
130 *                      accept.
131 *      Alan Cox      :      Kept the state trace facility since
132 *                      it's handy for debugging.
133 *      Alan Cox      :      More reset handler fixes.
134 *      Alan Cox      :      Started rewriting the code based on
135 *                      the RFC's for other useful protocol
136 *                      references see: Comer, KA9Q NOS, and
137 *                      for a reference on the difference
138 *                      between specifications and how BSD
139 *                      works see the 4.4Lite source.
140 *      A.N.Kuznetsov  :      Don't time wait on completion of tidy
141 *                      close.
142 *      Linus Torvalds :      Fin/Shutdown & copied_seq changes.
143 *      Linus Torvalds :      Fixed BSD port reuse to work first syn
144 *      Alan Cox      :      Reimplemented timers as per the RFC
145 *                      and using multiple timers for sanity.
146 *      Alan Cox      :      Small bug fixes, and a lot of new
147 *                      comments.
148 *      Alan Cox      :      Fixed dual reader crash by locking
149 *                      the buffers (much like datagram.c)
150 *      Alan Cox      :      Fixed stuck sockets in probe. A probe

```

```

151 *                               now gets fed up of retrying without
152 *                               (even a no space) answer.
153 *                               Alan Cox      :   Extracted closing code better
154 *                               Alan Cox      :   Fixed the closing state machine to
155 *                               :               resemble the RFC.
156 *                               Alan Cox      :   More 'per spec' fixes.
157 *                               Jorge Cwik    :   Even faster checksumming.
158 *                               Alan Cox      :   tcp_data() doesn't ack illegal PSH
159 *                               :               only frames. At least one pc tcp stack
160 *                               :               generates them.
161 *                               Alan Cox      :   Cache last socket.
162 *                               Alan Cox      :   Per route irtt.
163 *                               Matt Day      :   poll()->select() match BSD precisely on error
164 *                               Alan Cox      :   New buffers
165 *                               Marc Tamsky   :   Various sk->prot->retransmits and
166 *                               :               sk->retransmits misupdating fixed.
167 *                               :               Fixed tcp_write_timeout: stuck close,
168 *                               :               and TCP syn retries gets used now.
169 *                               Mark Yarvis   :   In tcp_read_wakeup(), don't send an
170 *                               :               ack if state is TCP_CLOSED.
171 *                               Alan Cox      :   Look up device on a retransmit - routes may
172 *                               :               change. Doesn't yet cope with MSS shrink right
173 *                               :               but it's a start!
174 *                               Marc Tamsky   :   Closing in closing fixes.
175 *                               Mike Shaver   :   RFC1122 verifications.
176 *                               Alan Cox      :   rcv_saddr errors.
177 *                               Alan Cox      :   Block double connect().
178 *                               Alan Cox      :   Small hooks for enSKIP.
179 *                               Alexey Kuznetsov:   Path MTU discovery.
180 *                               Alan Cox      :   Support soft errors.
181 *                               Alan Cox      :   Fix MTU discovery pathological case
182 *                               :               when the remote claims no mtu!
183 *                               Marc Tamsky   :   TCP_CLOSE fix.
184 *                               Colin (G3TNE) :   Send a reset on syn ack replies in
185 *                               :               window but wrong (fixes NT Lpd problems)
186 *                               Pedro Roque   :   Better TCP window handling, delayed ack.
187 *                               Joerg Reuter  :   No modification of locked buffers in
188 *                               :               tcp_do_retransmit()
189 *                               Eric Schenk    :   Changed receiver side silly window
190 *                               :               avoidance algorithm to BSD style
191 *                               :               algorithm. This doubles throughput
192 *                               :               against machines running Solaris,
193 *                               :               and seems to result in general
194 *                               :               improvement.
195 *                               Stefan Magdalinski :   adjusted tcp_readable() to fix FIONREAD
196 *                               Willy Konynenberg :   Transparent proxying support.
197 *                               Mike McLagan   :   Routing by source
198 *                               Keith Owens   :   Do proper merging with partial SKB's in
199 *                               :               tcp_do_sendmsg to avoid burstiness.
200 *                               Eric Schenk    :   Fix fast close down bug with
201 *                               :               shutdown() followed by close().
202 *                               Andi Kleen     :   Make poll agree with SIGIO
203 *                               Salvatore Sanfilippo :   Support SO_LINGER with linger == 1 and
204 *                               :               lingertime == 0 (RFC 793 ABORT Call)
205 *                               Hirokazu Takahashi :   Use copy_from_user() instead of
206 *                               :               csum_and_copy_from_user() if possible.
207 *
208 *   This program is free software; you can redistribute it and/or
209 *   modify it under the terms of the GNU General Public License
210 *   as published by the Free Software Foundation; either version
211 *   2 of the License, or(at your option) any later version.
212 *
213 * Description of States:
214 *
215 *   TCP_SYN_SENT      sent a connection request, waiting for ack
216 *
217 *   TCP_SYN_RECV      received a connection request, sent ack,
218 *                       waiting for final ack in three-way handshake.
219 *
220 *   TCP_ESTABLISHED   connection established
221 *
222 *   TCP_FIN_WAIT1     our side has shutdown, waiting to complete
223 *                       transmission of remaining buffered data
224 *
225 *   TCP_FIN_WAIT2     all buffered data sent, waiting for remote
226 *                       to shutdown
227 *
228 *   TCP_CLOSING       both sides have shutdown but we still have
229 *                       data we have to finish sending
230 *
231 *   TCP_TIME_WAIT     timeout to catch resent junk before entering
232 *                       closed, can only be entered from FIN_WAIT2
233 *                       or CLOSING. Required because the other end
234 *                       may not have gotten our last ACK causing it
235 *                       to retransmit the data packet (which we ignore)

```

```

236 *
237 *      TCP_CLOSE_WAIT      remote side has shutdown and is waiting for
238 *                          us to finish writing our data and to shutdown
239 *                          (we have to close() to move on to LAST_ACK)
240 *
241 *      TCP_LAST_ACK        out side has shutdown after remote has
242 *                          shutdown. There may still be data in our
243 *                          buffer that we have to finish sending
244 *
245 *      TCP_CLOSE           socket is finished
246 */
247
248 #define pr_fmt(fmt) "TCP: " fmt
249
250 #include <linux/kernel.h>
251 #include <linux/module.h>
252 #include <linux/types.h>
253 #include <linux/fcntl.h>
254 #include <linux/poll.h>
255 #include <linux/init.h>
256 #include <linux/fs.h>
257 #include <linux/skbuff.h>
258 #include <linux/scatterlist.h>
259 #include <linux/splice.h>
260 #include <linux/net.h>
261 #include <linux/socket.h>
262 #include <linux/random.h>
263 #include <linux/bootmem.h>
264 #include <linux/highmem.h>
265 #include <linux/swap.h>
266 #include <linux/cache.h>
267 #include <linux/err.h>
268 #include <linux/crypto.h>
269 #include <linux/time.h>
270 #include <linux/slab.h>
271
272 #include <net/icmp.h>
273 #include <net/inet_common.h>
274 #include <net/tcp.h>
275 #include <net/xfrm.h>
276 #include <net/ip.h>
277 #include <net/netdma.h>
278 #include <net/sock.h>
279
280 #include <asm/uaccess.h>
281 #include <asm/ioctls.h>
282 #include <net/busy_poll.h>
283
284 int sysctl_tcp_fin_timeout __read_mostly = TCP_FIN_TIMEOUT;
285
286 int sysctl_tcp_min_tso_segs __read_mostly = 2;
287
288 int sysctl_tcp_autocorking __read_mostly = 1;
289
290 struct percpu_counter tcp_orphan_count;
291 EXPORT_SYMBOL_GPL(tcp_orphan_count);
292
293 long sysctl_tcp_mem[3] __read_mostly;
294 int sysctl_tcp_wmem[3] __read_mostly;
295 int sysctl_tcp_rmem[3] __read_mostly;
296
297 EXPORT_SYMBOL(sysctl_tcp_mem);
298 EXPORT_SYMBOL(sysctl_tcp_rmem);
299 EXPORT_SYMBOL(sysctl_tcp_wmem);
300
301 atomic_long_t tcp_memory_allocated; /* Current allocated memory. */
302 EXPORT_SYMBOL(tcp_memory_allocated);
303
304 /*
305  * Current number of TCP sockets.
306  */
307 struct percpu_counter tcp_sockets_allocated;
308 EXPORT_SYMBOL(tcp_sockets_allocated);
309
310 /*
311  * TCP splice context
312  */
313 struct tcp_splice_state {
314     struct pipe_inode_info *pipe;
315     size_t len;
316     unsigned int flags;
317 };
318
319 /*
320  * Pressure flag: try to collapse.

```

```

321  * Technical note: it is used by multiple contexts non atomically.
322  * All the __sk_mem_schedule() is of this nature: accounting
323  * is strict, actions are advisory and have some latency.
324  */
325  int tcp_memory_pressure __read_mostly;
326  EXPORT_SYMBOL(tcp_memory_pressure);
327
328  void tcp_enter_memory_pressure(struct sock *sk)
329  {
330      if (!tcp_memory_pressure) {
331          NET_INC_STATS(sock_net(sk), LINUX_MIB_TCPMEMORYPRESSURES);
332          tcp_memory_pressure = 1;
333      }
334  }
335  EXPORT_SYMBOL(tcp_enter_memory_pressure);
336
337  /* Convert seconds to retransmits based on initial and max timeout */
338  static u8 secs_to_retrans(int seconds, int timeout, int rto_max)
339  {
340      u8 res = 0;
341
342      if (seconds > 0) {
343          int period = timeout;
344
345          res = 1;
346          while (seconds > period && res < 255) {
347              res++;
348              timeout <= 1;
349              if (timeout > rto_max)
350                  timeout = rto_max;
351              period += timeout;
352          }
353      }
354      return res;
355  }
356
357  /* Convert retransmits to seconds based on initial and max timeout */
358  static int retrans_to_secs(u8 retrans, int timeout, int rto_max)
359  {
360      int period = 0;
361
362      if (retrans > 0) {
363          period = timeout;
364          while (--retrans) {
365              timeout <= 1;
366              if (timeout > rto_max)
367                  timeout = rto_max;
368              period += timeout;
369          }
370      }
371      return period;
372  }
373
374  /* Address-family independent initialization for a tcp_sock.
375  *
376  * NOTE: A lot of things set to zero explicitly by call to
377  * sk_alloc() so need not be done here.
378  */
379  void tcp_init_sock(struct sock *sk)
380  {
381      struct inet_connection_sock *icsk = inet_csk(sk);
382      struct tcp_sock *tp = tcp_sk(sk);
383
384      __skb_queue_head_init(&tp->out_of_order_queue);
385      tcp_init_xmit_timers(sk);
386      tcp_prequeue_init(tp);
387      INIT_LIST_HEAD(&tp->tsq_node);
388
389      icsk->icsk_rto = TCP_TIMEOUT_INIT;
390      tp->mdev_us = jiffies_to_usecs(TCP_TIMEOUT_INIT);
391
392      /* So many TCP implementations out there (incorrectly) count the
393       * initial SYN frame in their delayed-ACK and congestion control
394       * algorithms that we must have the following bandaid to talk
395       * efficiently to them. -DaveM
396       */
397      tp->snd_cwnd = TCP_INIT_CWND;
398
399      /* See draft-stevens-tcpca-spec-01 for discussion of the
400       * initialization of these values.
401       */
402      tp->snd_ssthresh = TCP_INFINITE_SSTHRESH;
403      tp->snd_cwnd_clamp = ~0;
404      tp->mss_cache = TCP_MSS_DEFAULT;
405

```

```

406 tp->reordering = sysctl\_tcp\_reordering;
407 tcp\_enable\_early\_retrans(tp);
408 icsk->icsk_ca_ops = &tcp\_init\_congestion\_ops;
409
410 tp->tsoffset = 0;
411
412 sk->sk\_state = TCP_CLOSE;
413
414 sk->sk_write_space = sk\_stream\_write\_space;
415 sock\_set\_flag(sk, SOCK_USE_WRITE_QUEUE);
416
417 icsk->icsk_sync_mss = tcp\_sync\_mss;
418
419 sk->sk_sndbuf = sysctl\_tcp\_wmem[1];
420 sk->sk_rcvbuf = sysctl\_tcp\_rmem[1];
421
422 local\_bh\_disable();
423 sock\_update\_memcg(sk);
424 sk\_sockets\_allocated\_inc(sk);
425 local\_bh\_enable();
426 }
427 EXPORT_SYMBOL(tcp\_init\_sock);
428
429 static void tcp\_tx\_timestamp(struct sock *sk, struct sk\_buff *skb)
430 {
431     if (sk->sk_tsflags) {
432         struct skb\_shared\_info *shinfo = skb\_shinfo(skb);
433
434         sock\_tx\_timestamp(sk, &shinfo->tx_flags);
435         if (shinfo->tx_flags & SKBTX_ANY_TSTAMP)
436             shinfo->tskey = TCP\_SKB\_CB(skb)->seq + skb->len - 1;
437     }
438 }
439
440 /*
441  *      Wait for a TCP event.
442  *
443  *      Note that we don't need to lock the socket, as the upper poll layers
444  *      take care of normal races (between the test and the event) and we don't
445  *      go look at any of the socket buffers directly.
446  */
447 unsigned int tcp\_poll(struct file *file, struct socket *sock, poll\_table *wait)
448 {
449     unsigned int mask;
450     struct sock *sk = sock->sk;
451     const struct tcp\_sock *tp = tcp\_sk(sk);
452
453     sock\_rps\_record\_flow(sk);
454
455     sock\_poll\_wait(file, sk\_sleep(sk), wait);
456     if (sk->sk\_state == TCP_LISTEN)
457         return inet\_csk\_listen\_poll(sk);
458
459     /* Socket is not locked. We are protected from async events
460      * by poll logic and correct handling of state changes
461      * made by other threads is impossible in any case.
462      */
463
464     mask = 0;
465
466     /*
467      * POLLHUP is certainly not done right. But poll() doesn't
468      * have a notion of HUP in just one direction, and for a
469      * socket the read side is more interesting.
470      *
471      * Some poll() documentation says that POLLHUP is incompatible
472      * with the POLLOUT/POLLWR flags, so somebody should check this
473      * all. But careful, it tends to be safer to return too many
474      * bits than too few, and you can easily break real applications
475      * if you don't tell them that something has hung up!
476      *
477      * Check-me.
478      *
479      * Check number 1. POLLHUP is _UNMASKABLE_ event (see UNIX98 and
480      * our fs/select.c). It means that after we received EOF,
481      * poll always returns immediately, making impossible poll() on write()
482      * in state CLOSE_WAIT. One solution is evident --- to set POLLHUP
483      * if and only if shutdown has been made in both directions.
484      * Actually, it is interesting to look how Solaris and DUX
485      * solve this dilemma. I would prefer, if POLLHUP were maskable,
486      * then we could set it on SND_SHUTDOWN. BTW examples given
487      * in Stevens' books assume exactly this behaviour, it explains
488      * why POLLHUP is incompatible with POLLOUT. --ANK
489      *
490      * NOTE. Check for TCP_CLOSE is added. The goal is to prevent

```



```

491      * blocking on fresh not-connected or disconnected socket. --ANK
492      */
493      if (sk->sk_shutdown == SHUTDOWN_MASK || sk->sk_state == TCP_CLOSE)
494          mask |= POLLHUP;
495      if (sk->sk_shutdown & RCV_SHUTDOWN)
496          mask |= POLLIN | POLLRDNORM | POLLRDHUP;
497
498      /* Connected or passive Fast Open socket? */
499      if (sk->sk_state != TCP_SYN_SENT &&
500          (sk->sk_state != TCP_SYN_RECV || tp->fastopen_rsk != NULL)) {
501          int target = sock_rcvlowat(sk, 0, INT_MAX);
502
503          if (tp->urg_seq == tp->copied_seq &&
504              !sock_flag(sk, SOCK_URGINLINE) &&
505              tp->urg_data)
506              target++;
507
508          /* Potential race condition. If read of tp below will
509             * escape above sk->sk_state, we can be illegally awoken
510             * in SYN_* states. */
511          if (tp->rcv_nxt - tp->copied_seq >= target)
512              mask |= POLLIN | POLLRDNORM;
513
514          if (!(sk->sk_shutdown & SEND_SHUTDOWN)) {
515              if (sk_stream_is_writeable(sk)) {
516                  mask |= POLLOUT | POLLWRNORM;
517              } else { /* send SIGIO later */
518                  set_bit(SOCK_ASYNC_NOSPACE,
519                          &sk->sk_socket->flags);
520                  set_bit(SOCK_NOSPACE, &sk->sk_socket->flags);
521
522                  /* Race breaker. If space is freed after
523                     * wspace test but before the flags are set,
524                     * IO signal will be lost.
525                     */
526                  if (sk_stream_is_writeable(sk))
527                      mask |= POLLOUT | POLLWRNORM;
528              }
529          } else
530              mask |= POLLOUT | POLLWRNORM;
531
532          if (tp->urg_data & TCP_URG_VALID)
533              mask |= POLLPRI;
534      }
535      /* This barrier is coupled with smp_wmb() in tcp_reset() */
536      smp_rmb();
537      if (sk->sk_err || !skb_queue_empty(&sk->sk_error_queue))
538          mask |= POLLERR;
539
540      return mask;
541 }
542 EXPORT_SYMBOL(tcp_poll);
543
544 int tcp_ioctl(struct sock *sk, int cmd, unsigned long arg)
545 {
546     struct tcp_sock *tp = tcp_sk(sk);
547     int ans;
548     bool slow;
549
550     switch (cmd) {
551     case SIOCINQ:
552         if (sk->sk_state == TCP_LISTEN)
553             return -EINVAL;
554
555         slow = lock_sock_fast(sk);
556         if ((1 << sk->sk_state) & (TCPF_SYN_SENT | TCPF_SYN_RECV))
557             ans = 0;
558         else if (sock_flag(sk, SOCK_URGINLINE) ||
559                 !tp->urg_data ||
560                 before(tp->urg_seq, tp->copied_seq) ||
561                 !before(tp->urg_seq, tp->rcv_nxt)) {
562
563             ans = tp->rcv_nxt - tp->copied_seq;
564
565             /* Subtract 1, if FIN was received */
566             if (ans && sock_flag(sk, SOCK_DONE))
567                 ans--;
568         } else
569             ans = tp->urg_seq - tp->copied_seq;
570         unlock_sock_fast(sk, slow);
571         break;
572     case SIOCATMARK:
573         ans = tp->urg_data && tp->urg_seq == tp->copied_seq;
574         break;
575     case SIOCOUTQ:

```

```

576         if (sk->sk_state == TCP_LISTEN)
577             return -EINVAL;
578
579         if ((1 << sk->sk_state) & (TCPF_SYN_SENT | TCPF_SYN_RECV))
580             answ = 0;
581         else
582             answ = tp->write_seq - tp->snd_una;
583         break;
584     case SIOCOUTQSD:
585         if (sk->sk_state == TCP_LISTEN)
586             return -EINVAL;
587
588         if ((1 << sk->sk_state) & (TCPF_SYN_SENT | TCPF_SYN_RECV))
589             answ = 0;
590         else
591             answ = tp->write_seq - tp->snd_nxt;
592         break;
593     default:
594         return -ENOIOCTLCMD;
595     }
596
597     return put_user(answ, (int __user *)arg);
598 }
599 EXPORT_SYMBOL(tcp_ioctl);
600
601 static inline void tcp_mark_push(struct tcp_sock *tp, struct sk_buff *skb)
602 {
603     TCP_SKB_CB(skb)->tcp_flags |= TCPHDR_PSH;
604     tp->pushed_seq = tp->write_seq;
605 }
606
607 static inline bool forced_push(const struct tcp_sock *tp)
608 {
609     return after(tp->write_seq, tp->pushed_seq + (tp->max_window >> 1));
610 }
611
612 static inline void skb_entail(struct sock *sk, struct sk_buff *skb)
613 {
614     struct tcp_sock *tp = tcp_sk(sk);
615     struct tcp_skb_cb *tcb = TCP_SKB_CB(skb);
616
617     skb->csum = 0;
618     tcb->seq = tcb->end_seq = tp->write_seq;
619     tcb->tcp_flags = TCPHDR_ACK;
620     tcb->sacked = 0;
621     skb_header_release(skb);
622     tcp_add_write_queue_tail(sk, skb);
623     sk->sk_wmem_queued += skb->truesize;
624     sk_mem_charge(sk, skb->truesize);
625     if (tp->nonagle & TCP_NAGLE_PUSH)
626         tp->nonagle &= ~TCP_NAGLE_PUSH;
627 }
628
629 static inline void tcp_mark_urg(struct tcp_sock *tp, int flags)
630 {
631     if (flags & MSG_OOB)
632         tp->snd_up = tp->write_seq;
633 }
634
635 /* If a not yet filled skb is pushed, do not send it if
636  * we have data packets in Qdisc or NIC queues :
637  * Because TX completion will happen shortly, it gives a chance
638  * to coalesce future sendmsg() payload into this skb, without
639  * need for a timer, and with no latency trade off.
640  * As packets containing data payload have a bigger truesize
641  * than pure acks (dataless) packets, the last checks prevent
642  * autocorking if we only have an ACK in Qdisc/NIC queues,
643  * or if TX completion was delayed after we processed ACK packet.
644  */
645 static bool tcp_should_autocork(struct sock *sk, struct sk_buff *skb,
646                                int size_goal)
647 {
648     return skb->len < size_goal &&
649         sysctl_tcp_autocorking &&
650         skb != tcp_write_queue_head(sk) &&
651         atomic_read(&sk->sk_wmem_alloc) > skb->truesize;
652 }
653
654 static void tcp_push(struct sock *sk, int flags, int mss_now,
655                     int nonagle, int size_goal)
656 {
657     struct tcp_sock *tp = tcp_sk(sk);
658     struct sk_buff *skb;
659
660     if (!tcp_send_head(sk))

```



```

661         return;
662
663         skb = tcp_write_queue_tail(sk);
664         if (!(flags & MSG_MORE) || forced_push(tp))
665             tcp_mark_push(tp, skb);
666
667         tcp_mark_urg(tp, flags);
668
669         if (tcp_should_autocork(sk, skb, size_goal)) {
670
671             /* avoid atomic op if TSQ_THROTTLED bit is already set */
672             if (!test_bit(TSQ_THROTTLED, &tp->tsq_flags)) {
673                 NET_INC_STATS(sock_net(sk), LINUX_MIB_TCPAUTOCORKING);
674                 set_bit(TSQ_THROTTLED, &tp->tsq_flags);
675             }
676             /* It is possible TX completion already happened
677              * before we set TSQ_THROTTLED.
678              */
679             if (atomic_read(&sk->sk_wmem_alloc) > skb->truesize)
680                 return;
681         }
682
683         if (flags & MSG_MORE)
684             nonagle = TCP_NAGLE_CORK;
685
686         __tcp_push_pending_frames(sk, mss_now, nonagle);
687     }
688
689 static int tcp_splice_data_recv(read_descriptor_t *rd_desc, struct sk_buff *skb,
690                               unsigned int offset, size_t len)
691 {
692     struct tcp_splice_state *tss = rd_desc->arg.data;
693     int ret;
694
695     ret = skb_splice_bits(skb, offset, tss->pipe, min(rd_desc->count, len),
696                          tss->flags);
697     if (ret > 0)
698         rd_desc->count -= ret;
699     return ret;
700 }
701
702 static int __tcp_splice_read(struct sock *sk, struct tcp_splice_state *tss)
703 {
704     /* Store TCP splice context information in read_descriptor_t. */
705     read_descriptor_t rd_desc = {
706         .arg.data = tss,
707         .count = tss->len,
708     };
709
710     return tcp_read_sock(sk, &rd_desc, tcp_splice_data_recv);
711 }
712
713 /**
714  * tcp_splice_read - splice data from TCP socket to a pipe
715  * @sock:      socket to splice from
716  * @ppos:      position (not valid)
717  * @pipe:      pipe to splice to
718  * @len:       number of bytes to splice
719  * @flags:     splice modifier flags
720  *
721  * Description:
722  *   Will read pages from given socket and fill them into a pipe.
723  */
724
725 ssize_t tcp_splice_read(struct socket *sock, loff_t *ppos,
726                        struct pipe_inode_info *pipe, size_t len,
727                        unsigned int flags)
728 {
729     struct sock *sk = sock->sk;
730     struct tcp_splice_state tss = {
731         .pipe = pipe,
732         .len = len,
733         .flags = flags,
734     };
735     long timeo;
736     ssize_t spliced;
737     int ret;
738
739     sock_rps_record_flow(sk);
740     /*
741      * We can't seek on a socket input
742      */
743     if (unlikely(*ppos))
744         return -ESPIPE;
745

```

```

746     ret = spliced = 0;
747
748     lock_sock(sk);
749
750     timeo = sock_rcvtimeo(sk, sock->file->f_flags & O_NONBLOCK);
751     while (tss.len) {
752         ret = tcp_splice_read(sk, &tss);
753         if (ret < 0)
754             break;
755         else if (!ret) {
756             if (spliced)
757                 break;
758             if (sock_flag(sk, SOCK_DONE))
759                 break;
760             if (sk->sk_err) {
761                 ret = sock_error(sk);
762                 break;
763             }
764             if (sk->sk_shutdown & RCV_SHUTDOWN)
765                 break;
766             if (sk->sk_state == TCP_CLOSE) {
767                 /*
768                  * This occurs when user tries to read
769                  * from never connected socket.
770                  */
771                 if (!sock_flag(sk, SOCK_DONE))
772                     ret = -ENOTCONN;
773                 break;
774             }
775             if (!timeo) {
776                 ret = -EAGAIN;
777                 break;
778             }
779             sk_wait_data(sk, &timeo);
780             if (signal_pending(current)) {
781                 ret = sock_intr_errno(timeo);
782                 break;
783             }
784             continue;
785         }
786         tss.len -= ret;
787         spliced += ret;
788
789         if (!timeo)
790             break;
791         release_sock(sk);
792         lock_sock(sk);
793
794         if (sk->sk_err || sk->sk_state == TCP_CLOSE ||
795             (sk->sk_shutdown & RCV_SHUTDOWN) ||
796             signal_pending(current))
797             break;
798     }
799
800     release_sock(sk);
801
802     if (spliced)
803         return spliced;
804
805     return ret;
806 }
807 EXPORT_SYMBOL(tcp_splice_read);
808
809 struct sk_buff *sk_stream_alloc_skb(struct sock *sk, int size, gfp_t gfp)
810 {
811     struct sk_buff *skb;
812
813     /* The TCP header must be at least 32-bit aligned. */
814     size = ALIGN(size, 4);
815
816     skb = alloc_skb_fclone(size + sk->sk_prot->max_header, gfp);
817     if (skb) {
818         if (sk_wmem_schedule(sk, skb->truesize)) {
819             skb_reserve(skb, sk->sk_prot->max_header);
820             /*
821              * Make sure that we have exactly size bytes
822              * available to the caller, no more, no less.
823              */
824             skb->reserved_tailroom = skb->end - skb->tail - size;
825             return skb;
826         }
827         kfree_skb(skb);
828     } else {
829         sk->sk_prot->enter_memory_pressure(sk);
830         sk_stream_moderate_sndbuf(sk);

```

```

831     }
832     return NULL;
833 }
834
835 static unsigned int tcp_xmit_size_goal(struct sock *sk, u32 mss_now,
836                                         int large_allowed)
837 {
838     struct tcp_sock *tp = tcp_sk(sk);
839     u32 xmit_size_goal, old_size_goal;
840
841     xmit_size_goal = mss_now;
842
843     if (large_allowed && sk_can_gso(sk)) {
844         u32 gso_size, hlen;
845
846         /* Maybe we should/could use sk->sk_prot->max_header here ? */
847         hlen = inet_csk(sk)->icsk_af_ops->net_header_len +
848             inet_csk(sk)->icsk_ext_hdr_len +
849             tp->tcp_header_len;
850
851         /* Goal is to send at least one packet per ms,
852          * not one big TSO packet every 100 ms.
853          * This preserves ACK clocking and is consistent
854          * with tcp_tso_should_defer() heuristic.
855          */
856         gso_size = sk->sk_pacing_rate / (2 * MSEC_PER_SEC);
857         gso_size = max_t(u32, gso_size,
858                         sysctl_tcp_min_tso_segs * mss_now);
859
860         xmit_size_goal = min_t(u32, gso_size,
861                                sk->sk_gso_max_size - 1 - hlen);
862
863         xmit_size_goal = tcp_bound_to_half_wnd(tp, xmit_size_goal);
864
865         /* We try hard to avoid divides here */
866         old_size_goal = tp->xmit_size_goal_segs * mss_now;
867
868         if (likely(old_size_goal <= xmit_size_goal &&
869                  old_size_goal + mss_now > xmit_size_goal)) {
870             xmit_size_goal = old_size_goal;
871         } else {
872             tp->xmit_size_goal_segs =
873                 min_t(u16, xmit_size_goal / mss_now,
874                       sk->sk_gso_max_segs);
875             xmit_size_goal = tp->xmit_size_goal_segs * mss_now;
876         }
877     }
878
879     return max(xmit_size_goal, mss_now);
880 }
881
882 static int tcp_send_mss(struct sock *sk, int *size_goal, int flags)
883 {
884     int mss_now;
885
886     mss_now = tcp_current_mss(sk);
887     *size_goal = tcp_xmit_size_goal(sk, mss_now, !(flags & MSG_OOB));
888
889     return mss_now;
890 }
891
892 static ssize_t do_tcp_sendpages(struct sock *sk, struct page *page, int offset,
893                                 size_t size, int flags)
894 {
895     struct tcp_sock *tp = tcp_sk(sk);
896     int mss_now, size_goal;
897     int err;
898     ssize_t copied;
899     long timeo = sock_sndtimeo(sk, flags & MSG_DONTWAIT);
900
901     /* Wait for a connection to finish. One exception is TCP Fast Open
902      * (passive side) where data is allowed to be sent before a connection
903      * is fully established.
904      */
905     if (((1 << sk->sk_state) & ~(TCPF_ESTABLISHED | TCPF_CLOSE_WAIT)) &&
906         !tcp_passive_fastopen(sk)) {
907         if ((err = sk_stream_wait_connect(sk, &timeo)) != 0)
908             goto out_err;
909     }
910
911     clear_bit(SOCK_ASYNC_NOSPACE, &sk->sk_socket->flags);
912
913     mss_now = tcp_send_mss(sk, &size_goal, flags);
914     copied = 0;
915

```

```

916     err = -EPIPE;
917     if (sk->sk_err || (sk->sk_shutdown & SEND_SHUTDOWN))
918         goto out_err;
919
920     while (size > 0) {
921         struct sk_buff *skb = tcp_write_queue_tail(sk);
922         int copy, i;
923         bool can_coalesce;
924
925         if (!tcp_send_head(sk) || (copy = size_goal - skb->len) <= 0) {
926 new_segment:
927             if (!sk_stream_memory_free(sk))
928                 goto wait_for_sndbuf;
929
930             skb = sk_stream_alloc_skb(sk, 0, sk->sk_allocation);
931             if (!skb)
932                 goto wait_for_memory;
933
934             skb_entail(sk, skb);
935             copy = size_goal;
936         }
937
938         if (copy > size)
939             copy = size;
940
941         i = skb_shinfo(skb)->nr_frags;
942         can_coalesce = skb_can_coalesce(skb, i, page, offset);
943         if (!can_coalesce && i >= MAX_SKB_FRAGS) {
944             tcp_mark_push(tp, skb);
945             goto new_segment;
946         }
947         if (!sk_wmem_schedule(sk, copy))
948             goto wait_for_memory;
949
950         if (can_coalesce) {
951             skb_frag_size_add(&skb_shinfo(skb)->frags[i - 1], copy);
952         } else {
953             get_page(page);
954             skb_fill_page_desc(skb, i, page, offset, copy);
955         }
956         skb_shinfo(skb)->tx_flags |= SKBTX_SHARED_FRAG;
957
958         skb->len += copy;
959         skb->data_len += copy;
960         skb->truesize += copy;
961         sk->sk_wmem_queued += copy;
962         sk_mem_charge(sk, copy);
963         skb->ip_summed = CHECKSUM_PARTIAL;
964         tp->write_seq += copy;
965         TCP_SKB_CB(skb)->end_seq += copy;
966         skb_shinfo(skb)->gso_segs = 0;
967
968         if (!copied)
969             TCP_SKB_CB(skb)->tcp_flags &= ~TCPHDR_PSH;
970
971         copied += copy;
972         offset += copy;
973         if (!(size -= copy)) {
974             tcp_tx_timestamp(sk, skb);
975             goto out;
976         }
977
978         if (skb->len < size_goal || (flags & MSG_OOB))
979             continue;
980
981         if (forced_push(tp)) {
982             tcp_mark_push(tp, skb);
983             tcp_push_pending_frames(sk, mss_now, TCP_NAGLE_PUSH);
984         } else if (skb == tcp_send_head(sk))
985             tcp_push_one(sk, mss_now);
986         continue;
987
988 wait_for_sndbuf:
989     set_bit(SOCK_NOSPACE, &sk->sk_socket->flags);
990 wait_for_memory:
991     tcp_push(sk, flags & ~MSG_MORE, mss_now,
992             TCP_NAGLE_PUSH, size_goal);
993
994     if ((err = sk_stream_wait_memory(sk, &timeo)) != 0)
995         goto do_error;
996
997     mss_now = tcp_send_mss(sk, &size_goal, flags);
998 }
999
1000 out:

```

```

1001     if (copied && !(flags & MSG_SENDPAGE_NOTLAST))
1002         tcp_push(sk, flags, mss_now, tp->nonagle, size_goal);
1003     return copied;
1004
1005 do_error:
1006     if (copied)
1007         goto out;
1008 out_err:
1009     return sk_stream_error(sk, flags, err);
1010 }
1011
1012 int tcp_sendpage(struct sock *sk, struct page *page, int offset,
1013                 size_t size, int flags)
1014 {
1015     ssize_t res;
1016
1017     if (!(sk->sk_route_caps & NETIF_F_SG) ||
1018         !(sk->sk_route_caps & NETIF_F_ALL_CSUM))
1019         return sock_no_sendpage(sk->sk_socket, page, offset, size,
1020                                 flags);
1021
1022     lock_sock(sk);
1023     res = do_tcp_sendpages(sk, page, offset, size, flags);
1024     release_sock(sk);
1025     return res;
1026 }
1027 EXPORT_SYMBOL(tcp_sendpage);
1028
1029 static inline int select_size(const struct sock *sk, bool sg)
1030 {
1031     const struct tcp_sock *tp = tcp_sk(sk);
1032     int tmp = tp->mss_cache;
1033
1034     if (sg) {
1035         if (sk_can_gso(sk)) {
1036             /* Small frames wont use a full page:
1037              * Payload will immediately follow tcp header.
1038              */
1039             tmp = SKB_WITH_OVERHEAD(2048 - MAX_TCP_HEADER);
1040         } else {
1041             int pgbreak = SKB_MAX_HEAD(MAX_TCP_HEADER);
1042
1043             if (tmp >= pgbreak &&
1044                 tmp <= pgbreak + (MAX_SKB_FRAGS - 1) * PAGE_SIZE)
1045                 tmp = pgbreak;
1046         }
1047     }
1048
1049     return tmp;
1050 }
1051
1052 void tcp_free_fastopen_req(struct tcp_sock *tp)
1053 {
1054     if (tp->fastopen_req != NULL) {
1055         kfree(tp->fastopen_req);
1056         tp->fastopen_req = NULL;
1057     }
1058 }
1059
1060 static int tcp_sendmsg_fastopen(struct sock *sk, struct msghdr *msg,
1061                                int *copied, size_t size)
1062 {
1063     struct tcp_sock *tp = tcp_sk(sk);
1064     int err, flags;
1065
1066     if (!(sysctl_tcp_fastopen & TFO_CLIENT_ENABLE))
1067         return -EOPNOTSUPP;
1068     if (tp->fastopen_req != NULL)
1069         return -EALREADY; /* Another Fast Open is in progress */
1070
1071     tp->fastopen_req = kzalloc(sizeof(struct tcp_fastopen_request),
1072                               sk->sk_allocation);
1073     if (unlikely(tp->fastopen_req == NULL))
1074         return -ENOMEM;
1075     tp->fastopen_req->data = msg;
1076     tp->fastopen_req->size = size;
1077
1078     flags = (msg->msg_flags & MSG_DONTWAIT) ? 0_NONBLOCK : 0;
1079     err = __inet_stream_connect(sk->sk_socket, msg->msg_name,
1080                                msg->msg_namelen, flags);
1081     *copied = tp->fastopen_req->copied;
1082     tcp_free_fastopen_req(tp);
1083     return err;
1084 }
1085

```

```

1086 int tcp_sendmsg(struct kiocb *iocb, struct sock *sk, struct msghdr *msg,
1087                 size\_t size)
1088 {
1089     struct iovec *iov;
1090     struct tcp\_sock *tp = tcp\_sk(sk);
1091     struct sk\_buff *skb;
1092     int iovlen, flags, err, copied = 0;
1093     int mss_now = 0, size_goal, copied_syn = 0, offset = 0;
1094     bool sg;
1095     long timeo;
1096
1097     lock\_sock(sk);
1098
1099     flags = msg->msg_flags;
1100     if (flags & MSG_FASTOPEN) {
1101         err = tcp\_sendmsg\_fastopen(sk, msg, &copied_syn, size);
1102         if (err == -EINPROGRESS && copied_syn > 0)
1103             goto out;
1104         else if (err)
1105             goto out_err;
1106         offset = copied_syn;
1107     }
1108
1109     timeo = sock\_sndtimeo(sk, flags & MSG_DONTWAIT);
1110
1111     /* Wait for a connection to finish. One exception is TCP Fast Open
1112      * (passive side) where data is allowed to be sent before a connection
1113      * is fully established.
1114     */
1115     if (((1 << sk->sk_state) & ~(TCPF_ESTABLISHED | TCPF_CLOSE_WAIT)) &&
1116         !tcp\_passive\_fastopen(sk)) {
1117         if ((err = sk\_stream\_wait\_connect(sk, &timeo)) != 0)
1118             goto do\_error;
1119     }
1120
1121     if (unlikely(tp->repair)) {
1122         if (tp->repair_queue == TCP_RECV_QUEUE) {
1123             copied = tcp\_send\_rcvq(sk, msg, size);
1124             goto out_nopush;
1125         }
1126
1127         err = -EINVAL;
1128         if (tp->repair_queue == TCP_NO_QUEUE)
1129             goto out_err;
1130
1131         /* 'common' sending to sendq */
1132     }
1133
1134     /* This should be in poll */
1135     clear\_bit(SOCK_ASYNC_NOSPACE, &sk->sk_socket->flags);
1136
1137     mss_now = tcp\_send\_mss(sk, &size_goal, flags);
1138
1139     /* Ok commence sending. */
1140     iovlen = msg->msg_iovlen;
1141     iov = msg->msg_iov;
1142     copied = 0;
1143
1144     err = -EPIPE;
1145     if (sk->sk_err || (sk->sk_shutdown & SEND_SHUTDOWN))
1146         goto out_err;
1147
1148     sg = !(sk->sk_route_caps & NETIF_F_SG);
1149
1150     while (--iovlen >= 0) {
1151         size\_t seglen = iov->iov_len;
1152         unsigned char \_\_user *from = iov->iov_base;
1153
1154         iov++;
1155         if (unlikely(offset > 0)) { /* Skip bytes copied in SYN */
1156             if (offset >= seglen) {
1157                 offset -= seglen;
1158                 continue;
1159             }
1160             seglen -= offset;
1161             from += offset;
1162             offset = 0;
1163         }
1164
1165         while (seglen > 0) {
1166             int copy = 0;
1167             int max = size_goal;
1168
1169             skb = tcp\_write\_queue\_tail(sk);
1170             if (tcp\_send\_head(sk)) {

```



```

1171 if (skb->ip_summed == CHECKSUM_NONE)
1172     max = mss_now;
1173     copy = max - skb->len;
1174 }
1175
1176 if (copy <= 0) {
1177     new_segment:
1178         /* Allocate new segment. If the interface is SG,
1179          * allocate skb fitting to single page.
1180          */
1181         if (!sk_stream_memory_free(sk))
1182             goto wait_for_sndbuf;
1183
1184         skb = sk_stream_alloc_skb(sk,
1185                                 select_size(sk, sg),
1186                                 sk->sk_allocation);
1187
1188         if (!skb)
1189             goto wait_for_memory;
1190
1191         /*
1192          * Check whether we can use HW checksum.
1193          */
1194         if (sk->sk_route_caps & NETIF_F_ALL_CSUM)
1195             skb->ip_summed = CHECKSUM_PARTIAL;
1196
1197         skb_ensure_ownership(skb);
1198         copy = size_goal;
1199         max = size_goal;
1200
1201         /* All packets are restored as if they have
1202          * already been sent. skb_mstamp isn't set to
1203          * avoid wrong rtt estimation.
1204          */
1205         if (tp->repair)
1206             TCP_SKB_CB(skb)->sacked |= TCPCB_REPAIRED;
1207     }
1208
1209     /* Try to append data to the end of skb. */
1210     if (copy > seglen)
1211         copy = seglen;
1212
1213     /* Where to copy to? */
1214     if (skb_availroom(skb) > 0) {
1215         /* We have some space in skb head. Superb! */
1216         copy = min_t(int, copy, skb_availroom(skb));
1217         err = skb_add_data_nocache(sk, skb, from, copy);
1218         if (err)
1219             goto do_fault;
1220     } else {
1221         bool merge = true;
1222         int i = skb_shinfo(skb)->nr_frags;
1223         struct page_frag *pfrag = sk_page_frag(sk);
1224
1225         if (!sk_page_frag_refill(sk, pfrag))
1226             goto wait_for_memory;
1227
1228         if (!skb_can_coalesce(skb, i, pfrag->page,
1229                             pfrag->offset)) {
1230             if (i == MAX_SKB_FRAGS || !sg) {
1231                 tcp_mark_push(tp, skb);
1232                 goto new_segment;
1233             }
1234             merge = false;
1235         }
1236
1237         copy = min_t(int, copy, pfrag->size - pfrag->offset);
1238
1239         if (!sk_wmem_schedule(sk, copy))
1240             goto wait_for_memory;
1241
1242         err = skb_copy_to_page_nocache(sk, from, skb,
1243                                       pfrag->page,
1244                                       pfrag->offset,
1245                                       copy);
1246         if (err)
1247             goto do_error;
1248
1249         /* Update the skb. */
1250         if (merge) {
1251             skb_frag_size_add(&skb_shinfo(skb)->frags[i - 1], copy);
1252         } else {
1253             skb_fill_page_desc(skb, i, pfrag->page,
1254                               pfrag->offset, copy);
1255             get_page(pfrag->page);
1256         }
1257     }

```

```

1256         pfrag->offset += copy;
1257     }
1258
1259     if (!copied)
1260         TCP_SKB_CB(skb)->tcp_flags &= ~TCPHDR_PSH;
1261
1262     tp->write_seq += copy;
1263     TCP_SKB_CB(skb)->end_seq += copy;
1264     skb_shinfo(skb)->gso_segs = 0;
1265
1266     from += copy;
1267     copied += copy;
1268     if ((seglen -= copy) == 0 && iovlen == 0) {
1269         tcp_tx_timestamp(sk, skb);
1270         goto out;
1271     }
1272
1273     if (skb->len < max || (flags & MSG_OOB) || unlikely(tp->repair))
1274         continue;
1275
1276     if (forced_push(tp)) {
1277         tcp_mark_push(tp, skb);
1278         tcp_push_pending_frames(sk, mss_now, TCP_NAGLE_PUSH);
1279     } else if (skb == tcp_send_head(sk))
1280         tcp_push_one(sk, mss_now);
1281     continue;
1282
1283 wait_for_sndbuf:
1284     set_bit(SOCK_NOSPACE, &sk->sk_socket->flags);
1285 wait_for_memory:
1286     if (copied)
1287         tcp_push(sk, flags & ~MSG_MORE, mss_now,
1288                 TCP_NAGLE_PUSH, size_goal);
1289
1290     if ((err = sk_stream_wait_memory(sk, &timeo)) != 0)
1291         goto do_error;
1292
1293     mss_now = tcp_send_mss(sk, &size_goal, flags);
1294 }
1295 }
1296
1297 out:
1298     if (copied)
1299         tcp_push(sk, flags, mss_now, tp->nonagle, size_goal);
1300 out_nopush:
1301     release_sock(sk);
1302     return copied + copied_syn;
1303
1304 do_fault:
1305     if (!skb->len) {
1306         tcp_unlink_write_queue(skb, sk);
1307         /* It is the one place in all of TCP, except connection
1308          * reset, where we can be unlinking the send_head.
1309          */
1310         tcp_check_send_head(sk, skb);
1311         sk_wmem_free_skb(sk, skb);
1312     }
1313
1314 do_error:
1315     if (copied + copied_syn)
1316         goto out;
1317 out_err:
1318     err = sk_stream_error(sk, flags, err);
1319     release_sock(sk);
1320     return err;
1321 }
1322 EXPORT_SYMBOL(tcp_sendmsg);
1323
1324 /*
1325  *      Handle reading urgent data. BSD has very simple semantics for
1326  *      this, no blocking and very strange errors 8)
1327  */
1328
1329 static int tcp_recv_urg(struct sock *sk, struct msghdr *msg, int len, int flags)
1330 {
1331     struct tcp_sock *tp = tcp_sk(sk);
1332
1333     /* No URG data to read. */
1334     if (sock_flag(sk, SOCK_URGINLINE) || !tp->urg_data ||
1335         tp->urg_data == TCP_URG_READ)
1336         return -EINVAL; /* Yes this is right ! */
1337
1338     if (sk->sk_state == TCP_CLOSE && !sock_flag(sk, SOCK_DONE))
1339         return -ENOTCONN;
1340

```

```

1341 if (tp->urg_data & TCP_URG_VALID) {
1342     int err = 0;
1343     char c = tp->urg_data;
1344
1345     if (!(flags & MSG_PEEK))
1346         tp->urg_data = TCP_URG_READ;
1347
1348     /* Read urgent data. */
1349     msg->msg_flags |= MSG_OOB;
1350
1351     if (len > 0) {
1352         if (!(flags & MSG_TRUNC))
1353             err = memcpy_toiovec(msg->msg_iov, &c, 1);
1354         len = 1;
1355     } else
1356         msg->msg_flags |= MSG_TRUNC;
1357
1358     return err ? -EFAULT : len;
1359 }
1360
1361 if (sk->sk_state == TCP_CLOSE || (sk->sk_shutdown & RCV_SHUTDOWN))
1362     return 0;
1363
1364 /* Fixed the recv(..., MSG_OOB) behaviour. BSD docs and
1365  * the available implementations agree in this case:
1366  * this call should never block, independent of the
1367  * blocking state of the socket.
1368  * Mike <pall@rz.uni-karlsruhe.de>
1369  */
1370 return -EAGAIN;
1371 }
1372
1373 static int tcp_peek_sndq(struct sock *sk, struct msghdr *msg, int len)
1374 {
1375     struct sk_buff *skb;
1376     int copied = 0, err = 0;
1377
1378     /* XXX -- need to support SO_PEEK_OFF */
1379
1380     skb_queue_walk(&sk->sk_write_queue, skb) {
1381         err = skb_copy_datagram_iovec(skb, 0, msg->msg_iov, skb->len);
1382         if (err)
1383             break;
1384
1385         copied += skb->len;
1386     }
1387
1388     return err ? : copied;
1389 }
1390
1391 /* Clean up the receive buffer for full frames taken by the user,
1392  * then send an ACK if necessary. COPIED is the number of bytes
1393  * tcp_recvmmsg has given to the user so far, it speeds up the
1394  * calculation of whether or not we must ACK for the sake of
1395  * a window update.
1396  */
1397 void tcp_cleanup_rbuf(struct sock *sk, int copied)
1398 {
1399     struct tcp_sock *tp = tcp_sk(sk);
1400     bool time_to_ack = false;
1401
1402     struct sk_buff *skb = skb_peek(&sk->sk_receive_queue);
1403
1404     WARN(skb && !before(tp->copied_seq, TCP_SKB_CB(skb)->end_seq),
1405          "cleanup rbuf bug: copied %X seq %X rcvnext %X\n",
1406          tp->copied_seq, TCP_SKB_CB(skb)->end_seq, tp->rcv_nxt);
1407
1408     if (inet_csk_ack_scheduled(sk)) {
1409         const struct inet_connection_sock *icsk = inet_csk(sk);
1410         /* Delayed ACKs frequently hit locked sockets during bulk
1411          * receive. */
1412         if (icsk->icsk_ack.blocked ||
1413             /* Once-per-two-segments ACK was not sent by tcp_input.c */
1414             tp->rcv_nxt - tp->rcv_wup > icsk->icsk_ack.rcv_mss ||
1415             /*
1416              * If this read emptied read buffer, we send ACK, if
1417              * connection is not bidirectional, user drained
1418              * receive buffer and there was a small segment
1419              * in queue.
1420              */
1421             (copied > 0 &&
1422              ((icsk->icsk_ack.pending & ICSK_ACK_PUSHED2) ||
1423               (icsk->icsk_ack.pending & ICSK_ACK_PUSHED) &&
1424               !icsk->icsk_ack.pingpong)) &&
1425             !atomic_read(&sk->sk_rmem_alloc)))

```

```

1426         time_to_ack = true;
1427     }
1428
1429     /* We send an ACK if we can now advertise a non-zero window
1430      * which has been raised "significantly".
1431      *
1432      * Even if window raised up to infinity, do not send window open ACK
1433      * in states, where we will not receive more. It is useless.
1434      */
1435     if (copied > 0 && !time_to_ack && !(sk->sk_shutdown & RCV\_SHUTDOWN)) {
1436         u32 rcv_window_now = tcp\_receive\_window(tp);
1437
1438         /* Optimize, __tcp_select_window() is not cheap. */
1439         if (2*rcv_window_now <= tp->window_clamp) {
1440             u32 new_window = tcp\_select\_window(sk);
1441
1442             /* Send ACK now, if this read freed lots of space
1443              * in our buffer. Certainly, new_window is new window.
1444              * We can advertise it now, if it is not less than current one.
1445              * "Lots" means "at least twice" here.
1446              */
1447             if (new_window && new_window >= 2 * rcv_window_now)
1448                 time_to_ack = true;
1449         }
1450     }
1451     if (time_to_ack)
1452         tcp\_send\_ack(sk);
1453 }
1454
1455 static void tcp\_prequeue\_process(struct sock *sk)
1456 {
1457     struct sk\_buff *skb;
1458     struct tcp\_sock *tp = tcp\_sk(sk);
1459
1460     NET\_INC\_STATS\_USER(sock\_net(sk), LINUX\_MIB\_TCPPREQUEUED);
1461
1462     /* RX process wants to run with disabled BHs, though it is not
1463      * necessary */
1464     local\_bh\_disable();
1465     while ((skb = skb\_dequeue(&tp->ucopy.prequeue)) != NULL)
1466         sk\_backlog\_rcv(sk, skb);
1467     local\_bh\_enable();
1468
1469     /* Clear memory counter. */
1470     tp->ucopy.memory = 0;
1471 }
1472
1473 #ifdef CONFIG_NET_DMA
1474 static void tcp\_service\_net\_dma(struct sock *sk, bool wait)
1475 {
1476     dma\_cookie\_t done, used;
1477     dma\_cookie\_t last_issued;
1478     struct tcp\_sock *tp = tcp\_sk(sk);
1479
1480     if (!tp->ucopy.dma_chan)
1481         return;
1482
1483     last_issued = tp->ucopy.dma_cookie;
1484     dma\_async\_issue\_pending(tp->ucopy.dma_chan);
1485
1486     do {
1487         if (dma\_async\_is\_tx\_complete(tp->ucopy.dma_chan,
1488                                     last_issued, &done,
1489                                     &used) == DMA\_COMPLETE) {
1490             /* Safe to free early-copied skbs now */
1491             skb\_queue\_purge(&sk->sk_async_wait_queue);
1492             break;
1493         } else {
1494             struct sk\_buff *skb;
1495             while ((skb = skb\_peek(&sk->sk_async_wait_queue)) &&
1496                   (dma\_async\_is\_complete(skb->dma_cookie, done,
1497                                           used) == DMA\_COMPLETE)) {
1498                 skb\_dequeue(&sk->sk_async_wait_queue);
1499                 kfree\_skb(skb);
1500             }
1501         }
1502     } while (wait);
1503 }
1504 #endif
1505
1506 static struct sk\_buff *tcp\_rcv\_skb(struct sock *sk, u32 seq, u32 *off)
1507 {
1508     struct sk\_buff *skb;
1509     u32 offset;
1510

```

```

1511 while ((skb = skb_peek(&sk->sk_receive_queue)) != NULL) {
1512     offset = seq - TCP_SKB_CB(skb)->seq;
1513     if (tcp_hdr(skb)->syn)
1514         offset--;
1515     if (offset < skb->len || tcp_hdr(skb)->fin) {
1516         *off = offset;
1517         return skb;
1518     }
1519     /* This looks weird, but this can happen if TCP collapsing
1520      * splitted a fat GRO packet, while we released socket lock
1521      * in skb_splice_bits()
1522      */
1523     sk_eat_skb(sk, skb, false);
1524 }
1525 return NULL;
1526 }
1527
1528 /*
1529  * This routine provides an alternative to tcp_recvmg() for routines
1530  * that would like to handle copying from skbuffs directly in 'sendfile'
1531  * fashion.
1532  * Note:
1533  * - It is assumed that the socket was locked by the caller.
1534  * - The routine does not block.
1535  * - At present, there is no support for reading OOB data
1536  *   or for 'peeking' the socket using this routine
1537  *   (although both would be easy to implement).
1538  */
1539 int tcp_read_sock(struct sock *sk, read_descriptor_t *desc,
1540                   sk_read_actor_t recv_actor)
1541 {
1542     struct sk_buff *skb;
1543     struct tcp_sock *tp = tcp_sk(sk);
1544     u32 seq = tp->copied_seq;
1545     u32 offset;
1546     int copied = 0;
1547
1548     if (sk->sk_state == TCP_LISTEN)
1549         return -ENOTCONN;
1550     while ((skb = tcp_recv_skb(sk, seq, &offset)) != NULL) {
1551         if (offset < skb->len) {
1552             int used;
1553             size_t len;
1554
1555             len = skb->len - offset;
1556             /* Stop reading if we hit a patch of urgent data */
1557             if (tp->urg_data) {
1558                 u32 urg_offset = tp->urg_seq - seq;
1559                 if (urg_offset < len)
1560                     len = urg_offset;
1561                 if (!len)
1562                     break;
1563             }
1564             used = recv_actor(desc, skb, offset, len);
1565             if (used <= 0) {
1566                 if (!copied)
1567                     copied = used;
1568                 break;
1569             } else if (used <= len) {
1570                 seq += used;
1571                 copied += used;
1572                 offset += used;
1573             }
1574             /* If recv_actor drops the lock (e.g. TCP splice
1575              * receive) the skb pointer might be invalid when
1576              * getting here: tcp_collapse might have deleted it
1577              * while aggregating skbs from the socket queue.
1578              */
1579             skb = tcp_recv_skb(sk, seq - 1, &offset);
1580             if (!skb)
1581                 break;
1582             /* TCP coalescing might have appended data to the skb.
1583              * Try to splice more frags
1584              */
1585             if (offset + 1 != skb->len)
1586                 continue;
1587         }
1588         if (tcp_hdr(skb)->fin) {
1589             sk_eat_skb(sk, skb, false);
1590             ++seq;
1591             break;
1592         }
1593         sk_eat_skb(sk, skb, false);
1594         if (!desc->count)
1595             break;

```

```

1596         tp->copied_seq = seq;
1597     }
1598     tp->copied_seq = seq;
1599
1600     tcp_rcv_space_adjust(sk);
1601
1602     /* Clean up data we have read: This will do ACK frames. */
1603     if (copied > 0) {
1604         tcp_recv_skb(sk, seq, &offset);
1605         tcp_cleanup_rbuf(sk, copied);
1606     }
1607     return copied;
1608 }
1609 EXPORT_SYMBOL(tcp_read_sock);
1610
1611 /*
1612  * This routine copies from a sock struct into the user buffer.
1613  *
1614  * Technical note: in 2.3 we work on _locked_socket, so that
1615  * tricks with *seq access order and skb->users are not required.
1616  * Probably, code can be easily improved even more.
1617  */
1618
1619 int tcp_recvmmsg(struct kiocb *iocb, struct sock *sk, struct msghdr *msg,
1620                 size_t len, int nonblock, int flags, int *addr_len)
1621 {
1622     struct tcp_sock *tp = tcp_sk(sk);
1623     int copied = 0;
1624     u32 peek_seq;
1625     u32 *seq;
1626     unsigned long used;
1627     int err;
1628     int target;          /* Read at least this many bytes */
1629     long timeo;
1630     struct task_struct *user_recv = NULL;
1631     bool copied_early = false;
1632     struct sk_buff *skb;
1633     u32 urg_hole = 0;
1634
1635     if (unlikely(flags & MSG_ERRQUEUE))
1636         return ip_recv_error(sk, msg, len, addr_len);
1637
1638     if (sk_can_busy_loop(sk) && skb_queue_empty(&sk->sk_receive_queue) &&
1639         (sk->sk_state == TCP_ESTABLISHED))
1640         sk_busy_loop(sk, nonblock);
1641
1642     lock_sock(sk);
1643
1644     err = -ENOTCONN;
1645     if (sk->sk_state == TCP_LISTEN)
1646         goto out;
1647
1648     timeo = sock_rcvtimeo(sk, nonblock);
1649
1650     /* Urgent data needs to be handled specially. */
1651     if (flags & MSG_OOB)
1652         goto recv_urg;
1653
1654     if (unlikely(tp->repair)) {
1655         err = -EPERM;
1656         if (!(flags & MSG_PEEK))
1657             goto out;
1658
1659         if (tp->repair_queue == TCP_SEND_QUEUE)
1660             goto recv_sndq;
1661
1662         err = -EINVAL;
1663         if (tp->repair_queue == TCP_NO_QUEUE)
1664             goto out;
1665
1666         /* 'common' recv queue MSG_PEEK-ing */
1667     }
1668
1669     seq = &tp->copied_seq;
1670     if (flags & MSG_PEEK) {
1671         peek_seq = tp->copied_seq;
1672         seq = &peek_seq;
1673     }
1674
1675     target = sock_rcvlowat(sk, flags & MSG_WAITALL, len);
1676
1677 #ifdef CONFIG_NET_DMA
1678     tp->ucopy.dma_chan = NULL;
1679     preempt_disable();
1680     skb = skb_peek_tail(&sk->sk_receive_queue);

```



```

1681 {
1682     int available = 0;
1683
1684     if (skb)
1685         available = TCP_SKB_CB(skb)->seq + skb->len - (*seq);
1686     if ((available < target) &&
1687         (len > sysctl_tcp_dma_copybreak) && !(flags & MSG_PEEK) &&
1688         !sysctl_tcp_low_latency &&
1689         net_dma_find_channel()) {
1690         preempt_enable();
1691         tp->ucopy.pinned_list =
1692             dma_pin_iovec_pages(msg->msg_iov, len);
1693     } else {
1694         preempt_enable();
1695     }
1696 }
1697 #endif
1698
1699 do {
1700     u32 offset;
1701
1702     /* Are we at urgent data? Stop if we have read anything or have SIGURG pending. */
1703     if (tp->urg_data && tp->urg_seq == *seq) {
1704         if (copied)
1705             break;
1706         if (signal_pending(current)) {
1707             copied = timeo ? sock_intr_errno(timeo) : -EAGAIN;
1708             break;
1709         }
1710     }
1711
1712     /* Next get a buffer. */
1713
1714     skb_queue_walk(&sk->sk_receive_queue, skb) {
1715         /* Now that we have two receive queues this
1716          * shouldn't happen.
1717          */
1718         if (WARN(before(*seq, TCP_SKB_CB(skb)->seq),
1719                 "recvmsg bug: copied %X seq %X rcvnxt %X fl %X\n",
1720                 *seq, TCP_SKB_CB(skb)->seq, tp->rcv_nxt,
1721                 flags))
1722             break;
1723
1724         offset = *seq - TCP_SKB_CB(skb)->seq;
1725         if (tcp_hdr(skb)->syn)
1726             offset--;
1727         if (offset < skb->len)
1728             goto found_ok_skb;
1729         if (tcp_hdr(skb)->fin)
1730             goto found_fin_ok;
1731         WARN(!(flags & MSG_PEEK),
1732             "recvmsg bug 2: copied %X seq %X rcvnxt %X fl %X\n",
1733             *seq, TCP_SKB_CB(skb)->seq, tp->rcv_nxt, flags);
1734     }
1735
1736     /* Well, if we have backlog, try to process it now yet. */
1737
1738     if (copied >= target && !sk->sk_backlog.tail)
1739         break;
1740
1741     if (copied) {
1742         if (sk->sk_err ||
1743             sk->sk_state == TCP_CLOSE ||
1744             (sk->sk_shutdown & RCV_SHUTDOWN) ||
1745             !timeo ||
1746             signal_pending(current))
1747             break;
1748     } else {
1749         if (sock_flag(sk, SOCK_DONE))
1750             break;
1751
1752         if (sk->sk_err) {
1753             copied = sock_error(sk);
1754             break;
1755         }
1756
1757         if (sk->sk_shutdown & RCV_SHUTDOWN)
1758             break;
1759
1760         if (sk->sk_state == TCP_CLOSE) {
1761             if (!sock_flag(sk, SOCK_DONE)) {
1762                 /* This occurs when user tries to read
1763                  * from never connected socket.
1764                  */
1765                 copied = -ENOTCONN;

```

```

1766             break;
1767         }
1768         break;
1769     }
1770
1771     if (!timeo) {
1772         copied = -EAGAIN;
1773         break;
1774     }
1775
1776     if (signal_pending(current)) {
1777         copied = sock_intr_errno(timeo);
1778         break;
1779     }
1780 }
1781
1782 tcp_cleanup_rbuf(sk, copied);
1783
1784 if (!sysctl_tcp_low_latency && tp->ucopy.task == user_rcv) {
1785     /* Install new reader */
1786     if (!user_rcv && !(flags & (MSG_TRUNC | MSG_PEEK))) {
1787         user_rcv = current;
1788         tp->ucopy.task = user_rcv;
1789         tp->ucopy.iov = msg->msg_iov;
1790     }
1791
1792     tp->ucopy.len = len;
1793
1794     WARN_ON(tp->copied_seq != tp->rcv_nxt &&
1795             !(flags & (MSG_PEEK | MSG_TRUNC)));
1796
1797     /* Ugly... If prequeue is not empty, we have to
1798      * process it before releasing socket, otherwise
1799      * order will be broken at second iteration.
1800      * More elegant solution is required!!!
1801      *
1802      * Look: we have the following (pseudo)queues:
1803      *
1804      * 1. packets in flight
1805      * 2. backlog
1806      * 3. prequeue
1807      * 4. receive_queue
1808      *
1809      * Each queue can be processed only if the next ones
1810      * are empty. At this point we have empty receive_queue.
1811      * But prequeue _can_ be not empty after 2nd iteration,
1812      * when we jumped to start of loop because backlog
1813      * processing added something to receive_queue.
1814      * We cannot release_sock(), because backlog contains
1815      * packets arrived _after_ prequeued ones.
1816      *
1817      * Shortly, algorithm is clear --- to process all
1818      * the queues in order. We could make it more directly,
1819      * requeueing packets from backlog to prequeue, if
1820      * is not empty. It is more elegant, but eats cycles,
1821      * unfortunately.
1822      */
1823     if (!skb_queue_empty(&tp->ucopy.prequeue))
1824         goto do_prequeue;
1825
1826     /* __ Set realtime policy in scheduler __ */
1827 }
1828
1829 #ifdef CONFIG_NET_DMA
1830     if (tp->ucopy.dma_chan) {
1831         if (tp->rcv_wnd == 0 &&
1832             !skb_queue_empty(&sk->sk_async_wait_queue)) {
1833             tcp_service_net_dma(sk, true);
1834             tcp_cleanup_rbuf(sk, copied);
1835         } else
1836             dma_async_issue_pending(tp->ucopy.dma_chan);
1837     }
1838 #endif
1839
1840 if (copied >= target) {
1841     /* Do not sleep, just process backlog. */
1842     release_sock(sk);
1843     lock_sock(sk);
1844 } else
1845     sk_wait_data(sk, &timeo);
1846
1847 #ifdef CONFIG_NET_DMA
1848     tcp_service_net_dma(sk, false); /* Don't block */
1849     tp->ucopy.wakeup = 0;
1850 #endif

```

```

1851     if (user_rcv) {
1852         int chunk;
1853
1854         /* __ Restore normal policy in scheduler __ */
1855
1856         if ((chunk = len - tp->ucopy.len) != 0) {
1857             NET_ADD_STATS_USER(sock_net(sk), LINUX_MIB_TCPDIRECTCOPYFROMBACKLOG, chunk);
1858             len -= chunk;
1859             copied += chunk;
1860         }
1861
1862         if (tp->rcv_nxt == tp->copied_seq &&
1863             !skb_queue_empty(&tp->ucopy.prequeue)) {
1864 do_prequeue:
1865             tcp_prequeue_process(sk);
1866
1867             if ((chunk = len - tp->ucopy.len) != 0) {
1868                 NET_ADD_STATS_USER(sock_net(sk), LINUX_MIB_TCPDIRECTCOPYFROMPREQUEUE, chunk);
1869                 len -= chunk;
1870                 copied += chunk;
1871             }
1872         }
1873     }
1874     if ((flags & MSG_PEEK) &&
1875         (peek_seq - copied - urg_hole != tp->copied_seq)) {
1876         net_dbg_ratelimited("TCP(%s:%d): Application bug, race in MSG_PEEK\n",
1877             current->comm,
1878             task_pid_nr(current));
1879         peek_seq = tp->copied_seq;
1880     }
1881     continue;
1882
1883 found_ok_skb:
1884     /* Ok so how much can we use? */
1885     used = skb->len - offset;
1886     if (len < used)
1887         used = len;
1888
1889     /* Do we have urgent data here? */
1890     if (tp->urg_data) {
1891         u32 urg_offset = tp->urg_seq - *seq;
1892         if (urg_offset < used) {
1893             if (!urg_offset) {
1894                 if (!sock_flag(sk, SOCK_URGINLINE)) {
1895                     ++*seq;
1896                     urg_hole++;
1897                     offset++;
1898                     used--;
1899                     if (!used)
1900                         goto skip_copy;
1901                 }
1902             } else
1903                 used = urg_offset;
1904         }
1905     }
1906
1907     if (!(flags & MSG_TRUNC)) {
1908 #ifdef CONFIG_NET_DMA
1909         if (!tp->ucopy.dma_chan && tp->ucopy.pinned_list)
1910             tp->ucopy.dma_chan = net_dma_find_channel();
1911
1912         if (tp->ucopy.dma_chan) {
1913             tp->ucopy.dma_cookie = dma_skb_copy_datagram_iovec(
1914                 tp->ucopy.dma_chan, skb, offset,
1915                 msg->msg_iov, used,
1916                 tp->ucopy.pinned_list);
1917
1918             if (tp->ucopy.dma_cookie < 0) {
1919                 pr_alert("%s: dma_cookie < 0\n",
1920                     __func__);
1921
1922                 /* Exception. Bailout! */
1923                 if (!copied)
1924                     copied = -EFAULT;
1925                 break;
1926             }
1927
1928             dma_async_issue_pending(tp->ucopy.dma_chan);
1929
1930             if ((offset + used) == skb->len)
1931                 copied_early = true;
1932         } else
1933 #endif
1934     }
1935 #endif

```

```

1936         {
1937             err = skb_copy_datagram_iovec(skb, offset,
1938             msg->msg_iov, used);
1939             if (err) {
1940                 /* Exception. Bailout! */
1941                 if (!copied)
1942                     copied = -EFAULT;
1943                 break;
1944             }
1945         }
1946     }
1947
1948     *seq += used;
1949     copied += used;
1950     len -= used;
1951
1952     tcp_rcv_space_adjust(sk);
1953
1954 skip_copy:
1955     if (tp->urg_data && after(tp->copied_seq, tp->urg_seq)) {
1956         tp->urg_data = 0;
1957         tcp_fast_path_check(sk);
1958     }
1959     if (used + offset < skb->len)
1960         continue;
1961
1962     if (tcp_hdr(skb)->fin)
1963         goto found_fin_ok;
1964     if (!(flags & MSG_PEEK)) {
1965         sk_eat_skb(sk, skb, copied_early);
1966         copied_early = false;
1967     }
1968     continue;
1969
1970 found_fin_ok:
1971     /* Process the FIN. */
1972     ++*seq;
1973     if (!(flags & MSG_PEEK)) {
1974         sk_eat_skb(sk, skb, copied_early);
1975         copied_early = false;
1976     }
1977     break;
1978 } while (len > 0);
1979
1980 if (user_recv) {
1981     if (!skb_queue_empty(&tp->ucopy.prequeue)) {
1982         int chunk;
1983
1984         tp->ucopy.len = copied > 0 ? len : 0;
1985
1986         tcp_prequeue_process(sk);
1987
1988         if (copied > 0 && (chunk = len - tp->ucopy.len) != 0) {
1989             NET_ADD_STATS_USER(sock_net(sk), LINUX_MIB_TCPDIRECTCOPYFROMPREQUEUE, chunk);
1990             len -= chunk;
1991             copied += chunk;
1992         }
1993     }
1994
1995     tp->ucopy.task = NULL;
1996     tp->ucopy.len = 0;
1997 }
1998
1999 #ifdef CONFIG_NET_DMA
2000 tcp_service_net_dma(sk, true); /* Wait for queue to drain */
2001 tp->ucopy.dma_chan = NULL;
2002
2003 if (tp->ucopy.pinned_list) {
2004     dma_unpin_iovec_pages(tp->ucopy.pinned_list);
2005     tp->ucopy.pinned_list = NULL;
2006 }
2007 #endif
2008
2009 /* According to UNIX98, msg_name/msg_namelen are ignored
2010  * on connected socket. I was just happy when found this 8) --ANK
2011  */
2012
2013 /* Clean up data we have read: This will do ACK frames. */
2014 tcp_cleanup_rbuf(sk, copied);
2015
2016 release_sock(sk);
2017 return copied;
2018
2019 out:
2020 release_sock(sk);

```

```

2021         return err;
2022
2023     recv_urg:
2024         err = tcp\_recv\_urg(sk, msg, len, flags);
2025         goto out;
2026
2027     recv_sndq:
2028         err = tcp\_peek\_sndq(sk, msg, len);
2029         goto out;
2030 }
2031 EXPORT\_SYMBOL(tcp\_recvmmsg);
2032
2033 void tcp\_set\_state(struct sock *sk, int state)
2034 {
2035     int oldstate = sk->sk\_state;
2036
2037     switch (state) {
2038     case TCP_ESTABLISHED:
2039         if (oldstate != TCP_ESTABLISHED)
2040             TCP\_INC\_STATS(sock\_net(sk), TCP_MIB_CURRESTAB);
2041         break;
2042
2043     case TCP_CLOSE:
2044         if (oldstate == TCP_CLOSE_WAIT || oldstate == TCP_ESTABLISHED)
2045             TCP\_INC\_STATS(sock\_net(sk), TCP_MIB_ESTABRESETS);
2046
2047         sk->sk\_prot->unhash(sk);
2048         if (inet\_csk(sk)->icsk_bind_hash &&
2049             !(sk->sk\_userlocks & SOCK\_BINDPORT\_LOCK))
2050             inet\_put\_port(sk);
2051         /* fall through */
2052     default:
2053         if (oldstate == TCP_ESTABLISHED)
2054             TCP\_DEC\_STATS(sock\_net(sk), TCP_MIB_CURRESTAB);
2055     }
2056
2057     /* Change state AFTER socket is unhashed to avoid closed
2058      * socket sitting in hash tables.
2059      */
2060     sk->sk\_state = state;
2061
2062     #ifdef STATE\_TRACE
2063     SOCK\_DEBUG(sk, "TCP sk=%p, State %s -> %s\n", sk, statename[oldstate], statename[state]);
2064     #endif
2065 }
2066 EXPORT\_SYMBOL\_GPL(tcp\_set\_state);
2067
2068 /*
2069  *      State processing on a close. This implements the state shift for
2070  *      sending our FIN frame. Note that we only send a FIN for some
2071  *      states. A shutdown() may have already sent the FIN, or we may be
2072  *      closed.
2073  */
2074
2075 static const unsigned char new\_state[16] = {
2076     /* current state:      new state:      action:      */
2077     /* (Invalid)          */ TCP_CLOSE,
2078     /* TCP_ESTABLISHED    */ TCP_FIN_WAIT1 | TCP\_ACTION\_FIN,
2079     /* TCP_SYN_SENT       */ TCP_CLOSE,
2080     /* TCP_SYN_RECV       */ TCP_FIN_WAIT1 | TCP\_ACTION\_FIN,
2081     /* TCP_FIN_WAIT1      */ TCP_FIN_WAIT1,
2082     /* TCP_FIN_WAIT2      */ TCP_FIN_WAIT2,
2083     /* TCP_TIME_WAIT      */ TCP_CLOSE,
2084     /* TCP_CLOSE          */ TCP_CLOSE,
2085     /* TCP_CLOSE_WAIT     */ TCP_LAST_ACK | TCP\_ACTION\_FIN,
2086     /* TCP_LAST_ACK       */ TCP_LAST_ACK,
2087     /* TCP_LISTEN         */ TCP_CLOSE,
2088     /* TCP_CLOSING        */ TCP_CLOSING,
2089 };
2090
2091 static int tcp\_close\_state(struct sock *sk)
2092 {
2093     int next = (int)new\_state[sk->sk\_state];
2094     int ns = next & TCP\_STATE\_MASK;
2095
2096     tcp\_set\_state(sk, ns);
2097
2098     return next & TCP\_ACTION\_FIN;
2099 }
2100
2101 /*
2102  *      Shutdown the sending side of a connection. Much like close except
2103  *      that we don't receive shut down or sock_set_flag(sk, SOCK_DEAD).
2104  */
2105

```

```

2106 void tcp_shutdown(struct sock *sk, int how)
2107 {
2108     /*      We need to grab some memory, and put together a FIN,
2109     *      and then put it into the queue to be sent.
2110     *      Tim MacKenzie(tym@dibbler.cs.monash.edu.au) 4 Dec '92.
2111     */
2112     if (!(how & SEND_SHUTDOWN))
2113         return;
2114
2115     /* If we've already sent a FIN, or it's a closed state, skip this. */
2116     if ((1 << sk->sk_state) &
2117         (TCPF_ESTABLISHED | TCPF_SYN_SENT |
2118          TCPF_SYN_RECV | TCPF_CLOSE_WAIT)) {
2119         /* Clear out any half completed packets. FIN if needed. */
2120         if (tcp_close_state(sk))
2121             tcp_send_fin(sk);
2122     }
2123 }
2124 EXPORT_SYMBOL(tcp_shutdown);
2125
2126 bool tcp_check_oom(struct sock *sk, int shift)
2127 {
2128     bool too_many_orphans, out_of_socket_memory;
2129
2130     too_many_orphans = tcp_too_many_orphans(sk, shift);
2131     out_of_socket_memory = tcp_out_of_memory(sk);
2132
2133     if (too_many_orphans)
2134         net_info_ratelimited("too many orphaned sockets\n");
2135     if (out_of_socket_memory)
2136         net_info_ratelimited("out of memory -- consider tuning tcp_mem\n");
2137     return too_many_orphans || out_of_socket_memory;
2138 }
2139
2140 void tcp_close(struct sock *sk, long timeout)
2141 {
2142     struct sk_buff *skb;
2143     int data_was_unread = 0;
2144     int state;
2145
2146     lock_sock(sk);
2147     sk->sk_shutdown = SHUTDOWN_MASK;
2148
2149     if (sk->sk_state == TCP_LISTEN) {
2150         tcp_set_state(sk, TCP_CLOSE);
2151
2152         /* Special case. */
2153         inet_csk_listen_stop(sk);
2154
2155         goto adjudge_to_death;
2156     }
2157
2158     /* We need to flush the recv. buffs. We do this only on the
2159     * descriptor close, not protocol-sourced closes, because the
2160     * reader process may not have drained the data yet!
2161     */
2162     while ((skb = __skb_dequeue(&sk->sk_receive_queue)) != NULL) {
2163         u32 len = TCP_SKB_CB(skb)->end_seq - TCP_SKB_CB(skb)->seq -
2164             tcp_hdr(skb)->fin;
2165         data_was_unread += len;
2166         kfree_skb(skb);
2167     }
2168
2169     sk_mem_reclaim(sk);
2170
2171     /* If socket has been already reset (e.g. in tcp_reset()) - kill it. */
2172     if (sk->sk_state == TCP_CLOSE)
2173         goto adjudge_to_death;
2174
2175     /* As outlined in RFC 2525, section 2.17, we send a RST here because
2176     * data was lost. To witness the awful effects of the old behavior of
2177     * always doing a FIN, run an older 2.1.x kernel or 2.0.x, start a bulk
2178     * GET in an FTP client, suspend the process, wait for the client to
2179     * advertise a zero window, then kill -9 the FTP client, wheee...
2180     * Note: timeout is always zero in such a case.
2181     */
2182     if (unlikely(tcp_sk(sk)->repair)) {
2183         sk->sk_prot->disconnect(sk, 0);
2184     } else if (data_was_unread) {
2185         /* Unread data was tossed, zap the connection. */
2186         NET_INC_STATS_USER(sock_net(sk), LINUX_MIB_TCPABORTONCLOSE);
2187         tcp_set_state(sk, TCP_CLOSE);
2188         tcp_send_active_reset(sk, sk->sk_allocation);
2189     } else if (sock_flag(sk, SOCK_LINGER) && !sk->sk_lingertime) {
2190         /* Check zero linger _after_ checking for unread data. */

```



```

2191 sk->sk_prot->disconnect(sk, 0);
2192 NET_INC_STATS_USER(sock_net(sk), LINUX_MIB_TCPABORTONDATA);
2193 } else if (tcp_close_state(sk)) {
2194     /* We FIN if the application ate all the data before
2195      * zapping the connection.
2196      */
2197
2198     /* RED-PEN. Formally speaking, we have broken TCP state
2199      * machine. State transitions:
2200      *
2201      * TCP_ESTABLISHED -> TCP_FIN_WAIT1
2202      * TCP_SYN_RECV -> TCP_FIN_WAIT1 (forget it, it's impossible)
2203      * TCP_CLOSE_WAIT -> TCP_LAST_ACK
2204      *
2205      * are legal only when FIN has been sent (i.e. in window),
2206      * rather than queued out of window. Purists blame.
2207      *
2208      * F.e. "RFC state" is ESTABLISHED,
2209      * if Linux state is FIN-WAIT-1, but FIN is still not sent.
2210      *
2211      * The visible declinations are that sometimes
2212      * we enter time-wait state, when it is not required really
2213      * (harmless), do not send active resets, when they are
2214      * required by specs (TCP_ESTABLISHED, TCP_CLOSE_WAIT, when
2215      * they look as CLOSING or LAST_ACK for Linux)
2216      * Probably, I missed some more holelets.
2217      *
2218      * XXX (TFO) - To start off we don't support SYN+ACK+FIN
2219      * in a single packet! (May consider it later but will
2220      * probably need API support or TCP_CORK SYN-ACK until
2221      * data is written and socket is closed.)
2222      */
2223     tcp_send_fin(sk);
2224 }
2225
2226 sk_stream_wait_close(sk, timeout);
2227
2228 adjudge_to_death:
2229     state = sk->sk_state;
2230     sock_hold(sk);
2231     sock_orphan(sk);
2232
2233     /* It is the last release_sock in its life. It will remove backlog. */
2234     release_sock(sk);
2235
2236
2237     /* Now socket is owned by kernel and we acquire BH lock
2238      * to finish close. No need to check for user refs.
2239      */
2240     local_bh_disable();
2241     bh_lock_sock(sk);
2242     WARN_ON(sock_owned_by_user(sk));
2243
2244     percpu_counter_inc(sk->sk_prot->orphan_count);
2245
2246     /* Have we already been destroyed by a softirq or backlog? */
2247     if (state != TCP_CLOSE && sk->sk_state == TCP_CLOSE)
2248         goto out;
2249
2250     /* This is a (useful) BSD violating of the RFC. There is a
2251      * problem with TCP as specified in that the other end could
2252      * keep a socket open forever with no application left this end.
2253      * We use a 1 minute timeout (about the same as BSD) then kill
2254      * our end. If they send after that then tough - BUT: Long enough
2255      * that we won't make the old 4*rto = almost no time - whoops
2256      * reset mistake.
2257      *
2258      * Nope, it was not mistake. It is really desired behaviour
2259      * f.e. on http servers, when such sockets are useless, but
2260      * consume significant resources. Let's do it with special
2261      * Linger2 option.
2262      *
2263      */
2264     if (sk->sk_state == TCP_FIN_WAIT2) {
2265         struct tcp_sock *tp = tcp_sk(sk);
2266         if (tp->linger2 < 0) {
2267             tcp_set_state(sk, TCP_CLOSE);
2268             tcp_send_active_reset(sk, GFP_ATOMIC);
2269             NET_INC_STATS_BH(sock_net(sk),
2270                             LINUX_MIB_TCPABORTONLINGER);
2271         } else {
2272             const int tmo = tcp_fin_time(sk);
2273
2274             if (tmo > TCP_TIMEWAIT_LEN) {
2275                 inet_csk_reset_keepalive_timer(sk,

```

```

2276                                     tmo - TCP\_TIMEWAIT\_LEN);
2277     } else {
2278         tcp\_time\_wait(sk, TCP_FIN_WAIT2, tmo);
2279         goto out;
2280     }
2281 }
2282 }
2283 if (sk->sk\_state != TCP_CLOSE) {
2284     sk\_mem\_reclaim(sk);
2285     if (tcp\_check\_oom(sk, 0)) {
2286         tcp\_set\_state(sk, TCP_CLOSE);
2287         tcp\_send\_active\_reset(sk, GFP\_ATOMIC);
2288         NET\_INC\_STATS\_BH(sock\_net(sk),
2289                        LINUX_MIB_TCPABORTONMEMORY);
2290     }
2291 }
2292
2293 if (sk->sk\_state == TCP_CLOSE) {
2294     struct request\_sock *req = tcp\_sk(sk)->fastopen_rsk;
2295     /* We could get here with a non-NULL req if the socket is
2296      * aborted (e.g., closed with unread data) before 3WHS
2297      * finishes.
2298      */
2299     if (req != NULL)
2300         reqsk\_fastopen\_remove(sk, req, false);
2301     inet\_csk\_destroy\_sock(sk);
2302 }
2303 /* Otherwise, socket is reprieved until protocol close. */
2304
2305 out:
2306     bh\_unlock\_sock(sk);
2307     local\_bh\_enable();
2308     sock\_put(sk);
2309 }
2310 EXPORT\_SYMBOL(tcp\_close);
2311
2312 /* These states need RST on ABORT according to RFC793 */
2313
2314 static inline bool tcp\_need\_reset(int state)
2315 {
2316     return (1 << state) &
2317         (TCPF_ESTABLISHED | TCPF_CLOSE_WAIT | TCPF_FIN_WAIT1 |
2318          TCPF_FIN_WAIT2 | TCPF_SYN_RECV);
2319 }
2320
2321 int tcp\_disconnect(struct sock *sk, int flags)
2322 {
2323     struct inet\_sock *inet = inet\_sk(sk);
2324     struct inet\_connection\_sock *icsk = inet\_csk(sk);
2325     struct tcp\_sock *tp = tcp\_sk(sk);
2326     int err = 0;
2327     int old_state = sk->sk\_state;
2328
2329     if (old_state != TCP_CLOSE)
2330         tcp\_set\_state(sk, TCP_CLOSE);
2331
2332     /* ABORT function of RFC793 */
2333     if (old_state == TCP_LISTEN) {
2334         inet\_csk\_listen\_stop(sk);
2335     } else if (unlikely(tp->repair)) {
2336         sk->sk\_err = ECONNABORTED;
2337     } else if (tcp\_need\_reset(old_state) ||
2338               (tp->snd\_nxt != tp->write\_seq &&
2339                (1 << old_state) & (TCPF_CLOSING | TCPF_LAST_ACK))) {
2340         /* The last check adjusts for discrepancy of Linux wrt. RFC
2341          * states
2342          */
2343         tcp\_send\_active\_reset(sk, gfp\_any());
2344         sk->sk\_err = ECONNRESET;
2345     } else if (old_state == TCP_SYN_SENT)
2346         sk->sk\_err = ECONNRESET;
2347
2348     tcp\_clear\_xmit\_timers(sk);
2349     skb\_queue\_purge(&sk->sk\_receive\_queue);
2350     tcp\_write\_queue\_purge(sk);
2351     skb\_queue\_purge(&tp->out\_of\_order\_queue);
2352 #ifdef CONFIG_NET_DMA
2353     skb\_queue\_purge(&sk->sk\_async\_wait\_queue);
2354 #endif
2355
2356     inet->inet\_dport = 0;
2357
2358     if (!(sk->sk\_userlocks & SOCK\_BINDADDR\_LOCK))
2359         inet\_reset\_saddr(sk);
2360

```

```

2361     sk->sk_shutdown = 0;
2362     sock\_reset\_flag(sk, SOCK_DONE);
2363     tp->srtt\_us = 0;
2364     if ((tp->write\_seq += tp->max\_window + 2) == 0)
2365         tp->write\_seq = 1;
2366     icsk->icsk\_backoff = 0;
2367     tp->snd\_cwnd = 2;
2368     icsk->icsk\_probes\_out = 0;
2369     tp->packets\_out = 0;
2370     tp->snd\_ssthresh = TCP\_INFINITE\_SSTHRESH;
2371     tp->snd\_cwnd\_cnt = 0;
2372     tp->window\_clamp = 0;
2373     tcp\_set\_ca\_state(sk, TCP_CA_Open);
2374     tcp\_clear\_retrans(tp);
2375     inet\_csk\_delack\_init(sk);
2376     tcp\_init\_send\_head(sk);
2377     memset(&tp->rx\_opt, 0, sizeof(tp->rx\_opt));
2378     \_\_sk\_dst\_reset(sk);
2379
2380     WARN\_ON(inet->inet\_num && !icsk->icsk\_bind\_hash);
2381
2382     sk->sk_error_report(sk);
2383     return err;
2384 }
2385 EXPORT\_SYMBOL(tcp\_disconnect);
2386
2387 void tcp\_sock\_destruct(struct sock *sk)
2388 {
2389     inet\_sock\_destruct(sk);
2390
2391     kfree(inet\_csk(sk)->icsk\_accept\_queue.fastopenq);
2392 }
2393
2394 static inline bool tcp\_can\_repair\_sock(const struct sock *sk)
2395 {
2396     return ns\_capable(sock\_net(sk)->user\_ns, CAP\_NET\_ADMIN) &&
2397         ((1 << sk->sk\_state) & (TCPF_CLOSE | TCPF_ESTABLISHED));
2398 }
2399
2400 static int tcp\_repair\_options\_est(struct tcp\_sock *tp,
2401     struct tcp\_repair\_opt __user *optbuf, unsigned int len)
2402 {
2403     struct tcp\_repair\_opt opt;
2404
2405     while (len >= sizeof(opt)) {
2406         if (copy\_from\_user(&opt, optbuf, sizeof(opt)))
2407             return -EFAULT;
2408
2409         optbuf++;
2410         len -= sizeof(opt);
2411
2412         switch (opt.opt_code) {
2413             case TCPOPT\_MSS:
2414                 tp->rx\_opt.mss\_clamp = opt.opt_val;
2415                 break;
2416             case TCPOPT\_WINDOW:
2417                 {
2418                     u16 snd_wscale = opt.opt_val & 0xFFFF;
2419                     u16 rcv_wscale = opt.opt_val >> 16;
2420
2421                     if (snd_wscale > 14 || rcv_wscale > 14)
2422                         return -EFBIG;
2423
2424                     tp->rx\_opt.snd\_wscale = snd_wscale;
2425                     tp->rx\_opt.rcv\_wscale = rcv_wscale;
2426                     tp->rx\_opt.wscale\_ok = 1;
2427                 }
2428                 break;
2429             case TCPOPT\_SACK\_PERM:
2430                 if (opt.opt_val != 0)
2431                     return -EINVAL;
2432
2433                 tp->rx\_opt.sack\_ok |= TCP\_SACK\_SEEN;
2434                 if (sysctl\_tcp\_fack)
2435                     tcp\_enable\_fack(tp);
2436                 break;
2437             case TCPOPT\_TIMESTAMP:
2438                 if (opt.opt_val != 0)
2439                     return -EINVAL;
2440
2441                 tp->rx\_opt.timestamp\_ok = 1;
2442                 break;
2443         }
2444     }
2445 }

```

<http://lxr.free-electrons.com/source/net/ipv4/tcp.c>

```

2531     else {
2532         tp->thin_dupack = val;
2533         if (tp->thin_dupack)
2534             tcp_disable_early_retrans(tp);
2535     }
2536     break;
2537
2538 case TCP_REPAIR:
2539     if (!tcp_can_repair_sock(sk))
2540         err = -EPERM;
2541     else if (val == 1) {
2542         tp->repair = 1;
2543         sk->sk_reuse = SK_FORCE_REUSE;
2544         tp->repair_queue = TCP_NO_QUEUE;
2545     } else if (val == 0) {
2546         tp->repair = 0;
2547         sk->sk_reuse = SK_NO_REUSE;
2548         tcp_send_window_probe(sk);
2549     } else
2550         err = -EINVAL;
2551
2552     break;
2553
2554 case TCP_REPAIR_QUEUE:
2555     if (!tp->repair)
2556         err = -EPERM;
2557     else if (val < TCP_QUEUES_NR)
2558         tp->repair_queue = val;
2559     else
2560         err = -EINVAL;
2561     break;
2562
2563 case TCP_QUEUE_SEQ:
2564     if (sk->sk_state != TCP_CLOSE)
2565         err = -EPERM;
2566     else if (tp->repair_queue == TCP_SEND_QUEUE)
2567         tp->write_seq = val;
2568     else if (tp->repair_queue == TCP_RECV_QUEUE)
2569         tp->rcv_nxt = val;
2570     else
2571         err = -EINVAL;
2572     break;
2573
2574 case TCP_REPAIR_OPTIONS:
2575     if (!tp->repair)
2576         err = -EINVAL;
2577     else if (sk->sk_state == TCP_ESTABLISHED)
2578         err = tcp_repair_options_est(tp,
2579                                     (struct tcp_repair_opt_user *)optval,
2580                                     optlen);
2581     else
2582         err = -EPERM;
2583     break;
2584
2585 case TCP_CORK:
2586     /* When set indicates to always queue non-full frames.
2587      * Later the user clears this option and we transmit
2588      * any pending partial frames in the queue. This is
2589      * meant to be used alongside sendfile() to get properly
2590      * filled frames when the user (for example) must write
2591      * out headers with a write() call first and then use
2592      * sendfile to send out the data parts.
2593      *
2594      * TCP_CORK can be set together with TCP_NODELAY and it is
2595      * stronger than TCP_NODELAY.
2596      */
2597     if (val) {
2598         tp->nonagle |= TCP_NAGLE_CORK;
2599     } else {
2600         tp->nonagle &= ~TCP_NAGLE_CORK;
2601         if (tp->nonagle & TCP_NAGLE_OFF)
2602             tp->nonagle |= TCP_NAGLE_PUSH;
2603         tcp_push_pending_frames(sk);
2604     }
2605     break;
2606
2607 case TCP_KEEPIDLE:
2608     if (val < 1 || val > MAX_TCP_KEEPIDLE)
2609         err = -EINVAL;
2610     else {
2611         tp->keepalive_time = val * HZ;
2612         if (sock_flag(sk, SOCK_KEEPOPEN) &&
2613             !((1 << sk->sk_state) &
2614              (TCPF_CLOSE | TCPF_LISTEN))) {
2615             u32 elapsed = keepalive_time_elapsed(tp);

```

```

2616         if (tp->keepalive_time > elapsed)
2617             elapsed = tp->keepalive_time - elapsed;
2618         else
2619             elapsed = 0;
2620         inet_csk_reset_keepalive_timer(sk, elapsed);
2621     }
2622 }
2623 break;
2624 case TCP_KEEPINTVL:
2625     if (val < 1 || val > MAX_TCP_KEEPINTVL)
2626         err = -EINVAL;
2627     else
2628         tp->keepalive_intvl = val * HZ;
2629     break;
2630 case TCP_KEEPCNT:
2631     if (val < 1 || val > MAX_TCP_KEEPCNT)
2632         err = -EINVAL;
2633     else
2634         tp->keepalive_probes = val;
2635     break;
2636 case TCP_SYNCNT:
2637     if (val < 1 || val > MAX_TCP_SYNCNT)
2638         err = -EINVAL;
2639     else
2640         icsk->icsk_syn_retries = val;
2641     break;
2642 case TCP_LINGER2:
2643     if (val < 0)
2644         tp->linger2 = -1;
2645     else if (val > sysctl_tcp_fin_timeout / HZ)
2646         tp->linger2 = 0;
2647     else
2648         tp->linger2 = val * HZ;
2649     break;
2650 case TCP_DEFER_ACCEPT:
2651     /* Translate value in seconds to number of retransmits */
2652     icsk->icsk_accept_queue.rskq_defer_accept =
2653         secs_to_retrans(val, TCP_TIMEOUT_INIT / HZ,
2654             TCP_RTO_MAX / HZ);
2655     break;
2656 case TCP_WINDOW_CLAMP:
2657     if (!val) {
2658         if (sk->sk_state != TCP_CLOSE) {
2659             err = -EINVAL;
2660             break;
2661         }
2662         tp->window_clamp = 0;
2663     } else
2664         tp->window_clamp = val < SOCK_MIN_RCVBUF / 2 ?
2665             SOCK_MIN_RCVBUF / 2 : val;
2666     break;
2667 case TCP_QUICKACK:
2668     if (!val) {
2669         icsk->icsk_ack.pingpong = 1;
2670     } else {
2671         icsk->icsk_ack.pingpong = 0;
2672         if ((1 << sk->sk_state) &
2673             (TCPF_ESTABLISHED | TCPF_CLOSE_WAIT) &&
2674             inet_csk_ack_scheduled(sk)) {
2675             icsk->icsk_ack.pending |= ICSK_ACK_PUSHED;
2676             tcp_cleanup_rbuf(sk, 1);
2677             if (!(val & 1))
2678                 icsk->icsk_ack.pingpong = 1;
2679         }
2680     }
2681     break;
2682 #ifdef CONFIG_TCP_MD5SIG
2683 case TCP_MD5SIG:
2684     /* Read the IP->Key mappings from userspace */
2685     err = tp->af_specific->md5_parse(sk, optval, optlen);
2686     break;
2687 #endif
2688 case TCP_USER_TIMEOUT:
2689     /* Cap the max timeout in ms TCP will retry/retrans
2690      * before giving up and aborting (ETIMEDOUT) a connection.
2691      */
2692     if (val < 0)
2693         err = -EINVAL;
2694     else
2695         icsk->icsk_user_timeout = msecs_to_jiffies(val);
2696

```



```

2701         break;
2702
2703     case TCP_FASTOPEN:
2704         if (val >= 0 && ((1 << sk->sk_state) & (TCPF_CLOSE |
2705             TCPF_LISTEN)))
2706             err = fastopen_init_queue(sk, val);
2707         else
2708             err = -EINVAL;
2709         break;
2710     case TCP_TIMESTAMP:
2711         if (!tp->repair)
2712             err = -EPERM;
2713         else
2714             tp->tsoffset = val - tcp_time_stamp;
2715         break;
2716     case TCP_NOTSENT_LOWAT:
2717         tp->notsent_lowat = val;
2718         sk->sk_write_space(sk);
2719         break;
2720     default:
2721         err = -ENOPROTOOPT;
2722         break;
2723 }
2724
2725 release_sock(sk);
2726 return err;
2727 }
2728
2729 int tcp_setsockopt(struct sock *sk, int level, int optname, char __user *optval,
2730     unsigned int optlen)
2731 {
2732     const struct inet_connection_sock *icsk = inet_csk(sk);
2733
2734     if (level != SOL_TCP)
2735         return icsk->icsk_af_ops->setsockopt(sk, level, optname,
2736             optval, optlen);
2737     return do_tcp_setsockopt(sk, level, optname, optval, optlen);
2738 }
2739 EXPORT_SYMBOL(tcp_setsockopt);
2740
2741 #ifdef CONFIG_COMPAT
2742 int compat_tcp_setsockopt(struct sock *sk, int level, int optname,
2743     char __user *optval, unsigned int optlen)
2744 {
2745     if (level != SOL_TCP)
2746         return inet_csk_compat_setsockopt(sk, level, optname,
2747             optval, optlen);
2748     return do_tcp_setsockopt(sk, level, optname, optval, optlen);
2749 }
2750 EXPORT_SYMBOL(compat_tcp_setsockopt);
2751 #endif
2752
2753 /* Return information about state of tcp endpoint in API format. */
2754 void tcp_get_info(const struct sock *sk, struct tcp_info *info)
2755 {
2756     const struct tcp_sock *tp = tcp_sk(sk);
2757     const struct inet_connection_sock *icsk = inet_csk(sk);
2758     u32 now = tcp_time_stamp;
2759
2760     memset(info, 0, sizeof(*info));
2761
2762     info->tcpi_state = sk->sk_state;
2763     info->tcpi_ca_state = icsk->icsk_ca_state;
2764     info->tcpi_retransmits = icsk->icsk_retransmits;
2765     info->tcpi_probes = icsk->icsk_probes_out;
2766     info->tcpi_backoff = icsk->icsk_backoff;
2767
2768     if (tp->rx_opt.tstamp_ok)
2769         info->tcpi_options |= TCPI_OPT_TIMESTAMPS;
2770     if (tcp_is_sack(tp))
2771         info->tcpi_options |= TCPI_OPT_SACK;
2772     if (tp->rx_opt.wscale_ok) {
2773         info->tcpi_options |= TCPI_OPT_WSCALE;
2774         info->tcpi_snd_wscale = tp->rx_opt.snd_wscale;
2775         info->tcpi_rcv_wscale = tp->rx_opt.rcv_wscale;
2776     }
2777
2778     if (tp->ecn_flags & TCP_ECN_OK)
2779         info->tcpi_options |= TCPI_OPT_ECN;
2780     if (tp->ecn_flags & TCP_ECN_SEEN)
2781         info->tcpi_options |= TCPI_OPT_ECN_SEEN;
2782     if (tp->syn_data_acked)
2783         info->tcpi_options |= TCPI_OPT_SYN_DATA;
2784
2785     info->tcpi_rto = jiffies_to_usecs(icsk->icsk_rto);

```

```

2786 info->tcp_i_ato = jiffies_to_usecs(icsk->icsk_ack.ato);
2787 info->tcp_i_snd_mss = tp->mss_cache;
2788 info->tcp_i_rcv_mss = icsk->icsk_ack.rcv_mss;
2789
2790 if (sk->sk_state == TCP_LISTEN) {
2791     info->tcp_i_unacked = sk->sk_ack_backlog;
2792     info->tcp_i_sacked = sk->sk_max_ack_backlog;
2793 } else {
2794     info->tcp_i_unacked = tp->packets_out;
2795     info->tcp_i_sacked = tp->sacked_out;
2796 }
2797 info->tcp_i_lost = tp->lost_out;
2798 info->tcp_i_retrans = tp->retrans_out;
2799 info->tcp_i_fackets = tp->fackets_out;
2800
2801 info->tcp_i_last_data_sent = jiffies_to_msecs(now - tp->lsndtime);
2802 info->tcp_i_last_data_recv = jiffies_to_msecs(now - icsk->icsk_ack.lrcvtime);
2803 info->tcp_i_last_ack_recv = jiffies_to_msecs(now - tp->rcv_tstamp);
2804
2805 info->tcp_i_pmtu = icsk->icsk_pmtu_cookie;
2806 info->tcp_i_rcv_ssthresh = tp->rcv_ssthresh;
2807 info->tcp_i_rtt = tp->srtt_us >> 3;
2808 info->tcp_i_rttvar = tp->mdev_us >> 2;
2809 info->tcp_i_snd_ssthresh = tp->snd_ssthresh;
2810 info->tcp_i_snd_cwnd = tp->snd_cwnd;
2811 info->tcp_i_advmss = tp->advmss;
2812 info->tcp_i_reordering = tp->reordering;
2813
2814 info->tcp_i_rcv_rtt = jiffies_to_usecs(tp->rcv_rtt_est.rtt)>>3;
2815 info->tcp_i_rcv_space = tp->rcvq_space.space;
2816
2817 info->tcp_i_total_retrans = tp->total_retrans;
2818
2819 info->tcp_i_pacing_rate = sk->sk_pacing_rate != ~0U ?
2820     sk->sk_pacing_rate : ~0ULL;
2821 info->tcp_i_max_pacing_rate = sk->sk_max_pacing_rate != ~0U ?
2822     sk->sk_max_pacing_rate : ~0ULL;
2823 }
2824 EXPORT_SYMBOL_GPL(tcp_get_info);
2825
2826 static int do_tcp_getsockopt(struct sock *sk, int level,
2827     int optname, char __user *optval, int __user *optlen)
2828 {
2829     struct inet_connection_sock *icsk = inet_csk(sk);
2830     struct tcp_sock *tp = tcp_sk(sk);
2831     int val, len;
2832
2833     if (get_user(len, optlen))
2834         return -EFAULT;
2835
2836     len = min_t(unsigned int, len, sizeof(int));
2837
2838     if (len < 0)
2839         return -EINVAL;
2840
2841     switch (optname) {
2842     case TCP_MAXSEG:
2843         val = tp->mss_cache;
2844         if (!val && ((1 << sk->sk_state) & (TCPF_CLOSE | TCPF_LISTEN)))
2845             val = tp->rx_opt.user_mss;
2846         if (tp->repair)
2847             val = tp->rx_opt.mss_clamp;
2848         break;
2849     case TCP_NODELAY:
2850         val = !(tp->nonagle&TCP_NAGLE_OFF);
2851         break;
2852     case TCP_CORK:
2853         val = !(tp->nonagle&TCP_NAGLE_CORK);
2854         break;
2855     case TCP_KEEPIDL:
2856         val = keepalive_time_when(tp) / HZ;
2857         break;
2858     case TCP_KEEPINTVL:
2859         val = keepalive_intvl_when(tp) / HZ;
2860         break;
2861     case TCP_KEEPCNT:
2862         val = keepalive_probes(tp);
2863         break;
2864     case TCP_SYNCNT:
2865         val = icsk->icsk_syn_retries ? : sysctl_tcp_syn_retries;
2866         break;
2867     case TCP_LINGER2:
2868         val = tp->linger2;
2869         if (val >= 0)
2870             val = (val ? : sysctl_tcp_fin_timeout) / HZ;

```

```

2871         break;
2872     case TCP_DEFER_ACCEPT:
2873         val = retrans_to_secs(icsk->icsk_accept_queue.rskq_defer_accept,
2874                               TCP_TIMEOUT_INIT / HZ, TCP_RTO_MAX / HZ);
2875         break;
2876     case TCP_WINDOW_CLAMP:
2877         val = tp->window_clamp;
2878         break;
2879     case TCP_INFO: {
2880         struct tcp_info info;
2881
2882         if (get_user(len, optlen))
2883             return -EFAULT;
2884
2885         tcp_get_info(sk, &info);
2886
2887         len = min_t(unsigned int, len, sizeof(info));
2888         if (put_user(len, optlen))
2889             return -EFAULT;
2890         if (copy_to_user(optval, &info, len))
2891             return -EFAULT;
2892         return 0;
2893     }
2894     case TCP_QUICKACK:
2895         val = !icsk->icsk_ack.pingpong;
2896         break;
2897
2898     case TCP_CONGESTION:
2899         if (get_user(len, optlen))
2900             return -EFAULT;
2901         len = min_t(unsigned int, len, TCP_CA_NAME_MAX);
2902         if (put_user(len, optlen))
2903             return -EFAULT;
2904         if (copy_to_user(optval, icsk->icsk_ca_ops->name, len))
2905             return -EFAULT;
2906         return 0;
2907
2908     case TCP_THIN_LINEAR_TIMEOUTS:
2909         val = tp->thin_lto;
2910         break;
2911     case TCP_THIN_DUPACK:
2912         val = tp->thin_dupack;
2913         break;
2914
2915     case TCP_REPAIR:
2916         val = tp->repair;
2917         break;
2918
2919     case TCP_REPAIR_QUEUE:
2920         if (tp->repair)
2921             val = tp->repair_queue;
2922         else
2923             return -EINVAL;
2924         break;
2925
2926     case TCP_QUEUE_SEQ:
2927         if (tp->repair_queue == TCP_SEND_QUEUE)
2928             val = tp->write_seq;
2929         else if (tp->repair_queue == TCP_RECV_QUEUE)
2930             val = tp->rcv_nxt;
2931         else
2932             return -EINVAL;
2933         break;
2934
2935     case TCP_USER_TIMEOUT:
2936         val = jiffies_to_msecs(icsk->icsk_user_timeout);
2937         break;
2938
2939     case TCP_FASTOPEN:
2940         if (icsk->icsk_accept_queue.fastopenq != NULL)
2941             val = icsk->icsk_accept_queue.fastopenq->max_qlen;
2942         else
2943             val = 0;
2944         break;
2945
2946     case TCP_TIMESTAMP:
2947         val = tcp_time_stamp + tp->tsoffset;
2948         break;
2949     case TCP_NOTSENT_LOWAT:
2950         val = tp->notsent_lowat;
2951         break;
2952     default:
2953         return -ENOPROTOPT;
2954     }
2955

```

```

2956     if (put\_user(len, optlen))
2957         return -EFAULT;
2958     if (copy\_to\_user(optval, &val, len))
2959         return -EFAULT;
2960     return 0;
2961 }
2962
2963 int tcp\_getsockopt(struct sock *sk, int level, int optname, char \_\_user *optval,
2964                   int \_\_user *optlen)
2965 {
2966     struct inet\_connection\_sock *icsk = inet\_csk(sk);
2967
2968     if (level != SOL\_TCP)
2969         return icsk->icsk_af_ops->getsockopt(sk, level, optname,
2970                                             optval, optlen);
2971     return do\_tcp\_getsockopt(sk, level, optname, optval, optlen);
2972 }
2973 EXPORT\_SYMBOL(tcp\_getsockopt);
2974
2975 #ifdef CONFIG_COMPAT
2976 int compat\_tcp\_getsockopt(struct sock *sk, int level, int optname,
2977                           char \_\_user *optval, int \_\_user *optlen)
2978 {
2979     if (level != SOL\_TCP)
2980         return inet\_csk\_compat\_getsockopt(sk, level, optname,
2981                                           optval, optlen);
2982     return do\_tcp\_getsockopt(sk, level, optname, optval, optlen);
2983 }
2984 EXPORT\_SYMBOL(compat\_tcp\_getsockopt);
2985 #endif
2986
2987 #ifdef CONFIG_TCP_MD5SIG
2988 static struct tcp\_md5sig\_pool \_\_percpu *tcp\_md5sig\_pool \_\_read\_mostly;
2989 static DEFINE\_MUTEX(tcp\_md5sig\_mutex);
2990
2991 static void \_\_tcp\_free\_md5sig\_pool(struct tcp\_md5sig\_pool \_\_percpu *pool)
2992 {
2993     int cpu;
2994
2995     for\_each\_possible\_cpu(cpu) {
2996         struct tcp\_md5sig\_pool *p = per\_cpu\_ptr(pool, cpu);
2997
2998         if (p->md5_desc.tfm)
2999             crypto\_free\_hash(p->md5_desc.tfm);
3000     }
3001     free\_percpu(pool);
3002 }
3003
3004 static void \_\_tcp\_alloc\_md5sig\_pool(void)
3005 {
3006     int cpu;
3007     struct tcp\_md5sig\_pool \_\_percpu *pool;
3008
3009     pool = alloc\_percpu(struct tcp\_md5sig\_pool);
3010     if (!pool)
3011         return;
3012
3013     for\_each\_possible\_cpu(cpu) {
3014         struct crypto\_hash *hash;
3015
3016         hash = crypto\_alloc\_hash("md5", 0, CRYPTO\_ALG\_ASYNC);
3017         if (IS\_ERR\_OR\_NULL(hash))
3018             goto out_free;
3019
3020         per\_cpu\_ptr(pool, cpu)->md5_desc.tfm = hash;
3021     }
3022     /* before setting tcp_md5sig_pool, we must commit all writes
3023      * to memory. See ACCESS_ONCE() in tcp_get_md5sig_pool()
3024      */
3025     smp\_wmb();
3026     tcp\_md5sig\_pool = pool;
3027     return;
3028 out_free:
3029     \_\_tcp\_free\_md5sig\_pool(pool);
3030 }
3031
3032 bool tcp\_alloc\_md5sig\_pool(void)
3033 {
3034     if (unlikely(!tcp\_md5sig\_pool)) {
3035         mutex\_lock(&tcp\_md5sig\_mutex);
3036
3037         if (!tcp\_md5sig\_pool)
3038             \_\_tcp\_alloc\_md5sig\_pool();
3039
3040         mutex\_unlock(&tcp\_md5sig\_mutex);

```

```

3041     }
3042     return tcp\_md5sig\_pool != NULL;
3043 }
3044 EXPORT\_SYMBOL(tcp\_alloc\_md5sig\_pool);
3045
3046
3047 /**
3048  *      tcp\_get\_md5sig\_pool - get md5sig_pool for this user
3049  *
3050  *      We use percpu structure, so if we succeed, we exit with preemption
3051  *      and BH disabled, to make sure another thread or softirq handling
3052  *      wont try to get same context.
3053  */
3054 struct tcp\_md5sig\_pool *tcp\_get\_md5sig\_pool(void)
3055 {
3056     struct tcp\_md5sig\_pool __percpu *p;
3057
3058     local\_bh\_disable();
3059     p = ACCESS\_ONCE(tcp\_md5sig\_pool);
3060     if (p)
3061         return this\_cpu\_ptr(p);
3062
3063     local\_bh\_enable();
3064     return NULL;
3065 }
3066 EXPORT\_SYMBOL(tcp\_get\_md5sig\_pool);
3067
3068 int tcp\_md5\_hash\_header(struct tcp\_md5sig\_pool *hp,
3069                        const struct tcphdr *th)
3070 {
3071     struct scatterlist sg;
3072     struct tcphdr hdr;
3073     int err;
3074
3075     /* We are not allowed to change tcphdr, make a local copy */
3076     memcpy(&hdr, th, sizeof(hdr));
3077     hdr.check = 0;
3078
3079     /* options aren't included in the hash */
3080     sg\_init\_one(&sg, &hdr, sizeof(hdr));
3081     err = crypto\_hash\_update(&hp->md5_desc, &sg, sizeof(hdr));
3082     return err;
3083 }
3084 EXPORT\_SYMBOL(tcp\_md5\_hash\_header);
3085
3086 int tcp\_md5\_hash\_skb\_data(struct tcp\_md5sig\_pool *hp,
3087                          const struct sk\_buff *skb, unsigned int header_len)
3088 {
3089     struct scatterlist sg;
3090     const struct tcphdr *tp = tcp\_hdr(skb);
3091     struct hash\_desc *desc = &hp->md5_desc;
3092     unsigned int i;
3093     const unsigned int head_data_len = skb\_headlen(skb) > header_len ?
3094                                         skb\_headlen(skb) - header_len : 0;
3095     const struct skb\_shared\_info *shi = skb\_shinfo(skb);
3096     struct sk\_buff *frag_iter;
3097
3098     sg\_init\_table(&sg, 1);
3099
3100     sg\_set\_buf(&sg, ((u8 *) tp) + header_len, head_data_len);
3101     if (crypto\_hash\_update(desc, &sg, head_data_len))
3102         return 1;
3103
3104     for (i = 0; i < shi->nr_frags; ++i) {
3105         const struct skb\_frag\_struct *f = &shi->frags[i];
3106         unsigned int offset = f->page_offset;
3107         struct page *page = skb\_frag\_page(f) + (offset >> PAGE\_SHIFT);
3108
3109         sg\_set\_page(&sg, page, skb\_frag\_size(f),
3110                   offset\_in\_page(offset));
3111         if (crypto\_hash\_update(desc, &sg, skb\_frag\_size(f)))
3112             return 1;
3113     }
3114
3115     skb\_walk\_frags(skb, frag_iter)
3116         if (tcp\_md5\_hash\_skb\_data(hp, frag_iter, 0))
3117             return 1;
3118
3119     return 0;
3120 }
3121 EXPORT\_SYMBOL(tcp\_md5\_hash\_skb\_data);
3122
3123 int tcp\_md5\_hash\_key(struct tcp\_md5sig\_pool *hp, const struct tcp\_md5sig\_key *key)
3124 {
3125     struct scatterlist sg;

```

```

3126
3127     sg_init_one(&sg, key->key, key->keylen);
3128     return crypto_hash_update(&hp->md5_desc, &sg, key->keylen);
3129 }
3130 EXPORT_SYMBOL(tcp_md5_hash_key);
3131
3132 #endif
3133
3134 void tcp_done(struct sock *sk)
3135 {
3136     struct request_sock *req = tcp_sk(sk)->fastopen_rsk;
3137
3138     if (sk->sk_state == TCP_SYN_SENT || sk->sk_state == TCP_SYN_RECV)
3139         TCP_INC_STATS_BH(sock_net(sk), TCP_MIB_ATTEMPTFAILS);
3140
3141     tcp_set_state(sk, TCP_CLOSE);
3142     tcp_clear_xmit_timers(sk);
3143     if (req != NULL)
3144         reqsk_fastopen_remove(sk, req, false);
3145
3146     sk->sk_shutdown = SHUTDOWN_MASK;
3147
3148     if (!sock_flag(sk, SOCK_DEAD))
3149         sk->sk_state_change(sk);
3150     else
3151         inet_csk_destroy_sock(sk);
3152 }
3153 EXPORT_SYMBOL_GPL(tcp_done);
3154
3155 extern struct tcp_congestion_ops tcp_reno;
3156
3157 static __initdata unsigned long thash_entries;
3158 static int __init set_thash_entries(char *str)
3159 {
3160     ssize_t ret;
3161
3162     if (!str)
3163         return 0;
3164
3165     ret = kstrtoul(str, 0, &thash_entries);
3166     if (ret)
3167         return 0;
3168
3169     return 1;
3170 }
3171 __setup("thash_entries=", set_thash_entries);
3172
3173 static void tcp_init_mem(void)
3174 {
3175     unsigned long limit = nr_free_buffer_pages() / 8;
3176     limit = max(limit, 128UL);
3177     sysctl_tcp_mem[0] = limit / 4 * 3;
3178     sysctl_tcp_mem[1] = limit;
3179     sysctl_tcp_mem[2] = sysctl_tcp_mem[0] * 2;
3180 }
3181
3182 void __init tcp_init(void)
3183 {
3184     struct sk_buff *skb = NULL;
3185     unsigned long limit;
3186     int max_rshare, max_wshare, cnt;
3187     unsigned int i;
3188
3189     BUILD_BUG_ON(sizeof(struct tcp_skb_cb) > sizeof(skb->cb));
3190
3191     percpu_counter_init(&tcp_sockets_allocated, 0);
3192     percpu_counter_init(&tcp_orphan_count, 0);
3193     tcp_hashinfo.bind_bucket_cachep =
3194         kmem_cache_create("tcp_bind_bucket",
3195             sizeof(struct inet_bind_bucket), 0,
3196             SLAB_HWCACHE_ALIGN|SLAB_PANIC, NULL);
3197
3198     /* Size and allocate the main established and bind bucket
3199      * hash tables.
3200      *
3201      * The methodology is similar to that of the buffer cache.
3202      */
3203     tcp_hashinfo.ehash =
3204         alloc_large_system_hash("TCP established",
3205             sizeof(struct inet_ehash_bucket),
3206             thash_entries,
3207             17, /* one slot per 128 KB of memory */
3208             0,
3209             NULL,
3210             &tcp_hashinfo.ehash_mask,

```

```

3211         0,
3212         thash\_entries ? 0 : 512 * 1024);
3213     for (i = 0; i <= tcp\_hashinfo.ehash_mask; i++)
3214         INIT\_HLIST\_NULLS\_HEAD(&tcp\_hashinfo.ehash[i].chain, i);
3215
3216     if (inet\_ehash\_locks\_alloc(&tcp\_hashinfo))
3217         panic("TCP: failed to alloc ehash_locks");
3218     tcp\_hashinfo.bhash =
3219         alloc\_large\_system\_hash("TCP bind",
3220             sizeof(struct inet\_bind\_hashbucket),
3221             tcp\_hashinfo.ehash_mask + 1,
3222             17, /* one slot per 128 KB of memory */
3223             0,
3224             &tcp\_hashinfo.bhash_size,
3225             NULL,
3226             0,
3227             64 * 1024);
3228     tcp\_hashinfo.bhash_size = 1U << tcp\_hashinfo.bhash_size;
3229     for (i = 0; i < tcp\_hashinfo.bhash_size; i++) {
3230         spin\_lock\_init(&tcp\_hashinfo.bhash[i].lock);
3231         INIT\_HLIST\_HEAD(&tcp\_hashinfo.bhash[i].chain);
3232     }
3233
3234
3235     cnt = tcp\_hashinfo.ehash_mask + 1;
3236
3237     tcp\_death\_row.sysctl_max_tw_buckets = cnt / 2;
3238     sysctl\_tcp\_max\_orphans = cnt / 2;
3239     sysctl\_max\_syn\_backlog = max(128, cnt / 256);
3240
3241     tcp\_init\_mem();
3242     /* Set per-socket limits to no more than 1/128 the pressure threshold */
3243     limit = nr\_free\_buffer\_pages() << (PAGE\_SHIFT - 7);
3244     max\_wshare = min(4UL*1024*1024, limit);
3245     max\_rshare = min(6UL*1024*1024, limit);
3246
3247     sysctl\_tcp\_wmem[0] = SK\_MEM\_QUANTUM;
3248     sysctl\_tcp\_wmem[1] = 16*1024;
3249     sysctl\_tcp\_wmem[2] = max(64*1024, max\_wshare);
3250
3251     sysctl\_tcp\_rmem[0] = SK\_MEM\_QUANTUM;
3252     sysctl\_tcp\_rmem[1] = 87380;
3253     sysctl\_tcp\_rmem[2] = max(87380, max\_rshare);
3254
3255     pr\_info("Hash tables configured (established %u bind %u)\n",
3256         tcp\_hashinfo.ehash_mask + 1, tcp\_hashinfo.bhash_size);
3257
3258     tcp\_metrics\_init();
3259
3260     tcp\_register\_congestion\_control(&tcp\_reno);
3261
3262     tcp\_tasklet\_init();
3263 }
3264

```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)