

Linux Cross Reference

Free Electrons

Embedded Linux Experts

• [source navigation](#) • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version: [2.0.40](#) [2.2.26](#) [2.4.37](#) [3.1](#) [3.2](#) [3.3](#) [3.4](#) [3.5](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#)

Linux/net/ipv4/tcp_bic.c

```

1  /*
2  * Binary Increase Congestion control for TCP
3  * Home page:
4  * http://netsrv.csc.ncsu.edu/twiki/bin/view/Main/BIC
5  * This is from the implementation of BICTCP in
6  * Lison-Xu, Kahalel Harfoush, and Injong Rhee.
7  * "Binary Increase Congestion Control for Fast, Long Distance
8  * Networks" in InfoComm 2004
9  * Available from:
10 * http://netsrv.csc.ncsu.edu/export/bitcp.pdf
11 *
12 * Unless BIC is enabled and congestion window is large
13 * this behaves the same as the original Reno.
14 */
15
16 #include <linux/mm.h>
17 #include <linux/module.h>
18 #include <net/tcp.h>
19
20
21 #define BICTCP_BETA_SCALE    1024    /* Scale factor beta calculation
22                                     * max_cwnd = snd_cwnd * beta
23                                     */
24 #define BICTCP_B            4        /*
25                                     * In binary search,
26                                     * go to point (max+min)/N
27                                     */
28
29 static int fast_convergence = 1;
30 static int max_increment = 16;
31 static int low_window = 14;
32 static int beta = 819;             /* = 819/1024 (BICTCP_BETA_SCALE) */
33 static int initial_ssthresh;
34 static int smooth_part = 20;
35
36 module_param(fast_convergence, int, 0644);
37 MODULE_PARM_DESC(fast_convergence, "turn on/off fast convergence");
38 module_param(max_increment, int, 0644);
39 MODULE_PARM_DESC(max_increment, "Limit on increment allowed during binary search");
40 module_param(low_window, int, 0644);
41 MODULE_PARM_DESC(low_window, "Lower bound on congestion window (for TCP friendliness)");
42 module_param(beta, int, 0644);
43 MODULE_PARM_DESC(beta, "beta for multiplicative increase");
44 module_param(initial_ssthresh, int, 0644);
45 MODULE_PARM_DESC(initial_ssthresh, "initial value of slow start threshold");
46 module_param(smooth_part, int, 0644);
47 MODULE_PARM_DESC(smooth_part, "Log(B/(B*Smin))/Log(B/(B-1))+B, # of RTT from Wmax-B to Wmax");
48

```

```

49
50 /* BIC TCP Parameters */
51 struct bictcp {
52     u32      cnt;          /* increase cwnd by 1 after ACKs */
53     u32      last\_max\_cwnd; /* Last maximum snd_cwnd */
54     u32      loss\_cwnd;    /* congestion window at Last Loss */
55     u32      last\_cwnd;    /* the last snd_cwnd */
56     u32      last\_time;    /* time when updated last_cwnd */
57     u32      epoch\_start;  /* beginning of an epoch */
58 #define ACK\_RATIO\_SHIFT 4
59     u32      delayed\_ack;  /* estimate the ratio of Packets/ACKs << 4 */
60 };
61
62 static inline void bictcp\_reset(struct bictcp *ca)
63 {
64     ca->cnt = 0;
65     ca->last\_max\_cwnd = 0;
66     ca->last\_cwnd = 0;
67     ca->last\_time = 0;
68     ca->epoch\_start = 0;
69     ca->delayed\_ack = 2 << ACK\_RATIO\_SHIFT;
70 }
71
72 static void bictcp\_init(struct sock *sk)
73 {
74     struct bictcp *ca = inet\_csk\_ca(sk);
75
76     bictcp\_reset(ca);
77     ca->loss\_cwnd = 0;
78
79     if (initial\_ssthresh)
80         tcp\_sk(sk)->snd\_ssthresh = initial\_ssthresh;
81 }
82
83 /*
84  * Compute congestion window to use.
85  */
86 static inline void bictcp\_update(struct bictcp *ca, u32 cwnd)
87 {
88     if (ca->last\_cwnd == cwnd &&
89         (s32)(tcp\_time\_stamp - ca->last\_time) <= HZ / 32)
90         return;
91
92     ca->last\_cwnd = cwnd;
93     ca->last\_time = tcp\_time\_stamp;
94
95     if (ca->epoch\_start == 0) /* record the beginning of an epoch */
96         ca->epoch\_start = tcp\_time\_stamp;
97
98     /* start off normal */
99     if (cwnd <= low\_window) {
100         ca->cnt = cwnd;
101         return;
102     }
103
104     /* binary increase */
105     if (cwnd < ca->last\_max\_cwnd) {
106         u32 dist = (ca->last\_max\_cwnd - cwnd)
107             / BICTCP\_B;
108
109         if (dist > max\_increment)
110             /* linear increase */
111             ca->cnt = cwnd / max\_increment;
112         else if (dist <= 1U)
113             /* binary search increase */
114             ca->cnt = (cwnd * smooth\_part) / BICTCP\_B;
115         else
116             /* binary search increase */

```

```

117         ca->cnt = cwnd / dist;
118     } else {
119         /* slow start AMD linear increase */
120         if (cwnd < ca->last_max_cwnd + BICTCP_B)
121             /* slow start */
122             ca->cnt = (cwnd * smooth_part) / BICTCP_B;
123         else if (cwnd < ca->last_max_cwnd + max_increment*(BICTCP_B-1))
124             /* slow start */
125             ca->cnt = (cwnd * (BICTCP_B-1))
126                     / (cwnd - ca->last_max_cwnd);
127         else
128             /* linear increase */
129             ca->cnt = cwnd / max_increment;
130     }
131
132     /* if in slow start or link utilization is very low */
133     if (ca->last_max_cwnd == 0) {
134         if (ca->cnt > 20) /* increase cwnd 5% per RTT */
135             ca->cnt = 20;
136     }
137
138     ca->cnt = (ca->cnt << ACK_RATIO_SHIFT) / ca->delayed_ack;
139     if (ca->cnt == 0) /* cannot be zero */
140         ca->cnt = 1;
141 }
142
143 static void bictcp_cong_avoid(struct sock *sk, u32 ack, u32 acked)
144 {
145     struct tcp_sock *tp = tcp_sk(sk);
146     struct bictcp *ca = inet_csk_ca(sk);
147
148     if (!tcp_is_cwnd_limited(sk))
149         return;
150
151     if (tp->snd_cwnd <= tp->snd_ssthresh)
152         tcp_slow_start(tp, acked);
153     else {
154         bictcp_update(ca, tp->snd_cwnd);
155         tcp_cong_avoid_ai(tp, ca->cnt);
156     }
157 }
158
159 /*
160  * behave like Reno until low_window is reached,
161  * then increase congestion window slowly
162  */
163
164 static u32 bictcp_recalc_ssthresh(struct sock *sk)
165 {
166     const struct tcp_sock *tp = tcp_sk(sk);
167     struct bictcp *ca = inet_csk_ca(sk);
168
169     ca->epoch_start = 0; /* end of epoch */
170
171     /* Wmax and fast convergence */
172     if (tp->snd_cwnd < ca->last_max_cwnd && fast_convergence)
173         ca->last_max_cwnd = (tp->snd_cwnd * (BICTCP_BETA_SCALE + beta))
174                             / (2 * BICTCP_BETA_SCALE);
175     else
176         ca->last_max_cwnd = tp->snd_cwnd;
177
178     ca->loss_cwnd = tp->snd_cwnd;
179
180     if (tp->snd_cwnd <= low_window)
181         return max(tp->snd_cwnd >> 1U, 2U);
182     else
183         return max((tp->snd_cwnd * beta) / BICTCP_BETA_SCALE, 2U);
184 }

```

```

185 }
186
187 static u32 bictcp_undo_cwnd(struct sock *sk)
188 {
189     const struct tcp_sock *tp = tcp_sk(sk);
190     const struct bictcp *ca = inet_csk_ca(sk);
191     return max(tp->snd_cwnd, ca->loss_cwnd);
192 }
193
194 static void bictcp_state(struct sock *sk, u8 new_state)
195 {
196     if (new_state == TCP_CA_Loss)
197         bictcp_reset(inet_csk_ca(sk));
198 }
199
200 /* Track delayed acknowledgment ratio using sliding window
201  * ratio = (15*ratio + sample) / 16
202  */
203 static void bictcp_acked(struct sock *sk, u32 cnt, s32 rtt)
204 {
205     const struct inet_connection_sock *icsk = inet_csk(sk);
206
207     if (icsk->icsk_ca_state == TCP_CA_Open) {
208         struct bictcp *ca = inet_csk_ca(sk);
209         cnt -= ca->delayed_ack >> ACK_RATIO_SHIFT;
210         ca->delayed_ack += cnt;
211     }
212 }
213
214
215 static struct tcp_congestion_ops bictcp __read_mostly = {
216     .init          = bictcp_init,
217     .sssthresh     = bictcp_recalc_ssthresh,
218     .cong_avoid    = bictcp_cong_avoid,
219     .set_state     = bictcp_state,
220     .undo_cwnd     = bictcp_undo_cwnd,
221     .pkts_acked    = bictcp_acked,
222     .owner         = THIS_MODULE,
223     .name          = "bic",
224 };
225
226 static int __init bictcp_register(void)
227 {
228     BUILD_BUG_ON(sizeof(struct bictcp) > ICSK_CA_PRIV_SIZE);
229     return tcp_register_congestion_control(&bictcp);
230 }
231
232 static void __exit bictcp_unregister(void)
233 {
234     tcp_unregister_congestion_control(&bictcp);
235 }
236
237 module_init(bictcp_register);
238 module_exit(bictcp_unregister);
239
240 MODULE_AUTHOR("Stephen Hemminger");
241 MODULE_LICENSE("GPL");
242 MODULE_DESCRIPTION("BIC TCP");
243

```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)

- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)