

Linux Cross Reference

[Free Electrons](#)

Embedded Linux Experts

• [source navigation](#) • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version:

[2.0.40](#) [2.2.26](#) [2.4.37](#) [3.1](#) [3.2](#) [3.3](#) [3.4](#) [3.5](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#)

[Linux/net/ipv4/tcp_fastopen.c](#)

```

1 #include <linux/err.h>
2 #include <linux/init.h>
3 #include <linux/kernel.h>
4 #include <linux/list.h>
5 #include <linux/tcp.h>
6 #include <linux/rcupdate.h>
7 #include <linux/rculist.h>
8 #include <net/inetpeer.h>
9 #include <net/tcp.h>
10
11 int sysctl_tcp_fastopen __read_mostly = TFO_CLIENT_ENABLE;
12
13 struct tcp_fastopen_context __rcu *tcp_fastopen_ctx;
14
15 static DEFINE_SPINLOCK(tcp_fastopen_ctx_lock);
16
17 void tcp_fastopen_init_key_once(bool publish)
18 {
19     static u8 key[TCP_FASTOPEN_KEY_LENGTH];
20
21     /* tcp_fastopen_reset_cipher publishes the new context
22      * atomically, so we allow this race happening here.
23      *
24      * All call sites of tcp_fastopen_cookie_gen also check
25      * for a valid cookie, so this is an acceptable risk.
26      */
27     if (net_get_random_once(key, sizeof(key)) && publish)
28         tcp_fastopen_reset_cipher(key, sizeof(key));
29 }
30
31 static void tcp_fastopen_ctx_free(struct rcu_head *head)
32 {
33     struct tcp_fastopen_context *ctx =
34         container_of(head, struct tcp_fastopen_context, rcu);
35     crypto_free_cipher(ctx->tfm);
36     kfree(ctx);
37 }
38
39 int tcp_fastopen_reset_cipher(void *key, unsigned int len)
40 {
41     int err;
42     struct tcp_fastopen_context *ctx, *octx;
43 
```

```

44  ctx = kmalloc(sizeof(*ctx), GFP_KERNEL);
45  if (!ctx)
46      return -ENOMEM;
47  ctx->tfm = crypto_alloc_cipher("aes", 0, 0);
48
49  if (IS_ERR(ctx->tfm)) {
50      err = PTR_ERR(ctx->tfm);
51  error:
52      kfree(ctx);
53      pr_err("TCP: TFO aes cipher alloc error: %d\n", err);
54      return err;
55  }
56  err = crypto_cipher_setkey(ctx->tfm, key, len);
57  if (err) {
58      pr_err("TCP: TFO cipher key error: %d\n", err);
59      crypto_free_cipher(ctx->tfm);
60      goto error;
61  }
62  memcpy(ctx->key, key, len);
63
64  spin_lock(&tcp_fastopen_ctx_lock);
65
66  octx = rcu_dereference_protected(tcp_fastopen_ctx,
67                                  lockdep_is_held(&tcp_fastopen_ctx_lock));
68  rcu_assign_pointer(tcp_fastopen_ctx, ctx);
69  spin_unlock(&tcp_fastopen_ctx_lock);
70
71  if (octx)
72      call_rcu(&octx->rcu, tcp_fastopen_ctx_free);
73  return err;
74 }
75
76 static bool tcp_fastopen_cookie_gen(const void *path,
77                                     struct tcp_fastopen_cookie *foc)
78 {
79     struct tcp_fastopen_context *ctx;
80     bool ok = false;
81
82     tcp_fastopen_init_key_once(true);
83
84     rcu_read_lock();
85     ctx = rcu_dereference(tcp_fastopen_ctx);
86     if (ctx) {
87         crypto_cipher_encrypt_one(ctx->tfm, foc->val, path);
88         foc->len = TCP_FASTOPEN_COOKIE_SIZE;
89         ok = true;
90     }
91     rcu_read_unlock();
92     return ok;
93 }
94
95 /* Generate the fastopen cookie by doing aes128 encryption on both
96  * the source and destination addresses. Pad 0s for IPv4 or IPv4-mapped-IPv6
97  * addresses. For the longer IPv6 addresses use CBC-MAC.
98  *
99  * XXX (TFO) - refactor when TCP_FASTOPEN_COOKIE_SIZE != AES_BLOCK_SIZE.
100 */
101 static bool tcp_fastopen_cookie_gen(struct request_sock *req,
102                                     struct sk_buff *syn,
103                                     struct tcp_fastopen_cookie *foc)
104 {
105     if (req->rsk_ops->family == AF_INET) {
106         const struct iphdr *iph = ip_hdr(syn);
107
108         __be32 path[4] = { iph->saddr, iph->daddr, 0, 0 };
109         return tcp_fastopen_cookie_gen(path, foc);
110     }

```

```

109     }
110
111     #if IS_ENABLED(CONFIG_IPV6)
112     if (req->rsk_ops->family == AF_INET6) {
113         const struct ipv6hdr *ip6h = ipv6_hdr(syn);
114         struct tcp_fastopen_cookie tmp;
115
116         if (tcp_fastopen_cookie_gen(&ip6h->saddr, &tmp)) {
117             struct in6_addr *buf = (struct in6_addr *) tmp.val;
118             int i = 4;
119
120             for (i = 0; i < 4; i++)
121                 buf->s6_addr32[i] ^= ip6h->daddr.s6_addr32[i];
122             return tcp_fastopen_cookie_gen(buf, foc);
123         }
124     }
125 #endif
126     return false;
127 }
128
129 static bool tcp_fastopen_create_child(struct sock *sk,
130                                     struct sk_buff *skb,
131                                     struct dst_entry *dst,
132                                     struct request_sock *req)
133 {
134     struct tcp_sock *tp;
135     struct request_sock_queue *queue = &inet_csk(sk)->icsk_accept_queue;
136     struct sock *child;
137
138     req->num_retrans = 0;
139     req->num_timeout = 0;
140     req->sk = NULL;
141
142     child = inet_csk(sk)->icsk_af_ops->syn_recv_sock(sk, skb, req, NULL);
143     if (child == NULL)
144         return false;
145
146     spin_lock(&queue->fastopenq->lock);
147     queue->fastopenq->qlen++;
148     spin_unlock(&queue->fastopenq->lock);
149
150     /* Initialize the child socket. Have to fix some values to take
151      * into account the child is a Fast Open socket and is created
152      * only out of the bits carried in the SYN packet.
153      */
154     tp = tcp_sk(child);
155
156     tp->fastopen_rsk = req;
157     /* Do a hold on the listener sk so that if the listener is being
158      * closed, the child that has been accepted can live on and still
159      * access listen_lock.
160      */
161     sock_hold(sk);
162     tcp_rsk(req)->listener = sk;
163
164     /* RFC1323: The window in SYN & SYN/ACK segments is never
165      * scaled. So correct it appropriately.
166      */
167     tp->snd_wnd = ntohs(tcp_hdr(skb)->window);
168
169     /* Activate the retrans timer so that SYNACK can be retransmitted.
170      * The request socket is not added to the SYN table of the parent
171      * because it's been added to the accept queue directly.
172      */
173     inet_csk_reset_xmit_timer(child, ICSK_TIME_RETRANS,

```

```

174                                     TCP\_TIMEOUT\_INIT, TCP\_RTO\_MAX);
175
176 /* Add the child socket directly into the accept queue */
177 inet\_csk\_reqsk\_queue\_add(sk, req, child);
178
179 /* Now finish processing the fastopen child socket. */
180 inet\_csk(child)->icsk_af_ops->rebuild_header(child);
181 tcp\_init\_congestion\_control(child);
182 tcp\_mtup\_init(child);
183 tcp\_init\_metrics(child);
184 tcp\_init\_buffer\_space(child);
185
186 /* Queue the data carried in the SYN packet. We need to first
187 * bump skb's refcnt because the caller will attempt to free it.
188 *
189 * XXX (TFO) - we honor a zero-payload TFO request for now,
190 * (any reason not to?) but no need to queue the skb since
191 * there is no data. How about SYN+FIN?
192 */
193 if (TCP\_SKB\_CB(skb)->end_seq != TCP\_SKB\_CB(skb)->seq + 1) {
194     skb = skb\_get(skb);
195     skb\_dst\_drop(skb);
196     skb\_pull(skb, tcp\_hdr(skb)->doff * 4);
197     skb\_set\_owner\_r(skb, child);
198     skb\_queue\_tail(&child->sk_receive_queue, skb);
199     tp->syn_data_acked = 1;
200 }
201 tcp\_rsk(req)->rcv_nxt = tp->rcv_nxt = TCP\_SKB\_CB(skb)->end_seq;
202 sk->sk_data_ready(sk);
203 bh\_unlock\_sock(child);
204 sock\_put(child);
205 WARN\_ON(req->sk == NULL);
206 return true;
207 }
208 EXPORT\_SYMBOL(tcp\_fastopen\_create\_child);
209
210 static bool tcp\_fastopen\_queue\_check(struct sock *sk)
211 {
212     struct fastopen\_queue *fastopenq;
213
214     /* Make sure the listener has enabled fastopen, and we don't
215     * exceed the max # of pending TFO requests allowed before trying
216     * to validating the cookie in order to avoid burning CPU cycles
217     * unnecessarily.
218     *
219     * XXX (TFO) - The implication of checking the max_qlen before
220     * processing a cookie request is that clients can't differentiate
221     * between qlen overflow causing Fast Open to be disabled
222     * temporarily vs a server not supporting Fast Open at all.
223     */
224     fastopenq = inet\_csk(sk)->icsk_accept_queue.fastopenq;
225     if (fastopenq == NULL || fastopenq->max_qlen == 0)
226         return false;
227
228     if (fastopenq->qlen >= fastopenq->max_qlen) {
229         struct request\_sock *req1;
230         spin\_lock(&fastopenq->lock);
231         req1 = fastopenq->rskq_rst_head;
232         if ((req1 == NULL) || time\_after(req1->expires, jiffies)) {
233             spin\_unlock(&fastopenq->lock);
234             NET\_INC\_STATS\_BH(sock\_net(sk),
235                             LINUX\_MIB\_TCPFASTOPENLISTENOVERFLOW);
236             return false;
237         }
238         fastopenq->rskq_rst_head = req1->d1_next;

```

```

239         fastopenq->qlen--;
240         spin_unlock(&fastopenq->lock);
241         reqsk_free(req1);
242     }
243     return true;
244 }
245
246 /* Returns true if we should perform Fast Open on the SYN. The cookie (foc)
247  * may be updated and return the client in the SYN-ACK Later. E.g., Fast Open
248  * cookie request (foc->len == 0).
249  */
250 bool tcp_try_fastopen(struct sock *sk, struct sk_buff *skb,
251                      struct request_sock *req,
252                      struct tcp_fastopen_cookie *foc,
253                      struct dst_entry *dst)
254 {
255     struct tcp_fastopen_cookie valid_foc = { .len = -1 };
256     bool syn_data = TCP_SKB_CB(skb)->end_seq != TCP_SKB_CB(skb)->seq + 1;
257
258     if (!((sysctl_tcp_fastopen & TFO_SERVER_ENABLE) &&
259         (syn_data || foc->len >= 0) &&
260         tcp_fastopen_queue_check(sk))) {
261         foc->len = -1;
262         return false;
263     }
264
265     if (syn_data && (sysctl_tcp_fastopen & TFO_SERVER_COOKIE_NOT_REQD))
266         goto fastopen;
267
268     if (tcp_fastopen_cookie_gen(req, skb, &valid_foc) &&
269         foc->len == TCP_FASTOPEN_COOKIE_SIZE &&
270         foc->len == valid_foc.len &&
271         !memcmp(foc->val, valid_foc.val, foc->len)) {
272         /* Cookie is valid. Create a (full) child socket to accept
273          * the data in SYN before returning a SYN-ACK to ack the
274          * data. If we fail to create the socket, fall back and
275          * ack the ISN only but includes the same cookie.
276          *
277          * Note: Data-less SYN with valid cookie is allowed to send
278          * data in SYN_RECV state.
279          */
280     fastopen:
281         if (tcp_fastopen_create_child(sk, skb, dst, req)) {
282             foc->len = -1;
283             NET_INC_STATS_BH(sock_net(sk),
284                             LINUX_MIB_TCPFASTOPENPASSIVE);
285             return true;
286         }
287     }
288
289     NET_INC_STATS_BH(sock_net(sk), foc->len ?
290                     LINUX_MIB_TCPFASTOPENPASSIVEFAIL :
291                     LINUX_MIB_TCPFASTOPENCOOKIEREQD);
292     *foc = valid_foc;
293     return false;
294 }
295 EXPORT_SYMBOL(tcp_try_fastopen);
296

```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)

- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)