

# Linux Cross Reference

## Free Electrons

## Embedded Linux Experts

• [source navigation](#) • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version: [2.0.40](#) [2.2.26](#) [2.4.37](#) [3.1](#) [3.2](#) [3.3](#) [3.4](#) [3.5](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#)

## Linux/net/ipv4/tcp\_highspeed.c

```

1  /*
2   * Sally Floyd's High Speed TCP (RFC 3649) congestion control
3   *
4   * See http://www.icir.org/floyd/hstcp.html
5   *
6   * John Heffner <jheffner@psc.edu>
7   */
8
9  #include <linux/module.h>
10 #include <net/tcp.h>
11
12
13 /* From AIMD tables from RFC 3649 appendix B,
14  * with fixed-point MD scaled <<8.
15  */
16 static const struct hstcp\_aimd\_val {
17     unsigned int cwnd;
18     unsigned int md;
19 } hstcp\_aimd\_vals[] = {
20     { 38, 128, /* 0.50 */ },
21     { 118, 112, /* 0.44 */ },
22     { 221, 104, /* 0.41 */ },
23     { 347, 98, /* 0.38 */ },
24     { 495, 93, /* 0.37 */ },
25     { 663, 89, /* 0.35 */ },
26     { 851, 86, /* 0.34 */ },
27     { 1058, 83, /* 0.33 */ },
28     { 1284, 81, /* 0.32 */ },
29     { 1529, 78, /* 0.31 */ },
30     { 1793, 76, /* 0.30 */ },
31     { 2076, 74, /* 0.29 */ },
32     { 2378, 72, /* 0.28 */ },
33     { 2699, 71, /* 0.28 */ },
34     { 3039, 69, /* 0.27 */ },
35     { 3399, 68, /* 0.27 */ },
36     { 3778, 66, /* 0.26 */ },
37     { 4177, 65, /* 0.26 */ },
38     { 4596, 64, /* 0.25 */ },
39     { 5036, 62, /* 0.25 */ },
40     { 5497, 61, /* 0.24 */ },
41     { 5979, 60, /* 0.24 */ },
42     { 6483, 59, /* 0.23 */ },
43     { 7009, 58, /* 0.23 */ },
44     { 7558, 57, /* 0.22 */ },
45     { 8130, 56, /* 0.22 */ },

```

```

46 { 8726, 55, /* 0.22 */ },
47 { 9346, 54, /* 0.21 */ },
48 { 9991, 53, /* 0.21 */ },
49 { 10661, 52, /* 0.21 */ },
50 { 11358, 52, /* 0.20 */ },
51 { 12082, 51, /* 0.20 */ },
52 { 12834, 50, /* 0.20 */ },
53 { 13614, 49, /* 0.19 */ },
54 { 14424, 48, /* 0.19 */ },
55 { 15265, 48, /* 0.19 */ },
56 { 16137, 47, /* 0.19 */ },
57 { 17042, 46, /* 0.18 */ },
58 { 17981, 45, /* 0.18 */ },
59 { 18955, 45, /* 0.18 */ },
60 { 19965, 44, /* 0.17 */ },
61 { 21013, 43, /* 0.17 */ },
62 { 22101, 43, /* 0.17 */ },
63 { 23230, 42, /* 0.17 */ },
64 { 24402, 41, /* 0.16 */ },
65 { 25618, 41, /* 0.16 */ },
66 { 26881, 40, /* 0.16 */ },
67 { 28193, 39, /* 0.16 */ },
68 { 29557, 39, /* 0.15 */ },
69 { 30975, 38, /* 0.15 */ },
70 { 32450, 38, /* 0.15 */ },
71 { 33986, 37, /* 0.15 */ },
72 { 35586, 36, /* 0.14 */ },
73 { 37253, 36, /* 0.14 */ },
74 { 38992, 35, /* 0.14 */ },
75 { 40808, 35, /* 0.14 */ },
76 { 42707, 34, /* 0.13 */ },
77 { 44694, 33, /* 0.13 */ },
78 { 46776, 33, /* 0.13 */ },
79 { 48961, 32, /* 0.13 */ },
80 { 51258, 32, /* 0.13 */ },
81 { 53677, 31, /* 0.12 */ },
82 { 56230, 30, /* 0.12 */ },
83 { 58932, 30, /* 0.12 */ },
84 { 61799, 29, /* 0.12 */ },
85 { 64851, 28, /* 0.11 */ },
86 { 68113, 28, /* 0.11 */ },
87 { 71617, 27, /* 0.11 */ },
88 { 75401, 26, /* 0.10 */ },
89 { 79517, 26, /* 0.10 */ },
90 { 84035, 25, /* 0.10 */ },
91 { 89053, 24, /* 0.10 */ },
92 };
93
94 #define HSTCP_AIMD_MAX ARRAY_SIZE(hstcp_aimd_vals)
95
96 struct hstcp {
97     u32 ai;
98 };
99
100 static void hstcp_init(struct sock *sk)
101 {
102     struct tcp_sock *tp = tcp_sk(sk);
103     struct hstcp *ca = inet_csk_ca(sk);
104
105     ca->ai = 0;
106
107     /* Ensure the MD arithmetic works. This is somewhat pedantic,
108      * since I don't think we will see a cwnd this large. :) */
109     tp->snd_cwnd_clamp = min_t(u32, tp->snd_cwnd_clamp, 0xffffffff/128);
110 }
111

```

```

112 static void hstcp\_cong\_avoid(struct sock *sk, u32 ack, u32 acked)
113 {
114     struct tcp\_sock *tp = tcp\_sk(sk);
115     struct hstcp *ca = inet\_csk\_ca(sk);
116
117     if (!tcp\_is\_cwnd\_limited(sk))
118         return;
119
120     if (tp->snd_cwnd <= tp->snd_ssthresh)
121         tcp\_slow\_start(tp, acked);
122     else {
123         /* Update AIMD parameters.
124          *
125          * We want to guarantee that:
126          *   hstcp_aimd_vals[ca->ai-1].cwnd <
127          *   snd_cwnd <=
128          *   hstcp_aimd_vals[ca->ai].cwnd
129          */
130         if (tp->snd_cwnd > hstcp\_aimd\_vals[ca->ai].cwnd) {
131             while (tp->snd_cwnd > hstcp\_aimd\_vals[ca->ai].cwnd &&
132                   ca->ai < HSTCP\_AIMD\_MAX - 1)
133                 ca->ai++;
134         } else if (ca->ai && tp->snd_cwnd <= hstcp\_aimd\_vals[ca->ai-1].cwnd) {
135             while (ca->ai && tp->snd_cwnd <= hstcp\_aimd\_vals[ca->ai-1].cwnd)
136                 ca->ai--;
137         }
138
139         /* Do additive increase */
140         if (tp->snd_cwnd < tp->snd_cwnd_clamp) {
141             /* cwnd = cwnd + a(w) / cwnd */
142             tp->snd_cwnd_cnt += ca->ai + 1;
143             if (tp->snd_cwnd_cnt >= tp->snd_cwnd) {
144                 tp->snd_cwnd_cnt -= tp->snd_cwnd;
145                 tp->snd_cwnd++;
146             }
147         }
148     }
149 }
150
151 static u32 hstcp\_ssthresh(struct sock *sk)
152 {
153     const struct tcp\_sock *tp = tcp\_sk(sk);
154     const struct hstcp *ca = inet\_csk\_ca(sk);
155
156     /* Do multiplicative decrease */
157     return max(tp->snd_cwnd - ((tp->snd_cwnd * hstcp\_aimd\_vals[ca->ai].md) >> 8), 2U);
158 }
159
160
161 static struct tcp\_congestion\_ops tcp_highspeed __read_mostly = {
162     .init          = hstcp\_init,
163     .ssthresh      = hstcp\_ssthresh,
164     .cong_avoid    = hstcp\_cong\_avoid,
165
166     .owner         = THIS\_MODULE,
167     .name          = "highspeed"
168 };
169
170 static int __init hstcp\_register(void)
171 {
172     BUILD\_BUG\_ON(sizeof(struct hstcp) > ICSK\_CA\_PRIV\_SIZE);
173     return tcp\_register\_congestion\_control(&tcp_highspeed);
174 }
175
176 static void __exit hstcp\_unregister(void)

```

```
177 {  
178     tcp\_unregister\_congestion\_control(&tcp_highspeed);  
179 }  
180  
181 module\_init(hstcp\_register);  
182 module\_exit(hstcp\_unregister);  
183  
184 MODULE\_AUTHOR("John Heffner");  
185 MODULE\_LICENSE("GPL");  
186 MODULE\_DESCRIPTION("High Speed TCP");  
187
```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)