

Linux Cross Reference

[Free Electrons](#)

Embedded Linux Experts

• [source navigation](#) • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version:

[2.0.40](#) [2.2.26](#) [2.4.37](#) [3.1](#) [3.2](#) [3.3](#) [3.4](#) [3.5](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#)

[Linux/net/ipv4/tcp_hybla.c](#)

```

1  /*
2   * TCP HYBLA
3   *
4   * TCP-HYBLA Congestion control algorithm, based on:
5   *   C.Caini, R.Firrincieli, "TCP-Hybla: A TCP Enhancement
6   *   for Heterogeneous Networks",
7   *   International Journal on satellite Communications,
8   *   September 2004
9   *   Daniele Lacamera
10  *   root at danielinux.net
11  */
12
13 #include <linux/module.h>
14 #include <net/tcp.h>
15
16 /* Tcp Hybla structure. */
17 struct hybla {
18     bool    hybla_en;
19     u32     snd_cwnd_cents; /* Keeps increment values when it is <1, <<7 */
20     u32     rho;           /* Rho parameter, integer part */
21     u32     rho2;          /* Rho * Rho, integer part */
22     u32     rho_3ls;       /* Rho parameter, <<3 */
23     u32     rho2_7ls;      /* Rho^2, <<7 */
24     u32     minrtt_us;     /* Minimum smoothed round trip time value seen */
25 };
26
27 /* Hybla reference round trip time (default= 1/40 sec = 25 ms), in ms */
28 static int rtt0 = 25;
29 module_param(rtt0, int, 0644);
30 MODULE_PARM_DESC(rtt0, "reference rout trip time (ms)");
31
32
33 /* This is called to refresh values for hybla parameters */
34 static inline void hybla_recalc_param (struct sock *sk)
35 {
36     struct hybla *ca = inet_csk_ca(sk);
37
38     ca->rho_3ls = max_t(u32,
39                       tcp_sk(sk)->srtt_us / (rtt0 * USEC_PER_MSEC),
40                       8U);
41     ca->rho = ca->rho_3ls >> 3;
42     ca->rho2_7ls = (ca->rho_3ls * ca->rho_3ls) << 1;
43     ca->rho2 = ca->rho2_7ls >> 7;

```

```

44 }
45
46 static void hybla\_init(struct sock *sk)
47 {
48     struct tcp\_sock *tp = tcp\_sk(sk);
49     struct hybla *ca = inet\_csk\_ca(sk);
50
51     ca->rho = 0;
52     ca->rho2 = 0;
53     ca->rho_3ls = 0;
54     ca->rho2_7ls = 0;
55     ca->snd_cwnd_cents = 0;
56     ca->hybla_en = true;
57     tp->snd_cwnd = 2;
58     tp->snd_cwnd_clamp = 65535;
59
60     /* 1st Rho measurement based on initial srtt */
61     hybla\_recalc\_param(sk);
62
63     /* set minimum rtt as this is the 1st ever seen */
64     ca->minrtt_us = tp->srtt_us;
65     tp->snd_cwnd = ca->rho;
66 }
67
68 static void hybla\_state(struct sock *sk, u8 ca_state)
69 {
70     struct hybla *ca = inet\_csk\_ca(sk);
71
72     ca->hybla_en = (ca_state == TCP_CA_Open);
73 }
74
75 static inline u32 hybla\_fraction(u32 odds)
76 {
77     static const u32 fractions[] = {
78         128, 139, 152, 165, 181, 197, 215, 234,
79     };
80
81     return (odds < ARRAY\_SIZE(fractions)) ? fractions[odds] : 128;
82 }
83
84 /* TCP Hybla main routine.
85 * This is the algorithm behavior:
86 *   o Recalc Hybla parameters if min_rtt has changed
87 *   o Give cwnd a new value based on the model proposed
88 *   o remember increments <1
89 */
90 static void hybla\_cong\_avoid(struct sock *sk, u32 ack, u32 acked)
91 {
92     struct tcp\_sock *tp = tcp\_sk(sk);
93     struct hybla *ca = inet\_csk\_ca(sk);
94     u32 increment, odd, rho_fractions;
95     int is_slowstart = 0;
96
97     /* Recalculate rho only if this srtt is the lowest */
98     if (tp->srtt_us < ca->minrtt_us) {
99         hybla\_recalc\_param(sk);
100         ca->minrtt_us = tp->srtt_us;
101     }
102
103     if (!tcp\_is\_cwnd\_limited(sk))
104         return;
105
106     if (!ca->hybla_en) {
107         tcp\_reno\_cong\_avoid(sk, ack, acked);
108         return;

```

```

109     }
110
111     if (ca->rho == 0)
112         hybla_recalc_param(sk);
113
114     rho_fractions = ca->rho_3ls - (ca->rho << 3);
115
116     if (tp->snd_cwnd < tp->snd_ssthresh) {
117         /*
118          * slow start
119          * INC = 2^RHO - 1
120          * This is done by splitting the rho parameter
121          * into 2 parts: an integer part and a fraction part.
122          * Inrement<<7 is estimated by doing:
123          * [2^(int+fract)]<<7
124          * that is equal to:
125          * (2^int) * [(2^fract) <<7]
126          * 2^int is straightly computed as 1<<int,
127          * while we will use hybla_slowstart_fraction_increment() to
128          * calculate 2^fract in a <<7 value.
129          */
130         is_slowstart = 1;
131         increment = ((1 << min(ca->rho, 16U)) *
132                     hybla_fraction(rho_fractions)) - 128;
133     } else {
134         /*
135          * congestion avoidance
136          * INC = RHO^2 / W
137          * as long as increment is estimated as (rho<<7)/window
138          * it already is <<7 and we can easily count its fractions.
139          */
140         increment = ca->rho2_7ls / tp->snd_cwnd;
141         if (increment < 128)
142             tp->snd_cwnd_cnt++;
143     }
144
145     odd = increment % 128;
146     tp->snd_cwnd += increment >> 7;
147     ca->snd_cwnd_cents += odd;
148
149     /* check when fractions goes >=128 and increase cwnd by 1. */
150     while (ca->snd_cwnd_cents >= 128) {
151         tp->snd_cwnd++;
152         ca->snd_cwnd_cents -= 128;
153         tp->snd_cwnd_cnt = 0;
154     }
155     /* check when cwnd has not been incremented for a while */
156     if (increment == 0 && odd == 0 && tp->snd_cwnd_cnt >= tp->snd_cwnd) {
157         tp->snd_cwnd++;
158         tp->snd_cwnd_cnt = 0;
159     }
160     /* clamp down slowstart cwnd to ssthresh value. */
161     if (is_slowstart)
162         tp->snd_cwnd = min(tp->snd_cwnd, tp->snd_ssthresh);
163
164     tp->snd_cwnd = min_t(u32, tp->snd_cwnd, tp->snd_cwnd_clamp);
165 }
166
167 static struct tcp_congestion_ops tcp_hybla_read_mostly = {
168     .init = hybla_init,
169     .ssthresh = tcp_reno_ssthresh,
170     .cong_avoid = hybla_cong_avoid,
171     .set_state = hybla_state,
172
173     .owner = THIS_MODULE,

```

```
174     .name          = "hybla"
175 };
176
177 static int __init hybla_register(void)
178 {
179     BUILD_BUG_ON(sizeof(struct hybla) > ICSK_CA_PRIV_SIZE);
180     return tcp_register_congestion_control(&tcp_hybla);
181 }
182
183 static void __exit hybla_unregister(void)
184 {
185     tcp_unregister_congestion_control(&tcp_hybla);
186 }
187
188 module_init(hybla_register);
189 module_exit(hybla_unregister);
190
191 MODULE_AUTHOR("Daniele Lacamera");
192 MODULE_LICENSE("GPL");
193 MODULE_DESCRIPTION("TCP Hybla");
194
```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)