

# Linux Cross Reference

## Free Electrons

## Embedded Linux Experts

• *source navigation* • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version:

[2.0.40](#) [2.2.26](#) [2.4.37](#) [3.1](#) [3.2](#) [3.3](#) [3.4](#) [3.5](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#)

## Linux/net/ipv4/tcp\_lp.c

```

1  /*
2   * TCP Low Priority (TCP-LP)
3   *
4   * TCP Low Priority is a distributed algorithm whose goal is to utilize only
5   * the excess network bandwidth as compared to the ``fair share`` of
6   * bandwidth as targeted by TCP.
7   *
8   * As of 2.6.13, Linux supports pluggable congestion control algorithms.
9   * Due to the limitation of the API, we take the following changes from
10  * the original TCP-LP implementation:
11  *   o We use newReno in most core CA handling. Only add some checking
12  *     within cong_avoid.
13  *   o Error correcting in remote HZ, therefore remote HZ will be kepted
14  *     on checking and updating.
15  *   o Handling calculation of One-Way-Delay (OWD) within rtt_sample, since
16  *     OWD have a similar meaning as RTT. Also correct the buggy formular.
17  *   o Handle reaction for Early Congestion Indication (ECI) within
18  *     pkts_acked, as mentioned within pseudo code.
19  *   o OWD is handled in relative format, where local time stamp will in
20  *     tcp_time_stamp format.
21  *
22  * Original Author:
23  * Aleksandar Kuzmanovic <akuzma@northwestern.edu>
24  * Available from:
25  * http://www.ece.rice.edu/~akuzma/Doc/akuzma/TCP-LP.pdf
26  * Original implementation for 2.4.19:
27  * http://www-ece.rice.edu/networks/TCP-LP/
28  *
29  * 2.6.x module Authors:
30  * Wong Hoi Sing, Edison <hswong3i@gmail.com>
31  * Hung Hing Lun, Mike <hlhung3i@gmail.com>
32  * SourceForge project page:
33  * http://tcp-lp-mod.sourceforge.net/
34  */
35
36 #include <linux/module.h>
37 #include <net/tcp.h>
38
39 /* resolution of owd */
40 #define LP_RESOL 1000
41
42 /**
43  * enum tcp_lp_state

```

```

44 * @LP_VALID_RHZ: is remote HZ valid?
45 * @LP_VALID_OWD: is OWD valid?
46 * @LP_WITHIN_THR: are we within threshold?
47 * @LP_WITHIN_INF: are we within inference?
48 *
49 * TCP-LP's state flags.
50 * We create this set of state flag mainly for debugging.
51 */
52 enum tcp_lp_state {
53     LP_VALID_RHZ = (1 << 0),
54     LP_VALID_OWD = (1 << 1),
55     LP_WITHIN_THR = (1 << 3),
56     LP_WITHIN_INF = (1 << 4),
57 };
58
59 /**
60 * struct lp
61 * @flag: TCP-LP state flag
62 * @sowd: smoothed OWD << 3
63 * @owd_min: min OWD
64 * @owd_max: max OWD
65 * @owd_max_rsv: resrved max owd
66 * @remote_hz: estimated remote HZ
67 * @remote_ref_time: remote reference time
68 * @local_ref_time: local reference time
69 * @last_drop: time for last active drop
70 * @inference: current inference
71 *
72 * TCP-LP's private struct.
73 * We get the idea from original TCP-LP implementation where only left those we
74 * found are really useful.
75 */
76 struct lp {
77     u32 flag;
78     u32 sowd;
79     u32 owd_min;
80     u32 owd_max;
81     u32 owd_max_rsv;
82     u32 remote_hz;
83     u32 remote_ref_time;
84     u32 local_ref_time;
85     u32 last_drop;
86     u32 inference;
87 };
88
89 /**
90 * tcp_lp_init
91 *
92 * Init all required variables.
93 * Clone the handling from Vegas module implementation.
94 */
95 static void tcp_lp_init(struct sock *sk)
96 {
97     struct lp *lp = inet_csk_ca(sk);
98
99     lp->flag = 0;
100    lp->sowd = 0;
101    lp->owd_min = 0xffffffff;
102    lp->owd_max = 0;
103    lp->owd_max_rsv = 0;
104    lp->remote_hz = 0;
105    lp->remote_ref_time = 0;
106    lp->local_ref_time = 0;
107    lp->last_drop = 0;
108    lp->inference = 0;

```

```

109 }
110
111 /**
112  * tcp_lp_cong_avoid
113  *
114  * Implementation of cong_avoid.
115  * Will only call newReno CA when away from inference.
116  * From TCP-LP's paper, this will be handled in additive increasement.
117  */
118 static void tcp_lp_cong_avoid(struct sock *sk, u32 ack, u32 acked)
119 {
120     struct lp *lp = inet_csk_ca(sk);
121
122     if (!(lp->flag & LP_WITHIN_INF))
123         tcp_reno_cong_avoid(sk, ack, acked);
124 }
125
126 /**
127  * tcp_lp_remote_hz_estimator
128  *
129  * Estimate remote HZ.
130  * We keep on updating the estimated value, where original TCP-LP
131  * implementation only guest it for once and use forever.
132  */
133 static u32 tcp_lp_remote_hz_estimator(struct sock *sk)
134 {
135     struct tcp_sock *tp = tcp_sk(sk);
136     struct lp *lp = inet_csk_ca(sk);
137     s64 rhz = lp->remote_hz << 6; /* remote HZ << 6 */
138     s64 m = 0;
139
140     /* not yet record reference time
141      * go away!! record it before come back!! */
142     if (lp->remote_ref_time == 0 || lp->local_ref_time == 0)
143         goto out;
144
145     /* we can't calc remote HZ with no different!! */
146     if (tp->rx_opt.rcv_tsval == lp->remote_ref_time ||
147         tp->rx_opt.rcv_tsecr == lp->local_ref_time)
148         goto out;
149
150     m = HZ * (tp->rx_opt.rcv_tsval -
151              lp->remote_ref_time) / (tp->rx_opt.rcv_tsecr -
152                                     lp->local_ref_time);
153     if (m < 0)
154         m = -m;
155
156     if (rhz > 0) {
157         m -= rhz >> 6; /* m is now error in remote HZ est */
158         rhz += m; /* 63/64 old + 1/64 new */
159     } else
160         rhz = m << 6;
161
162 out:
163     /* record time for successful remote HZ calc */
164     if ((rhz >> 6) > 0)
165         lp->flag |= LP_VALID_RHZ;
166     else
167         lp->flag &= ~LP_VALID_RHZ;
168
169     /* record reference time stamp */
170     lp->remote_ref_time = tp->rx_opt.rcv_tsval;
171     lp->local_ref_time = tp->rx_opt.rcv_tsecr;
172
173     return rhz >> 6;

```

```

174 }
175
176 /**
177  * tcp_lp_owd_calculator
178  *
179  * Calculate one way delay (in relative format).
180  * Original implement OWD as minus of remote time difference to local time
181  * difference directly. As this time difference just simply equal to RTT, when
182  * the network status is stable, remote RTT will equal to local RTT, and result
183  * OWD into zero.
184  * It seems to be a bug and so we fixed it.
185  */
186 static u32 tcp_lp_owd_calculator(struct sock *sk)
187 {
188     struct tcp_sock *tp = tcp_sk(sk);
189     struct lp *lp = inet_csk_ca(sk);
190     s64 owd = 0;
191
192     lp->remote_hz = tcp_lp_remote_hz_estimator(sk);
193
194     if (lp->flag & LP_VALID_RHZ) {
195         owd =
196             tp->rx_opt.rcv_tsval * (LP_RESOL / lp->remote_hz) -
197             tp->rx_opt.rcv_tsecr * (LP_RESOL / HZ);
198         if (owd < 0)
199             owd = -owd;
200     }
201
202     if (owd > 0)
203         lp->flag |= LP_VALID_OWD;
204     else
205         lp->flag &= ~LP_VALID_OWD;
206
207     return owd;
208 }
209
210 /**
211  * tcp_lp_rtt_sample
212  *
213  * Implementation or rtt_sample.
214  * Will take the following action,
215  * 1. calc OWD,
216  * 2. record the min/max OWD,
217  * 3. calc smoothed OWD (SOWD).
218  * Most ideas come from the original TCP-LP implementation.
219  */
220 static void tcp_lp_rtt_sample(struct sock *sk, u32 rtt)
221 {
222     struct lp *lp = inet_csk_ca(sk);
223     s64 mowd = tcp_lp_owd_calculator(sk);
224
225     /* sorry that we don't have valid data */
226     if (!(lp->flag & LP_VALID_RHZ) || !(lp->flag & LP_VALID_OWD))
227         return;
228
229     /* record the next min owd */
230     if (mowd < lp->owd_min)
231         lp->owd_min = mowd;
232
233     /* always forget the max of the max
234      * we just set owd_max as one below it */
235     if (mowd > lp->owd_max) {
236         if (mowd > lp->owd_max_rsv) {
237             if (lp->owd_max_rsv == 0)
238                 lp->owd_max = mowd;

```

```

239         else
240             lp->owd_max = lp->owd_max_rsv;
241             lp->owd_max_rsv = mowd;
242     } else
243         lp->owd_max = mowd;
244 }
245
246 /* calc for smoothed owd */
247 if (lp->sowd != 0) {
248     mowd -= lp->sowd >> 3; /* m is now error in owd est */
249     lp->sowd += mowd;      /* owd = 7/8 owd + 1/8 new */
250 } else
251     lp->sowd = mowd << 3; /* take the measured time be owd */
252 }
253
254 /**
255  * tcp_lp_pkts_acked
256  *
257  * Implementation of pkts_acked.
258  * Deal with active drop under Early Congestion Indication.
259  * Only drop to half and 1 will be handle, because we hope to use back
260  * newReno in increase case.
261  * We work it out by following the idea from TCP-LP's paper directly
262  */
263 static void tcp_lp_pkts_acked(struct sock *sk, u32 num_acked, s32 rtt_us)
264 {
265     struct tcp_sock *tp = tcp_sk(sk);
266     struct lp *lp = inet_csk_ca(sk);
267
268     if (rtt_us > 0)
269         tcp_lp_rtt_sample(sk, rtt_us);
270
271     /* calc inference */
272     if (tcp_time_stamp > tp->rx_opt.rcv_tsecr)
273         lp->inference = 3 * (tcp_time_stamp - tp->rx_opt.rcv_tsecr);
274
275     /* test if within inference */
276     if (lp->last_drop && (tcp_time_stamp - lp->last_drop < lp->inference))
277         lp->flag |= LP_WITHIN_INF;
278     else
279         lp->flag &= ~LP_WITHIN_INF;
280
281     /* test if within threshold */
282     if (lp->sowd >> 3 <
283         lp->owd_min + 15 * (lp->owd_max - lp->owd_min) / 100)
284         lp->flag |= LP_WITHIN_THR;
285     else
286         lp->flag &= ~LP_WITHIN_THR;
287
288     pr_debug("TCP-LP: %05o/%5u/%5u/%15u/%15u/%15u\n", lp->flag,
289             tp->snd_cwnd, lp->remote_hz, lp->owd_min, lp->owd_max,
290             lp->sowd >> 3);
291
292     if (lp->flag & LP_WITHIN_THR)
293         return;
294
295     /* FIXME: try to reset owd_min and owd_max here
296      * so decrease the chance the min/max is no longer suitable
297      * and will usually within threshold when whithin inference */
298     lp->owd_min = lp->sowd >> 3;
299     lp->owd_max = lp->sowd >> 2;
300     lp->owd_max_rsv = lp->sowd >> 2;
301
302     /* happened within inference
303      * drop snd_cwnd into 1 */

```

```

304     if (lp->flag & LP_WITHIN_INF)
305         tp->snd_cwnd = 1U;
306
307     /* happened after inference
308     * cut snd_cwnd into half */
309     else
310         tp->snd_cwnd = max(tp->snd_cwnd >> 1U, 1U);
311
312     /* record this drop time */
313     lp->last_drop = tcp_time_stamp;
314 }
315
316 static struct tcp_congestion_ops tcp_lp __read_mostly = {
317     .init = tcp_lp_init,
318     .ssthresh = tcp_reno_ssthresh,
319     .cong_avoid = tcp_lp_cong_avoid,
320     .pkts_acked = tcp_lp_pkts_acked,
321
322     .owner = THIS_MODULE,
323     .name = "lp"
324 };
325
326 static int __init tcp_lp_register(void)
327 {
328     BUILD_BUG_ON(sizeof(struct lp) > ICSK_CA_PRIV_SIZE);
329     return tcp_register_congestion_control(&tcp_lp);
330 }
331
332 static void __exit tcp_lp_unregister(void)
333 {
334     tcp_unregister_congestion_control(&tcp_lp);
335 }
336
337 module_init(tcp_lp_register);
338 module_exit(tcp_lp_unregister);
339
340 MODULE_AUTHOR("Wong Hoi Sing Edison, Hung Hing Lun Mike");
341 MODULE_LICENSE("GPL");
342 MODULE_DESCRIPTION("TCP Low Priority");
343

```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)