

Linux Cross Reference

[Free Electrons](#)

Embedded Linux Experts

• [source navigation](#) • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version:

[2.0.40](#) [2.2.26](#) [2.4.37](#) [3.1](#) [3.2](#) [3.3](#) [3.4](#) [3.5](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#)

[Linux/net/ipv4/tcp_memcontrol.c](#)

```

1 #include <net/tcp.h>
2 #include <net/tcp_memcontrol.h>
3 #include <net/sock.h>
4 #include <net/ip.h>
5 #include <linux/nsproxy.h>
6 #include <linux/memcontrol.h>
7 #include <linux/module.h>
8
9 int tcp_init_cgroup(struct mem_cgroup *memcg, struct cgroup_subsys *ss)
10 {
11     /*
12      * The root cgroup does not use res_counters, but rather,
13      * rely on the data already collected by the network
14      * subsystem
15      */
16     struct res_counter *res_parent = NULL;
17     struct cg_proto *cg_proto, *parent_cg;
18     struct mem_cgroup *parent = parent_mem_cgroup(memcg);
19
20     cg_proto = tcp_prot.proto_cgroup(memcg);
21     if (!cg_proto)
22         return 0;
23
24     cg_proto->sysctl_mem[0] = sysctl_tcp_mem[0];
25     cg_proto->sysctl_mem[1] = sysctl_tcp_mem[1];
26     cg_proto->sysctl_mem[2] = sysctl_tcp_mem[2];
27     cg_proto->memory_pressure = 0;
28     cg_proto->memcg = memcg;
29
30     parent_cg = tcp_prot.proto_cgroup(parent);
31     if (parent_cg)
32         res_parent = &parent_cg->memory_allocated;
33
34     res_counter_init(&cg_proto->memory_allocated, res_parent);
35     percpu_counter_init(&cg_proto->sockets_allocated, 0);
36
37     return 0;
38 }
39 EXPORT_SYMBOL(tcp_init_cgroup);
40
41 void tcp_destroy_cgroup(struct mem_cgroup *memcg)
42 {
43     struct cg_proto *cg_proto;

```

```

44
45     cg_proto = tcp_prot.proto_cgroup(memcg);
46     if (!cg_proto)
47         return;
48
49     percpu_counter_destroy(&cg_proto->sockets_allocated);
50 }
51 EXPORT_SYMBOL(tcp_destroy_cgroup);
52
53 static int tcp_update_limit(struct mem_cgroup *memcg, u64 val)
54 {
55     struct cg_proto *cg_proto;
56     int i;
57     int ret;
58
59     cg_proto = tcp_prot.proto_cgroup(memcg);
60     if (!cg_proto)
61         return -EINVAL;
62
63     if (val > RES_COUNTER_MAX)
64         val = RES_COUNTER_MAX;
65
66     ret = res_counter_set_limit(&cg_proto->memory_allocated, val);
67     if (ret)
68         return ret;
69
70     for (i = 0; i < 3; i++)
71         cg_proto->sysctl_mem[i] = min_t(long, val >> PAGE_SHIFT,
72                                         sysctl_tcp_mem[i]);
73
74     if (val == RES_COUNTER_MAX)
75         clear_bit(MEMCG_SOCK_ACTIVE, &cg_proto->flags);
76     else if (val != RES_COUNTER_MAX) {
77         /*
78          * The active bit needs to be written after the static_key
79          * update. This is what guarantees that the socket activation
80          * function is the last one to run. See sock_update_memcg() for
81          * details, and note that we don't mark any socket as belonging
82          * to this memcg until that flag is up.
83          *
84          * We need to do this, because static_keys will span multiple
85          * sites, but we can't control their order. If we mark a socket
86          * as accounted, but the accounting functions are not patched in
87          * yet, we'll lose accounting.
88          *
89          * We never race with the readers in sock_update_memcg(),
90          * because when this value change, the code to process it is not
91          * patched in yet.
92          *
93          * The activated bit is used to guarantee that no two writers
94          * will do the update in the same memcg. Without that, we can't
95          * properly shutdown the static key.
96          */
97         if (!test_and_set_bit(MEMCG_SOCK_ACTIVATED, &cg_proto->flags))
98             static_key_slow_inc(&memcg_socket_limit_enabled);
99         set_bit(MEMCG_SOCK_ACTIVE, &cg_proto->flags);
100     }
101
102     return 0;
103 }
104
105 static ssize_t tcp_cgroup_write(struct kernfs_open_file *of,
106                                char *buf, size_t nbytes, loff_t off)
107 {
108     struct mem_cgroup *memcg = mem_cgroup_from_css(of_css(of));

```

```

109 unsigned long long val;
110 int ret = 0;
111
112 buf = rstrip(buf);
113
114 switch (of cft(of)->private) {
115 case RES_LIMIT:
116     /* see memcontrol.c */
117     ret = res_counter_memparse_write_strategy(buf, &val);
118     if (ret)
119         break;
120     ret = tcp_update_limit(memcg, val);
121     break;
122 default:
123     ret = -EINVAL;
124     break;
125 }
126 return ret ?: nbytes;
127 }
128
129 static u64 tcp_read_stat(struct mem_cgroup *memcg, int type, u64 default_val)
130 {
131     struct cg_proto *cg_proto;
132
133     cg_proto = tcp_prot.proto_cgroup(memcg);
134     if (!cg_proto)
135         return default_val;
136
137     return res_counter_read_u64(&cg_proto->memory_allocated, type);
138 }
139
140 static u64 tcp_read_usage(struct mem_cgroup *memcg)
141 {
142     struct cg_proto *cg_proto;
143
144     cg_proto = tcp_prot.proto_cgroup(memcg);
145     if (!cg_proto)
146         return atomic_long_read(&tcp_memory_allocated) << PAGE_SHIFT;
147
148     return res_counter_read_u64(&cg_proto->memory_allocated, RES_USAGE);
149 }
150
151 static u64 tcp_cgroup_read(struct cgroup_subsys_state *css, struct cftype *cft)
152 {
153     struct mem_cgroup *memcg = mem_cgroup_from_css(css);
154     u64 val;
155
156     switch (cft->private) {
157     case RES_LIMIT:
158         val = tcp_read_stat(memcg, RES_LIMIT, RES_COUNTER_MAX);
159         break;
160     case RES_USAGE:
161         val = tcp_read_usage(memcg);
162         break;
163     case RES_FAILCNT:
164     case RES_MAX_USAGE:
165         val = tcp_read_stat(memcg, cft->private, 0);
166         break;
167     default:
168         BUG();
169     }
170     return val;
171 }
172
173 static ssize_t tcp_cgroup_reset(struct kernfs_open_file *of,

```

```

174         char *buf, size_t nbytes, loff_t off)
175     {
176         struct mem_cgroup *memcg;
177         struct cg_proto *cg_proto;
178
179         memcg = mem_cgroup_from_css(of_css(of));
180         cg_proto = tcp_prot.proto_cgroup(memcg);
181         if (!cg_proto)
182             return nbytes;
183
184         switch (of_cft(of)->private) {
185         case RES_MAX_USAGE:
186             res_counter_reset_max(&cg_proto->memory_allocated);
187             break;
188         case RES_FAILCNT:
189             res_counter_reset_failcnt(&cg_proto->memory_allocated);
190             break;
191         }
192
193         return nbytes;
194     }
195
196 static struct cftype tcp_files[] = {
197     {
198         .name = "kmem.tcp.limit_in_bytes",
199         .write = tcp_cgroup_write,
200         .read_u64 = tcp_cgroup_read,
201         .private = RES_LIMIT,
202     },
203     {
204         .name = "kmem.tcp.usage_in_bytes",
205         .read_u64 = tcp_cgroup_read,
206         .private = RES_USAGE,
207     },
208     {
209         .name = "kmem.tcp.failcnt",
210         .private = RES_FAILCNT,
211         .write = tcp_cgroup_reset,
212         .read_u64 = tcp_cgroup_read,
213     },
214     {
215         .name = "kmem.tcp.max_usage_in_bytes",
216         .private = RES_MAX_USAGE,
217         .write = tcp_cgroup_reset,
218         .read_u64 = tcp_cgroup_read,
219     },
220     { } /* terminate */
221 };
222
223 static int __init tcp_memcontrol_init(void)
224 {
225     WARN_ON(cgroup_add_legacy_cftypes(&memory_cgrp_subsys, tcp_files));
226     return 0;
227 }
228 __initcall(tcp_memcontrol_init);
229

```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)

- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)