

# Linux Cross Reference

## [Free Electrons](#)

## Embedded Linux Experts

• [source navigation](#) • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version: [2.0.40](#) [2.2.26](#) [2.4.37](#) [3.1](#) [3.2](#) [3.3](#) [3.4](#) [3.5](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#)

## [Linux/net/ipv4/tcp\\_minisocks.c](#)

```

1  /*
2  * INET          An implementation of the TCP/IP protocol suite for the LINUX
3  *              operating system. INET is implemented using the BSD Socket
4  *              interface as the means of communication with the user level.
5  *
6  *              Implementation of the Transmission Control Protocol(TCP).
7  *
8  * Authors:      Ross Biro
9  *              Fred N. van Kempen, <waltje@uWalt.NL.Mugnet.ORG>
10 *              Mark Evans, <evansmp@uhura.aston.ac.uk>
11 *              Corey Minyard <wf-rch!minyard@relay.EU.net>
12 *              Florian La Roche, <flla@stud.uni-sb.de>
13 *              Charles Hedrick, <hedrick@rlinzhai.rutgers.edu>
14 *              Linus Torvalds, <torvalds@cs.helsinki.fi>
15 *              Alan Cox, <gw4pts@gw4pts.ampr.org>
16 *              Matthew Dillon, <dillon@apollo.west.oic.com>
17 *              Arnt Gulbrandsen, <agulbra@nvg.unit.no>
18 *              Jorge Cwik, <jorge@laser.satlink.net>
19 */
20
21 #include <linux/mm.h>
22 #include <linux/module.h>
23 #include <linux/slab.h>
24 #include <linux/sysctl.h>
25 #include <linux/workqueue.h>
26 #include <net/tcp.h>
27 #include <net/inet_common.h>
28 #include <net/xfrm.h>
29
30 int sysctl_tcp_syncookies __read_mostly = 1;
31 EXPORT_SYMBOL(sysctl_tcp_syncookies);
32
33 int sysctl_tcp_abort_on_overflow __read_mostly;
34
35 struct inet_timewait_death_row tcp_death_row = {
36     .sysctl_max_tw_buckets = NR_FILE * 2,
37     .period = TCP_TIMEWAIT_LEN / INET_TWDR_TWKILL_SLOTS,
38     .death_lock = SPIN_LOCK_UNLOCKED(tcp_death_row.death_lock),
39     .hashinfo = &tcp_hashinfo,
40     .tw_timer = TIMER_INITIALIZER(inet_twdr_hangman, 0,
41                                   (unsigned long)&tcp_death_row),
42     .twkill_work = WORK_INITIALIZER(tcp_death_row.twkill_work,
43                                     inet_twdr_twkill_work),
44     /* Short-time timewait calendar */
45
46     .twcal_hand = -1,
47     .twcal_timer = TIMER_INITIALIZER(inet_twdr_twcal_tick, 0,
48                                     (unsigned long)&tcp_death_row),
49 };
50 EXPORT_SYMBOL_GPL(tcp_death_row);
51
52 static bool tcp_in_window(u32 seq, u32 end_seq, u32 s_win, u32 e_win)
53 {
54     if (seq == s_win)
55         return true;
56     if (after(end_seq, s_win) && before(seq, e_win))
57         return true;
58     return seq == e_win && seq == end_seq;

```

```

59 }
60
61 /*
62  * * Main purpose of TIME-WAIT state is to close connection gracefully,
63  * * when one of ends sits in LAST-ACK or CLOSING retransmitting FIN
64  * * (and, probably, tail of data) and one or more our ACKs are lost.
65  * * What is TIME-WAIT timeout? It is associated with maximal packet
66  * * lifetime in the internet, which results in wrong conclusion, that
67  * * it is set to catch "old duplicate segments" wandering out of their path.
68  * * It is not quite correct. This timeout is calculated so that it exceeds
69  * * maximal retransmission timeout enough to allow to lose one (or more)
70  * * segments sent by peer and our ACKs. This time may be calculated from RTO.
71  * * When TIME-WAIT socket receives RST, it means that another end
72  * * finally closed and we are allowed to kill TIME-WAIT too.
73  * * Second purpose of TIME-WAIT is catching old duplicate segments.
74  * * Well, certainly it is pure paranoia, but if we load TIME-WAIT
75  * * with this semantics, we MUST NOT kill TIME-WAIT state with RSTs.
76  * * If we invented some more clever way to catch duplicates
77  * * (f.e. based on PAWS), we could truncate TIME-WAIT to several RTOs.
78  *
79  * The algorithm below is based on FORMAL INTERPRETATION of RFCs.
80  * When you compare it to RFCs, please, read section SEGMENT ARRIVES
81  * from the very beginning.
82  *
83  * NOTE. With recycling (and later with fin-wait-2) TW bucket
84  * is _not_ stateless. It means, that strictly speaking we must
85  * spinlock it. I do not want! Well, probability of misbehaviour
86  * is ridiculously low and, seems, we could use some mb() tricks
87  * to avoid misread sequence numbers, states etc. --ANK
88  *
89  * We don't need to initialize tmp_out.sack_ok as we don't use the results
90  */
91 enum tcp_tw_status
92 tcp_timewait_state_process(struct inet_timewait_sock *tw, struct sk_buff *skb,
93                           const struct tcphdr *th)
94 {
95     struct tcp_options_received tmp_opt;
96     struct tcp_timewait_sock *tcptw = tcp_tws((struct sock *)tw);
97     bool paws_reject = false;
98
99     tmp_opt.saw_tstamp = 0;
100     if (th->doff > (sizeof(*th) >> 2) && tcptw->tw_ts_recent_stamp) {
101         tcp_parse_options(skb, &tmp_opt, 0, NULL);
102
103         if (tmp_opt.saw_tstamp) {
104             tmp_opt.rcv_tsecr      -= tcptw->tw_ts_offset;
105             tmp_opt.ts_recent      = tcptw->tw_ts_recent;
106             tmp_opt.ts_recent_stamp = tcptw->tw_ts_recent_stamp;
107             paws_reject = tcp_paws_reject(&tmp_opt, th->rst);
108         }
109     }
110
111     if (tw->tw_substate == TCP_FIN_WAIT2) {
112         /* Just repeat all the checks of tcp_rcv_state_process() */
113
114         /* Out of window, send ACK */
115         if (paws_reject ||
116             !tcp_in_window(TCP_SKB_CB(skb)->seq, TCP_SKB_CB(skb)->end_seq,
117                           tcptw->tw_rcv_nxt,
118                           tcptw->tw_rcv_nxt + tcptw->tw_rcv_wnd))
119             return TCP_TW_ACK;
120
121         if (th->rst)
122             goto kill;
123
124         if (th->syn && !before(TCP_SKB_CB(skb)->seq, tcptw->tw_rcv_nxt))
125             goto kill_with_rst;
126
127         /* Dup ACK? */
128         if (!th->ack ||
129             !after(TCP_SKB_CB(skb)->end_seq, tcptw->tw_rcv_nxt) ||
130             TCP_SKB_CB(skb)->end_seq == TCP_SKB_CB(skb)->seq) {
131             inet_tws_put(tw);
132             return TCP_TW_SUCCESS;
133         }
134
135         /* New data or FIN. If new data arrive after half-duplex close,
136          * reset.

```

```

137 */
138 if (!th->fin ||
139     TCP_SKB_CB(skb)->end_seq != tcptw->tw_rcv_nxt + 1) {
140 kill_with_rst:
141     inet_tsk_deschedule(tw, &tcp_death_row);
142     inet_tsk_put(tw);
143     return TCP_TW_RST;
144 }
145
146 /* FIN arrived, enter true time-wait state. */
147 tw->tw_substate = TCP_TIME_WAIT;
148 tcptw->tw_rcv_nxt = TCP_SKB_CB(skb)->end_seq;
149 if (tmp_opt.saw_tstamp) {
150     tcptw->tw_ts_recent_stamp = get_seconds();
151     tcptw->tw_ts_recent = tmp_opt.rcv_tsval;
152 }
153
154 if (tcp_death_row.sysctl_tw_recycle &&
155     tcptw->tw_ts_recent_stamp &&
156     tcp_tw_remember_stamp(tw))
157     inet_tsk_schedule(tw, &tcp_death_row, tw->tw_timeout,
158                      TCP_TIMEWAIT_LEN);
159 else
160     inet_tsk_schedule(tw, &tcp_death_row, TCP_TIMEWAIT_LEN,
161                      TCP_TIMEWAIT_LEN);
162 return TCP_TW_ACK;
163 }
164
165 /*
166  * Now real TIME-WAIT state.
167  *
168  * RFC 1122:
169  * "When a connection is [...] on TIME-WAIT state [...]
170  * [a TCP] MAY accept a new SYN from the remote TCP to
171  * reopen the connection directly, if it:
172  *
173  * (1) assigns its initial sequence number for the new
174  * connection to be larger than the largest sequence
175  * number it used on the previous connection incarnation,
176  * and
177  *
178  * (2) returns to TIME-WAIT state if the SYN turns out
179  * to be an old duplicate".
180  */
181
182 if (!paws_reject &&
183     (TCP_SKB_CB(skb)->seq == tcptw->tw_rcv_nxt &&
184     (TCP_SKB_CB(skb)->seq == TCP_SKB_CB(skb)->end_seq || th->rst))) {
185     /* In window segment, it may be only reset or bare ack. */
186
187     if (th->rst) {
188         /* This is TIME_WAIT assassination, in two flavors.
189          * Oh well... nobody has a sufficient solution to this
190          * protocol bug yet.
191          */
192         if (sysctl_tcp_rfc1337 == 0) {
193 kill:
194             inet_tsk_deschedule(tw, &tcp_death_row);
195             inet_tsk_put(tw);
196             return TCP_TW_SUCCESS;
197         }
198     }
199     inet_tsk_schedule(tw, &tcp_death_row, TCP_TIMEWAIT_LEN,
200                      TCP_TIMEWAIT_LEN);
201
202     if (tmp_opt.saw_tstamp) {
203         tcptw->tw_ts_recent = tmp_opt.rcv_tsval;
204         tcptw->tw_ts_recent_stamp = get_seconds();
205     }
206
207     inet_tsk_put(tw);
208     return TCP_TW_SUCCESS;
209 }
210
211 /* Out of window segment.
212
213 ALL the segments are ACKed immediately.
214

```

```

215 The only exception is new SYN. We accept it, if it is
216 not old duplicate and we are not in danger to be killed
217 by delayed old duplicates. RFC check is that it has
218 newer sequence number works at rates <40Mbit/sec.
219 However, if paws works, it is reliable AND even more,
220 we even may relax silly seq space cutoff.
221
222 RED-PEN: we violate main RFC requirement, if this SYN will appear
223 old duplicate (i.e. we receive RST in reply to SYN-ACK),
224 we must return socket to time-wait state. It is not good,
225 but not fatal yet.
226 */
227
228 if (th->syn && !th->rst && !th->ack && !paws_reject &&
229     (after(TCP_SKB_CB(skb))->seq, tcptw->tw_rcv_nxt) ||
230     (tmp_opt.saw_tstamp &&
231      (s32)(tcptw->tw_ts_recent - tmp_opt.rcv_tsval) < 0))) {
232     u32 isn = tcptw->tw_snd_nxt + 65535 + 2;
233     if (isn == 0)
234         isn++;
235     TCP_SKB_CB(skb)->when = isn;
236     return TCP_TW_SYN;
237 }
238
239 if (paws_reject)
240     NET_INC_STATS_BH(twsk_net(tw), LINUX_MIB_PAWSESTABREJECTED);
241
242 if (!th->rst) {
243     /* In this case we must reset the TIMEWAIT timer.
244     */
245     * If it is ACKless SYN it may be both old duplicate
246     * and new good SYN with random sequence number <rcv_nxt.
247     * Do not reschedule in the last case.
248     */
249     if (paws_reject || th->ack)
250         inet_tsk_schedule(tw, &tcp_death_row, TCP_TIMEWAIT_LEN,
251                          TCP_TIMEWAIT_LEN);
252
253     /* Send ACK. Note, we do not put the bucket,
254     * it will be released by caller.
255     */
256     return TCP_TW_ACK;
257 }
258 inet_tsk_put(tw);
259 return TCP_TW_SUCCESS;
260 }
261 EXPORT_SYMBOL(tcp_timewait_state_process);
262
263 /*
264 * Move a socket to time-wait or dead fin-wait-2 state.
265 */
266 void tcp_time_wait(struct sock *sk, int state, int timeo)
267 {
268     struct inet_timewait_sock *tw = NULL;
269     const struct inet_connection_sock *icsk = inet_csk(sk);
270     const struct tcp_sock *tp = tcp_sk(sk);
271     bool recycle_ok = false;
272
273     if (tcp_death_row.sysctl_tw_recycle && tp->rx_opt.ts_recent_stamp)
274         recycle_ok = tcp_remember_stamp(sk);
275
276     if (tcp_death_row.tw_count < tcp_death_row.sysctl_max_tw_buckets)
277         tw = inet_tsk_alloc(sk, state);
278
279     if (tw != NULL) {
280         struct tcp_timewait_sock *tcptw = tcp_tsk((struct sock *)tw);
281         const int rto = (icsk->icsk_rto << 2) - (icsk->icsk_rto >> 1);
282         struct inet_sock *inet = inet_sk(sk);
283
284         tw->tw_transparent = inet->transparent;
285         tw->tw_rcv_wscale = tp->rx_opt.rcv_wscale;
286         tcptw->tw_rcv_nxt = tp->rcv_nxt;
287         tcptw->tw_snd_nxt = tp->snd_nxt;
288         tcptw->tw_rcv_wnd = tcp_receive_window(tp);
289         tcptw->tw_ts_recent = tp->rx_opt.ts_recent;
290         tcptw->tw_ts_recent_stamp = tp->rx_opt.ts_recent_stamp;
291         tcptw->tw_ts_offset = tp->tsoffset;
292

```

```

293 #if IS_ENABLED(CONFIG_IPV6)
294     if (tw->tw_family == PF_INET6) {
295         struct ipv6_pininfo *np = inet6_sk(sk);
296
297         tw->tw_v6_daddr = sk->sk_v6_daddr;
298         tw->tw_v6_rcv_saddr = sk->sk_v6_rcv_saddr;
299         tw->tw_tclass = np->tclass;
300         tw->tw_flowlabel = np->flow_label >> 12;
301         tw->tw_ipv6only = sk->sk_ipv6only;
302     }
303 #endif
304
305 #ifdef CONFIG_TCP_MD5SIG
306     /*
307      * The timewait bucket does not have the key DB from the
308      * sock structure. We just make a quick copy of the
309      * md5 key being used (if indeed we are using one)
310      * so the timewait ack generating code has the key.
311      */
312     do {
313         struct tcp_md5sig_key *key;
314         tcptw->tw_md5_key = NULL;
315         key = tp->af_specific->md5_lookup(sk, sk);
316         if (key != NULL) {
317             tcptw->tw_md5_key = kmemdup(key, sizeof(*key), GFP_ATOMIC);
318             if (tcptw->tw_md5_key && !tcp_alloc_md5sig_pool())
319                 BUG();
320         }
321     } while (0);
322 #endif
323
324     /* Linkage updates. */
325     inet_twsk_hashdance(tw, sk, &tcp_hashinfo);
326
327     /* Get the TIME_WAIT timeout firing. */
328     if (timeo < rto)
329         timeo = rto;
330
331     if (recycle_ok) {
332         tw->tw_timeout = rto;
333     } else {
334         tw->tw_timeout = TCP_TIMEWAIT_LEN;
335         if (state == TCP_TIME_WAIT)
336             timeo = TCP_TIMEWAIT_LEN;
337     }
338
339     inet_twsk_schedule(tw, &tcp_death_row, timeo,
340                       TCP_TIMEWAIT_LEN);
341     inet_twsk_put(tw);
342 } else {
343     /* Sorry, if we're out of memory, just CLOSE this
344      * socket up. We've got bigger problems than
345      * non-graceful socket closings.
346      */
347     NET_INC_STATS_BH(sock_net(sk), LINUX_MIB_TCPTIMEWAITOVERFLOW);
348 }
349
350 tcp_update_metrics(sk);
351 tcp_done(sk);
352 }
353
354 void tcp_twsk_destructor(struct sock *sk)
355 {
356     #ifdef CONFIG_TCP_MD5SIG
357         struct tcp_timewait_sock *twsk = tcp_twsk(sk);
358
359         if (twsk->tw_md5_key)
360             kfree_rcu(twsk->tw_md5_key, rcu);
361     #endif
362 }
363 EXPORT_SYMBOL_GPL(tcp_twsk_destructor);
364
365 void tcp_openreq_init_rwin(struct request_sock *req,
366                           struct sock *sk, struct dst_entry *dst)
367 {
368     struct inet_request_sock *ireq = inet_rsk(req);
369     struct tcp_sock *tp = tcp_sk(sk);
370     __u8 rcv_wscale;

```

```

371 int mss = dst\_metric\_advms(dst);
372
373 if (tp->rx_opt.user_mss && tp->rx_opt.user_mss < mss)
374     mss = tp->rx_opt.user_mss;
375
376 /* Set this up on the first call only */
377 req->window_clamp = tp->window_clamp ? : dst\_metric(dst, RTAX\_WINDOW);
378
379 /* Limit the window selection if the user enforce a smaller rx buffer */
380 if (sk->sk_userlocks & SOCK\_RCVBUF\_LOCK &&
381     (req->window_clamp > tcp\_full\_space(sk) || req->window_clamp == 0))
382     req->window_clamp = tcp\_full\_space(sk);
383
384 /* tcp_full_space because it is guaranteed to be the first packet */
385 tcp\_select\_initial\_window(tcp\_full\_space(sk),
386     mss - (ireq->tstamp_ok ? TCPOLEN\_TSTAMP\_ALIGNED : 0),
387     &req->rcv_wnd,
388     &req->window_clamp,
389     ireq->wscale_ok,
390     &rcv\_wscale,
391     dst\_metric(dst, RTAX\_INITRWND));
392 ireq->rcv_wscale = rcv\_wscale;
393 }
394 EXPORT\_SYMBOL(tcp\_openreq\_init\_rwin);
395
396 static inline void TCP\_ECN\_openreq\_child(struct tcp\_sock *tp,
397     struct request\_sock *req)
398 {
399     tp->ecn_flags = inet\_rsk(req)->ecn_ok ? TCP\_ECN\_OK : 0;
400 }
401
402 /* This is not only more efficient than what we used to do, it eliminates
403 * a lot of code duplication between IPv4/IPv6 SYN recv processing. -DaveM
404 *
405 * Actually, we could lots of memory writes here. tp of Listening
406 * socket contains all necessary default parameters.
407 */
408 struct sock *tcp\_create\_openreq\_child(struct sock *sk, struct request\_sock *req, struct sk\_buff *skb)
409 {
410     struct sock *newsk = inet\_csk\_clone\_lock(sk, req, GFP\_ATOMIC);
411
412     if (newsk != NULL) {
413         const struct inet\_request\_sock *ireq = inet\_rsk(req);
414         struct tcp\_request\_sock *treq = tcp\_rsk(req);
415         struct inet\_connection\_sock *newicsk = inet\_csk(newsk);
416         struct tcp\_sock *newtp = tcp\_sk(newsk);
417
418         /* Now setup tcp_sock */
419         newtp->pred_flags = 0;
420
421         newtp->rcv_wup = newtp->copied_seq =
422         newtp->rcv_nxt = treq->rcv_isn + 1;
423
424         newtp->snd_sml = newtp->snd_una =
425         newtp->snd_nxt = newtp->snd_up = treq->snt_isn + 1;
426
427         tcp\_prequeue\_init(newtp);
428         INIT\_LIST\_HEAD(&newtp->tsq_node);
429
430         tcp\_init\_wl(newtp, treq->rcv_isn);
431
432         newtp->srtt_us = 0;
433         newtp->mdev_us = jiffies\_to\_usecs(TCP\_TIMEOUT\_INIT);
434         newicsk->icsk_rto = TCP\_TIMEOUT\_INIT;
435
436         newtp->packets_out = 0;
437         newtp->retrans_out = 0;
438         newtp->sacked_out = 0;
439         newtp->fackets_out = 0;
440         newtp->snd_ssthresh = TCP\_INFINITE\_SSTHRESH;
441         tcp\_enable\_early\_retrans(newtp);
442         newtp->tlp_high_seq = 0;
443         newtp->lsndtime = treq->snt_synack;
444         newtp->total_retrans = req->num_retrans;
445
446         /* So many TCP implementations out there (incorrectly) count the
447         * initial SYN frame in their delayed-ACK and congestion control
448         * algorithms that we must have the following bandaid to talk

```

```

449     * efficiently to them.  -DaveM
450 */
451 newtp->snd_cwnd = TCP_INIT_CWND;
452 newtp->snd_cwnd_cnt = 0;
453
454 if (newicsk->icsk_ca_ops != &tcp_init_congestion_ops &&
455     !try_module_get(newicsk->icsk_ca_ops->owner))
456     newicsk->icsk_ca_ops = &tcp_init_congestion_ops;
457
458 tcp_set_ca_state(newsk, TCP_CA_Open);
459 tcp_init_xmit_timers(newsk);
460 skb_queue_head_init(&newtp->out_of_order_queue);
461 newtp->write_seq = newtp->pushed_seq = treq->snt_isn + 1;
462
463 newtp->rx_opt.saw_tstamp = 0;
464
465 newtp->rx_opt.dsack = 0;
466 newtp->rx_opt.num_sacks = 0;
467
468 newtp->urg_data = 0;
469
470 if (sock_flag(newsk, SOCK_KEEPOPEN))
471     inet_csk_reset_keepalive_timer(newsk,
472                                     keepalive_time_when(newtp));
473
474 newtp->rx_opt.tstamp_ok = ireq->tstamp_ok;
475 if ((newtp->rx_opt.sack_ok = ireq->sack_ok) != 0) {
476     if (sysctl_tcp_fack)
477         tcp_enable_fack(newtp);
478 }
479 newtp->window_clamp = req->>window_clamp;
480 newtp->rcv_ssthresh = req->rcv_wnd;
481 newtp->rcv_wnd = req->rcv_wnd;
482 newtp->rx_opt.wscale_ok = ireq->wscale_ok;
483 if (newtp->rx_opt.wscale_ok) {
484     newtp->rx_opt.snd_wscale = ireq->snd_wscale;
485     newtp->rx_opt.rcv_wscale = ireq->rcv_wscale;
486 } else {
487     newtp->rx_opt.snd_wscale = newtp->rx_opt.rcv_wscale = 0;
488     newtp->window_clamp = min(newtp->window_clamp, 65535U);
489 }
490 newtp->snd_wnd = (ntohs(tcp_hdr(skb)->window) <<
491                 newtp->rx_opt.snd_wscale);
492 newtp->max_window = newtp->snd_wnd;
493
494 if (newtp->rx_opt.tstamp_ok) {
495     newtp->rx_opt.ts_recent = req->ts_recent;
496     newtp->rx_opt.ts_recent_stamp = get_seconds();
497     newtp->tcp_header_len = sizeof(struct tcphdr) + TCPOLEN_TSTAMP_ALIGNED;
498 } else {
499     newtp->rx_opt.ts_recent_stamp = 0;
500     newtp->tcp_header_len = sizeof(struct tcphdr);
501 }
502 newtp->tsoffset = 0;
503 #ifdef CONFIG_TCP_MD5SIG
504 newtp->md5sig_info = NULL; /*XXX*/
505 if (newtp->af_specific->md5_lookup(sk, newsk))
506     newtp->tcp_header_len += TCPOLEN_MD5SIG_ALIGNED;
507 #endif
508
509 if (skb->len >= TCP_MSS_DEFAULT + newtp->tcp_header_len)
510     newicsk->icsk_ack.last_seg_size = skb->len - newtp->tcp_header_len;
511 newtp->rx_opt.mss_clamp = req->mss;
512 TCP_ECN_openreq_child(newtp, req);
513 newtp->fastopen_rsk = NULL;
514 newtp->syn_data_acked = 0;
515
516 TCP_INC_STATS_BH(sock_net(sk), TCP_MIB_PASSIVEOPENS);
517 }
518 return newsk;
519 }
520 EXPORT_SYMBOL(tcp_create_openreq_child);
521 /*
522  * Process an incoming packet for SYN_RECV sockets represented as a
523  * request_sock. Normally sk is the listener socket but for TFO it
524  * points to the child socket.
525  *
526  * XXX (TFO) - The current impl contains a special check for ack

```



```

527 * validation and inside tcp_v4_reqsk_send_ack(). Can we do better?
528 *
529 * We don't need to initialize tmp_opt.sack_ok as we don't use the results
530 */
531
532 struct sock *tcp_check_req(struct sock *sk, struct sk_buff *skb,
533                          struct request_sock *req,
534                          struct request_sock **prev,
535                          bool fastopen)
536 {
537     struct tcp_options_received tmp_opt;
538     struct sock *child;
539     const struct tcphdr *th = tcp_hdr(skb);
540     __be32 flg = tcp_flag_word(th) & (TCP_FLAG_RST|TCP_FLAG_SYN|TCP_FLAG_ACK);
541     bool paws_reject = false;
542
543     BUG_ON(fastopen == (sk->sk_state == TCP_LISTEN));
544
545     tmp_opt.saw_timestamp = 0;
546     if (th->doff > (sizeof(struct tcphdr)>>2)) {
547         tcp_parse_options(skb, &tmp_opt, 0, NULL);
548
549         if (tmp_opt.saw_timestamp) {
550             tmp_opt.ts_recent = req->ts_recent;
551             /* We do not store true stamp, but it is not required,
552              * it can be estimated (approximately)
553              * from another data.
554              */
555             tmp_opt.ts_recent_stamp = get_seconds() - ((TCP_TIMEOUT_INIT/HZ)<<req->num_timeout);
556             paws_reject = tcp_paws_reject(&tmp_opt, th->rst);
557         }
558     }
559
560     /* Check for pure retransmitted SYN. */
561     if (TCP_SKB_CB(skb)->seq == tcp_rsk(req)->rcv_isn &&
562         flg == TCP_FLAG_SYN &&
563         !paws_reject) {
564         /*
565          * RFC793 draws (Incorrectly! It was fixed in RFC1122)
566          * this case on figure 6 and figure 8, but formal
567          * protocol description says NOTHING.
568          * To be more exact, it says that we should send ACK,
569          * because this segment (at least, if it has no data)
570          * is out of window.
571          *
572          * CONCLUSION: RFC793 (even with RFC1122) DOES NOT
573          * describe SYN-RECV state. All the description
574          * is wrong, we cannot believe to it and should
575          * rely only on common sense and implementation
576          * experience.
577          *
578          * Enforce "SYN-ACK" according to figure 8, figure 6
579          * of RFC793, fixed by RFC1122.
580          *
581          * Note that even if there is new data in the SYN packet
582          * they will be thrown away too.
583          *
584          * Reset timer after retransmitting SYNACK, similar to
585          * the idea of fast retransmit in recovery.
586          */
587         if (!inet_rtx_syn_ack(sk, req))
588             req->expires = min(TCP_TIMEOUT_INIT << req->num_timeout,
589                               TCP_RTO_MAX) + jiffies;
590         return NULL;
591     }
592
593     /* Further reproduces section "SEGMENT ARRIVES"
594     for state SYN-RECEIVED of RFC793.
595     It is broken, however, it does not work only
596     when SYNs are crossed.
597
598     You would think that SYN crossing is impossible here, since
599     we should have a SYN_SENT socket (from connect()) on our end,
600     but this is not true if the crossed SYNs were sent to both
601     ends by a malicious third party. We must defend against this,
602     and to do that we first verify the ACK (as per RFC793, page
603     36) and reset if it is invalid. Is this a true full defense?
604     To convince ourselves, let us consider a way in which the ACK

```



*test can still pass in this 'malicious crossed SYNs' case.  
 Malicious sender sends identical SYNs (and thus identical sequence  
 numbers) to both A and B:*

*A: gets SYN, seq=7  
 B: gets SYN, seq=7*

*By our good fortune, both A and B select the same initial  
 send sequence number of seven :-)*

*A: sends SYN/ACK, seq=7, ack\_seq=8  
 B: sends SYN/ACK, seq=7, ack\_seq=8*

*So we are now A eating this SYN/ACK, ACK test passes. So  
 does sequence test, SYN is truncated, and thus we consider  
 it a bare ACK.*

*If icsk->icsk\_accept\_queue.rskq\_defer\_accept, we silently drop this  
 bare ACK. Otherwise, we create an established connection. Both  
 ends (listening sockets) accept the new incoming connection and try  
 to talk to each other. 8-)*

*Note: This case is both harmless, and rare. Possibility is about the  
 same as us discovering intelligent life on another planet tomorrow.*

*But generally, we should (RFC lies!) to accept ACK  
 from SYNACK both here and in tcp\_rcv\_state\_process().  
 tcp\_rcv\_state\_process() does not, hence, we do not too.*

*Note that the case is absolutely generic:  
 we cannot optimize anything here without  
 violating protocol. All the checks must be made  
 before attempt to create socket.*

```

*/
/* RFC793 page 36: "If the connection is in any non-synchronized state ...
 * and the incoming segment acknowledges something not yet
 * sent (the segment carries an unacceptable ACK) ...
 * a reset is sent."
 * Invalid ACK: reset will be sent by listening socket.
 * Note that the ACK validity check for a Fast Open socket is done
 * elsewhere and is checked directly against the child socket rather
 * than req because user data may have been sent out.
 */
if ((flg & TCP_FLAG_ACK) && !fastopen &&
    (TCP_SKB_CB(skb)->ack_seq !=
     tcp_rsk(req)->snt_isn + 1))
    return sk;

/* Also, it would be not so bad idea to check rcv_tsecr, which
 * is essentially ACK extension and too early or too late values
 * should cause reset in unsynchronized states.
 */

/* RFC793: "first check sequence number". */

if (paws_reject || !tcp_in_window(TCP_SKB_CB(skb)->seq, TCP_SKB_CB(skb)->end_seq,
    tcp_rsk(req)->rcv_nxt, tcp_rsk(req)->rcv_nxt + req->rcv_wnd)) {
    /* Out of window: send ACK and drop. */
    if (!(flg & TCP_FLAG_RST))
        req->rsk_ops->send_ack(sk, skb, req);
    if (paws_reject)
        NET_INC_STATS_BH(sock_net(sk), LINUX_MIB_PAWESESTABREJECTED);
    return NULL;
}

/* In sequence, PAWS is OK. */

if (tmp_opt.saw_timestamp && !after(TCP_SKB_CB(skb)->seq, tcp_rsk(req)->rcv_nxt))
    req->ts_recent = tmp_opt.rcv_tsval;

if (TCP_SKB_CB(skb)->seq == tcp_rsk(req)->rcv_isn) {
    /* Truncate SYN, it is out of window starting
     * at tcp_rsk(req)->rcv_isn + 1. */
    flg &= ~TCP_FLAG_SYN;
}

```

```

683  /* RFC793: "second check the RST bit" and
684  *          "fourth, check the SYN bit"
685  */
686  if (flg & (TCP_FLAG_RST|TCP_FLAG_SYN)) {
687      TCP_INC_STATS_BH(sock_net(sk), TCP_MIB_ATTEMPTFAILS);
688      goto embryonic_reset;
689  }
690
691  /* ACK sequence verified above, just make sure ACK is
692  * set. If ACK not set, just silently drop the packet.
693  *
694  * XXX (TFO) - if we ever allow "data after SYN", the
695  * following check needs to be removed.
696  */
697  if (!(flg & TCP_FLAG_ACK))
698      return NULL;
699
700  /* For Fast Open no more processing is needed (sk is the
701  * child socket).
702  */
703  if (fastopen)
704      return sk;
705
706  /* While TCP_DEFER_ACCEPT is active, drop bare ACK. */
707  if (req->num_timeout < inet_csk(sk)->icsk_accept_queue.rskq_defer_accept &&
708      TCP_SKB_CB(skb)->end_seq == tcp_rsk(req)->rcv_isn + 1) {
709      inet_rsk(req)->acked = 1;
710      NET_INC_STATS_BH(sock_net(sk), LINUX_MIB_TCPDEFERACCEPTDROP);
711      return NULL;
712  }
713
714  /* OK, ACK is valid, create big socket and
715  * feed this segment to it. It will repeat all
716  * the tests. THIS SEGMENT MUST MOVE SOCKET TO
717  * ESTABLISHED STATE. If it will be dropped after
718  * socket is created, wait for troubles.
719  */
720  child = inet_csk(sk)->icsk_af_ops->syn_rcv_sock(sk, skb, req, NULL);
721  if (child == NULL)
722      goto listen_overflow;
723
724  inet_csk_reqsk_queue_unlink(sk, req, prev);
725  inet_csk_reqsk_queue_removed(sk, req);
726
727  inet_csk_reqsk_queue_add(sk, req, child);
728  return child;
729
730 listen_overflow:
731  if (!sysctl_tcp_abort_on_overflow) {
732      inet_rsk(req)->acked = 1;
733      return NULL;
734  }
735
736 embryonic_reset:
737  if (!(flg & TCP_FLAG_RST)) {
738      /* Received a bad SYN pkt - for TFO We try not to reset
739      * the local connection unless it's really necessary to
740      * avoid becoming vulnerable to outside attack aiming at
741      * resetting legit local connections.
742      */
743      req->rsk_ops->send_reset(sk, skb);
744  } else if (fastopen) { /* received a valid RST pkt */
745      reqsk_fastopen_remove(sk, req, true);
746      tcp_reset(sk);
747  }
748  if (!fastopen) {
749      inet_csk_reqsk_queue_drop(sk, req, prev);
750      NET_INC_STATS_BH(sock_net(sk), LINUX_MIB_EMBRYONICRSTS);
751  }
752  return NULL;
753 }
754 EXPORT_SYMBOL(tcp_check_req);
755
756 /*
757  * Queue segment on the new socket if the new socket is active,
758  * otherwise we just shortcircuit this and continue with
759  * the new socket.
760  */

```

```

761  * For the vast majority of cases child->sk_state will be TCP_SYN_RECV
762  * when entering. But other states are possible due to a race condition
763  * where after __inet_lookup_established() fails but before the listener
764  * locked is obtained, other packets cause the same connection to
765  * be created.
766  */
767
768 int tcp_child_process(struct sock *parent, struct sock *child,
769                      struct sk_buff *skb)
770 {
771     int ret = 0;
772     int state = child->sk_state;
773
774     if (!sock_owned_by_user(child)) {
775         ret = tcp_rcv_state_process(child, skb, tcp_hdr(skb),
776                                     skb->len);
777         /* Wakeup parent, send SIGIO */
778         if (state == TCP_SYN_RECV && child->sk_state != state)
779             parent->sk_data_ready(parent);
780     } else {
781         /* Alas, it is possible again, because we do Lookup
782          * in main socket hash table and lock on listening
783          * socket does not protect us more.
784          */
785         __sk_add_backlog(child, skb);
786     }
787
788     bh_unlock_sock(child);
789     sock_put(child);
790     return ret;
791 }
792 EXPORT_SYMBOL(tcp_child_process);
793

```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)