

# Linux Cross Reference

## [Free Electrons](#)

## Embedded Linux Experts

• [source navigation](#) • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version: [2.0.40](#) [2.2.26](#) [2.4.37](#) [3.1](#) [3.2](#) [3.3](#) [3.4](#) [3.5](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#)

## [Linux/net/ipv4/tcp\\_output.c](#)

```

1  /*
2  * INET      An implementation of the TCP/IP protocol suite for the LINUX
3  *           operating system. INET is implemented using the BSD Socket
4  *           interface as the means of communication with the user level.
5  *
6  *           Implementation of the Transmission Control Protocol(TCP).
7  *
8  * Authors:  Ross Biro
9  *           Fred N. van Kempen, <waltje@uWalt.NL.Mugnet.ORG>
10 *           Mark Evans, <evansmp@uhura.aston.ac.uk>
11 *           Corey Minyard <wf-rch!minyard@relay.EU.net>
12 *           Florian La Roche, <flla@stud.uni-sb.de>
13 *           Charles Hedrick, <hedrick@klinzhai.rutgers.edu>
14 *           Linus Torvalds, <torvalds@cs.helsinki.fi>
15 *           Alan Cox, <gw4pts@gw4pts.ampr.org>
16 *           Matthew Dillon, <dillon@apollo.west.oic.com>
17 *           Arnt Gulbrandsen, <agulbra@nvg.unit.no>
18 *           Jorge Cwik, <jorge@Laser.satlink.net>
19 */
20
21 /*
22 * Changes:  Pedro Roque      :      Retransmit queue handled by TCP.
23 *           :                  :      Fragmentation on mtu decrease
24 *           :                  :      Segment collapse on retransmit
25 *           :                  :      AF independence
26 *
27 *           Linus Torvalds   :      send_delayed_ack
28 *           David S. Miller  :      Charge memory using the right skb
29 *           :                  :      during syn/ack processing.
30 *           David S. Miller  :      Output engine completely rewritten.
31 *           Andrea Arcangeli :      SYNACK carry ts_recent in tsecr.
32 *           Cacophonix Gaul  :      draft-minshall-nagle-01
33 *           J Hadi Salim    :      ECN support
34 *
35 */
36
37 #define pr_fmt(fmt) "TCP: " fmt
38
39 #include <net/tcp.h>
40
41 #include <linux/compiler.h>
42 #include <linux/gfp.h>
43 #include <linux/module.h>
44
45 /* People can turn this off for buggy TCP's found in printers etc. */
46 int sysctl_tcp_retrans_collapse __read_mostly = 1;
47
48 /* People can turn this on to work with those rare, broken TCPs that
49 * interpret the window field as a signed quantity.
50 */
51 int sysctl_tcp_workaround_signed_windows __read_mostly = 0;
52
53 /* Default TSQ Limit of two TSO segments */
54 int sysctl_tcp_limit_output_bytes __read_mostly = 131072;
55
56 /* This limits the percentage of the congestion window which we
57 * will allow a single TSO frame to consume. Building TSO frames
58 * which are too large can cause TCP streams to be bursty.
59 */
60 int sysctl_tcp_tso_win_divisor __read_mostly = 3;
61
62 int sysctl_tcp_mtu_probing __read_mostly = 0;
63 int sysctl_tcp_base_mss __read_mostly = TCP_BASE_MSS;
64
65 /* By default, RFC2861 behavior. */
66 int sysctl_tcp_slow_start_after_idle __read_mostly = 1;
67
68 unsigned int sysctl_tcp_notsent_lowat __read_mostly = UINT_MAX;
69 EXPORT_SYMBOL(sysctl_tcp_notsent_lowat);
70

```

```

71 static bool tcp_write_xmit(struct sock *sk, unsigned int mss_now, int nonagle,
72                             int push_one, gfp_t gfp);
73
74 /* Account for new data that has been sent to the network. */
75 static void tcp_event_new_data_sent(struct sock *sk, const struct sk_buff *skb)
76 {
77     struct inet_connection_sock *icsk = inet_csk(sk);
78     struct tcp_sock *tp = tcp_sk(sk);
79     unsigned int prior_packets = tp->packets_out;
80
81     tcp_advance_send_head(sk, skb);
82     tp->snd_nxt = TCP_SKB_CB(skb)->end_seq;
83
84     tp->packets_out += tcp_skb_pcount(skb);
85     if (!prior_packets || icsk->icsk_pending == ICSK_TIME_EARLY_RETRANS ||
86         icsk->icsk_pending == ICSK_TIME_LOSS_PROBE) {
87         tcp_rearm_rto(sk);
88     }
89
90     NET_ADD_STATS(sock_net(sk), LINUX_MIB_TCPORIGDATASENT,
91                  tcp_skb_pcount(skb));
92 }
93
94 /* SND.NXT, if window was not shrunk.
95  * If window has been shrunk, what should we make? It is not clear at all.
96  * Using SND.UNA we will fail to open window, SND.NXT is out of window. :-(
97  * Anything in between SND.UNA...SND.UNA+SND.WND also can be already
98  * invalid. OK, let's make this for now:
99  */
100 static inline __u32 tcp_acceptable_seq(const struct sock *sk)
101 {
102     const struct tcp_sock *tp = tcp_sk(sk);
103
104     if (!before(tcp_wnd_end(tp), tp->snd_nxt))
105         return tp->snd_nxt;
106     else
107         return tcp_wnd_end(tp);
108 }
109
110 /* Calculate mss to advertise in SYN segment.
111  * RFC1122, RFC1063, draft-ietf-tcpimpl-pmtud-01 state that:
112  *
113  * 1. It is independent of path mtu.
114  * 2. Ideally, it is maximal possible segment size i.e. 65535-40.
115  * 3. For IPv4 it is reasonable to calculate it from maximal MTU of
116  *    attached devices, because some buggy hosts are confused by
117  *    large MSS.
118  * 4. We do not make 3, we advertise MSS, calculated from first
119  *    hop device mtu, but allow to raise it to ip_rt_min_advmss.
120  *    This may be overridden via information stored in routing table.
121  * 5. Value 65535 for MSS is valid in IPv6 and means "as large as possible,
122  *    probably even Jumbo".
123  */
124 static __u16 tcp_advertise_mss(struct sock *sk)
125 {
126     struct tcp_sock *tp = tcp_sk(sk);
127     const struct dst_entry *dst = __sk_dst_get(sk);
128     int mss = tp->advms;
129
130     if (dst) {
131         unsigned int metric = dst_metric_advms(dst);
132
133         if (metric < mss) {
134             mss = metric;
135             tp->advms = mss;
136         }
137     }
138
139     return (__u16)mss;
140 }
141
142 /* RFC2861. Reset CWND after idle period longer RTO to "restart window".
143  * This is the first part of cwnd validation mechanism. */
144 static void tcp_cwnd_restart(struct sock *sk, const struct dst_entry *dst)
145 {
146     struct tcp_sock *tp = tcp_sk(sk);
147     s32 delta = tcp_time_stamp - tp->lsndtime;
148     __u32 restart_cwnd = tcp_init_cwnd(tp, dst);
149     __u32 cwnd = tp->snd_cwnd;
150
151     tcp_ca_event(sk, CA_EVENT_CWND_RESTART);
152
153     tp->snd_ssthresh = tcp_current_ssthresh(sk);
154     restart_cwnd = min(restart_cwnd, cwnd);
155
156     while ((delta -= inet_csk(sk)->icsk_rto) > 0 && cwnd > restart_cwnd)
157         cwnd >>= 1;
158     tp->snd_cwnd = max(cwnd, restart_cwnd);
159     tp->snd_cwnd_stamp = tcp_time_stamp;
160     tp->snd_cwnd_used = 0;

```

```

161 }
162
163 /* Congestion state accounting after a packet has been sent. */
164 static void tcp_event_data_sent(struct tcp_sock *tp,
165                               struct sock *sk)
166 {
167     struct inet_connection_sock *icsk = inet_csk(sk);
168     const u32 now = tcp_time_stamp;
169     const struct dst_entry *dst = __sk_dst_get(sk);
170
171     if (sysctl_tcp_slow_start_after_idle &&
172         (!tp->packets_out && (s32)(now - tp->lsndtime) > icsk->icsk_rto))
173         tcp_cwnd_restart(sk, __sk_dst_get(sk));
174
175     tp->lsndtime = now;
176
177     /* If it is a reply for ato after last received
178      * packet, enter pingpong mode.
179      */
180     if ((u32)(now - icsk->icsk_ack.lrcvtime) < icsk->icsk_ack.ato &&
181         (!dst || !dst_metric(dst, RTAX_QUICKACK)))
182         icsk->icsk_ack.pingpong = 1;
183 }
184
185 /* Account for an ACK we sent. */
186 static inline void tcp_event_ack_sent(struct sock *sk, unsigned int pkts)
187 {
188     tcp_dec_quickack_mode(sk, pkts);
189     inet_csk_clear_xmit_timer(sk, ICSK_TIME_DACK);
190 }
191
192
193 u32 tcp_default_init_rwnd(u32 mss)
194 {
195     /* Initial receive window should be twice of TCP_INIT_CWND to
196      * enable proper sending of new unsent data during fast recovery
197      * (RFC 3517, Section 4, NextSeg() rule (2)). Further place a
198      * limit when mss is larger than 1460.
199      */
200     u32 init_rwnd = TCP_INIT_CWND * 2;
201
202     if (mss > 1460)
203         init_rwnd = max((1460 * init_rwnd) / mss, 2U);
204     return init_rwnd;
205 }
206
207 /* Determine a window scaling and initial window to offer.
208  * Based on the assumption that the given amount of space
209  * will be offered. Store the results in the tp structure.
210  * NOTE: for smooth operation initial space offering should
211  * be a multiple of mss if possible. We assume here that mss >= 1.
212  * This MUST be enforced by all callers.
213  */
214 void tcp_select_initial_window(int __space, u32 mss,
215                               u32 *rcv_wnd, u32 *window_clamp,
216                               int wscale_ok, u8 *rcv_wscale,
217                               u32 init_rcv_wnd)
218 {
219     unsigned int space = (__space < 0 ? 0 : __space);
220
221     /* If no clamp set the clamp to the max possible scaled window */
222     if (*window_clamp == 0)
223         (*window_clamp) = (65535 << 14);
224     space = min(*window_clamp, space);
225
226     /* Quantize space offering to a multiple of mss if possible. */
227     if (space > mss)
228         space = (space / mss) * mss;
229
230     /* NOTE: offering an initial window larger than 32767
231      * will break some buggy TCP stacks. If the admin tells us
232      * it is likely we could be speaking with such a buggy stack
233      * we will truncate our initial window offering to 32K-1
234      * unless the remote has sent us a window scaling option,
235      * which we interpret as a sign the remote TCP is not
236      * misinterpreting the window field as a signed quantity.
237      */
238     if (sysctl_tcp_workaround_signed_windows)
239         (*rcv_wnd) = min(space, MAX_TCP_WINDOW);
240     else
241         (*rcv_wnd) = space;
242
243     (*rcv_wscale) = 0;
244     if (wscale_ok) {
245         /* Set window scaling on max possible window
246          * See RFC1323 for an explanation of the limit to 14
247          */
248         space = max_t(u32, sysctl_tcp_rmem[2], sysctl_rmem_max);
249         space = min_t(u32, space, *window_clamp);
250         while (space > 65535 && (*rcv_wscale) < 14) {

```

```

251         space >= 1;
252         (*rcv_wscale)++;
253     }
254 }
255
256 if (mss > (1 << *rcv_wscale)) {
257     if (!init_rcv_wnd) /* Use default unless specified otherwise */
258         init_rcv_wnd = tcp_default_init_rwnd(mss);
259     *rcv_wnd = min(*rcv_wnd, init_rcv_wnd * mss);
260 }
261
262 /* Set the clamp no higher than max representable value */
263 (*window_clamp) = min(65535U << (*rcv_wscale), *window_clamp);
264 }
265 EXPORT_SYMBOL(tcp_select_initial_window);
266
267 /* Chose a new window to advertise, update state in tcp_sock for the
268  * socket, and return result with RFC1323 scaling applied. The return
269  * value can be stuffed directly into th->window for an outgoing
270  * frame.
271  */
272 static u16 tcp_select_window(struct sock *sk)
273 {
274     struct tcp_sock *tp = tcp_sk(sk);
275     u32 old_win = tp->rcv_wnd;
276     u32 cur_win = tcp_receive_window(tp);
277     u32 new_win = tcp_select_window(sk);
278
279     /* Never shrink the offered window */
280     if (new_win < cur_win) {
281         /* Danger Will Robinson!
282          * Don't update rcv_wup/rcv_wnd here or else
283          * we will not be able to advertise a zero
284          * window in time. --DaveM
285          *
286          * Relax Will Robinson.
287          */
288         if (new_win == 0)
289             NET_INC_STATS(sock_net(sk),
290                           LINUX_MIB_TCPWANTZEROWINDOWADV);
291         new_win = ALIGN(cur_win, 1 << tp->rx_opt.rcv_wscale);
292     }
293     tp->rcv_wnd = new_win;
294     tp->rcv_wup = tp->rcv_nxt;
295
296     /* Make sure we do not exceed the maximum possible
297      * scaled window.
298      */
299     if (!tp->rx_opt.rcv_wscale && sysctl_tcp_workaround_signed_windows)
300         new_win = min(new_win, MAX_TCP_WINDOW);
301     else
302         new_win = min(new_win, (65535U << tp->rx_opt.rcv_wscale));
303
304     /* RFC1323 scaling applied */
305     new_win >= tp->rx_opt.rcv_wscale;
306
307     /* If we advertise zero window, disable fast path. */
308     if (new_win == 0) {
309         tp->pred_flags = 0;
310         if (old_win)
311             NET_INC_STATS(sock_net(sk),
312                           LINUX_MIB_TCPTOZEROWINDOWADV);
313     } else if (old_win == 0) {
314         NET_INC_STATS(sock_net(sk), LINUX_MIB_TCPFROMZEROWINDOWADV);
315     }
316
317     return new_win;
318 }
319
320 /* Packet ECN state for a SYN-ACK */
321 static inline void TCP_ECN_send_synack(const struct tcp_sock *tp, struct sk_buff *skb)
322 {
323     TCP_SKB_CB(skb)->tcp_flags &= ~TCPHDR_CWR;
324     if (!(tp->ecn_flags & TCP_ECN_OK))
325         TCP_SKB_CB(skb)->tcp_flags &= ~TCPHDR_ECE;
326 }
327
328 /* Packet ECN state for a SYN. */
329 static inline void TCP_ECN_send_syn(struct sock *sk, struct sk_buff *skb)
330 {
331     struct tcp_sock *tp = tcp_sk(sk);
332
333     tp->ecn_flags = 0;
334     if (sock_net(sk)->ipv4.sysctl_tcp_ecn == 1) {
335         TCP_SKB_CB(skb)->tcp_flags |= TCPHDR_ECE | TCPHDR_CWR;
336         tp->ecn_flags = TCP_ECN_OK;
337     }
338 }
339
340 static __inline__ void

```

```

341 TCP_ECN_make_synack(const struct request_sock *req, struct tcphdr *th)
342 {
343     if (inet_rsk(req)->ecn_ok)
344         th->ece = 1;
345 }
346
347 /* Set up ECN state for a packet on a ESTABLISHED socket that is about to
348  * be sent.
349  */
350 static inline void TCP_ECN_send(struct sock *sk, struct sk_buff *skb,
351                                int tcp_header_len)
352 {
353     struct tcp_sock *tp = tcp_sk(sk);
354
355     if (tp->ecn_flags & TCP_ECN_OK) {
356         /* Not-retransmitted data segment: set ECT and inject CWR. */
357         if (skb->len != tcp_header_len &&
358             !before(TCP_SKB_CB(skb)->seq, tp->snd_nxt)) {
359             INET_ECN_xmit(sk);
360             if (tp->ecn_flags & TCP_ECN_QUEUE_CWR) {
361                 tp->ecn_flags &= ~TCP_ECN_QUEUE_CWR;
362                 tcp_hdr(skb)->cwr = 1;
363                 skb_shinfo(skb)->gso_type |= SKB_GSO_TCP_ECN;
364             }
365         } else {
366             /* ACK or retransmitted segment: clear ECT/CE */
367             INET_ECN_dontxmit(sk);
368         }
369         if (tp->ecn_flags & TCP_ECN_DEMAND_CWR)
370             tcp_hdr(skb)->ece = 1;
371     }
372 }
373
374 /* Constructs common control bits of non-data skb. If SYN/FIN is present,
375  * auto increment end seqno.
376  */
377 static void tcp_init_nondata_skb(struct sk_buff *skb, u32 seq, u8 flags)
378 {
379     struct skb_shared_info *shinfo = skb_shinfo(skb);
380
381     skb->ip_summed = CHECKSUM_PARTIAL;
382     skb->csum = 0;
383
384     TCP_SKB_CB(skb)->tcp_flags = flags;
385     TCP_SKB_CB(skb)->sacked = 0;
386
387     shinfo->gso_segs = 1;
388     shinfo->gso_size = 0;
389     shinfo->gso_type = 0;
390
391     TCP_SKB_CB(skb)->seq = seq;
392     if (flags & (TCPHDR_SYN | TCPhdr_FIN))
393         seq++;
394     TCP_SKB_CB(skb)->end_seq = seq;
395 }
396
397 static inline bool tcp_urg_mode(const struct tcp_sock *tp)
398 {
399     return tp->snd_una != tp->snd_up;
400 }
401
402 #define OPTION_SACK_ADVERTISE (1 << 0)
403 #define OPTION_TS (1 << 1)
404 #define OPTION_MD5 (1 << 2)
405 #define OPTION_WSCALE (1 << 3)
406 #define OPTION_FAST_OPEN_COOKIE (1 << 8)
407
408 struct tcp_out_options {
409     u16 options; /* bit field of OPTION_* */
410     u16 mss; /* 0 to disable */
411     u8 ws; /* window scale, 0 to disable */
412     u8 num_sack_blocks; /* number of SACK blocks to include */
413     u8 hash_size; /* bytes in hash_location */
414     u8 *hash_location; /* temporary pointer, overloaded */
415     u32 tsval, tsecr; /* need to include OPTION_TS */
416     struct tcp_fastopen_cookie *fastopen_cookie; /* Fast open cookie */
417 };
418
419 /* Write previously computed TCP options to the packet.
420  *
421  * Beware: Something in the Internet is very sensitive to the ordering of
422  * TCP options, we learned this through the hard way, so be careful here.
423  * Luckily we can at least blame others for their non-compliance but from
424  * inter-operability perspective it seems that we're somewhat stuck with
425  * the ordering which we have been using if we want to keep working with
426  * those broken things (not that it currently hurts anybody as there isn't
427  * particular reason why the ordering would need to be changed).
428  *
429  * At Least SACK_PERM as the first option is known to lead to a disaster
430  * (but it may well be that other scenarios fail similarly).

```

```

431 */
432 static void tcp_options_write(__be32 *ptr, struct tcp_sock *tp,
433                             struct tcp_out_options *opts)
434 {
435     u16 options = opts->options;    /* mungable copy */
436
437     if (unlikely(OPTION_MD5 & options)) {
438         *ptr++ = htonl((TCPOPT_NOP << 24) | (TCPOPT_NOP << 16) |
439                       (TCPOPT_MD5SIG << 8) | TCPOLEN_MD5SIG);
440         /* overLoad cookie hash Location */
441         opts->hash_location = (__u8 *)ptr;
442         ptr += 4;
443     }
444
445     if (unlikely(opts->mss)) {
446         *ptr++ = htonl((TCPOPT_MSS << 24) |
447                       (TCPOLEN_MSS << 16) |
448                       opts->mss);
449     }
450
451     if (likely(OPTION_TS & options)) {
452         if (unlikely(OPTION_SACK_ADVERTISE & options)) {
453             *ptr++ = htonl((TCPOPT_SACK_PERM << 24) |
454                           (TCPOLEN_SACK_PERM << 16) |
455                           (TCPOPT_TIMESTAMP << 8) |
456                           TCPOLEN_TIMESTAMP);
457             options &= ~OPTION_SACK_ADVERTISE;
458         } else {
459             *ptr++ = htonl((TCPOPT_NOP << 24) |
460                           (TCPOPT_NOP << 16) |
461                           (TCPOPT_TIMESTAMP << 8) |
462                           TCPOLEN_TIMESTAMP);
463         }
464         *ptr++ = htonl(opts->tsval);
465         *ptr++ = htonl(opts->tsecr);
466     }
467
468     if (unlikely(OPTION_SACK_ADVERTISE & options)) {
469         *ptr++ = htonl((TCPOPT_NOP << 24) |
470                       (TCPOPT_NOP << 16) |
471                       (TCPOPT_SACK_PERM << 8) |
472                       TCPOLEN_SACK_PERM);
473     }
474
475     if (unlikely(OPTION_WSCALE & options)) {
476         *ptr++ = htonl((TCPOPT_NOP << 24) |
477                       (TCPOPT_WINDOW << 16) |
478                       (TCPOLEN_WINDOW << 8) |
479                       opts->ws);
480     }
481
482     if (unlikely(opts->num_sack_blocks)) {
483         struct tcp_sack_block *sp = tp->rx_opt.dsack ?
484             tp->duplicate_sack : tp->selective_acks;
485         int this_sack;
486
487         *ptr++ = htonl((TCPOPT_NOP << 24) |
488                       (TCPOPT_NOP << 16) |
489                       (TCPOPT_SACK << 8) |
490                       (TCPOLEN_SACK_BASE + (opts->num_sack_blocks *
491                                           TCPOLEN_SACK_PERBLOCK)));
492
493         for (this_sack = 0; this_sack < opts->num_sack_blocks;
494             ++this_sack) {
495             *ptr++ = htonl(sp[this_sack].start_seq);
496             *ptr++ = htonl(sp[this_sack].end_seq);
497         }
498
499         tp->rx_opt.dsack = 0;
500     }
501
502     if (unlikely(OPTION_FAST_OPEN_COOKIE & options)) {
503         struct tcp_fastopen_cookie *foc = opts->fastopen_cookie;
504
505         *ptr++ = htonl((TCPOPT_EXP << 24) |
506                       ((TCPOLEN_EXP_FASTOPEN_BASE + foc->len) << 16) |
507                       TCPOPT_FASTOPEN_MAGIC);
508
509         memcpy(ptr, foc->val, foc->len);
510         if ((foc->len & 3) == 2) {
511             u8 *align = ((__u8 *)ptr) + foc->len;
512             align[0] = align[1] = TCPOPT_NOP;
513         }
514         ptr += (foc->len + 3) >> 2;
515     }
516 }
517
518 /* Compute TCP options for SYN packets. This is not the final
519  * network wire format yet.
520 */

```

```

521 static unsigned int tcp_syn_options(struct sock *sk, struct sk_buff *skb,
522                                     struct tcp_out_options *opts,
523                                     struct tcp_md5sig_key **md5)
524 {
525     struct tcp_sock *tp = tcp_sk(sk);
526     unsigned int remaining = MAX_TCP_OPTION_SPACE;
527     struct tcp_fastopen_request *fastopen = tp->fastopen_req;
528
529     #ifdef CONFIG_TCP_MD5SIG
530     *md5 = tp->af_specific->md5_lookup(sk, sk);
531     if (*md5) {
532         opts->options |= OPTION_MD5;
533         remaining -= TCPOLEN_MD5SIG_ALIGNED;
534     }
535     #else
536     *md5 = NULL;
537     #endif
538
539     /* We always get an MSS option. The option bytes which will be seen in
540     * normal data packets should timestamps be used, must be in the MSS
541     * advertised. But we subtract them from tp->mss_cache so that
542     * calculations in tcp_sendmsg are simpler etc. So account for this
543     * fact here if necessary. If we don't do this correctly, as a
544     * receiver we won't recognize data packets as being full sized when we
545     * should, and thus we won't abide by the delayed ACK rules correctly.
546     * SACKs don't matter, we never delay an ACK when we have any of those
547     * going out. */
548     opts->mss = tcp_advertise_mss(sk);
549     remaining -= TCPOLEN_MSS_ALIGNED;
550
551     if (likely(sysctl_tcp_timestamps && *md5 == NULL)) {
552         opts->options |= OPTION_TS;
553         opts->tsval = TCP_SKB_CB(skb)->when + tp->tsoffset;
554         opts->tsecr = tp->rx_opt.ts_recent;
555         remaining -= TCPOLEN_TSTAMP_ALIGNED;
556     }
557     if (likely(sysctl_tcp_window_scaling)) {
558         opts->ws = tp->rx_opt.rcv_wscale;
559         opts->options |= OPTION_WSCALE;
560         remaining -= TCPOLEN_WSCALE_ALIGNED;
561     }
562     if (likely(sysctl_tcp_sack)) {
563         opts->options |= OPTION_SACK_ADVERTISE;
564         if (unlikely(!(OPTION_TS & opts->options)))
565             remaining -= TCPOLEN_SACKPERM_ALIGNED;
566     }
567
568     if (fastopen && fastopen->cookie.len >= 0) {
569         u32 need = TCPOLEN_EXP_FASTOPEN_BASE + fastopen->cookie.len;
570         need = (need + 3) & ~3U; /* Align to 32 bits */
571         if (remaining >= need) {
572             opts->options |= OPTION_FAST_OPEN_COOKIE;
573             opts->fastopen_cookie = &fastopen->cookie;
574             remaining -= need;
575             tp->syn_fastopen = 1;
576         }
577     }
578
579     return MAX_TCP_OPTION_SPACE - remaining;
580 }
581
582 /* Set up TCP options for SYN-ACKs. */
583 static unsigned int tcp_synack_options(struct sock *sk,
584                                         struct request_sock *req,
585                                         unsigned int mss, struct sk_buff *skb,
586                                         struct tcp_out_options *opts,
587                                         struct tcp_md5sig_key **md5,
588                                         struct tcp_fastopen_cookie *foc)
589 {
590     struct inet_request_sock *ireq = inet_rsk(req);
591     unsigned int remaining = MAX_TCP_OPTION_SPACE;
592
593     #ifdef CONFIG_TCP_MD5SIG
594     *md5 = tcp_rsk(req)->af_specific->md5_lookup(sk, req);
595     if (*md5) {
596         opts->options |= OPTION_MD5;
597         remaining -= TCPOLEN_MD5SIG_ALIGNED;
598
599         /* We can't fit any SACK blocks in a packet with MD5 + TS
600         * options. There was discussion about disabling SACK
601         * rather than TS in order to fit in better with old,
602         * buggy kernels, but that was deemed to be unnecessary.
603         */
604         ireq->tstamp_ok &= !ireq->sack_ok;
605     }
606     #else
607     *md5 = NULL;
608     #endif
609
610     /* We always send an MSS option. */

```



```

611     opts->mss = mss;
612     remaining -= TCPOLEN_MSS_ALIGNED;
613
614     if (likely(ireq->wscale_ok)) {
615         opts->ws = ireq->rcv_wscale;
616         opts->options |= OPTION_WSCALE;
617         remaining -= TCPOLEN_WSCALE_ALIGNED;
618     }
619     if (likely(ireq->tstamp_ok)) {
620         opts->options |= OPTION_TS;
621         opts->tsval = TCP_SKB_CB(skb)->when;
622         opts->tsecr = req->ts_recent;
623         remaining -= TCPOLEN_TSTAMP_ALIGNED;
624     }
625     if (likely(ireq->sack_ok)) {
626         opts->options |= OPTION_SACK_ADVERTISE;
627         if (unlikely(!ireq->tstamp_ok))
628             remaining -= TCPOLEN_SACKPERM_ALIGNED;
629     }
630     if (foc != NULL && foc->len >= 0) {
631         u32 need = TCPOLEN_EXP_FASTOPEN_BASE + foc->len;
632         need = (need + 3) & ~3U; /* Align to 32 bits */
633         if (remaining >= need) {
634             opts->options |= OPTION_FAST_OPEN_COOKIE;
635             opts->fastopen_cookie = foc;
636             remaining -= need;
637         }
638     }
639     return MAX_TCP_OPTION_SPACE - remaining;
640 }
641
642 /* Compute TCP options for ESTABLISHED sockets. This is not the
643  * final wire format yet.
644  */
645 static unsigned int tcp_established_options(struct sock *sk, struct sk_buff *skb,
646                                           struct tcp_out_options *opts,
647                                           struct tcp_md5sig_key **md5)
648 {
649     struct tcp_skb_cb *tcb = skb ? TCP_SKB_CB(skb) : NULL;
650     struct tcp_sock *tp = tcp_sk(sk);
651     unsigned int size = 0;
652     unsigned int eff_sacks;
653
654     opts->options = 0;
655
656 #ifdef CONFIG_TCP_MD5SIG
657     *md5 = tp->af_specific->md5_lookup(sk, sk);
658     if (unlikely(*md5)) {
659         opts->options |= OPTION_MD5;
660         size += TCPOLEN_MD5SIG_ALIGNED;
661     }
662 #else
663     *md5 = NULL;
664 #endif
665
666     if (likely(tp->rx_opt.tstamp_ok)) {
667         opts->options |= OPTION_TS;
668         opts->tsval = tcb ? tcb->when + tp->tsoffset : 0;
669         opts->tsecr = tp->rx_opt.ts_recent;
670         size += TCPOLEN_TSTAMP_ALIGNED;
671     }
672
673     eff_sacks = tp->rx_opt.num_sacks + tp->rx_opt.dsack;
674     if (unlikely(eff_sacks)) {
675         const unsigned int remaining = MAX_TCP_OPTION_SPACE - size;
676         opts->num_sack_blocks =
677             min_t(unsigned int, eff_sacks,
678                 (remaining - TCPOLEN_SACK_BASE_ALIGNED) /
679                 TCPOLEN_SACK_PERBLOCK);
680         size += TCPOLEN_SACK_BASE_ALIGNED +
681             opts->num_sack_blocks * TCPOLEN_SACK_PERBLOCK;
682     }
683
684     return size;
685 }
686
687 /* TCP SMALL QUEUES (TSQ)
688  *
689  * TSQ goal is to keep small amount of skbs per tcp flow in tx queues (qdisc+dev)
690  * to reduce RTT and bufferbloat.
691  * We do this using a special skb destructor (tcp_wfree).
692  *
693  * Its important tcp_wfree() can be replaced by sock_wfree() in the event skb
694  * needs to be reallocated in a driver.
695  * The invariant being skb->truesize subtracted from sk->sk_wmem_alloc
696  *
697  * Since transmit from skb destructor is forbidden, we use a tasklet
698  * to process all sockets that eventually need to send more skbs.
699
700

```



```

701  * We use one tasklet per cpu, with its own queue of sockets.
702  */
703  struct tsq_tasklet {
704      struct tasklet_struct tasklet;
705      struct list_head head; /* queue of tcp sockets */
706  };
707  static DEFINE_PER_CPU(struct tsq_tasklet, tsq_tasklet);
708
709  static void tcp_tsq_handler(struct sock *sk)
710  {
711      if ((1 << sk->sk_state) &
712          (TCPF_ESTABLISHED | TCPF_FIN_WAIT1 | TCPF_CLOSING |
713           TCPF_CLOSE_WAIT | TCPF_LAST_ACK))
714          tcp_write_xmit(sk, tcp_current_mss(sk), tcp_sk(sk)->nonagle,
715                        0, GFP_ATOMIC);
716  }
717  /*
718   * One tasklet per cpu tries to send more skbs.
719   * We run in tasklet context but need to disable irqs when
720   * transferring tsq->head because tcp_wfree() might
721   * interrupt us (non NAPI drivers)
722   */
723  static void tcp_tasklet_func(unsigned long data)
724  {
725      struct tsq_tasklet *tsq = (struct tsq_tasklet *)data;
726      LIST_HEAD(list);
727      unsigned long flags;
728      struct list_head *q, *n;
729      struct tcp_sock *tp;
730      struct sock *sk;
731
732      local_irq_save(flags);
733      list_splice_init(&tsq->head, &list);
734      local_irq_restore(flags);
735
736      list_for_each_safe(q, n, &list) {
737          tp = list_entry(q, struct tcp_sock, tsq_node);
738          list_del(&tp->tsq_node);
739
740          sk = (struct sock *)tp;
741          bh_lock_sock(sk);
742
743          if (!sock_owned_by_user(sk)) {
744              tcp_tsq_handler(sk);
745          } else {
746              /* defer the work to tcp_release_cb() */
747              set_bit(TCP_TSQ_DEFERRED, &tp->tsq_flags);
748          }
749          bh_unlock_sock(sk);
750
751          clear_bit(TSQ_QUEUED, &tp->tsq_flags);
752          sk_free(sk);
753      }
754  }
755
756  #define TCP_DEFERRED_ALL ((1UL << TCP_TSQ_DEFERRED) | \
757                          (1UL << TCP_WRITE_TIMER_DEFERRED) | \
758                          (1UL << TCP_DELACK_TIMER_DEFERRED) | \
759                          (1UL << TCP_MTU_REDUCED_DEFERRED))
760  /**
761   * tcp_release_cb - tcp release_sock() callback
762   * @sk: socket
763   *
764   * called from release_sock() to perform protocol dependent
765   * actions before socket release.
766   */
767  void tcp_release_cb(struct sock *sk)
768  {
769      struct tcp_sock *tp = tcp_sk(sk);
770      unsigned long flags, nflags;
771
772      /* perform an atomic operation only if at least one flag is set */
773      do {
774          flags = tp->tsq_flags;
775          if (!(flags & TCP_DEFERRED_ALL))
776              return;
777          nflags = flags & ~TCP_DEFERRED_ALL;
778      } while (cmpxchg(&tp->tsq_flags, flags, nflags) != flags);
779
780      if (flags & (1UL << TCP_TSQ_DEFERRED))
781          tcp_tsq_handler(sk);
782
783      /* Here begins the tricky part :
784       * We are called from release_sock() with :
785       * 1) BH disabled
786       * 2) sk_lock.slock spinlock held
787       * 3) socket owned by us (sk->sk_lock.owned == 1)
788       *
789       * But following code is meant to be called from BH handlers,
790       * so we should keep BH disabled, but early release socket ownership

```

```

791     */
792     sock_release_ownership(sk);
793
794     if (flags & (1UL << TCP_WRITE_TIMER_DEFERRED)) {
795         tcp_write_timer_handler(sk);
796         sock_put(sk);
797     }
798     if (flags & (1UL << TCP_DELACK_TIMER_DEFERRED)) {
799         tcp_delack_timer_handler(sk);
800         sock_put(sk);
801     }
802     if (flags & (1UL << TCP_MTU_REDUCED_DEFERRED)) {
803         inet_csk(sk)->icsk_af_ops->mtu_reduced(sk);
804         sock_put(sk);
805     }
806 }
807 EXPORT_SYMBOL(tcp_release_cb);
808
809 void __init tcp_tasklet_init(void)
810 {
811     int i;
812
813     for_each_possible_cpu(i) {
814         struct tsq_tasklet *tsq = &per_cpu(tsq_tasklet, i);
815
816         INIT_LIST_HEAD(&tsq->head);
817         tasklet_init(&tsq->tasklet,
818                     tcp_tasklet_func,
819                     (unsigned long)tsq);
820     }
821 }
822
823 /*
824  * Write buffer destructor automatically called from kfree_skb.
825  * We can't xmit new skbs from this context, as we might already
826  * hold qdisc lock.
827  */
828 void tcp_wfree(struct sk_buff *skb)
829 {
830     struct sock *sk = skb->sk;
831     struct tcp_sock *tp = tcp_sk(sk);
832
833     if (test_and_clear_bit(TSQ_THROTTLED, &tp->tsq_flags) &&
834         !test_and_set_bit(TSQ_QUEUED, &tp->tsq_flags)) {
835         unsigned long flags;
836         struct tsq_tasklet *tsq;
837
838         /* Keep a ref on socket.
839          * This last ref will be released in tcp_tasklet_func()
840          */
841         atomic_sub(skb->truesize - 1, &sk->sk_wmem_alloc);
842
843         /* queue this socket to tasklet queue */
844         local_irq_save(flags);
845         tsq = &__get_cpu_var(tsq_tasklet);
846         list_add(&tp->tsq_node, &tsq->head);
847         tasklet_schedule(&tsq->tasklet);
848         local_irq_restore(flags);
849     } else {
850         sock_wfree(skb);
851     }
852 }
853
854 /* This routine actually transmits TCP packets queued in by
855  * tcp_do_sendmsg(). This is used by both the initial
856  * transmission and possible later retransmissions.
857  * All SKB's seen here are completely headerless. It is our
858  * job to build the TCP header, and pass the packet down to
859  * IP so it can do the same plus pass the packet off to the
860  * device.
861  *
862  * We are working here with either a clone of the original
863  * SKB, or a fresh unique copy made by the retransmit engine.
864  */
865 static int tcp_transmit_skb(struct sock *sk, struct sk_buff *skb, int clone_it,
866                            gfp_t gfp_mask)
867 {
868     const struct inet_connection_sock *icsk = inet_csk(sk);
869     struct inet_sock *inet;
870     struct tcp_sock *tp;
871     struct tcp_skb_cb *tcb;
872     struct tcp_out_options opts;
873     unsigned int tcp_options_size, tcp_header_size;
874     struct tcp_md5sig_key *md5;
875     struct tcphdr *th;
876     int err;
877
878     BUG_ON(!skb || !tcp_skb_pcount(skb));
879
880     if (clone_it) {

```

```

881     skb_mstamp_get(&skb->skb_mstamp);
882
883     if (unlikely(skb_cloned(skb)))
884         skb = pskb_copy(skb, gfp_mask);
885     else
886         skb = skb_clone(skb, gfp_mask);
887     if (unlikely(!skb))
888         return -ENOBUFS;
889     /* Our usage of tstamp should remain private */
890     skb->tstamp.tv64 = 0;
891 }
892
893 inet = inet_sk(sk);
894 tp = tcp_sk(sk);
895 tcb = TCP_SKB_CB(skb);
896 memset(&opts, 0, sizeof(opts));
897
898 if (unlikely(tcb->tcp_flags & TCPHDR_SYN))
899     tcp_options_size = tcp_syn_options(sk, skb, &opts, &md5);
900 else
901     tcp_options_size = tcp_established_options(sk, skb, &opts,
902                                               &md5);
903 tcp_header_size = tcp_options_size + sizeof(struct tcphdr);
904
905 if (tcp_packets_in_flight(tp) == 0)
906     tcp_ca_event(sk, CA_EVENT_TX_START);
907
908 /* if no packet is in qdisc/device queue, then allow XPS to select
909  * another queue.
910  */
911 skb->ooo_okay = sk_wmem_alloc_get(sk) == 0;
912
913 skb_push(skb, tcp_header_size);
914 skb_reset_transport_header(skb);
915
916 skb_orphan(skb);
917 skb->sk = sk;
918 skb->destructor = tcp_wfree;
919 skb_set_hash_from_sk(skb, sk);
920 atomic_add(skb->truesize, &sk->sk_wmem_alloc);
921
922 /* Build TCP header and checksum it. */
923 th = tcp_hdr(skb);
924 th->source      = inet->inet_sport;
925 th->dest        = inet->inet_dport;
926 th->seq         = htonl(tcb->seq);
927 th->ack_seq     = htonl(tp->rcv_nxt);
928 *((__be16 *)th) + 6 = htons(((tcp_header_size >> 2) << 12) |
929                             tcb->tcp_flags);
930
931 if (unlikely(tcb->tcp_flags & TCPHDR_SYN)) {
932     /* RFC1323: The window in SYN & SYN/ACK segments
933      * is never scaled.
934      */
935     th->window = htons(min(tp->rcv_wnd, 65535U));
936 } else {
937     th->window = htons(tcp_select_window(sk));
938 }
939 th->check = 0;
940 th->urg_ptr = 0;
941
942 /* The urg_mode check is necessary during a below snd_una win probe */
943 if (unlikely(tcp_urg_mode(tp) && before(tcb->seq, tp->snd_up))) {
944     if (before(tp->snd_up, tcb->seq + 0x10000)) {
945         th->urg_ptr = htons(tp->snd_up - tcb->seq);
946         th->urg = 1;
947     } else if (after(tcb->seq + 0xFFFF, tp->snd_nxt)) {
948         th->urg_ptr = htons(0xFFFF);
949         th->urg = 1;
950     }
951 }
952
953 tcp_options_write((__be32 *) (th + 1), tp, &opts);
954 if (likely((tcb->tcp_flags & TCPHDR_SYN) == 0))
955     TCP_ECN_send(sk, skb, tcp_header_size);
956
957 #ifdef CONFIG_TCP_MD5SIG
958 /* Calculate the MD5 hash, as we have all we need now */
959 if (md5) {
960     sk_nocaps_add(sk, NETIF_F_GSO_MASK);
961     tp->af_specific->calc_md5_hash(opts.hash_location,
962                                   md5, sk, NULL, skb);
963 }
964 #endif
965
966 icsk->icsk_af_ops->send_check(sk, skb);
967
968 if (likely(tcb->tcp_flags & TCPHDR_ACK))
969     tcp_event_ack_sent(sk, tcp_skb_pcount(skb));
970

```

```

971     if (skb->len != tcp_header_size)
972         tcp_event_data_sent(tp, sk);
973
974     if (after(tcb->end_seq, tp->snd_nxt) || tcb->seq == tcb->end_seq)
975         TCP_ADD_STATS(sock_net(sk), TCP_MIB_OUTSEGS,
976             tcp_skb_pcount(skb));
977
978     err = icsk->icsk_af_ops->queue_xmit(sk, skb, &inet->cork.fl);
979     if (likely(err <= 0))
980         return err;
981
982     tcp_enter_cwr(sk);
983
984     return net_xmit_eval(err);
985 }
986
987 /* This routine just queues the buffer for sending.
988  *
989  * NOTE: probe0 timer is not checked, do not forget tcp_push_pending_frames,
990  * otherwise socket can stall.
991  */
992 static void tcp_queue_skb(struct sock *sk, struct sk_buff *skb)
993 {
994     struct tcp_sock *tp = tcp_sk(sk);
995
996     /* Advance write_seq and place onto the write_queue. */
997     tp->write_seq = TCP_SKB_CB(skb)->end_seq;
998     skb_header_release(skb);
999     tcp_add_write_queue_tail(sk, skb);
1000     sk->sk_wmem_queued += skb->truesize;
1001     sk_mem_charge(sk, skb->truesize);
1002 }
1003
1004 /* Initialize TSO segments for a packet. */
1005 static void tcp_set_skb_tso_segs(const struct sock *sk, struct sk_buff *skb,
1006     unsigned int mss_now)
1007 {
1008     struct skb_shared_info *shinfo = skb_shinfo(skb);
1009
1010     /* Make sure we own this skb before messing gso_size/gso_segs */
1011     WARN_ON_ONCE(skb_cloned(skb));
1012
1013     if (skb->len <= mss_now || skb->ip_summed == CHECKSUM_NONE) {
1014         /* Avoid the costly divide in the normal
1015          * non-TSO case.
1016          */
1017         shinfo->gso_segs = 1;
1018         shinfo->gso_size = 0;
1019         shinfo->gso_type = 0;
1020     } else {
1021         shinfo->gso_segs = DIV_ROUND_UP(skb->len, mss_now);
1022         shinfo->gso_size = mss_now;
1023         shinfo->gso_type = sk->sk_gso_type;
1024     }
1025 }
1026
1027 /* When a modification to fackets_out becomes necessary, we need to check
1028  * skb is counted to fackets_out or not.
1029  */
1030 static void tcp_adjust_fackets_out(struct sock *sk, const struct sk_buff *skb,
1031     int decr)
1032 {
1033     struct tcp_sock *tp = tcp_sk(sk);
1034
1035     if (!tp->sacked_out || tcp_is_reno(tp))
1036         return;
1037
1038     if (after(tcp_highest_sack_seq(tp), TCP_SKB_CB(skb)->seq))
1039         tp->fackets_out -= decr;
1040 }
1041
1042 /* Pcount in the middle of the write queue got changed, we need to do various
1043  * tweaks to fix counters
1044  */
1045 static void tcp_adjust_pcount(struct sock *sk, const struct sk_buff *skb, int decr)
1046 {
1047     struct tcp_sock *tp = tcp_sk(sk);
1048
1049     tp->packets_out -= decr;
1050
1051     if (TCP_SKB_CB(skb)->sacked & TCPCB_SACKED_ACKED)
1052         tp->sacked_out -= decr;
1053     if (TCP_SKB_CB(skb)->sacked & TCPCB_SACKED_RETRANS)
1054         tp->retrans_out -= decr;
1055     if (TCP_SKB_CB(skb)->sacked & TCPCB_LOST)
1056         tp->lost_out -= decr;
1057
1058     /* Reno case is special. Sigh... */
1059     if (tcp_is_reno(tp) && decr > 0)
1060         tp->sacked_out -= min_t(u32, tp->sacked_out, decr);

```

```

1061
1062     tcp_adjust_fackets_out(sk, skb, decr);
1063
1064     if (tp->lost_skb_hint &&
1065         before(TCP_SKB_CB(skb)->seq, TCP_SKB_CB(tp->lost_skb_hint)->seq) &&
1066         (tcp_is_fack(tp) || (TCP_SKB_CB(skb)->sacked & TCPCB_SACKED_ACKED)))
1067         tp->lost_cnt_hint -= decr;
1068
1069     tcp_verify_left_out(tp);
1070 }
1071
1072 static void tcp_fragment_tstamp(struct sk_buff *skb, struct sk_buff *skb2)
1073 {
1074     struct skb_shared_info *shinfo = skb_shinfo(skb);
1075
1076     if (unlikely(shinfo->tx_flags & SKBTX_ANY_TSTAMP) &&
1077         !before(shinfo->tskey, TCP_SKB_CB(skb2)->seq)) {
1078         struct skb_shared_info *shinfo2 = skb_shinfo(skb2);
1079         u8 tsflags = shinfo->tx_flags & SKBTX_ANY_TSTAMP;
1080
1081         shinfo->tx_flags &= ~tsflags;
1082         shinfo2->tx_flags |= tsflags;
1083         swap(shinfo->tskey, shinfo2->tskey);
1084     }
1085 }
1086
1087 /* Function to create two new TCP segments. Shrinks the given segment
1088  * to the specified size and appends a new segment with the rest of the
1089  * packet to the list. This won't be called frequently, I hope.
1090  * Remember, these are still headerless SKBs at this point.
1091  */
1092 int tcp_fragment(struct sock *sk, struct sk_buff *skb, u32 len,
1093                 unsigned int mss_now, gfp_t gfp)
1094 {
1095     struct tcp_sock *tp = tcp_sk(sk);
1096     struct sk_buff *buff;
1097     int nsize, old_factor;
1098     int nlen;
1099     u8 flags;
1100
1101     if (WARN_ON(len > skb->len))
1102         return -EINVAL;
1103
1104     nsize = skb_headlen(skb) - len;
1105     if (nsize < 0)
1106         nsize = 0;
1107
1108     if (skb_unclone(skb, gfp))
1109         return -ENOMEM;
1110
1111     /* Get a new skb... force flag on. */
1112     buff = sk_stream_alloc_skb(sk, nsize, gfp);
1113     if (buff == NULL)
1114         return -ENOMEM; /* We'll just try again later. */
1115
1116     sk->sk_wmem_queued += buff->truesize;
1117     sk_mem_charge(sk, buff->truesize);
1118     nlen = skb->len - len - nsize;
1119     buff->truesize += nlen;
1120     skb->truesize -= nlen;
1121
1122     /* Correct the sequence numbers. */
1123     TCP_SKB_CB(buff)->seq = TCP_SKB_CB(skb)->seq + len;
1124     TCP_SKB_CB(buff)->end_seq = TCP_SKB_CB(skb)->end_seq;
1125     TCP_SKB_CB(skb)->end_seq = TCP_SKB_CB(buff)->seq;
1126
1127     /* PSH and FIN should only be set in the second packet. */
1128     flags = TCP_SKB_CB(skb)->tcp_flags;
1129     TCP_SKB_CB(skb)->tcp_flags = flags & ~(TCPHDR_FIN | TCPHDR_PSH);
1130     TCP_SKB_CB(buff)->tcp_flags = flags;
1131     TCP_SKB_CB(buff)->sacked = TCP_SKB_CB(skb)->sacked;
1132
1133     if (!skb_shinfo(skb)->nr_frags && skb->ip_summed != CHECKSUM_PARTIAL) {
1134         /* Copy and checksum data tail into the new buffer. */
1135         buff->csum = csum_partial_copy_nocheck(skb->data + len,
1136                                                skb_put(buff, nsize),
1137                                                nsize, 0);
1138
1139         skb_trim(skb, len);
1140
1141         skb->csum = csum_block_sub(skb->csum, buff->csum, len);
1142     } else {
1143         skb->ip_summed = CHECKSUM_PARTIAL;
1144         skb_split(skb, buff, len);
1145     }
1146
1147     buff->ip_summed = skb->ip_summed;
1148
1149     /* Looks stupid, but our code really uses when of
1150      * skbs, which it never sent before. --ANK

```

```

1151     */
1152     TCP_SKB_CB(buff)->when = TCP_SKB_CB(skb)->when;
1153     buff->tstamp = skb->tstamp;
1154     tcp_fragment_tstamp(skb, buff);
1155
1156     old_factor = tcp_skb_pcount(skb);
1157
1158     /* Fix up tso_factor for both original and new SKB. */
1159     tcp_set_skb_tso_segs(sk, skb, mss_now);
1160     tcp_set_skb_tso_segs(sk, buff, mss_now);
1161
1162     /* If this packet has been sent out already, we must
1163      * adjust the various packet counters.
1164      */
1165     if (!before(tp->snd_nxt, TCP_SKB_CB(buff)->end_seq)) {
1166         int diff = old_factor - tcp_skb_pcount(skb) -
1167             tcp_skb_pcount(buff);
1168
1169         if (diff)
1170             tcp_adjust_pcount(sk, skb, diff);
1171     }
1172
1173     /* Link BUFF into the send queue. */
1174     skb_header_release(buff);
1175     tcp_insert_write_queue_after(skb, buff, sk);
1176
1177     return 0;
1178 }
1179
1180 /* This is similar to __pskb_pull_head() (it will go to core/skbuff.c
1181  * eventually). The difference is that pulled data not copied, but
1182  * immediately discarded.
1183  */
1184 static void __pskb_trim_head(struct sk_buff *skb, int len)
1185 {
1186     struct skb_shared_info *shinfo;
1187     int i, k, eat;
1188
1189     eat = min_t(int, len, skb_headlen(skb));
1190     if (eat) {
1191         __skb_pull(skb, eat);
1192         len -= eat;
1193         if (!len)
1194             return;
1195     }
1196     eat = len;
1197     k = 0;
1198     shinfo = skb_shinfo(skb);
1199     for (i = 0; i < shinfo->nr_frags; i++) {
1200         int size = skb_frag_size(&shinfo->frags[i]);
1201
1202         if (size <= eat) {
1203             skb_frag_unref(skb, i);
1204             eat -= size;
1205         } else {
1206             shinfo->frags[k] = shinfo->frags[i];
1207             if (eat) {
1208                 shinfo->frags[k].page_offset += eat;
1209                 skb_frag_size_sub(&shinfo->frags[k], eat);
1210                 eat = 0;
1211             }
1212             k++;
1213         }
1214     }
1215     shinfo->nr_frags = k;
1216
1217     skb_reset_tail_pointer(skb);
1218     skb->data_len -= len;
1219     skb->len = skb->data_len;
1220 }
1221
1222 /* Remove acked data from a packet in the transmit queue. */
1223 int tcp_trim_head(struct sock *sk, struct sk_buff *skb, u32 len)
1224 {
1225     if (skb_unclone(skb, GFP_ATOMIC))
1226         return -ENOMEM;
1227
1228     __pskb_trim_head(skb, len);
1229
1230     TCP_SKB_CB(skb)->seq += len;
1231     skb->ip_summed = CHECKSUM_PARTIAL;
1232
1233     skb->truesize -= len;
1234     sk->sk_wmem_queued -= len;
1235     sk_mem_uncharge(sk, len);
1236     sock_set_flag(sk, SOCK_QUEUE_SHRUNK);
1237
1238     /* Any change of skb->len requires recalculation of tso factor. */
1239     if (tcp_skb_pcount(skb) > 1)
1240         tcp_set_skb_tso_segs(sk, skb, tcp_skb_mss(skb));

```

```

1241
1242     return 0;
1243 }
1244
1245 /* Calculate MSS not accounting any TCP options. */
1246 static inline int tcp_mtu_to_mss(struct sock *sk, int pmtu)
1247 {
1248     const struct tcp_sock *tp = tcp_sk(sk);
1249     const struct inet_connection_sock *icsk = inet_csk(sk);
1250     int mss_now;
1251
1252     /* Calculate base mss without TCP options:
1253      * It is MSS_S - sizeof(tcphdr) of rfc1122
1254      */
1255     mss_now = pmtu - icsk->icsk_af_ops->net_header_len - sizeof(struct tcphdr);
1256
1257     /* IPv6 adds a frag_hdr in case RTAX_FEATURE_ALLFRAG is set */
1258     if (icsk->icsk_af_ops->net_frag_header_len) {
1259         const struct dst_entry *dst = __sk_dst_get(sk);
1260
1261         if (dst && dst_allfrag(dst))
1262             mss_now -= icsk->icsk_af_ops->net_frag_header_len;
1263     }
1264
1265     /* Clamp it (mss_clamp does not include tcp options) */
1266     if (mss_now > tp->rx_opt.mss_clamp)
1267         mss_now = tp->rx_opt.mss_clamp;
1268
1269     /* Now subtract optional transport overhead */
1270     mss_now -= icsk->icsk_ext_hdr_len;
1271
1272     /* Then reserve room for full set of TCP options and 8 bytes of data */
1273     if (mss_now < 48)
1274         mss_now = 48;
1275     return mss_now;
1276 }
1277
1278 /* Calculate MSS. Not accounting for SACKs here. */
1279 int tcp_mtu_to_mss(struct sock *sk, int pmtu)
1280 {
1281     /* Subtract TCP options size, not including SACKs */
1282     return tcp_mtu_to_mss(sk, pmtu) -
1283         (tcp_sk(sk)->tcp_header_len - sizeof(struct tcphdr));
1284 }
1285
1286 /* Inverse of above */
1287 int tcp_mss_to_mtu(struct sock *sk, int mss)
1288 {
1289     const struct tcp_sock *tp = tcp_sk(sk);
1290     const struct inet_connection_sock *icsk = inet_csk(sk);
1291     int mtu;
1292
1293     mtu = mss +
1294         tp->tcp_header_len +
1295         icsk->icsk_ext_hdr_len +
1296         icsk->icsk_af_ops->net_header_len;
1297
1298     /* IPv6 adds a frag_hdr in case RTAX_FEATURE_ALLFRAG is set */
1299     if (icsk->icsk_af_ops->net_frag_header_len) {
1300         const struct dst_entry *dst = __sk_dst_get(sk);
1301
1302         if (dst && dst_allfrag(dst))
1303             mtu += icsk->icsk_af_ops->net_frag_header_len;
1304     }
1305     return mtu;
1306 }
1307
1308 /* MTU probing init per socket */
1309 void tcp_mtup_init(struct sock *sk)
1310 {
1311     struct tcp_sock *tp = tcp_sk(sk);
1312     struct inet_connection_sock *icsk = inet_csk(sk);
1313
1314     icsk->icsk_mtup.enabled = sysctl_tcp_mtu_probing > 1;
1315     icsk->icsk_mtup.search_high = tp->rx_opt.mss_clamp + sizeof(struct tcphdr) +
1316         icsk->icsk_af_ops->net_header_len;
1317     icsk->icsk_mtup.search_low = tcp_mss_to_mtu(sk, sysctl_tcp_base_mss);
1318     icsk->icsk_mtup.probe_size = 0;
1319 }
1320 EXPORT_SYMBOL(tcp_mtup_init);
1321
1322 /* This function synchronize snd mss to current pmtu/exthdr set.
1323
1324 tp->rx_opt.user_mss is mss set by user by TCP_MAXSEG. It does NOT counts
1325 for TCP options, but includes only bare TCP header.
1326
1327 tp->rx_opt.mss_clamp is mss negotiated at connection setup.
1328 It is minimum of user_mss and mss received with SYN.
1329 It also does not include TCP options.
1330

```



```

1331 inet_csk(sk)->icsk_pmtu_cookie is last pmtu, seen by this function.
1332
1333 tp->mss_cache is current effective sending mss, including
1334 all tcp options except for SACKs. It is evaluated,
1335 taking into account current pmtu, but never exceeds
1336 tp->rx_opt.mss_clamp.
1337
1338 NOTE1. rfc1122 clearly states that advertised MSS
1339 DOES NOT include either tcp or ip options.
1340
1341 NOTE2. inet_csk(sk)->icsk_pmtu_cookie and tp->mss_cache
1342 are READ ONLY outside this function. --ANK (980731)
1343 */
1344 unsigned int tcp_sync_mss(struct sock *sk, u32 pmtu)
1345 {
1346     struct tcp_sock *tp = tcp_sk(sk);
1347     struct inet_connection_sock *icsk = inet_csk(sk);
1348     int mss_now;
1349
1350     if (icsk->icsk_mtup.search_high > pmtu)
1351         icsk->icsk_mtup.search_high = pmtu;
1352
1353     mss_now = tcp_mtu_to_mss(sk, pmtu);
1354     mss_now = tcp_bound_to_half_wnd(tp, mss_now);
1355
1356     /* And store cached results */
1357     icsk->icsk_pmtu_cookie = pmtu;
1358     if (icsk->icsk_mtup.enabled)
1359         mss_now = min(mss_now, tcp_mtu_to_mss(sk, icsk->icsk_mtup.search_low));
1360     tp->mss_cache = mss_now;
1361
1362     return mss_now;
1363 }
1364 EXPORT_SYMBOL(tcp_sync_mss);
1365
1366 /* Compute the current effective MSS, taking SACKs and IP options,
1367 * and even PMTU discovery events into account.
1368 */
1369 unsigned int tcp_current_mss(struct sock *sk)
1370 {
1371     const struct tcp_sock *tp = tcp_sk(sk);
1372     const struct dst_entry *dst = __sk_dst_get(sk);
1373     u32 mss_now;
1374     unsigned int header_len;
1375     struct tcp_out_options opts;
1376     struct tcp_md5sig_key *md5;
1377
1378     mss_now = tp->mss_cache;
1379
1380     if (dst) {
1381         u32 mtu = dst_mtu(dst);
1382         if (mtu != inet_csk(sk)->icsk_pmtu_cookie)
1383             mss_now = tcp_sync_mss(sk, mtu);
1384     }
1385
1386     header_len = tcp_established_options(sk, NULL, &opts, &md5) +
1387         sizeof(struct tcphdr);
1388     /* The mss_cache is sized based on tp->tcp_header_len, which assumes
1389     * some common options. If this is an odd packet (because we have SACK
1390     * blocks etc) then our calculated header_len will be different, and
1391     * we have to adjust mss_now correspondingly */
1392     if (header_len != tp->tcp_header_len) {
1393         int delta = (int) header_len - tp->tcp_header_len;
1394         mss_now -= delta;
1395     }
1396
1397     return mss_now;
1398 }
1399
1400 /* RFC2861, slow part. Adjust cwnd, after it was not full during one rto.
1401 * As additional protections, we do not touch cwnd in retransmission phases,
1402 * and if application hit its sndbuf limit recently.
1403 */
1404 static void tcp_cwnd_application_limited(struct sock *sk)
1405 {
1406     struct tcp_sock *tp = tcp_sk(sk);
1407
1408     if (inet_csk(sk)->icsk_ca_state == TCP_CA_Open &&
1409         sk->sk_socket && !test_bit(SOCK_NOSPACE, &sk->sk_socket->flags)) {
1410         /* Limited by application or receiver window. */
1411         u32 init_win = tcp_init_cwnd(tp, __sk_dst_get(sk));
1412         u32 win_used = max(tp->snd_cwnd_used, init_win);
1413         if (win_used < tp->snd_cwnd) {
1414             tp->snd_ssthresh = tcp_current_ssthresh(sk);
1415             tp->snd_cwnd = (tp->snd_cwnd + win_used) >> 1;
1416         }
1417         tp->snd_cwnd_used = 0;
1418     }
1419     tp->snd_cwnd_stamp = tcp_time_stamp;
1420 }

```

```

1421
1422 static void tcp_cwnd_validate(struct sock *sk, bool is_cwnd_limited)
1423 {
1424     struct tcp_sock *tp = tcp_sk(sk);
1425
1426     /* Track the maximum number of outstanding packets in each
1427      * window, and remember whether we were cwnd-limited then.
1428      */
1429     if (!before(tp->snd_una, tp->max_packets_seq) ||
1430         tp->packets_out > tp->max_packets_out) {
1431         tp->max_packets_out = tp->packets_out;
1432         tp->max_packets_seq = tp->snd_nxt;
1433         tp->is_cwnd_limited = is_cwnd_limited;
1434     }
1435
1436     if (tcp_is_cwnd_limited(sk)) {
1437         /* Network is feed fully. */
1438         tp->snd_cwnd_used = 0;
1439         tp->snd_cwnd_stamp = tcp_time_stamp;
1440     } else {
1441         /* Network starves. */
1442         if (tp->packets_out > tp->snd_cwnd_used)
1443             tp->snd_cwnd_used = tp->packets_out;
1444
1445         if (sysctl_tcp_slow_start_after_idle &&
1446             (s32)(tcp_time_stamp - tp->snd_cwnd_stamp) >= inet_csk(sk)->icsk_rto)
1447             tcp_cwnd_application_limited(sk);
1448     }
1449 }
1450
1451 /* Minshall's variant of the Nagle send check. */
1452 static bool tcp_minshall_check(const struct tcp_sock *tp)
1453 {
1454     return after(tp->snd_sml, tp->snd_una) &&
1455         !after(tp->snd_sml, tp->snd_nxt);
1456 }
1457
1458 /* Update snd_sml if this skb is under mss
1459  * Note that a TSO packet might end with a sub-mss segment
1460  * The test is really :
1461  * if ((skb->len % mss) != 0)
1462  *     tp->snd_sml = TCP_SKB_CB(skb)->end_seq;
1463  * But we can avoid doing the divide again given we already have
1464  * skb_pcount = skb->len / mss_now
1465  */
1466 static void tcp_minshall_update(struct tcp_sock *tp, unsigned int mss_now,
1467                                const struct sk_buff *skb)
1468 {
1469     if (skb->len < tcp_skb_pcount(skb) * mss_now)
1470         tp->snd_sml = TCP_SKB_CB(skb)->end_seq;
1471 }
1472
1473 /* Return false, if packet can be sent now without violation Nagle's rules:
1474  * 1. It is full sized. (provided by caller in %partial bool)
1475  * 2. Or it contains FIN. (already checked by caller)
1476  * 3. Or TCP_CORK is not set, and TCP_NODELAY is set.
1477  * 4. Or TCP_CORK is not set, and all sent packets are ACKed.
1478  * With Minshall's modification: all sent small packets are ACKed.
1479  */
1480 static bool tcp_nagle_check(bool partial, const struct tcp_sock *tp,
1481                             int nonagle)
1482 {
1483     return partial &&
1484         ((nonagle & TCP_NAGLE_CORK) ||
1485          (!nonagle && tp->packets_out && tcp_minshall_check(tp)));
1486 }
1487
1488 /* Returns the portion of skb which can be sent right away */
1489 static unsigned int tcp_mss_split_point(const struct sock *sk,
1490                                         const struct sk_buff *skb,
1491                                         unsigned int mss_now,
1492                                         unsigned int max_segs,
1493                                         int nonagle)
1494 {
1495     const struct tcp_sock *tp = tcp_sk(sk);
1496     u32 partial, needed, window, max_len;
1497
1498     window = tcp_wnd_end(tp) - TCP_SKB_CB(skb)->seq;
1499     max_len = mss_now * max_segs;
1500
1501     if (likely(max_len <= window && skb != tcp_write_queue_tail(sk)))
1502         return max_len;
1503
1504     needed = min(skb->len, window);
1505
1506     if (max_len <= needed)
1507         return max_len;
1508
1509     partial = needed % mss_now;
1510     /* If last segment is not a full MSS, check if Nagle rules allow us
1511      * to include this last segment in this skb.
1512     */

```

```

1511     * Otherwise, we'll split the skb at last MSS boundary
1512     */
1513     if (tcp_nagle_check(partial != 0, tp, nonagle))
1514         return needed - partial;
1515
1516     return needed;
1517 }
1518
1519 /* Can at least one segment of SKB be sent right now, according to the
1520  * congestion window rules? If so, return how many segments are allowed.
1521  */
1522 static inline unsigned int tcp_cwnd_test(const struct tcp_sock *tp,
1523                                         const struct sk_buff *skb)
1524 {
1525     u32 in_flight, cwnd;
1526
1527     /* Don't be strict about the congestion window for the final FIN. */
1528     if ((TCP_SKB_CB(skb)->tcp_flags & TCPHDR_FIN) &&
1529         tcp_skb_pcount(skb) == 1)
1530         return 1;
1531
1532     in_flight = tcp_packets_in_flight(tp);
1533     cwnd = tp->snd_cwnd;
1534     if (in_flight < cwnd)
1535         return (cwnd - in_flight);
1536
1537     return 0;
1538 }
1539
1540 /* Initialize TSO state of a skb.
1541  * This must be invoked the first time we consider transmitting
1542  * SKB onto the wire.
1543  */
1544 static int tcp_init_tso_segs(const struct sock *sk, struct sk_buff *skb,
1545                             unsigned int mss_now)
1546 {
1547     int tso_segs = tcp_skb_pcount(skb);
1548
1549     if (!tso_segs || (tso_segs > 1 && tcp_skb_mss(skb) != mss_now)) {
1550         tcp_set_skb_tso_segs(sk, skb, mss_now);
1551         tso_segs = tcp_skb_pcount(skb);
1552     }
1553     return tso_segs;
1554 }
1555
1556 /* Return true if the Nagle test allows this packet to be
1557  * sent now.
1558  */
1559 static inline bool tcp_nagle_test(const struct tcp_sock *tp, const struct sk_buff *skb,
1560                                  unsigned int cur_mss, int nonagle)
1561 {
1562     /* Nagle rule does not apply to frames, which sit in the middle of the
1563      * write_queue (they have no chances to get new data).
1564      *
1565      * This is implemented in the callers, where they modify the 'nonagle'
1566      * argument based upon the location of SKB in the send queue.
1567      */
1568     if (nonagle & TCP_NAGLE_PUSH)
1569         return true;
1570
1571     /* Don't use the nagle rule for urgent data (or for the final FIN). */
1572     if (tcp_urg_mode(tp) || (TCP_SKB_CB(skb)->tcp_flags & TCPHDR_FIN))
1573         return true;
1574
1575     if (!tcp_nagle_check(skb->len < cur_mss, tp, nonagle))
1576         return true;
1577
1578     return false;
1579 }
1580
1581 /* Does at least the first segment of SKB fit into the send window? */
1582 static bool tcp_snd_wnd_test(const struct tcp_sock *tp,
1583                             const struct sk_buff *skb,
1584                             unsigned int cur_mss)
1585 {
1586     u32 end_seq = TCP_SKB_CB(skb)->end_seq;
1587
1588     if (skb->len > cur_mss)
1589         end_seq = TCP_SKB_CB(skb)->seq + cur_mss;
1590
1591     return !after(end_seq, tcp_wnd_end(tp));
1592 }
1593
1594 /* This checks if the data bearing packet SKB (usually tcp_send_head(sk))
1595  * should be put on the wire right now. If so, it returns the number of
1596  * packets allowed by the congestion window.
1597  */
1598 static unsigned int tcp_snd_test(const struct sock *sk, struct sk_buff *skb,
1599                                 unsigned int cur_mss, int nonagle)
1600

```

```

1601 {
1602     const struct tcp\_sock *tp = tcp\_sk(sk);
1603     unsigned int cwnd_quota;
1604
1605     tcp\_init\_tso\_segs(sk, skb, cur_mss);
1606
1607     if (!tcp\_nagle\_test(tp, skb, cur_mss, nonagle))
1608         return 0;
1609
1610     cwnd_quota = tcp\_cwnd\_test(tp, skb);
1611     if (cwnd_quota && !tcp\_snd\_wnd\_test(tp, skb, cur_mss))
1612         cwnd_quota = 0;
1613
1614     return cwnd_quota;
1615 }
1616
1617 /* Test if sending is allowed right now. */
1618 bool tcp\_may\_send\_now(struct sock *sk)
1619 {
1620     const struct tcp\_sock *tp = tcp\_sk(sk);
1621     struct sk\_buff *skb = tcp\_send\_head(sk);
1622
1623     return skb &&
1624         tcp\_snd\_test(sk, skb, tcp\_current\_mss(sk),
1625                     (tcp\_skb\_is\_last(sk, skb) ?
1626                      tp->nonagle : TCP\_NAGLE\_PUSH));
1627 }
1628
1629 /* Trim TSO SKB to LEN bytes, put the remaining data into a new packet
1630  * which is put after SKB on the list. It is very much like
1631  * tcp_fragment() except that it may make several kinds of assumptions
1632  * in order to speed up the splitting operation. In particular, we
1633  * know that all the data is in scatter-gather pages, and that the
1634  * packet has never been sent out before (and thus is not cloned).
1635  */
1636 static int tso\_fragment(struct sock *sk, struct sk\_buff *skb, unsigned int len,
1637                        unsigned int mss_now, gfp\_t gfp)
1638 {
1639     struct sk\_buff *buff;
1640     int nlen = skb->len - len;
1641     u8 flags;
1642
1643     /* ALL of a TSO frame must be composed of paged data. */
1644     if (skb->len != skb->data_len)
1645         return tcp\_fragment(sk, skb, len, mss_now, gfp);
1646
1647     buff = sk\_stream\_alloc\_skb(sk, 0, gfp);
1648     if (unlikely(buff == NULL))
1649         return -ENOMEM;
1650
1651     sk->sk_wmem_queued += buff->truesize;
1652     sk\_mem\_charge(sk, buff->truesize);
1653     buff->truesize += nlen;
1654     skb->truesize -= nlen;
1655
1656     /* Correct the sequence numbers. */
1657     TCP\_SKB\_CB(buff)->seq = TCP\_SKB\_CB(skb)->seq + len;
1658     TCP\_SKB\_CB(buff)->end_seq = TCP\_SKB\_CB(skb)->end_seq;
1659     TCP\_SKB\_CB(skb)->end_seq = TCP\_SKB\_CB(buff)->seq;
1660
1661     /* PSH and FIN should only be set in the second packet. */
1662     flags = TCP\_SKB\_CB(skb)->tcp_flags;
1663     TCP\_SKB\_CB(skb)->tcp_flags = flags & ~(TCPHDR\_FIN | TCPHDR\_PSH);
1664     TCP\_SKB\_CB(buff)->tcp_flags = flags;
1665
1666     /* This packet was never sent out yet, so no SACK bits. */
1667     TCP\_SKB\_CB(buff)->sacked = 0;
1668
1669     buff->ip_summed = skb->ip_summed = CHECKSUM\_PARTIAL;
1670     skb\_split(skb, buff, len);
1671     tcp\_fragment\_tstamp(skb, buff);
1672
1673     /* Fix up tso_factor for both original and new SKB. */
1674     tcp\_set\_skb\_tso\_segs(sk, skb, mss_now);
1675     tcp\_set\_skb\_tso\_segs(sk, buff, mss_now);
1676
1677     /* Link BUFF into the send queue. */
1678     skb\_header\_release(buff);
1679     tcp\_insert\_write\_queue\_after(skb, buff, sk);
1680
1681     return 0;
1682 }
1683
1684 /* Try to defer sending, if possible, in order to minimize the amount
1685  * of TSO splitting we do. View it as a kind of TSO Nagle test.
1686  */
1687 /* This algorithm is from John Heffner.
1688  */
1689 static bool tcp\_tso\_should\_defer(struct sock *sk, struct sk\_buff *skb,
1690                                  bool *is_cwnd_limited)

```

```

1691 {
1692     struct tcp\_sock *tp = tcp\_sk(sk);
1693     const struct inet\_connection\_sock *icsk = inet\_csk(sk);
1694     u32 send_win, cong_win, limit, in\_flight;
1695     int win_divisor;
1696
1697     if (TCP\_SKB\_CB(skb)->tcp_flags & TCPHDR\_FIN)
1698         goto send_now;
1699
1700     if (icsk->icsk_ca_state != TCP_CA_Open)
1701         goto send_now;
1702
1703     /* Defer for less than two clock ticks. */
1704     if (tp->tso_deferred &&
1705         (((u32)jiffies << 1) >> 1) - (tp->tso_deferred >> 1) > 1)
1706         goto send_now;
1707
1708     in\_flight = tcp\_packets\_in\_flight(tp);
1709
1710     BUG\_ON(tcp\_skb\_pcount(skb) <= 1 || (tp->snd_cwnd <= in\_flight));
1711
1712     send_win = tcp\_wnd\_end(tp) - TCP\_SKB\_CB(skb)->seq;
1713
1714     /* From in_flight test above, we know that cwnd > in_flight. */
1715     cong_win = (tp->snd_cwnd - in\_flight) * tp->mss_cache;
1716
1717     limit = min(send_win, cong_win);
1718
1719     /* If a full-sized TSO skb can be sent, do it. */
1720     if (limit >= min\_t(unsigned int, sk->sk_gso_max_size,
1721         tp->xmit_size_goal_segs * tp->mss_cache))
1722         goto send_now;
1723
1724     /* Middle in queue won't get any more data, full sendable already? */
1725     if ((skb != tcp\_write\_queue\_tail(sk)) && (limit >= skb->len))
1726         goto send_now;
1727
1728     win_divisor = ACCESS\_ONCE(sysctl\_tcp\_tso\_win\_divisor);
1729     if (win_divisor) {
1730         u32 chunk = min(tp->snd_wnd, tp->snd_cwnd * tp->mss_cache);
1731
1732         /* If at least some fraction of a window is available,
1733          * just use it.
1734          */
1735         chunk /= win_divisor;
1736         if (limit >= chunk)
1737             goto send_now;
1738     } else {
1739         /* Different approach, try not to defer past a single
1740          * ACK. Receiver should ACK every other full sized
1741          * frame, so if we have space for more than 3 frames
1742          * then send now.
1743          */
1744         if (limit > tcp\_max\_tso\_deferred\_mss(tp) * tp->mss_cache)
1745             goto send_now;
1746     }
1747
1748     /* Ok, it looks like it is advisable to defer.
1749     * Do not rearm the timer if already set to not break TCP ACK clocking.
1750     */
1751     if (!tp->tso_deferred)
1752         tp->tso_deferred = 1 | (jiffies << 1);
1753
1754     if (cong_win < send_win && cong_win < skb->len)
1755         *is_cwnd_limited = true;
1756
1757     return true;
1758
1759 send_now:
1760     tp->tso_deferred = 0;
1761     return false;
1762 }
1763
1764 /* Create a new MTU probe if we are ready.
1765  * MTU probe is regularly attempting to increase the path MTU by
1766  * deliberately sending larger packets. This discovers routing
1767  * changes resulting in larger path MTUs.
1768  *
1769  * Returns 0 if we should wait to probe (no cwnd available),
1770  * 1 if a probe was sent,
1771  * -1 otherwise
1772  */
1773 static int tcp\_mtu\_probe(struct sock *sk)
1774 {
1775     struct tcp\_sock *tp = tcp\_sk(sk);
1776     struct inet\_connection\_sock *icsk = inet\_csk(sk);
1777     struct sk\_buff *skb, *nskb, *next;
1778     int len;
1779     int probe_size;
1780     int size_needed;

```

```

1781 int copy;
1782 int mss_now;
1783
1784 /* Not currently probing/verifying,
1785  * not in recovery,
1786  * have enough cwnd, and
1787  * not SACKing (the variable headers throw things off) */
1788 if (!icsk->icsk_mtup.enabled ||
1789     icsk->icsk_mtup.probe_size ||
1790     inet_csk(sk)->icsk_ca_state != TCP_CA_Open ||
1791     tp->snd_cwnd < 11 ||
1792     tp->rx_opt.num_sacks || tp->rx_opt.dsack)
1793     return -1;
1794
1795 /* Very simple search strategy: just double the MSS. */
1796 mss_now = tcp_current_mss(sk);
1797 probe_size = 2 * tp->mss_cache;
1798 size_needed = probe_size + (tp->reordering + 1) * tp->mss_cache;
1799 if (probe_size > tcp_mtu_to_mss(sk, icsk->icsk_mtup.search_high)) {
1800     /* TODO: set timer for probe_converge_event */
1801     return -1;
1802 }
1803
1804 /* Have enough data in the send queue to probe? */
1805 if (tp->write_seq - tp->snd_nxt < size_needed)
1806     return -1;
1807
1808 if (tp->snd_wnd < size_needed)
1809     return -1;
1810 if (after(tp->snd_nxt + size_needed, tcp_wnd_end(tp)))
1811     return 0;
1812
1813 /* Do we need to wait to drain cwnd? With none in flight, don't stall */
1814 if (tcp_packets_in_flight(tp) + 2 > tp->snd_cwnd) {
1815     if (!tcp_packets_in_flight(tp))
1816         return -1;
1817     else
1818         return 0;
1819 }
1820
1821 /* We're allowed to probe. Build it now. */
1822 if ((nskb = sk_stream_alloc_skb(sk, probe_size, GFP_ATOMIC)) == NULL)
1823     return -1;
1824 sk->sk_wmem_queued += nskb->truesize;
1825 sk_mem_charge(sk, nskb->truesize);
1826
1827 skb = tcp_send_head(sk);
1828
1829 TCP_SKB_CB(nskb)->seq = TCP_SKB_CB(skb)->seq;
1830 TCP_SKB_CB(nskb)->end_seq = TCP_SKB_CB(skb)->seq + probe_size;
1831 TCP_SKB_CB(nskb)->tcp_flags = TCPHDR_ACK;
1832 TCP_SKB_CB(nskb)->sacked = 0;
1833 nskb->csum = 0;
1834 nskb->ip_summed = skb->ip_summed;
1835
1836 tcp_insert_write_queue_before(nskb, skb, sk);
1837
1838 len = 0;
1839 tcp_for_write_queue_from_safe(skb, next, sk) {
1840     copy = min_t(int, skb->len, probe_size - len);
1841     if (nskb->ip_summed)
1842         skb_copy_bits(skb, 0, skb_put(nskb, copy), copy);
1843     else
1844         nskb->csum = skb_copy_and_csum_bits(skb, 0,
1845                                             skb_put(nskb, copy),
1846                                             copy, nskb->csum);
1847
1848     if (skb->len <= copy) {
1849         /* We've eaten all the data from this skb.
1850          * Throw it away. */
1851         TCP_SKB_CB(nskb)->tcp_flags |= TCP_SKB_CB(skb)->tcp_flags;
1852         tcp_unlink_write_queue(skb, sk);
1853         sk_wmem_free_skb(sk, skb);
1854     } else {
1855         TCP_SKB_CB(nskb)->tcp_flags |= TCP_SKB_CB(skb)->tcp_flags &
1856                                         ~(TCPHDR_FIN|TCPHDR_PSH);
1857         if (!skb_shinfo(skb)->nr_frags) {
1858             skb_pull(skb, copy);
1859             if (skb->ip_summed != CHECKSUM_PARTIAL)
1860                 skb->csum = csum_partial(skb->data,
1861                                           skb->len, 0);
1862         } else {
1863             pskb_trim_head(skb, copy);
1864             tcp_set_skb_tso_segs(sk, skb, mss_now);
1865         }
1866         TCP_SKB_CB(skb)->seq += copy;
1867     }
1868     len += copy;
1869 }
1870

```

```

1871         if (len >= probe_size)
1872             break;
1873     }
1874     tcp_init_tso_segs(sk, nskb, nskb->len);
1875
1876     /* We're ready to send. If this fails, the probe will
1877      * be resegmented into mss-sized pieces by tcp_write_xmit(). */
1878     TCP_SKB_CB(nskb)->when = tcp_time_stamp;
1879     if (!tcp_transmit_skb(sk, nskb, 1, GFP_ATOMIC)) {
1880         /* Decrement cwnd here because we are sending
1881          * effectively two packets. */
1882         tp->snd_cwnd--;
1883         tcp_event_new_data_sent(sk, nskb);
1884
1885         icsk->icsk_mtup.probe_size = tcp_mss_to_mtu(sk, nskb->len);
1886         tp->mtu_probe.probe_seq_start = TCP_SKB_CB(nskb)->seq;
1887         tp->mtu_probe.probe_seq_end = TCP_SKB_CB(nskb)->end_seq;
1888
1889         return 1;
1890     }
1891     return -1;
1892 }
1893
1894 /* This routine writes packets to the network. It advances the
1895  * send_head. This happens as incoming acks open up the remote
1896  * window for us.
1897  *
1898  * LARGESSEND note: !tcp_urg_mode is overkill, only frames between
1899  * snd_up-64k-mss .. snd_up cannot be large. However, taking into
1900  * account rare use of URG, this is not a big flaw.
1901  *
1902  * Send at most one packet when push_one > 0. Temporarily ignore
1903  * cwnd limit to force at most one packet out when push_one == 2.
1904  *
1905  * Returns true, if no segments are in flight and we have queued segments,
1906  * but cannot send anything now because of SWS or another problem.
1907  */
1908 static bool tcp_write_xmit(struct sock *sk, unsigned int mss_now, int nonagle,
1909                          int push_one, gfp_t gfp)
1910 {
1911     struct tcp_sock *tp = tcp_sk(sk);
1912     struct sk_buff *skb;
1913     unsigned int tso_segs, sent_pkts;
1914     int cwnd_quota;
1915     int result;
1916     bool is_cwnd_limited = false;
1917
1918     sent_pkts = 0;
1919
1920     if (!push_one) {
1921         /* Do MTU probing. */
1922         result = tcp_mtu_probe(sk);
1923         if (!result) {
1924             return false;
1925         } else if (result > 0) {
1926             sent_pkts = 1;
1927         }
1928     }
1929
1930     while ((skb = tcp_send_head(sk))) {
1931         unsigned int limit;
1932
1933         tso_segs = tcp_init_tso_segs(sk, skb, mss_now);
1934         BUG_ON(!tso_segs);
1935
1936         if (unlikely(tp->repair) && tp->repair_queue == TCP_SEND_QUEUE) {
1937             /* "when" is used as a start point for the retransmit timer */
1938             TCP_SKB_CB(skb)->when = tcp_time_stamp;
1939             goto repair; /* Skip network transmission */
1940         }
1941
1942         cwnd_quota = tcp_cwnd_test(tp, skb);
1943         if (!cwnd_quota) {
1944             is_cwnd_limited = true;
1945             if (push_one == 2)
1946                 /* Force out a loss probe pkt. */
1947                 cwnd_quota = 1;
1948             else
1949                 break;
1950         }
1951
1952         if (unlikely(!tcp_snd_wnd_test(tp, skb, mss_now)))
1953             break;
1954
1955         if (tso_segs == 1) {
1956             if (unlikely(!tcp_nagle_test(tp, skb, mss_now,
1957                                         (tcp_skb_is_last(sk, skb) ?
1958                                          nonagle : TCP_NAGLE_PUSH))))
1959                 break;
1960         }

```



```

1961     } else {
1962         if (!push_one &&
1963             tcp_tso_should_defer(sk, skb, &is_cwnd_limited))
1964             break;
1965     }
1966
1967     /* TCP Small Queues :
1968     * Control number of packets in qdisc/devices to two packets / or ~1 ms.
1969     * This allows for :
1970     * - better RTT estimation and ACK scheduling
1971     * - faster recovery
1972     * - high rates
1973     * Alas, some drivers / subsystems require a fair amount
1974     * of queued bytes to ensure line rate.
1975     * One example is wifi aggregation (802.11 AMPDU)
1976     */
1977     limit = max_t(unsigned int, sysctl_tcp_limit_output_bytes,
1978                   sk->sk_pacing_rate >> 10);
1979
1980     if (atomic_read(&sk->sk_wmem_alloc) > limit) {
1981         set_bit(TSQ_THROTTLED, &tp->tsq_flags);
1982         /* It is possible TX completion already happened
1983          * before we set TSQ_THROTTLED, so we must
1984          * test again the condition.
1985          */
1986         smp_mb_after_atomic();
1987         if (atomic_read(&sk->sk_wmem_alloc) > limit)
1988             break;
1989     }
1990
1991     limit = mss_now;
1992     if (tso_segs > 1 && !tcp_urg_mode(tp))
1993         limit = tcp_mss_split_point(sk, skb, mss_now,
1994                                     min_t(unsigned int,
1995                                             cwnd_quota,
1996                                             sk->sk_gso_max_segs),
1997                                     nonagle);
1998
1999     if (skb->len > limit &&
2000         unlikely(tso_fragment(sk, skb, limit, mss_now, gfp)))
2001         break;
2002
2003     TCP_SKB_CB(skb)->when = tcp_time_stamp;
2004
2005     if (unlikely(tcp_transmit_skb(sk, skb, 1, gfp)))
2006         break;
2007
2008     repair:
2009     /* Advance the send_head. This one is sent out.
2010     * This call will increment packets_out.
2011     */
2012     tcp_event_new_data_sent(sk, skb);
2013
2014     tcp_minshall_update(tp, mss_now, skb);
2015     sent_pkts += tcp_skb_pcount(skb);
2016
2017     if (push_one)
2018         break;
2019 }
2020
2021 if (likely(sent_pkts)) {
2022     if (tcp_in_cwnd_reduction(sk))
2023         tp->prrr_out += sent_pkts;
2024
2025     /* Send one loss probe per tail loss episode. */
2026     if (push_one != 2)
2027         tcp_schedule_loss_probe(sk);
2028     tcp_cwnd_validate(sk, is_cwnd_limited);
2029     return false;
2030 }
2031 return (push_one == 2) || (!tp->packets_out && tcp_send_head(sk));
2032 }
2033
2034 bool tcp_schedule_loss_probe(struct sock *sk)
2035 {
2036     struct inet_connection_sock *icsk = inet_csk(sk);
2037     struct tcp_sock *tp = tcp_sk(sk);
2038     u32 timeout, tlp_time_stamp, rto_time_stamp;
2039     u32 rtt = usecs_to_jiffies(tp->srtt_us >> 3);
2040
2041     if (WARN_ON(icsk->icsk_pending == ICSK_TIME_EARLY_RETRANS))
2042         return false;
2043     /* No consecutive loss probes. */
2044     if (WARN_ON(icsk->icsk_pending == ICSK_TIME_LOSS_PROBE)) {
2045         tcp_rearm_rto(sk);
2046         return false;
2047     }
2048     /* Don't do any loss probe on a Fast Open connection before 3WHS
2049     * finishes.
2050     */

```

```

2051 if (sk->sk_state == TCP_SYN_RECV)
2052     return false;
2053
2054 /* TLP is only scheduled when next timer event is RTO. */
2055 if (icsk->icsk_pending != ICSK_TIME_RETRANS)
2056     return false;
2057
2058 /* Schedule a Loss probe in 2*RTT for SACK capable connections
2059  * in Open state, that are either limited by cwnd or application.
2060  */
2061 if (sysctl_tcp_early_retrans < 3 || !tp->srtt_us || !tp->packets_out ||
2062     !tcp_is_sack(tp) || inet_csk(sk)->icsk_ca_state != TCP_CA_Open)
2063     return false;
2064
2065 if ((tp->snd_cwnd > tcp_packets_in_flight(tp)) &&
2066     tcp_send_head(sk))
2067     return false;
2068
2069 /* Probe timeout is at least 1.5*rtt + TCP_DELACK_MAX to account
2070  * for delayed ack when there's one outstanding packet.
2071  */
2072 timeout = rtt << 1;
2073 if (tp->packets_out == 1)
2074     timeout = max_t(u32, timeout,
2075                     (rtt + (rtt >> 1) + TCP_DELACK_MAX));
2076 timeout = max_t(u32, timeout, msecs_to_jiffies(10));
2077
2078 /* If RTO is shorter, just schedule TLP in its place. */
2079 tlp_time_stamp = tcp_time_stamp + timeout;
2080 rto_time_stamp = (u32)inet_csk(sk)->icsk_timeout;
2081 if ((s32)(tlp_time_stamp - rto_time_stamp) > 0) {
2082     s32 delta = rto_time_stamp - tcp_time_stamp;
2083     if (delta > 0)
2084         timeout = delta;
2085 }
2086
2087 inet_csk_reset_xmit_timer(sk, ICSK_TIME_LOSS_PROBE, timeout,
2088                           TCP_RTO_MAX);
2089 return true;
2090 }
2091
2092 /* Thanks to skb fast clones, we can detect if a prior transmit of
2093  * a packet is still in a qdisc or driver queue.
2094  * In this case, there is very little point doing a retransmit !
2095  * Note: This is called from BH context only.
2096  */
2097 static bool skb_still_in_host_queue(const struct sock *sk,
2098                                    const struct sk_buff *skb)
2099 {
2100     const struct sk_buff *fclone = skb + 1;
2101
2102     if (unlikely(skb->fclone == SKB_FCLONE_ORIG &&
2103                 fclone->fclone == SKB_FCLONE_CLONE)) {
2104         NET_INC_STATS_BH(sock_net(sk),
2105                          LINUX_MIB_TCPSPURIOUS_RTX_HOSTQUEUES);
2106         return true;
2107     }
2108     return false;
2109 }
2110
2111 /* When probe timeout (PTO) fires, send a new segment if one exists, else
2112  * retransmit the last segment.
2113  */
2114 void tcp_send_loss_probe(struct sock *sk)
2115 {
2116     struct tcp_sock *tp = tcp_sk(sk);
2117     struct sk_buff *skb;
2118     int pcount;
2119     int mss = tcp_current_mss(sk);
2120     int err = -1;
2121
2122     if (tcp_send_head(sk) != NULL) {
2123         err = tcp_write_xmit(sk, mss, TCP_NAGLE_OFF, 2, GFP_ATOMIC);
2124         goto rearm_timer;
2125     }
2126
2127     /* At most one outstanding TLP retransmission. */
2128     if (tp->tlp_high_seq)
2129         goto rearm_timer;
2130
2131     /* Retransmit last segment. */
2132     skb = tcp_write_queue_tail(sk);
2133     if (WARN_ON(!skb))
2134         goto rearm_timer;
2135
2136     if (skb_still_in_host_queue(sk, skb))
2137         goto rearm_timer;
2138
2139     pcount = tcp_skb_pcount(skb);
2140     if (WARN_ON(!pcount))

```

```

2141         goto rearm_timer;
2142
2143     if ((pcount > 1) && (skb->len > (pcount - 1) * mss)) {
2144         if (unlikely(tcp_fragment(sk, skb, (pcount - 1) * mss, mss,
2145                                GFP_ATOMIC)))
2146             goto rearm_timer;
2147         skb = tcp_write_queue_tail(sk);
2148     }
2149
2150     if (WARN_ON(!skb || !tcp_skb_pcount(skb)))
2151         goto rearm_timer;
2152
2153     err = __tcp_retransmit_skb(sk, skb);
2154
2155     /* Record snd_nxt for loss detection. */
2156     if (likely(!err))
2157         tp->tlp_high_seq = tp->snd_nxt;
2158
2159 rearm_timer:
2160     inet_csk_reset_xmit_timer(sk, ICSK_TIME_RETRANS,
2161                              inet_csk(sk)->icsk_rto,
2162                              TCP_RTO_MAX);
2163
2164     if (likely(!err))
2165         NET_INC_STATS_BH(sock_net(sk),
2166                          LINUX_MIB_TCPLOSSPROBES);
2167 }
2168
2169 /* Push out any pending frames which were held back due to
2170  * TCP_CORK or attempt at coalescing tiny packets.
2171  * The socket must be locked by the caller.
2172  */
2173 void __tcp_push_pending_frames(struct sock *sk, unsigned int cur_mss,
2174                               int nonagle)
2175 {
2176     /* If we are closed, the bytes will have to remain here.
2177      * In time closedown will finish, we empty the write queue and
2178      * all will be happy.
2179      */
2180     if (unlikely(sk->sk_state == TCP_CLOSE))
2181         return;
2182
2183     if (tcp_write_xmit(sk, cur_mss, nonagle, 0,
2184                       sk_gfp_atomic(sk, GFP_ATOMIC)))
2185         tcp_check_probe_timer(sk);
2186 }
2187
2188 /* Send_single_skb sitting at the send head. This function requires
2189  * true push pending frames to setup probe timer etc.
2190  */
2191 void tcp_push_one(struct sock *sk, unsigned int mss_now)
2192 {
2193     struct sk_buff *skb = tcp_send_head(sk);
2194
2195     BUG_ON(!skb || skb->len < mss_now);
2196
2197     tcp_write_xmit(sk, mss_now, TCP_NAGLE_PUSH, 1, sk->sk_allocation);
2198 }
2199
2200 /* This function returns the amount that we can raise the
2201  * usable window based on the following constraints
2202  *
2203  * 1. The window can never be shrunk once it is offered (RFC 793)
2204  * 2. We limit memory per socket
2205  *
2206  * RFC 1122:
2207  * "the suggested [SWS] avoidance algorithm for the receiver is to keep
2208  *  RECV.NEXT + RCV.WIN fixed until:
2209  *  RCV.BUFF - RCV.USER - RCV.WINDOW >= min(1/2 RCV.BUFF, MSS)"
2210  *
2211  * i.e. don't raise the right edge of the window until you can raise
2212  * it at least MSS bytes.
2213  *
2214  * Unfortunately, the recommended algorithm breaks header prediction,
2215  * since header prediction assumes th->window stays fixed.
2216  *
2217  * Strictly speaking, keeping th->window fixed violates the receiver
2218  * side SWS prevention criteria. The problem is that under this rule
2219  * a stream of single byte packets will cause the right side of the
2220  * window to always advance by a single byte.
2221  *
2222  * Of course, if the sender implements sender side SWS prevention
2223  * then this will not be a problem.
2224  *
2225  * BSD seems to make the following compromise:
2226  *
2227  * If the free space is less than the 1/4 of the maximum
2228  * space available and the free space is less than 1/2 mss,
2229  * then set the window to 0.
2230  * [ Actually, bsd uses MSS and 1/4 of maximal _window_ ]

```

```

2231 *      Otherwise, just prevent the window from shrinking
2232 *      and from being larger than the largest representable value.
2233 *
2234 * This prevents incremental opening of the window in the regime
2235 * where TCP is limited by the speed of the reader side taking
2236 * data out of the TCP receive queue. It does nothing about
2237 * those cases where the window is constrained on the sender side
2238 * because the pipeline is full.
2239 *
2240 * BSD also seems to "accidentally" limit itself to windows that are a
2241 * multiple of MSS, at least until the free space gets quite small.
2242 * This would appear to be a side effect of the mbuf implementation.
2243 * Combining these two algorithms results in the observed behavior
2244 * of having a fixed window size at almost all times.
2245 *
2246 * Below we obtain similar behavior by forcing the offered window to
2247 * a multiple of the mss when it is feasible to do so.
2248 *
2249 * Note, we don't "adjust" for TIMESTAMP or SACK option bytes.
2250 * Regular options like TIMESTAMP are taken into account.
2251 */
2252 u32 tcp_select_window(struct sock *sk)
2253 {
2254     struct inet_connection_sock *icsk = inet_csk(sk);
2255     struct tcp_sock *tp = tcp_sk(sk);
2256     /* MSS for the peer's data. Previous versions used mss_clamp
2257      * here. I don't know if the value based on our guesses
2258      * of peer's MSS is better for the performance. It's more correct
2259      * but may be worse for the performance because of rcv_mss
2260      * fluctuations. --SAW 1998/11/1
2261      */
2262     int mss = icsk->icsk_ack.rcv_mss;
2263     int free_space = tcp_space(sk);
2264     int allowed_space = tcp_full_space(sk);
2265     int full_space = min_t(int, tp->window_clamp, allowed_space);
2266     int window;
2267
2268     if (mss > full_space)
2269         mss = full_space;
2270
2271     if (free_space < (full_space >> 1)) {
2272         icsk->icsk_ack.quick = 0;
2273
2274         if (sk_under_memory_pressure(sk))
2275             tp->rcv_ssthresh = min(tp->rcv_ssthresh,
2276                                     4U * tp->advmss);
2277
2278         /* free_space might become our new window, make sure we don't
2279          * increase it due to wscale.
2280          */
2281         free_space = round_down(free_space, 1 << tp->rx_opt.rcv_wscale);
2282
2283         /* if free space is less than mss estimate, or is below 1/16th
2284          * of the maximum allowed, try to move to zero-window, else
2285          * tcp_clamp_window() will grow rcv buf up to tcp_rmem[2], and
2286          * new incoming data is dropped due to memory limits.
2287          * With large window, mss test triggers way too late in order
2288          * to announce zero window in time before rmem limit kicks in.
2289          */
2290         if (free_space < (allowed_space >> 4) || free_space < mss)
2291             return 0;
2292     }
2293
2294     if (free_space > tp->rcv_ssthresh)
2295         free_space = tp->rcv_ssthresh;
2296
2297     /* Don't do rounding if we are using window scaling, since the
2298      * scaled window will not line up with the MSS boundary anyway.
2299      */
2300     window = tp->rcv_wnd;
2301     if (tp->rx_opt.rcv_wscale) {
2302         window = free_space;
2303
2304         /* Advertise enough space so that it won't get scaled away.
2305          * Import case: prevent zero window announcement if
2306          * 1<rcv_wscale > mss.
2307          */
2308         if (((window >> tp->rx_opt.rcv_wscale) << tp->rx_opt.rcv_wscale) != window)
2309             window = (((window >> tp->rx_opt.rcv_wscale) + 1)
2310                       << tp->rx_opt.rcv_wscale);
2311     } else {
2312         /* Get the largest window that is a nice multiple of mss.
2313          * Window clamp already applied above.
2314          * If our current window offering is within 1 mss of the
2315          * free space we just keep it. This prevents the divide
2316          * and multiply from happening most of the time.
2317          * We also don't do any window rounding when the free space
2318          * is too small.
2319          */
2320         if (window <= free_space - mss || window > free_space)

```

```

2321         window = (free_space / mss) * mss;
2322     else if (mss == full_space &&
2323             free_space > window + (full_space >> 1))
2324         window = free_space;
2325 }
2326
2327     return window;
2328 }
2329
2330 /* Collapses two adjacent SKB's during retransmission. */
2331 static void tcp_collapse_retrans(struct sock *sk, struct sk_buff *skb)
2332 {
2333     struct tcp_sock *tp = tcp_sk(sk);
2334     struct sk_buff *next_skb = tcp_write_queue_next(sk, skb);
2335     int skb_size, next_skb_size;
2336
2337     skb_size = skb->len;
2338     next_skb_size = next_skb->len;
2339
2340     BUG_ON(tcp_skb_pcount(skb) != 1 || tcp_skb_pcount(next_skb) != 1);
2341
2342     tcp_highest_sack_combine(sk, next_skb, skb);
2343
2344     tcp_unlink_write_queue(next_skb, sk);
2345
2346     skb_copy_from_linear_data(next_skb, skb_put(skb, next_skb_size),
2347                               next_skb_size);
2348
2349     if (next_skb->ip_summed == CHECKSUM_PARTIAL)
2350         skb->ip_summed = CHECKSUM_PARTIAL;
2351
2352     if (skb->ip_summed != CHECKSUM_PARTIAL)
2353         skb->csum = csum_block_add(skb->csum, next_skb->csum, skb_size);
2354
2355     /* Update sequence range on original skb. */
2356     TCP_SKB_CB(skb)->end_seq = TCP_SKB_CB(next_skb)->end_seq;
2357
2358     /* Merge over control information. This moves PSH/FIN etc. over */
2359     TCP_SKB_CB(skb)->tcp_flags |= TCP_SKB_CB(next_skb)->tcp_flags;
2360
2361     /* ALL done, get rid of second SKB and account for it so
2362      * packet counting does not break.
2363      */
2364     TCP_SKB_CB(skb)->sacked |= TCP_SKB_CB(next_skb)->sacked & TCPCB_EVER_RETRANS;
2365
2366     /* changed transmit queue under us so clear hints */
2367     tcp_clear_retrans_hints_partial(tp);
2368     if (next_skb == tp->retransmit_skb_hint)
2369         tp->retransmit_skb_hint = skb;
2370
2371     tcp_adjust_pcount(sk, next_skb, tcp_skb_pcount(next_skb));
2372
2373     sk_wmem_free_skb(sk, next_skb);
2374 }
2375
2376 /* Check if coalescing SKBs is legal. */
2377 static bool tcp_can_collapse(const struct sock *sk, const struct sk_buff *skb)
2378 {
2379     if (tcp_skb_pcount(skb) > 1)
2380         return false;
2381     /* TODO: SACK collapsing could be used to remove this condition */
2382     if (skb_shinfo(skb)->nr_frags != 0)
2383         return false;
2384     if (skb_cloned(skb))
2385         return false;
2386     if (skb == tcp_send_head(sk))
2387         return false;
2388     /* Some heuristics for collapsing over SACK'd could be invented */
2389     if (TCP_SKB_CB(skb)->sacked & TCPCB_SACKED_ACKED)
2390         return false;
2391
2392     return true;
2393 }
2394
2395 /* Collapse packets in the retransmit queue to make to create
2396  * less packets on the wire. This is only done on retransmission.
2397  */
2398 static void tcp_retrans_try_collapse(struct sock *sk, struct sk_buff *to,
2399                                     int space)
2400 {
2401     struct tcp_sock *tp = tcp_sk(sk);
2402     struct sk_buff *skb = to, *tmp;
2403     bool first = true;
2404
2405     if (!sysctl_tcp_retrans_collapse)
2406         return;
2407     if (TCP_SKB_CB(skb)->tcp_flags & TCPHDR_SYN)
2408         return;
2409
2410     tcp_for_write_queue_from_safe(skb, tmp, sk) {

```

```

2411         if (!tcp_can_collapse(sk, skb))
2412             break;
2413
2414         space -= skb->len;
2415
2416         if (first) {
2417             first = false;
2418             continue;
2419         }
2420
2421         if (space < 0)
2422             break;
2423         /* Punt if not enough space exists in the first SKB for
2424            * the data in the second
2425            */
2426         if (skb->len > skb_availroom(to))
2427             break;
2428
2429         if (after(TCP_SKB_CB(skb)->end_seq, tcp_wnd_end(tp)))
2430             break;
2431
2432         tcp_collapse_retrans(sk, to);
2433     }
2434 }
2435
2436 /* This retransmits one SKB. Policy decisions and retransmit queue
2437  * state updates are done by the caller. Returns non-zero if an
2438  * error occurred which prevented the send.
2439  */
2440 int tcp_retransmit_skb(struct sock *sk, struct sk_buff *skb)
2441 {
2442     struct tcp_sock *tp = tcp_sk(sk);
2443     struct inet_connection_sock *icsk = inet_csk(sk);
2444     unsigned int cur_mss;
2445     int err;
2446
2447     /* Inconclusive MTU probe */
2448     if (icsk->icsk_mtup.probe_size) {
2449         icsk->icsk_mtup.probe_size = 0;
2450     }
2451
2452     /* Do not sent more than we queued. 1/4 is reserved for possible
2453        * copying overhead: fragmentation, tunneling, mangling etc.
2454        */
2455     if (atomic_read(&sk->sk_wmem_alloc) >
2456         min(sk->sk_wmem_queued + (sk->sk_wmem_queued >> 2), sk->sk_sndbuf))
2457         return -EAGAIN;
2458
2459     if (skb_still_in_host_queue(sk, skb))
2460         return -EBUSY;
2461
2462     if (before(TCP_SKB_CB(skb)->seq, tp->snd_una)) {
2463         if (before(TCP_SKB_CB(skb)->end_seq, tp->snd_una))
2464             BUG();
2465         if (tcp_trim_head(sk, skb, tp->snd_una - TCP_SKB_CB(skb)->seq))
2466             return -ENOMEM;
2467     }
2468
2469     if (inet_csk(sk)->icsk_af_ops->rebuild_header(sk))
2470         return -EHOSTUNREACH; /* Routing failure or similar. */
2471
2472     cur_mss = tcp_current_mss(sk);
2473
2474     /* If receiver has shrunk his window, and skb is out of
2475        * new window, do not retransmit it. The exception is the
2476        * case, when window is shrunk to zero. In this case
2477        * our retransmit serves as a zero window probe.
2478        */
2479     if (!before(TCP_SKB_CB(skb)->seq, tcp_wnd_end(tp)) &&
2480         TCP_SKB_CB(skb)->seq != tp->snd_una)
2481         return -EAGAIN;
2482
2483     if (skb->len > cur_mss) {
2484         if (tcp_fragment(sk, skb, cur_mss, cur_mss, GFP_ATOMIC))
2485             return -ENOMEM; /* We'll try again later. */
2486     } else {
2487         int oldpcount = tcp_skb_pcount(skb);
2488
2489         if (unlikely(oldpcount > 1)) {
2490             if (skb_unclone(skb, GFP_ATOMIC))
2491                 return -ENOMEM;
2492             tcp_init_tso_segs(sk, skb, cur_mss);
2493             tcp_adjust_pcount(sk, skb, oldpcount - tcp_skb_pcount(skb));
2494         }
2495     }
2496
2497     tcp_retrans_try_collapse(sk, skb, cur_mss);
2498
2499     /* Make a copy, if the first transmission SKB clone we made
2500        * is still in somebody's hands, else make a clone.

```

```

2501     */
2502     TCP_SKB_CB(skb)->when = tcp_time_stamp;
2503
2504     /* make sure skb->data is aligned on arches that require it
2505      * and check if ack-trimming & collapsing extended the headroom
2506      * beyond what csum_start can cover.
2507     */
2508     if (unlikely((NET_IP_ALIGN && ((unsigned long)skb->data & 3)) ||
2509                 skb_headroom(skb) >= 0xFFFF)) {
2510         struct sk_buff *nskb = __pskb_copy(skb, MAX_TCP_HEADER,
2511                                             GFP_ATOMIC);
2512         err = nskb ? tcp_transmit_skb(sk, nskb, 0, GFP_ATOMIC) :
2513                 -ENOBUFS;
2514     } else {
2515         err = tcp_transmit_skb(sk, skb, 1, GFP_ATOMIC);
2516     }
2517
2518     if (likely(!err)) {
2519         TCP_SKB_CB(skb)->sacked |= TCPCB_EVER_RETRANS;
2520         /* Update global TCP statistics. */
2521         TCP_INC_STATS(sock_net(sk), TCP_MIB_RETRANSSEGS);
2522         if (TCP_SKB_CB(skb)->tcp_flags & TCPHDR_SYN)
2523             NET_INC_STATS_BH(sock_net(sk), LINUX_MIB_TCPSYNRETRANS);
2524         tp->total_retrans++;
2525     }
2526     return err;
2527 }
2528
2529 int tcp_retransmit_skb(struct sock *sk, struct sk_buff *skb)
2530 {
2531     struct tcp_sock *tp = tcp_sk(sk);
2532     int err = __tcp_retransmit_skb(sk, skb);
2533
2534     if (err == 0) {
2535         #if FASTRETRANS_DEBUG > 0
2536         if (TCP_SKB_CB(skb)->sacked & TCPCB_SACKED_RETRANS) {
2537             net_dbg_ratelimited("retrans_out leaked\n");
2538         }
2539         #endif
2540         if (!tp->retrans_out)
2541             tp->lost_retrans_low = tp->snd_next;
2542         TCP_SKB_CB(skb)->sacked |= TCPCB_RETRANS;
2543         tp->retrans_out += tcp_skb_pcount(skb);
2544
2545         /* Save stamp of the first retransmit. */
2546         if (!tp->retrans_stamp)
2547             tp->retrans_stamp = TCP_SKB_CB(skb)->when;
2548
2549         /* snd_next is stored to detect loss of retransmitted segment,
2550          * see tcp_input.c tcp_sacktag_write_queue().
2551          */
2552         TCP_SKB_CB(skb)->ack_seq = tp->snd_next;
2553     } else if (err != -EBUSY) {
2554         NET_INC_STATS_BH(sock_net(sk), LINUX_MIB_TCPRETRANSFAIL);
2555     }
2556
2557     if (tp->undo_retrans < 0)
2558         tp->undo_retrans = 0;
2559     tp->undo_retrans += tcp_skb_pcount(skb);
2560     return err;
2561 }
2562
2563 /* Check if we forward retransmits are possible in the current
2564  * window/congestion state.
2565  */
2566 static bool tcp_can_forward_retransmit(struct sock *sk)
2567 {
2568     const struct inet_connection_sock *icsk = inet_csk(sk);
2569     const struct tcp_sock *tp = tcp_sk(sk);
2570
2571     /* Forward retransmissions are possible only during Recovery. */
2572     if (icsk->icsk_ca_state != TCP_CA_Recovery)
2573         return false;
2574
2575     /* No forward retransmissions in Reno are possible. */
2576     if (tcp_is_reno(tp))
2577         return false;
2578
2579     /* Yeah, we have to make difficult choice between forward transmission
2580      * and retransmission... Both ways have their merits...
2581      *
2582      * For now we do not retransmit anything, while we have some new
2583      * segments to send. In the other cases, follow rule 3 for
2584      * NextSeg() specified in RFC3517.
2585      */
2586
2587     if (tcp_may_send_now(sk))
2588         return false;
2589
2590     return true;

```



```

2591 }
2592
2593 /* This gets called after a retransmit timeout, and the initially
2594  * retransmitted data is acknowledged. It tries to continue
2595  * resending the rest of the retransmit queue, until either
2596  * we've sent it all or the congestion window limit is reached.
2597  * If doing SACK, the first ACK which comes back for a timeout
2598  * based retransmit packet might feed us FACK information again.
2599  * If so, we use it to avoid unnecessarily retransmissions.
2600  */
2601 void tcp_xmit_retransmit_queue(struct sock *sk)
2602 {
2603     const struct inet_connection_sock *icsk = inet_csk(sk);
2604     struct tcp_sock *tp = tcp_sk(sk);
2605     struct sk_buff *skb;
2606     struct sk_buff *hole = NULL;
2607     u32 last_lost;
2608     int mib_idx;
2609     int fwd_rexmitting = 0;
2610
2611     if (!tp->packets_out)
2612         return;
2613
2614     if (!tp->lost_out)
2615         tp->retransmit_high = tp->snd_una;
2616
2617     if (tp->retransmit_skb_hint) {
2618         skb = tp->retransmit_skb_hint;
2619         last_lost = TCP_SKB_CB(skb)->end_seq;
2620         if (after(last_lost, tp->retransmit_high))
2621             last_lost = tp->retransmit_high;
2622     } else {
2623         skb = tcp_write_queue_head(sk);
2624         last_lost = tp->snd_una;
2625     }
2626
2627     tcp_for_write_queue_from(skb, sk) {
2628         u8 sacked = TCP_SKB_CB(skb)->sacked;
2629
2630         if (skb == tcp_send_head(sk))
2631             break;
2632         /* we could do better than to assign each time */
2633         if (hole == NULL)
2634             tp->retransmit_skb_hint = skb;
2635
2636         /* Assume this retransmit will generate
2637          * only one packet for congestion window
2638          * calculation purposes. This works because
2639          * tcp_retransmit_skb() will chop up the
2640          * packet to be MSS sized and all the
2641          * packet counting works out.
2642          */
2643         if (tcp_packets_in_flight(tp) >= tp->snd_cwnd)
2644             return;
2645
2646         if (fwd_rexmitting) {
2647             begin_fwd:
2648                 if (!before(TCP_SKB_CB(skb)->seq, tcp_highest_sack_seq(tp)))
2649                     break;
2650                 mib_idx = LINUX_MIB_TCPFORWARDRETRANS;
2651
2652             } else if (!before(TCP_SKB_CB(skb)->seq, tp->retransmit_high)) {
2653                 tp->retransmit_high = last_lost;
2654                 if (!tcp_can_forward_retransmit(sk))
2655                     break;
2656                 /* Backtrack if necessary to non-L'ed skb */
2657                 if (hole != NULL) {
2658                     skb = hole;
2659                     hole = NULL;
2660                 }
2661                 fwd_rexmitting = 1;
2662                 goto begin_fwd;
2663             } else if (!(sacked & TCPCB_LOST)) {
2664                 if (hole == NULL && !(sacked & (TCPCB_SACKED_RETRANS|TCPCB_SACKED_ACKED)))
2665                     hole = skb;
2666                 continue;
2667             } else {
2668                 last_lost = TCP_SKB_CB(skb)->end_seq;
2669                 if (icsk->icsk_ca_state != TCP_CA_Loss)
2670                     mib_idx = LINUX_MIB_TCPFASTRETRANS;
2671                 else
2672                     mib_idx = LINUX_MIB_TCPSLOWSTARTRETRANS;
2673             }
2674         }
2675
2676         if (sacked & (TCPCB_SACKED_ACKED|TCPCB_SACKED_RETRANS))
2677             continue;
2678
2679         if (tcp_retransmit_skb(sk, skb))

```

```

2681         return;
2682
2683         NET_INC_STATS_BH(sock_net(sk), mib_idx);
2684
2685         if (tcp_in_cwnd_reduction(sk))
2686             tp->prrr_out += tcp_skb_pcount(skb);
2687
2688         if (skb == tcp_write_queue_head(sk))
2689             inet_csk_reset_xmit_timer(sk, ICSK_TIME_RETRANS,
2690                                     inet_csk(sk)->icsk_rto,
2691                                     TCP_RTO_MAX);
2692     }
2693 }
2694
2695 /* Send a fin. The caller locks the socket for us. This cannot be
2696  * allowed to fail queueing a FIN frame under any circumstances.
2697  */
2698 void tcp_send_fin(struct sock *sk)
2699 {
2700     struct tcp_sock *tp = tcp_sk(sk);
2701     struct sk_buff *skb = tcp_write_queue_tail(sk);
2702     int mss_now;
2703
2704     /* Optimization, tack on the FIN if we have a queue of
2705      * unsent frames. But be careful about outgoing SACKS
2706      * and IP options.
2707      */
2708     mss_now = tcp_current_mss(sk);
2709
2710     if (tcp_send_head(sk) != NULL) {
2711         TCP_SKB_CB(skb)->tcp_flags |= TCPHDR_FIN;
2712         TCP_SKB_CB(skb)->end_seq++;
2713         tp->write_seq++;
2714     } else {
2715         /* Socket is Locked, keep trying until memory is available. */
2716         for (;;) {
2717             skb = alloc_skb_fclone(MAX_TCP_HEADER,
2718                                   sk->sk_allocation);
2719             if (skb)
2720                 break;
2721             yield();
2722         }
2723
2724         /* Reserve space for headers and prepare control bits. */
2725         skb_reserve(skb, MAX_TCP_HEADER);
2726         /* FIN eats a sequence byte, write_seq advanced by tcp_queue_skb(). */
2727         tcp_init_nondata_skb(skb, tp->write_seq,
2728                             TCPHDR_ACK | TCPHDR_FIN);
2729         tcp_queue_skb(sk, skb);
2730     }
2731     tcp_push_pending_frames(sk, mss_now, TCP_NAGLE_OFF);
2732 }
2733
2734 /* We get here when a process closes a file descriptor (either due to
2735  * an explicit close() or as a byproduct of exit()'ing) and there
2736  * was unread data in the receive queue. This behavior is recommended
2737  * by RFC 2525, section 2.17. -DaveM
2738  */
2739 void tcp_send_active_reset(struct sock *sk, gfp_t priority)
2740 {
2741     struct sk_buff *skb;
2742
2743     /* NOTE: No TCP options attached and we never retransmit this. */
2744     skb = alloc_skb(MAX_TCP_HEADER, priority);
2745     if (!skb) {
2746         NET_INC_STATS(sock_net(sk), LINUX_MIB_TCPABORTFAILED);
2747         return;
2748     }
2749
2750     /* Reserve space for headers and prepare control bits. */
2751     skb_reserve(skb, MAX_TCP_HEADER);
2752     tcp_init_nondata_skb(skb, tcp_acceptable_seq(sk),
2753                         TCPHDR_ACK | TCPHDR_RST);
2754
2755     /* Send it off. */
2756     TCP_SKB_CB(skb)->when = tcp_time_stamp;
2757     if (tcp_transmit_skb(sk, skb, 0, priority))
2758         NET_INC_STATS(sock_net(sk), LINUX_MIB_TCPABORTFAILED);
2759
2760     TCP_INC_STATS(sock_net(sk), TCP_MIB_OUTRSTS);
2761 }
2762
2763 /* Send a crossed SYN-ACK during socket establishment.
2764  * WARNING: This routine must only be called when we have already sent
2765  * a SYN packet that crossed the incoming SYN that caused this routine
2766  * to get called. If this assumption fails then the initial rcv_wnd
2767  * and rcv_wscale values will not be correct.
2768  */
2769 int tcp_send_synack(struct sock *sk)
2770 {
2771     struct sk_buff *skb;

```

```

2771     skb = tcp_write_queue_head(skb);
2772     if (skb == NULL || !(TCP_SKB_CB(skb)->tcp_flags & TCPHDR_SYN)) {
2773         pr_debug("%s: wrong queue state\n", __func__);
2774         return -EFAULT;
2775     }
2776     if (!(TCP_SKB_CB(skb)->tcp_flags & TCPHDR_ACK)) {
2777         if (skb_cloned(skb)) {
2778             struct sk_buff *nskb = skb_copy(skb, GFP_ATOMIC);
2779             if (nskb == NULL)
2780                 return -ENOMEM;
2781             tcp_unlink_write_queue(skb, sk);
2782             skb_header_release(nskb);
2783             tcp_add_write_queue_head(sk, nskb);
2784             sk_wmem_free_skb(sk, skb);
2785             sk->sk_wmem_queued += nskb->truesize;
2786             sk_mem_charge(sk, nskb->truesize);
2787             skb = nskb;
2788         }
2789         TCP_SKB_CB(skb)->tcp_flags |= TCPHDR_ACK;
2790         TCP_ECN_send_synack(tcp_sk(sk), skb);
2791     }
2792     TCP_SKB_CB(skb)->when = tcp_time_stamp;
2793     return tcp_transmit_skb(sk, skb, 1, GFP_ATOMIC);
2794 }
2795
2796 /**
2797  * tcp_make_synack - Prepare a SYN-ACK.
2798  * sk: Listener socket
2799  * dst: dst entry attached to the SYNACK
2800  * req: request_sock pointer
2801  *
2802  * Allocate one skb and build a SYNACK packet.
2803  * @dst is consumed : Caller should not use it again.
2804  */
2805 struct sk_buff *tcp_make_synack(struct sock *sk, struct dst_entry *dst,
2806                                struct request_sock *req,
2807                                struct tcp_fastopen_cookie *foc)
2808 {
2809     struct tcp_out_options opts;
2810     struct inet_request_sock *ireq = inet_rsk(req);
2811     struct tcp_sock *tp = tcp_sk(sk);
2812     struct tcphdr *th;
2813     struct sk_buff *skb;
2814     struct tcp_md5sig_key *md5;
2815     int tcp_header_size;
2816     int mss;
2817
2818     skb = sock_wmalloc(sk, MAX_TCP_HEADER, 1, GFP_ATOMIC);
2819     if (unlikely(!skb)) {
2820         dst_release(dst);
2821         return NULL;
2822     }
2823     /* Reserve space for headers. */
2824     skb_reserve(skb, MAX_TCP_HEADER);
2825
2826     skb_dst_set(skb, dst);
2827     security_skb_owned_by(skb, sk);
2828
2829     mss = dst_metric_advms(dst);
2830     if (tp->rx_opt.user_mss && tp->rx_opt.user_mss < mss)
2831         mss = tp->rx_opt.user_mss;
2832
2833     memset(&opts, 0, sizeof(opts));
2834 #ifdef CONFIG_SYN_COOKIES
2835     if (unlikely(req->cookie_ts))
2836         TCP_SKB_CB(skb)->when = cookie_init_timestamp(req);
2837     else
2838 #endif
2839         TCP_SKB_CB(skb)->when = tcp_time_stamp;
2840
2841     tcp_header_size = tcp_synack_options(sk, req, mss, skb, &opts, &md5,
2842                                         foc) + sizeof(*th);
2843
2844     skb_push(skb, tcp_header_size);
2845     skb_reset_transport_header(skb);
2846
2847     th = tcp_hdr(skb);
2848     memset(th, 0, sizeof(struct tcphdr));
2849     th->syn = 1;
2850     th->ack = 1;
2851     TCP_ECN_make_synack(req, th);
2852     th->source = htons(ireq->ir_num);
2853     th->dest = ireq->ir_rmt_port;
2854     /* Setting of flags are superfluous here for callers (and ECE is
2855      * not even correctly set)
2856      */
2857     tcp_init_nondata_skb(skb, tcp_rsk(req)->snt_isn,
2858                          TCPHDR_SYN | TCPHDR_ACK);
2859
2860

```

```

2861     th->seq = htonl(TCP_SKB_CB(skb)->seq);
2862     /* XXX data is queued and acked as is. No buffer/window check */
2863     th->ack_seq = htonl(tcp_rsk(req)->rcv_nxt);
2864
2865     /* RFC1323: The window in SYN & SYN/ACK segments is never scaled. */
2866     th->window = htons(min(req->rcv_wnd, 65535U));
2867     tcp_options_write((__be32*)(th + 1), tp, &opts);
2868     th->doff = (tcp_header_size >> 2);
2869     TCP_INC_STATS_BH(sock_net(sk), TCP_MIB_OUTSEGS);
2870
2871 #ifdef CONFIG_TCP_MD5SIG
2872     /* Okay, we have all we need - do the md5 hash if needed */
2873     if (md5) {
2874         tcp_rsk(req)->af_specific->calc_md5_hash(opts.hash_location,
2875         md5, NULL, req, skb);
2876     }
2877 #endif
2878
2879     return skb;
2880 }
2881 EXPORT_SYMBOL(tcp_make_synack);
2882
2883 /* Do all connect socket setups that can be done AF independent. */
2884 static void tcp_connect_init(struct sock *sk)
2885 {
2886     const struct dst_entry *dst = __sk_dst_get(sk);
2887     struct tcp_sock *tp = tcp_sk(sk);
2888     u8 rcv_wscale;
2889
2890     /* We'll fix this up when we get a response from the other end.
2891      * See tcp_input.c:tcp_rcv_state_process case TCP_SYN_SENT.
2892      */
2893     tp->tcp_header_len = sizeof(struct tcphdr) +
2894         (sysctl_tcp_timestamps ? TCPOLEN_TSTAMP_ALIGNED : 0);
2895
2896 #ifdef CONFIG_TCP_MD5SIG
2897     if (tp->af_specific->md5_lookup(sk, sk) != NULL)
2898         tp->tcp_header_len += TCPOLEN_MD5SIG_ALIGNED;
2899 #endif
2900
2901     /* If user gave his TCP_MAXSEG, record it to clamp */
2902     if (tp->rx_opt.user_mss)
2903         tp->rx_opt.mss_clamp = tp->rx_opt.user_mss;
2904     tp->max_window = 0;
2905     tcp_mtup_init(sk);
2906     tcp_sync_mss(sk, dst_mtu(dst));
2907
2908     if (!tp->window_clamp)
2909         tp->window_clamp = dst_metric(dst, RTAX_WINDOW);
2910     tp->advms = dst_metric_advms(dst);
2911     if (tp->rx_opt.user_mss && tp->rx_opt.user_mss < tp->advms)
2912         tp->advms = tp->rx_opt.user_mss;
2913
2914     tcp_initialize_rcv_mss(sk);
2915
2916     /* Limit the window selection if the user enforce a smaller rx buffer */
2917     if (sk->sk_userlocks & SOCK_RCVBUF_LOCK &&
2918         (tp->window_clamp > tcp_full_space(sk) || tp->window_clamp == 0))
2919         tp->window_clamp = tcp_full_space(sk);
2920
2921     tcp_select_initial_window(tcp_full_space(sk),
2922         tp->advms - (tp->rx_opt.ts_recent_stamp ? tp->tcp_header_len - sizeof(struct tcphdr) : 0),
2923         &tp->rcv_wnd,
2924         &tp->window_clamp,
2925         sysctl_tcp_window_scaling,
2926         &rcv_wscale,
2927         dst_metric(dst, RTAX_INITRWND));
2928
2929     tp->rx_opt.rcv_wscale = rcv_wscale;
2930     tp->rcv_ssthresh = tp->rcv_wnd;
2931
2932     sk->sk_err = 0;
2933     sock_reset_flag(sk, SOCK_DONE);
2934     tp->snd_wnd = 0;
2935     tcp_init_wl(tp, 0);
2936     tp->snd_una = tp->write_seq;
2937     tp->snd_sml = tp->write_seq;
2938     tp->snd_up = tp->write_seq;
2939     tp->snd_nxt = tp->write_seq;
2940
2941     if (likely(!tp->repair))
2942         tp->rcv_nxt = 0;
2943     else
2944         tp->rcv_tstamp = tcp_time_stamp;
2945     tp->rcv_wup = tp->rcv_nxt;
2946     tp->copied_seq = tp->rcv_nxt;
2947
2948     inet_csk(sk)->icsk_rto = TCP_TIMEOUT_INIT;
2949     inet_csk(sk)->icsk_retransmits = 0;
2950     tcp_clear_retrans(tp);

```

```

2951 }
2952
2953 static void tcp_connect_queue_skb(struct sock *sk, struct sk_buff *skb)
2954 {
2955     struct tcp_sock *tp = tcp_sk(sk);
2956     struct tcp_skb_cb *tcb = TCP_SKB_CB(skb);
2957
2958     tcb->end_seq += skb->len;
2959     skb_header_release(skb);
2960     tcp_add_write_queue_tail(sk, skb);
2961     sk->sk_wmem_queued += skb->truesize;
2962     sk_mem_charge(sk, skb->truesize);
2963     tp->write_seq = tcb->end_seq;
2964     tp->packets_out += tcp_skb_pcount(skb);
2965 }
2966
2967 /* Build and send a SYN with data and (cached) Fast Open cookie. However,
2968  * queue a data-only packet after the regular SYN, such that regular SYNs
2969  * are retransmitted on timeouts. Also if the remote SYN-ACK acknowledges
2970  * only the SYN sequence, the data are retransmitted in the first ACK.
2971  * If cookie is not cached or other error occurs, falls back to send a
2972  * regular SYN with Fast Open cookie request option.
2973  */
2974 static int tcp_send_syn_data(struct sock *sk, struct sk_buff *syn)
2975 {
2976     struct tcp_sock *tp = tcp_sk(sk);
2977     struct tcp_fastopen_request *fo = tp->fastopen_req;
2978     int syn_loss = 0, space, i, err = 0, iovlen = fo->data->msg_iovlen;
2979     struct sk_buff *syn_data = NULL, *data;
2980     unsigned long last_syn_loss = 0;
2981
2982     tp->rx_opt.mss_clamp = tp->advmss; /* If MSS is not cached */
2983     tcp_fastopen_cache_get(sk, &tp->rx_opt.mss_clamp, &fo->cookie,
2984                           &syn_loss, &last_syn_loss);
2985     /* Recurring FO SYN losses: revert to regular handshake temporarily */
2986     if (syn_loss > 1 &&
2987         time_before(jiffies, last_syn_loss + (60*HZ << syn_loss))) {
2988         fo->cookie.len = -1;
2989         goto fallback;
2990     }
2991
2992     if (sysctl_tcp_fastopen & TFO_CLIENT_NO_COOKIE)
2993         fo->cookie.len = -1;
2994     else if (fo->cookie.len <= 0)
2995         goto fallback;
2996
2997     /* MSS for SYN-data is based on cached MSS and bounded by PMTU and
2998      * user-MSS. Reserve maximum option space for middleboxes that add
2999      * private TCP options. The cost is reduced data space in SYN :(
3000      */
3001     if (tp->rx_opt.user_mss && tp->rx_opt.user_mss < tp->rx_opt.mss_clamp)
3002         tp->rx_opt.mss_clamp = tp->rx_opt.user_mss;
3003     space = tcp_mtu_to_mss(sk, inet_csk(sk)->icsk_pmtu_cookie) -
3004             MAX_TCP_OPTION_SPACE;
3005
3006     space = min_t(size_t, space, fo->size);
3007
3008     /* Limit to order-0 allocations */
3009     space = min_t(size_t, space, SKB_MAX_HEAD(MAX_TCP_HEADER));
3010
3011     syn_data = skb_copy_expand(syn, MAX_TCP_HEADER, space,
3012                               sk->sk_allocation);
3013     if (syn_data == NULL)
3014         goto fallback;
3015
3016     for (i = 0; i < iovlen && syn_data->len < space; ++i) {
3017         struct iovec *iov = &fo->data->msg_iov[i];
3018         unsigned char __user *from = iov->iov_base;
3019         int len = iov->iov_len;
3020
3021         if (syn_data->len + len > space)
3022             len = space - syn_data->len;
3023         else if (i + 1 == iovlen)
3024             /* No more data pending in inet_wait_for_connect() */
3025             fo->data = NULL;
3026
3027         if (skb_add_data(syn_data, from, len))
3028             goto fallback;
3029     }
3030
3031     /* Queue a data-only packet after the regular SYN for retransmission */
3032     data = pskb_copy(syn_data, sk->sk_allocation);
3033     if (data == NULL)
3034         goto fallback;
3035     TCP_SKB_CB(data)->seq++;
3036     TCP_SKB_CB(data)->tcp_flags &= ~TCPHDR_SYN;
3037     TCP_SKB_CB(data)->tcp_flags = (TCPHDR_ACK|TCPHDR_PSH);
3038     tcp_connect_queue_skb(sk, data);
3039     fo->copied = data->len;
3040

```

```

3041  /* syn_data is about to be sent, we need to take current time stamps
3042  * for the packets that are in write queue : SYN packet and DATA
3043  */
3044  skb_mstamp_get(&syn->skb_mstamp);
3045  data->skb_mstamp = syn->skb_mstamp;
3046
3047  if (tcp_transmit_skb(sk, syn_data, 0, sk->sk_allocation) == 0) {
3048      tp->syn_data = (fo->copied > 0);
3049      NET_INC_STATS(sock_net(sk), LINUX_MIB_TCPORIGDATASENT);
3050      goto done;
3051  }
3052  syn_data = NULL;
3053
3054  fallback:
3055  /* Send a regular SYN with Fast Open cookie request option */
3056  if (fo->cookie.len > 0)
3057      fo->cookie.len = 0;
3058  err = tcp_transmit_skb(sk, syn, 1, sk->sk_allocation);
3059  if (err)
3060      tp->syn_fastopen = 0;
3061  kfree_skb(syn_data);
3062  done:
3063  fo->cookie.len = -1; /* Exclude Fast Open option for SYN retries */
3064  return err;
3065  }
3066
3067  /* Build a SYN and send it off. */
3068  int tcp_connect(struct sock *sk)
3069  {
3070      struct tcp_sock *tp = tcp_sk(sk);
3071      struct sk_buff *buff;
3072      int err;
3073
3074      tcp_connect_init(sk);
3075
3076      if (unlikely(tp->repair)) {
3077          tcp_finish_connect(sk, NULL);
3078          return 0;
3079      }
3080
3081      buff = alloc_skb_fclone(MAX_TCP_HEADER + 15, sk->sk_allocation);
3082      if (unlikely(buff == NULL))
3083          return -ENOMEM;
3084
3085      /* Reserve space for headers. */
3086      skb_reserve(buff, MAX_TCP_HEADER);
3087
3088      tcp_init_nondata_skb(buff, tp->write_seq++, TCPHDR_SYN);
3089      tp->retrans_stamp = TCP_SKB_CB(buff)->when = tcp_time_stamp;
3090      tcp_connect_queue_skb(sk, buff);
3091      TCP_ECN_send_syn(sk, buff);
3092
3093      /* Send off SYN; include data in Fast Open. */
3094      err = tp->fastopen_req ? tcp_send_syn_data(sk, buff) :
3095          tcp_transmit_skb(sk, buff, 1, sk->sk_allocation);
3096      if (err == -ECONNREFUSED)
3097          return err;
3098
3099      /* We change tp->snd_nxt after the tcp_transmit_skb() call
3100       * in order to make this packet get counted in tcpOutSegs.
3101       */
3102      tp->snd_nxt = tp->write_seq;
3103      tp->pushed_seq = tp->write_seq;
3104      TCP_INC_STATS(sock_net(sk), TCP_MIB_ACTIVEOPENS);
3105
3106      /* Timer for repeating the SYN until an answer. */
3107      inet_csk_reset_xmit_timer(sk, ICSK_TIME_RETRANS,
3108                              inet_csk(sk)->icsk_rto, TCP_RTO_MAX);
3109      return 0;
3110  }
3111  EXPORT_SYMBOL(tcp_connect);
3112
3113  /* Send out a delayed ack, the caller does the policy checking
3114  * to see if we should even be here. See tcp_input.c:tcp_ack_snd_check()
3115  * for details.
3116  */
3117  void tcp_send_delayed_ack(struct sock *sk)
3118  {
3119      struct inet_connection_sock *icsk = inet_csk(sk);
3120      int ato = icsk->icsk_ack.ato;
3121      unsigned long timeout;
3122
3123      if (ato > TCP_DELACK_MIN) {
3124          const struct tcp_sock *tp = tcp_sk(sk);
3125          int max_ato = HZ / 2;
3126
3127          if (icsk->icsk_ack.pingpong ||
3128              (icsk->icsk_ack.pending & ICSK_ACK_PUSHED))
3129              max_ato = TCP_DELACK_MAX;
3130      }

```

```

3131      /* Slow path, intersegment interval is "high". */
3132
3133      /* If some rtt estimate is known, use it to bound delayed ack.
3134       * Do not use inet_csk(sk)->icsk_rto here, use results of rtt measurements
3135       * directly.
3136       */
3137      if (tp->srtt_us) {
3138          int rtt = max_t(int, usecs_to_jiffies(tp->srtt_us >> 3),
3139                          TCP_DELACK_MIN);
3140
3141          if (rtt < max_ato)
3142              max_ato = rtt;
3143      }
3144
3145      ato = min(ato, max_ato);
3146  }
3147
3148  /* Stay within the limit we were given */
3149  timeout = jiffies + ato;
3150
3151  /* Use new timeout only if there wasn't a older one earlier. */
3152  if (icsk->icsk_ack.pending & ICSK_ACK_TIMER) {
3153      /* If delack timer was blocked or is about to expire,
3154       * send ACK now.
3155       */
3156      if (icsk->icsk_ack.blocked ||
3157          time_before_eq(icsk->icsk_ack.timeout, jiffies + (ato >> 2))) {
3158          tcp_send_ack(sk);
3159          return;
3160      }
3161
3162      if (!time_before(timeout, icsk->icsk_ack.timeout))
3163          timeout = icsk->icsk_ack.timeout;
3164  }
3165  icsk->icsk_ack.pending |= ICSK_ACK_SCHED | ICSK_ACK_TIMER;
3166  icsk->icsk_ack.timeout = timeout;
3167  sk_reset_timer(sk, &icsk->icsk_delack_timer, timeout);
3168 }
3169
3170 /* This routine sends an ack and also updates the window. */
3171 void tcp_send_ack(struct sock *sk)
3172 {
3173     struct sk_buff *buff;
3174
3175     /* If we have been reset, we may not send again. */
3176     if (sk->sk_state == TCP_CLOSE)
3177         return;
3178
3179     /* We are not putting this on the write queue, so
3180      * tcp_transmit_skb() will set the ownership to this
3181      * sock.
3182      */
3183     buff = alloc_skb(MAX_TCP_HEADER, sk_gfp_atomic(sk, GFP_ATOMIC));
3184     if (buff == NULL) {
3185         inet_csk_schedule_ack(sk);
3186         inet_csk(sk)->icsk_ack.ato = TCP_ATO_MIN;
3187         inet_csk_reset_xmit_timer(sk, ICSK_TIME_DACK,
3188                                   TCP_DELACK_MAX, TCP_RTO_MAX);
3189         return;
3190     }
3191
3192     /* Reserve space for headers and prepare control bits. */
3193     skb_reserve(buff, MAX_TCP_HEADER);
3194     tcp_init_nondata_skb(buff, tcp_acceptable_seq(sk), TCPHDR_ACK);
3195
3196     /* Send it off, this clears delayed acks for us. */
3197     TCP_SKB_CB(buff)->when = tcp_time_stamp;
3198     tcp_transmit_skb(sk, buff, 0, sk_gfp_atomic(sk, GFP_ATOMIC));
3199 }
3200
3201 /* This routine sends a packet with an out of date sequence
3202  * number. It assumes the other end will try to ack it.
3203  */
3204 /* Question: what should we make while urgent mode?
3205  * 4.4BSD forces sending single byte of data. We cannot send
3206  * out of window data, because we have SND.NXT==SND.MAX...
3207  */
3208 /* Current solution: to send TWO zero-length segments in urgent mode:
3209  * one is with SEG.SEQ=SND.UNA to deliver urgent pointer, another is
3210  * out-of-date with SND.UNA-1 to probe window.
3211  */
3212 static int tcp_xmit_probe_skb(struct sock *sk, int urgent)
3213 {
3214     struct tcp_sock *tp = tcp_sk(sk);
3215     struct sk_buff *skb;
3216
3217     /* We don't queue it, tcp_transmit_skb() sets ownership. */
3218     skb = alloc_skb(MAX_TCP_HEADER, sk_gfp_atomic(sk, GFP_ATOMIC));
3219     if (skb == NULL)
3220         return -1;

```



```

3221
3222 /* Reserve space for headers and set control bits. */
3223 skb_reserve(skb, MAX_TCP_HEADER);
3224 /* Use a previous sequence. This should cause the other
3225  * end to send an ack. Don't queue or clone SKB, just
3226  * send it.
3227  */
3228 tcp_init_nondata_skb(skb, tp->snd_una - !urgent, TCPHDR_ACK);
3229 TCP_SKB_CB(skb)->when = tcp_time_stamp;
3230 return tcp_transmit_skb(sk, skb, 0, GFP_ATOMIC);
3231 }
3232
3233 void tcp_send_window_probe(struct sock *sk)
3234 {
3235     if (sk->sk_state == TCP_ESTABLISHED) {
3236         tcp_sk(sk)->snd_wll = tcp_sk(sk)->rcv_nxt - 1;
3237         tcp_xmit_probe_skb(sk, 0);
3238     }
3239 }
3240
3241 /* Initiate keepalive or window probe from timer. */
3242 int tcp_write_wakeup(struct sock *sk)
3243 {
3244     struct tcp_sock *tp = tcp_sk(sk);
3245     struct sk_buff *skb;
3246
3247     if (sk->sk_state == TCP_CLOSE)
3248         return -1;
3249
3250     if ((skb = tcp_send_head(sk)) != NULL &&
3251         before(TCP_SKB_CB(skb)->seq, tcp_wnd_end(tp))) {
3252         int err;
3253         unsigned int mss = tcp_current_mss(sk);
3254         unsigned int seg_size = tcp_wnd_end(tp) - TCP_SKB_CB(skb)->seq;
3255
3256         if (before(tp->pushed_seq, TCP_SKB_CB(skb)->end_seq))
3257             tp->pushed_seq = TCP_SKB_CB(skb)->end_seq;
3258
3259         /* We are probing the opening of a window
3260          * but the window size is != 0
3261          * must have been a result SWS avoidance ( sender )
3262          */
3263         if (seg_size < TCP_SKB_CB(skb)->end_seq - TCP_SKB_CB(skb)->seq ||
3264             skb->len > mss) {
3265             seg_size = min(seg_size, mss);
3266             TCP_SKB_CB(skb)->tcp_flags |= TCPHDR_PSH;
3267             if (tcp_fragment(sk, skb, seg_size, mss, GFP_ATOMIC))
3268                 return -1;
3269         } else if (!tcp_skb_pcount(skb))
3270             tcp_set_skb_tso_segs(sk, skb, mss);
3271
3272         TCP_SKB_CB(skb)->tcp_flags |= TCPHDR_PSH;
3273         TCP_SKB_CB(skb)->when = tcp_time_stamp;
3274         err = tcp_transmit_skb(sk, skb, 1, GFP_ATOMIC);
3275         if (!err)
3276             tcp_event_new_data_sent(sk, skb);
3277         return err;
3278     } else {
3279         if (between(tp->snd_up, tp->snd_una + 1, tp->snd_una + 0xFFFF))
3280             tcp_xmit_probe_skb(sk, 1);
3281         return tcp_xmit_probe_skb(sk, 0);
3282     }
3283 }
3284
3285 /* A window probe timeout has occurred. If window is not closed send
3286  * a partial packet else a zero probe.
3287  */
3288 void tcp_send_probe0(struct sock *sk)
3289 {
3290     struct inet_connection_sock *icsk = inet_csk(sk);
3291     struct tcp_sock *tp = tcp_sk(sk);
3292     int err;
3293
3294     err = tcp_write_wakeup(sk);
3295
3296     if (tp->packets_out || !tcp_send_head(sk)) {
3297         /* Cancel probe timer, if it is not required. */
3298         icsk->icsk_probes_out = 0;
3299         icsk->icsk_backoff = 0;
3300         return;
3301     }
3302
3303     if (err <= 0) {
3304         if (icsk->icsk_backoff < sysctl_tcp_retries2)
3305             icsk->icsk_backoff++;
3306         icsk->icsk_probes_out++;
3307         inet_csk_reset_xmit_timer(sk, ICSK_TIME_PROBE0,
3308                                 min(icsk->icsk_rto << icsk->icsk_backoff, TCP_RTO_MAX),
3309                                 TCP_RTO_MAX);
3310     } else {

```

```

3311      /* If packet was not sent due to local congestion,
3312       * do not backoff and do not remember icsk_probes_out.
3313       * Let local senders to fight for local resources.
3314       *
3315       * Use accumulated backoff yet.
3316       */
3317      if (!icsk->icsk_probes_out)
3318          icsk->icsk_probes_out = 1;
3319      inet_csk_reset_xmit_timer(sk, ICSK_TIME_PROBE0,
3320                              min(icsk->icsk_rto << icsk->icsk_backoff,
3321                                  TCP_RESOURCE_PROBE_INTERVAL),
3322                              TCP_RTO_MAX);
3323  }
3324  }
3325
3326  int tcp_rtx_synack(struct sock *sk, struct request_sock *req)
3327  {
3328      const struct tcp_request_sock_ops *af_ops = tcp_rsk(req)->af_specific;
3329      struct flowi fl;
3330      int res;
3331
3332      res = af_ops->send_synack(sk, NULL, &fl, req, 0, NULL);
3333      if (!res) {
3334          TCP_INC_STATS_BH(sock_net(sk), TCP_MIB_RETRANSSEGS);
3335          NET_INC_STATS_BH(sock_net(sk), LINUX_MIB_TCPSYNRETRANS);
3336      }
3337      return res;
3338  }
3339  EXPORT_SYMBOL(tcp_rtx_synack);
3340

```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)