

# Linux Cross Reference

## [Free Electrons](#)

## Embedded Linux Experts

• *source navigation* • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version:

[2.0.40](#) [2.2.26](#) [2.4.37](#) [3.1](#) [3.2](#) [3.3](#) [3.4](#) [3.5](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#)

## [Linux/net/ipv4/tcp\\_probe.c](#)

```

1  /*
2   * tcpprobe - Observe the TCP flow with kprobes.
3   *
4   * The idea for this came from Werner Almesberger's umlsim
5   * Copyright (C) 2004, Stephen Hemminger <shemminger@osdl.org>
6   *
7   * This program is free software; you can redistribute it and/or modify
8   * it under the terms of the GNU General Public License as published by
9   * the Free Software Foundation; either version 2 of the License.
10  *
11  * This program is distributed in the hope that it will be useful,
12  * but WITHOUT ANY WARRANTY; without even the implied warranty of
13  * MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
14  * GNU General Public License for more details.
15  *
16  * You should have received a copy of the GNU General Public License
17  * along with this program; if not, write to the Free Software
18  * Foundation, Inc., 675 Mass Ave, Cambridge, MA 02139, USA.
19  */
20
21 #define pr_fmt(fmt) KBUILD_MODNAME ": " fmt
22
23 #include <linux/kernel.h>
24 #include <linux/kprobes.h>
25 #include <linux/socket.h>
26 #include <linux/tcp.h>
27 #include <linux/slab.h>
28 #include <linux/proc_fs.h>
29 #include <linux/module.h>
30 #include <linux/ktime.h>
31 #include <linux/time.h>
32 #include <net/net_namespace.h>
33
34 #include <net/tcp.h>
35
36 MODULE_AUTHOR("Stephen Hemminger <shemminger@linux-foundation.org>");
37 MODULE_DESCRIPTION("TCP cwnd snooper");
38 MODULE_LICENSE("GPL");
39 MODULE_VERSION("1.1");
40
41 static int port_read_mostly;
42 MODULE_PARM_DESC(port, "Port to match (0=all)");
43 module_param(port, int, 0);

```

```

44
45 static unsigned int bufsize __read_mostly = 4096;
46 MODULE_PARM_DESC(bufsize, "Log buffer size in packets (4096)");
47 module_param(bufsize, uint, 0);
48
49 static unsigned int fwmark __read_mostly;
50 MODULE_PARM_DESC(fwmark, "skb mark to match (0=no mark)");
51 module_param(fwmark, uint, 0);
52
53 static int full __read_mostly;
54 MODULE_PARM_DESC(full, "Full log (1=every ack packet received, 0=only cwnd changes)");
55 module_param(full, int, 0);
56
57 static const char procname[] = "tcpprobe";
58
59 struct tcp_log {
60     ktime_t timestamp;
61     union {
62         struct sockaddr raw;
63         struct sockaddr_in v4;
64         struct sockaddr_in6 v6;
65     } src, dst;
66     u16 length;
67     u32 snd_nxt;
68     u32 snd_una;
69     u32 snd_wnd;
70     u32 rcv_wnd;
71     u32 snd_cwnd;
72     u32 ssthresh;
73     u32 srtt;
74 };
75
76 static struct {
77     spinlock_t lock;
78     wait_queue_head_t wait;
79     ktime_t start;
80     u32 lastcwnd;
81
82     unsigned long head, tail;
83     struct tcp_log *log;
84 } tcp_probe;
85
86
87 static inline int tcp_probe_used(void)
88 {
89     return (tcp_probe.head - tcp_probe.tail) & (bufsize - 1);
90 }
91
92 static inline int tcp_probe_avail(void)
93 {
94     return bufsize - tcp_probe_used() - 1;
95 }
96
97 #define tcp_probe_copy_fl_to_si4(inet, si4, mem) \
98     do { \
99         si4.sin_family = AF_INET; \
100         si4.sin_port = inet->inet_##mem##port; \
101         si4.sin_addr.s_addr = inet->inet_##mem##addr; \
102     } while (0) \
103
104
105 /*
106  * Hook inserted to be called before each receive packet.
107  * Note: arguments must match tcp_rcv_established()!
108  */

```

```

109 static void jtcp_rcv_established(struct sock *sk, struct sk_buff *skb,
110                                 const struct tcphdr *th, unsigned int len)
111 {
112     const struct tcp_sock *tp = tcp_sk(sk);
113     const struct inet_sock *inet = inet_sk(sk);
114
115     /* Only update if port or skb mark matches */
116     if (((port == 0 && fwmark == 0) ||
117         ntohs(inet->inet_dport) == port ||
118         ntohs(inet->inet_sport) == port ||
119         (fwmark > 0 && skb->mark == fwmark)) &&
120         (full || tp->snd_cwnd != tcp_probe.lastcwnd)) {
121
122         spin_lock(&tcp_probe.lock);
123         /* If log fills, just silently drop */
124         if (tcp_probe.avail() > 1) {
125             struct tcp_log *p = tcp_probe.log + tcp_probe.head;
126
127             p->tstamp = ktime_get();
128             switch (sk->sk_family) {
129             case AF_INET:
130                 tcp_probe_copy_f1_to_si4(inet, p->src.v4, s);
131                 tcp_probe_copy_f1_to_si4(inet, p->dst.v4, d);
132                 break;
133             case AF_INET6:
134                 memset(&p->src.v6, 0, sizeof(p->src.v6));
135                 memset(&p->dst.v6, 0, sizeof(p->dst.v6));
136 #if IS_ENABLED(CONFIG_IPV6)
137                 p->src.v6.sin6_family = AF_INET6;
138                 p->src.v6.sin6_port = inet->inet_sport;
139                 p->src.v6.sin6_addr = inet6_sk(sk)->saddr;
140
141                 p->dst.v6.sin6_family = AF_INET6;
142                 p->dst.v6.sin6_port = inet->inet_dport;
143                 p->dst.v6.sin6_addr = sk->sk_v6_daddr;
144 #endif
145                 break;
146             default:
147                 BUG();
148             }
149
150             p->length = skb->len;
151             p->snd_nxt = tp->snd_nxt;
152             p->snd_una = tp->snd_una;
153             p->snd_cwnd = tp->snd_cwnd;
154             p->snd_wnd = tp->snd_wnd;
155             p->rcv_wnd = tp->rcv_wnd;
156             p->ssthresh = tcp_current_ssthresh(sk);
157             p->srtt = tp->srtt_us >> 3;
158
159             tcp_probe.head = (tcp_probe.head + 1) & (bufsize - 1);
160         }
161         tcp_probe.lastcwnd = tp->snd_cwnd;
162         spin_unlock(&tcp_probe.lock);
163
164         wake_up(&tcp_probe.wait);
165     }
166
167     jprobe_return();
168 }
169
170 static struct jprobe tcp_jprobe = {
171     .kp = {
172         .symbol_name = "tcp_rcv_established",
173     },

```

```

174     .entry = jtcp_rcv_established,
175 };
176
177 static int tcpprobe_open(struct inode *inode, struct file *file)
178 {
179     /* Reset (empty) Log */
180     spin_lock_bh(&tcp_probe.lock);
181     tcp_probe.head = tcp_probe.tail = 0;
182     tcp_probe.start = ktime_get();
183     spin_unlock_bh(&tcp_probe.lock);
184
185     return 0;
186 }
187
188 static int tcpprobe_sprint(char *tbuf, int n)
189 {
190     const struct tcp_log *p
191         = tcp_probe.log + tcp_probe.tail;
192     struct timespec tv
193         = ktime_to_timespec(ktime_sub(p->tstamp, tcp_probe.start));
194
195     return scnprintf(tbuf, n,
196                     "%lu.%09lu %pISpc %pISpc %d %#x %#x %u %u %u %u %u\n",
197                     (unsigned long) tv.tv_sec,
198                     (unsigned long) tv.tv_nsec,
199                     &p->src, &p->dst, p->length, p->snd_nxt, p->snd_una,
200                     p->snd_cwnd, p->ssthresh, p->snd_wnd, p->srtt, p->rcv_wnd);
201 }
202
203 static ssize_t tcpprobe_read(struct file *file, char __user *buf,
204                             size_t len, loff_t *ppos)
205 {
206     int error = 0;
207     size_t cnt = 0;
208
209     if (!buf)
210         return -EINVAL;
211
212     while (cnt < len) {
213         char tbuf[256];
214         int width;
215
216         /* Wait for data in buffer */
217         error = wait_event_interruptible(tcp_probe.wait,
218                                         tcp_probe_used() > 0);
219         if (error)
220             break;
221
222         spin_lock_bh(&tcp_probe.lock);
223         if (tcp_probe.head == tcp_probe.tail) {
224             /* multiple readers race? */
225             spin_unlock_bh(&tcp_probe.lock);
226             continue;
227         }
228
229         width = tcpprobe_sprint(tbuf, sizeof(tbuf));
230
231         if (cnt + width < len)
232             tcp_probe.tail = (tcp_probe.tail + 1) & (bufsize - 1);
233
234         spin_unlock_bh(&tcp_probe.lock);
235
236         /* if record greater than space available
237            return partial buffer (so far) */
238         if (cnt + width >= len)

```

```

239         break;
240
241         if (copy_to_user(buf + cnt, tbuf, width))
242             return -EFAULT;
243         cnt += width;
244     }
245
246     return cnt == 0 ? error : cnt;
247 }
248
249 static const struct file_operations tcpprobe_fops = {
250     .owner    = THIS_MODULE,
251     .open     = tcpprobe_open,
252     .read     = tcpprobe_read,
253     .llseek   = noop_llseek,
254 };
255
256 static __init int tcpprobe_init(void)
257 {
258     int ret = -ENOMEM;
259
260     /* Warning: if the function signature of tcp_rcv_established,
261      * has been changed, you also have to change the signature of
262      * jtcp_rcv_established, otherwise you end up right here!
263      */
264     BUILD_BUG_ON(__same_type(tcp_rcv_established,
265                               jtcp_rcv_established) == 0);
266
267     init_waitqueue_head(&tcp_probe.wait);
268     spin_lock_init(&tcp_probe.lock);
269
270     if (bufsize == 0)
271         return -EINVAL;
272
273     bufsize = roundup_pow_of_two(bufsize);
274     tcp_probe.log = kcalloc(bufsize, sizeof(struct tcp_log), GFP_KERNEL);
275     if (!tcp_probe.log)
276         goto err0;
277
278     if (!proc_create(procname, S_IRUSR, init_net.proc_net, &tcpprobe_fops))
279         goto err0;
280
281     ret = register_jprobe(&tcp_jprobe);
282     if (ret)
283         goto err1;
284
285     pr_info("probe registered (port=%d/fwmark=%u) bufsize=%u\n",
286            port, fwmark, bufsize);
287     return 0;
288 err1:
289     remove_proc_entry(procname, init_net.proc_net);
290 err0:
291     kfree(tcp_probe.log);
292     return ret;
293 }
294 module_init(tcpprobe_init);
295
296 static __exit void tcpprobe_exit(void)
297 {
298     remove_proc_entry(procname, init_net.proc_net);
299     unregister_jprobe(&tcp_jprobe);
300     kfree(tcp_probe.log);
301 }
302 module_exit(tcpprobe_exit);
303

```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)