

# Linux Cross Reference

## Free Electrons

## Embedded Linux Experts

• *source navigation* • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version:

[2.0.40](#) [2.2.26](#) [2.4.37](#) [3.1](#) [3.2](#) [3.3](#) [3.4](#) [3.5](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#)

## Linux/net/ipv4/tcp\_westwood.c

```

1  /*
2   * TCP Westwood+: end-to-end bandwidth estimation for TCP
3   *
4   *      Angelo Dell'Aera: author of the first version of TCP Westwood+ in Linux 2.4
5   *
6   * Support at http://c3lab.poliba.it/index.php/Westwood
7   * Main references in literature:
8   *
9   * - Mascolo S, Casetti, M. Gerla et al.
10  *   "TCP Westwood: bandwidth estimation for TCP" Proc. ACM Mobicom 2001
11  *
12  * - A. Grieco, s. Mascolo
13  *   "Performance evaluation of New Reno, Vegas, Westwood+ TCP" ACM Computer
14  *   Comm. Review, 2004
15  *
16  * - A. Dell'Aera, L. Grieco, S. Mascolo.
17  *   "Linux 2.4 Implementation of Westwood+ TCP with Rate-Halving :
18  *   A Performance Evaluation Over the Internet" (ICC 2004), Paris, June 2004
19  *
20  * Westwood+ employs end-to-end bandwidth measurement to set cwnd and
21  * ssthresh after packet loss. The probing phase is as the original Reno.
22  */
23
24 #include <linux/mm.h>
25 #include <linux/module.h>
26 #include <linux/skbuff.h>
27 #include <linux/inet_diag.h>
28 #include <net/tcp.h>
29
30 /* TCP Westwood structure */
31 struct westwood {
32     u32    bw_ns_est;    /* first bandwidth estimation..not too smoothed 8) */
33     u32    bw_est;      /* bandwidth estimate */
34     u32    rtt_win_sx;   /* here starts a new evaluation... */
35     u32    bk;
36     u32    snd_una;      /* used for evaluating the number of acked bytes */
37     u32    cumul_ack;
38     u32    accounted;
39     u32    rtt;
40     u32    rtt_min;      /* minimum observed RTT */
41     u8     first_ack;    /* flag which infers that this is the first ack */
42     u8     reset_rtt_min; /* Reset RTT min to next RTT sample*/
43 };

```

```

44
45
46 /* TCP Westwood functions and constants */
47 #define TCP_WESTWOOD_RTT_MIN    (HZ/20) /* 50ms */
48 #define TCP_WESTWOOD_INIT_RTT  (20*HZ) /* maybe too conservative?! */
49
50 /*
51  * @tcp_westwood_create
52  * This function initializes fields used in TCP Westwood+,
53  * it is called after the initial SYN, so the sequence numbers
54  * are correct but new passive connections we have no
55  * information about RTTmin at this time so we simply set it to
56  * TCP_WESTWOOD_INIT_RTT. This value was chosen to be too conservative
57  * since in this way we're sure it will be updated in a consistent
58  * way as soon as possible. It will reasonably happen within the first
59  * RTT period of the connection lifetime.
60  */
61 static void tcp_westwood_init(struct sock *sk)
62 {
63     struct westwood *w = inet_csk_ca(sk);
64
65     w->bk = 0;
66     w->bw_ns_est = 0;
67     w->bw_est = 0;
68     w->accounted = 0;
69     w->cumul_ack = 0;
70     w->reset_rtt_min = 1;
71     w->rtt_min = w->rtt = TCP_WESTWOOD_INIT_RTT;
72     w->rtt_win_sx = tcp_time_stamp;
73     w->snd_una = tcp_sk(sk)->snd_una;
74     w->first_ack = 1;
75 }
76
77 /*
78  * @westwood_do_filter
79  * Low-pass filter. Implemented using constant coefficients.
80  */
81 static inline u32 westwood_do_filter(u32 a, u32 b)
82 {
83     return ((7 * a) + b) >> 3;
84 }
85
86 static void westwood_filter(struct westwood *w, u32 delta)
87 {
88     /* If the filter is empty fill it with the first sample of bandwidth */
89     if (w->bw_ns_est == 0 && w->bw_est == 0) {
90         w->bw_ns_est = w->bk / delta;
91         w->bw_est = w->bw_ns_est;
92     } else {
93         w->bw_ns_est = westwood_do_filter(w->bw_ns_est, w->bk / delta);
94         w->bw_est = westwood_do_filter(w->bw_est, w->bw_ns_est);
95     }
96 }
97
98 /*
99  * @westwood_pkts_acked
100  * Called after processing group of packets.
101  * but all westwood needs is the last sample of srtt.
102  */
103 static void tcp_westwood_pkts_acked(struct sock *sk, u32 cnt, s32 rtt)
104 {
105     struct westwood *w = inet_csk_ca(sk);
106
107     if (rtt > 0)
108         w->rtt = usecs_to_jiffies(rtt);

```

```

109 }
110
111 /*
112  * @westwood_update_window
113  * It updates RTT evaluation window if it is the right moment to do
114  * it. If so it calls filter for evaluating bandwidth.
115  */
116 static void westwood_update_window(struct sock *sk)
117 {
118     struct westwood *w = inet_csk_ca(sk);
119     s32 delta = tcp_time_stamp - w->rtt_win_sx;
120
121     /* Initialize w->snd_una with the first acked sequence number in order
122      * to fix mismatch between tp->snd_una and w->snd_una for the first
123      * bandwidth sample
124      */
125     if (w->first_ack) {
126         w->snd_una = tcp_sk(sk)->snd_una;
127         w->first_ack = 0;
128     }
129
130     /*
131      * See if a RTT-window has passed.
132      * Be careful since if RTT is less than
133      * 50ms we don't filter but we continue 'building the sample'.
134      * This minimum limit was chosen since an estimation on small
135      * time intervals is better to avoid...
136      * Obviously on a LAN we reasonably will always have
137      * right_bound = left_bound + WESTWOOD_RTT_MIN
138      */
139     if (w->rtt && delta > max_t(u32, w->rtt, TCP_WESTWOOD_RTT_MIN)) {
140         westwood_filter(w, delta);
141
142         w->bk = 0;
143         w->rtt_win_sx = tcp_time_stamp;
144     }
145 }
146
147 static inline void update_rtt_min(struct westwood *w)
148 {
149     if (w->reset_rtt_min) {
150         w->rtt_min = w->rtt;
151         w->reset_rtt_min = 0;
152     } else
153         w->rtt_min = min(w->rtt, w->rtt_min);
154 }
155
156
157 /*
158  * @westwood_fast_bw
159  * It is called when we are in fast path. In particular it is called when
160  * header prediction is successful. In such case in fact update is
161  * straight forward and doesn't need any particular care.
162  */
163 static inline void westwood_fast_bw(struct sock *sk)
164 {
165     const struct tcp_sock *tp = tcp_sk(sk);
166     struct westwood *w = inet_csk_ca(sk);
167
168     westwood_update_window(sk);
169
170     w->bk += tp->snd_una - w->snd_una;
171     w->snd_una = tp->snd_una;
172     update_rtt_min(w);
173 }

```

```

174
175 /*
176  * @westwood_acked_count
177  * This function evaluates cumul_ack for evaluating bk in case of
178  * delayed or partial acks.
179  */
180 static inline u32 westwood_acked_count(struct sock *sk)
181 {
182     const struct tcp_sock *tp = tcp_sk(sk);
183     struct westwood *w = inet_csk_ca(sk);
184
185     w->cumul_ack = tp->snd_una - w->snd_una;
186
187     /* If cumul_ack is 0 this is a dupack since it's not moving
188      * tp->snd_una.
189      */
190     if (!w->cumul_ack) {
191         w->accounted += tp->mss_cache;
192         w->cumul_ack = tp->mss_cache;
193     }
194
195     if (w->cumul_ack > tp->mss_cache) {
196         /* Partial or delayed ack */
197         if (w->accounted >= w->cumul_ack) {
198             w->accounted -= w->cumul_ack;
199             w->cumul_ack = tp->mss_cache;
200         } else {
201             w->cumul_ack -= w->accounted;
202             w->accounted = 0;
203         }
204     }
205
206     w->snd_una = tp->snd_una;
207
208     return w->cumul_ack;
209 }
210
211
212 /*
213  * TCP Westwood
214  * Here limit is evaluated as Bw estimation*RTTmin (for obtaining it
215  * in packets we use mss_cache). Rttmin is guaranteed to be >= 2
216  * so avoids ever returning 0.
217  */
218 static u32 tcp_westwood_bw_rttmin(const struct sock *sk)
219 {
220     const struct tcp_sock *tp = tcp_sk(sk);
221     const struct westwood *w = inet_csk_ca(sk);
222     return max_t(u32, (w->bw_est * w->rtt_min) / tp->mss_cache, 2);
223 }
224
225 static void tcp_westwood_event(struct sock *sk, enum tcp_ca_event event)
226 {
227     struct tcp_sock *tp = tcp_sk(sk);
228     struct westwood *w = inet_csk_ca(sk);
229
230     switch (event) {
231     case CA_EVENT_FAST_ACK:
232         westwood_fast_bw(sk);
233         break;
234
235     case CA_EVENT_COMPLETE_CWR:
236         tp->snd_cwnd = tp->snd_ssthresh = tcp_westwood_bw_rttmin(sk);
237         break;
238

```

```

239 case CA_EVENT_LOSS:
240     tp->snd\_ssthresh = tcp\_westwood\_bw\_rttmin(sk);
241     /* Update RTT_min when next ack arrives */
242     w->reset\_rtt\_min = 1;
243     break;
244
245 case CA_EVENT_SLOW_ACK:
246     westwood\_update\_window(sk);
247     w->bk += westwood\_acked\_count(sk);
248     update\_rtt\_min(w);
249     break;
250
251 default:
252     /* don't care */
253     break;
254 }
255 }
256
257 /* Extract info for Tcp socket info provided via netLink. */
258 static void tcp\_westwood\_info(struct sock *sk, u32 ext,
259                               struct sk\_buff *skb)
260 {
261     const struct westwood *ca = inet\_csk\_ca(sk);
262     if (ext & (1 << (INET_DIAG_VEGASINFO - 1))) {
263         struct tcpvegas\_info info = {
264             .tcpv_enabled = 1,
265             .tcpv_rtt = jiffies\_to\_usecs(ca->rtt),
266             .tcpv_minrtt = jiffies\_to\_usecs(ca->rtt_min),
267         };
268         nla\_put(skb, INET_DIAG_VEGASINFO, sizeof(info), &info);
269     }
270 }
271
272
273
274
275 static struct tcp\_congestion\_ops tcp_westwood __read_mostly = {
276     .init = tcp\_westwood\_init,
277     .ssthresh = tcp\_reno\_ssthresh,
278     .cong_avoid = tcp\_reno\_cong\_avoid,
279     .cwnd_event = tcp\_westwood\_event,
280     .get_info = tcp\_westwood\_info,
281     .pkts_acked = tcp\_westwood\_pkts\_acked,
282
283     .owner = THIS_MODULE,
284     .name = "westwood"
285 };
286
287 static int __init tcp\_westwood\_register(void)
288 {
289     BUILD_BUG_ON(sizeof(struct westwood) > ICSK\_CA\_PRIV\_SIZE);
290     return tcp\_register\_congestion\_control(&tcp_westwood);
291 }
292
293 static void __exit tcp\_westwood\_unregister(void)
294 {
295     tcp\_unregister\_congestion\_control(&tcp_westwood);
296 }
297
298 module_init(tcp\_westwood\_register);
299 module_exit(tcp\_westwood\_unregister);
300
301 MODULE_AUTHOR("Stephen Hemminger, Angelo DeLL'Aera");
302 MODULE_LICENSE("GPL");
303 MODULE_DESCRIPTION("TCP Westwood+");

```

[304](#)

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)