

Linux Cross Reference

[Free Electrons](#)

Embedded Linux Experts

• [source navigation](#) • [diff markup](#) • [identifier search](#) • [freetext search](#) •

Version: [2.0.40](#) [2.2.26](#) [2.4.37](#) [3.1](#) [3.2](#) [3.3](#) [3.4](#) [3.5](#) [3.6](#) [3.7](#) [3.8](#) [3.9](#) [3.10](#) [3.11](#) [3.12](#) [3.13](#) [3.14](#) [3.15](#) [3.16](#) [3.17](#)

[Linux/net/ipv4/tcp_yeah.c](#)

```

1  /*
2  *
3  *   YeAH TCP
4  *
5  *   For further details look at:
6  *   https://web.archive.org/web/20080316215752/http://wil.cs.caltech.edu/pfldnet2007/paper/YeAH_TCP.pdf
7  *
8  */
9  #include <linux/mm.h>
10 #include <linux/module.h>
11 #include <linux/skbuff.h>
12 #include <linux/inet_diag.h>
13
14 #include <net/tcp.h>
15
16 #include "tcp_vegas.h"
17
18 #define TCP_YEAH_ALPHA      80 /* number of packets queued at the bottleneck */
19 #define TCP_YEAH_GAMMA      1 /* fraction of queue to be removed per rtt */
20 #define TCP_YEAH_DELTA      3 /* log minimum fraction of cwnd to be removed on loss */
21 #define TCP_YEAH_EPSILON    1 /* log maximum fraction to be removed on early decongestion */
22 #define TCP_YEAH_PHY        8 /* maximum delta from base */
23 #define TCP_YEAH_RHO        16 /* minimum number of consecutive rtt to consider competition on loss */
24 #define TCP_YEAH_ZETA       50 /* minimum number of state switches to reset reno_count */
25
26 #define TCP_SCALABLE_AI_CNT 100U
27
28 /* YeAH variables */
29 struct yeah {
30     struct vegas vegas;      /* must be first */
31
32     /* YeAH */
33     u32 lastQ;
34     u32 doing_reno_now;
35
36     u32 reno_count;
37     u32 fast_count;
38
39     u32 pkts_acked;
40 };
41
42 static void tcp_yeah_init(struct sock *sk)
43 {
44     struct tcp_sock *tp = tcp_sk(sk);
45     struct yeah *yeah = inet_csk_ca(sk);
46
47     tcp_vegas_init(sk);
48
49     yeah->doing_reno_now = 0;
50     yeah->lastQ = 0;
51
52     yeah->reno_count = 2;
53
54     /* Ensure the MD arithmetic works. This is somewhat pedantic,
55      * since I don't think we will see a cwnd this large. :) */

```

```

56     tp->snd_cwnd_clamp = min_t(u32, tp->snd_cwnd_clamp, 0xffffffff/128);
57 }
58 }
59
60
61 static void tcp_yeah_pkts_acked(struct sock *sk, u32 pkts_acked, s32 rtt_us)
62 {
63     const struct inet_connection_sock *icsk = inet_csk(sk);
64     struct yeah *yeah = inet_csk_ca(sk);
65
66     if (icsk->icsk_ca_state == TCP_CA_Open)
67         yeah->pkts_acked = pkts_acked;
68
69     tcp_vegas_pkts_acked(sk, pkts_acked, rtt_us);
70 }
71
72 static void tcp_yeah_cong_avoid(struct sock *sk, u32 ack, u32 acked)
73 {
74     struct tcp_sock *tp = tcp_sk(sk);
75     struct yeah *yeah = inet_csk_ca(sk);
76
77     if (!tcp_is_cwnd_limited(sk))
78         return;
79
80     if (tp->snd_cwnd <= tp->snd_ssthresh)
81         tcp_slow_start(tp, acked);
82
83     else if (!yeah->doing_reno_now) {
84         /* Scalable */
85
86         tp->snd_cwnd_cnt += yeah->pkts_acked;
87         if (tp->snd_cwnd_cnt > min(tp->snd_cwnd, TCP_SCALABLE_AI_CNT)){
88             if (tp->snd_cwnd < tp->snd_cwnd_clamp)
89                 tp->snd_cwnd++;
90             tp->snd_cwnd_cnt = 0;
91         }
92
93         yeah->pkts_acked = 1;
94
95     } else {
96         /* Reno */
97         tcp_cong_avoid_ai(tp, tp->snd_cwnd);
98     }
99
100     /* The key players are v_vegas.beg_snd_una and v_beg_snd_nxt.
101     *
102     * These are so named because they represent the approximate values
103     * of snd_una and snd_nxt at the beginning of the current RTT. More
104     * precisely, they represent the amount of data sent during the RTT.
105     * At the end of the RTT, when we receive an ACK for v_beg_snd_nxt,
106     * we will calculate that (v_beg_snd_nxt - v_vegas.beg_snd_una) outstanding
107     * bytes of data have been ACKed during the course of the RTT, giving
108     * an "actual" rate of:
109     *
110     * (v_beg_snd_nxt - v_vegas.beg_snd_una) / (rtt duration)
111     *
112     * Unfortunately, v_vegas.beg_snd_una is not exactly equal to snd_una,
113     * because delayed ACKs can cover more than one segment, so they
114     * don't line up yeahly with the boundaries of RTTs.
115     *
116     * Another unfortunate fact of life is that delayed ACKs delay the
117     * advance of the left edge of our send window, so that the number
118     * of bytes we send in an RTT is often less than our cwnd will allow.
119     * So we keep track of our cwnd separately, in v_beg_snd_cwnd.
120     */
121
122     if (after(ack, yeah->vegas.beg_snd_nxt)) {
123
124         /* We do the Vegas calculations only if we got enough RTT
125         * samples that we can be reasonably sure that we got
126         * at least one RTT sample that wasn't from a delayed ACK.
127         * If we only had 2 samples total,
128         * then that means we're getting only 1 ACK per RTT, which
129         * means they're almost certainly delayed ACKs.
130         * If we have 3 samples, we should be OK.

```

```

131 */
132
133 if (yeah->vegas.cntRTT > 2) {
134     u32 rtt, queue;
135     u64 bw;
136
137     /* We have enough RTT samples, so, using the Vegas
138      * algorithm, we determine if we should increase or
139      * decrease cwnd, and by how much.
140      */
141
142     /* Pluck out the RTT we are using for the Vegas
143      * calculations. This is the min RTT seen during the
144      * last RTT. Taking the min filters out the effects
145      * of delayed ACKs, at the cost of noticing congestion
146      * a bit later.
147      */
148     rtt = yeah->vegas.minRTT;
149
150     /* Compute excess number of packets above bandwidth
151      * Avoid doing full 64 bit divide.
152      */
153     bw = tp->snd_cwnd;
154     bw *= rtt - yeah->vegas.baseRTT;
155     do_div(bw, rtt);
156     queue = bw;
157
158     if (queue > TCP_YEAH_ALPHA ||
159         rtt - yeah->vegas.baseRTT > (yeah->vegas.baseRTT / TCP_YEAH_PHY)) {
160         if (queue > TCP_YEAH_ALPHA &&
161             tp->snd_cwnd > yeah->reno_count) {
162             u32 reduction = min(queue / TCP_YEAH_GAMMA,
163                                 tp->snd_cwnd >> TCP_YEAH_EPSILON);
164
165             tp->snd_cwnd -= reduction;
166
167             tp->snd_cwnd = max(tp->snd_cwnd,
168                             yeah->reno_count);
169
170             tp->snd_ssthresh = tp->snd_cwnd;
171         }
172
173         if (yeah->reno_count <= 2)
174             yeah->reno_count = max(tp->snd_cwnd>>1, 2U);
175         else
176             yeah->reno_count++;
177
178         yeah->doing_reno_now = min(yeah->doing_reno_now + 1,
179                                 0xffffffffU);
180     } else {
181         yeah->fast_count++;
182
183         if (yeah->fast_count > TCP_YEAH_ZETA) {
184             yeah->reno_count = 2;
185             yeah->fast_count = 0;
186         }
187
188         yeah->doing_reno_now = 0;
189     }
190
191     yeah->lastQ = queue;
192 }
193
194 /* Save the extent of the current window so we can use this
195  * at the end of the next RTT.
196  */
197 yeah->vegas.beg_snd_una = yeah->vegas.beg_snd_nxt;
198 yeah->vegas.beg_snd_nxt = tp->snd_nxt;
199 yeah->vegas.beg_snd_cwnd = tp->snd_cwnd;
200
201 /* Wipe the slate clean for the next RTT. */
202 yeah->vegas.cntRTT = 0;
203 yeah->vegas.minRTT = 0xffffffff;
204
205 }

```

```

206 }
207
208 static u32 tcp_yeah_ssthresh(struct sock *sk) {
209     const struct tcp_sock *tp = tcp_sk(sk);
210     struct yeah *yeah = inet_csk_ca(sk);
211     u32 reduction;
212
213     if (yeah->doing_reno_now < TCP_YEAH_RHO) {
214         reduction = yeah->lastQ;
215
216         reduction = min(reduction, max(tp->snd_cwnd>>1, 2U));
217
218         reduction = max(reduction, tp->snd_cwnd >> TCP_YEAH_DELTA);
219     } else
220         reduction = max(tp->snd_cwnd>>1, 2U);
221
222     yeah->fast_count = 0;
223     yeah->reno_count = max(yeah->reno_count>>1, 2U);
224
225     return tp->snd_cwnd - reduction;
226 }
227
228 static struct tcp_congestion_ops tcp_yeah __read_mostly = {
229     .init          = tcp_yeah_init,
230     .ssthresh      = tcp_yeah_ssthresh,
231     .cong_avoid    = tcp_yeah_cong_avoid,
232     .set_state     = tcp_vegas_state,
233     .cwnd_event    = tcp_vegas_cwnd_event,
234     .get_info      = tcp_vegas_get_info,
235     .pkts_acked    = tcp_yeah_pkts_acked,
236
237     .owner         = THIS_MODULE,
238     .name          = "yeah",
239 };
240
241 static int __init tcp_yeah_register(void)
242 {
243     BUG_ON(sizeof(struct yeah) > ICSK_CA_PRIV_SIZE);
244     tcp_register_congestion_control(&tcp_yeah);
245     return 0;
246 }
247
248 static void __exit tcp_yeah_unregister(void)
249 {
250     tcp_unregister_congestion_control(&tcp_yeah);
251 }
252
253 module_init(tcp_yeah_register);
254 module_exit(tcp_yeah_unregister);
255
256 MODULE_AUTHOR("Angelo P. Castellani");
257 MODULE_LICENSE("GPL");
258 MODULE_DESCRIPTION("YeAH TCP");
259

```

This page was automatically generated by [LXR](#) 0.3.1 ([source](#)). • Linux is a registered trademark of Linus Torvalds • [Contact us](#)

- [Home](#)
- [Development](#)
- [Services](#)
- [Training](#)
- [Docs](#)
- [Community](#)
- [Company](#)
- [Blog](#)