

# skbuff.h

[Go to the documentation of this file.](#)

```

00001 /*
00002  *      Definitions for the 'struct sk_buff' memory handlers.
00003  *
00004  *      Authors:
00005  *          Alan Cox, <gw4pts@gw4pts.ampr.org>
00006  *          Florian La Roche, <rzsf1@rz.uni-sb.de>
00007  *
00008  *      This program is free software; you can redistribute it and/or
00009  *      modify it under the terms of the GNU General Public License
00010  *      as published by the Free Software Foundation; either version
00011  *      2 of the License, or (at your option) any later version.
00012  */
00013
00014 #ifndef _LINUX_SKBUFF_H
00015 #define _LINUX_SKBUFF_H
00016
00017 #include <linux/config.h>
00018 #include <linux/kernel.h>
00019 #include <linux/compiler.h>
00020 #include <linux/time.h>
00021 #include <linux/cache.h>
00022
00023 #include <asm/atomic.h>
00024 #include <asm/types.h>
00025 #include <linux/spinlock.h>
00026 #include <linux/mm.h>
00027 #include <linux/highmem.h>
00028 #include <linux/poll.h>
00029 #include <linux/net.h>
00030
00031 #define HAVE_ALLOC_SKB          /* For the drivers to know */
00032 #define HAVE_ALIGNABLE_SKB     /* Ditto 8) */
00033 #define SLAB_SKB               /* Slabified skbuffs */
00034
00035 #define CHECKSUM_NONE 0
00036 #define CHECKSUM_HW 1
00037 #define CHECKSUM_UNNECESSARY 2
00038
00039 #define SKB_DATA_ALIGN(X)      (((X) + (SMP_CACHE_BYTES - 1)) & \
00040                                ~(SMP_CACHE_BYTES - 1))
00041 #define SKB_MAX_ORDER(X, ORDER) (((PAGE_SIZE << (ORDER)) - (X) - \
00042                                sizeof(struct skb_shared_info)) & \
00043                                ~(SMP_CACHE_BYTES - 1))
00044 #define SKB_MAX_HEAD(X)       (SKB_MAX_ORDER((X), 0))
00045 #define SKB_MAX_ALLOC         (SKB_MAX_ORDER(0, 2))
00046
00047 /* A. Checksumming of received packets by device.
00048  *
00049  *      NONE: device failed to checksum this packet.
00050  *          skb->csum is undefined.
00051  *
00052  *      UNNECESSARY: device parsed packet and wouldbe verified checksum.
00053  *          skb->csum is undefined.
00054  *          It is bad option, but, unfortunately, many of vendors do this.
00055  *          Apparently with secret goal to sell you new device, when you
00056  *          will add new protocol to your host. F.e. IPv6. 8)
00057  *
00058  *      HW: the most generic way. Device supplied checksum of _all_
00059  *          the packet as seen by netif_rx in skb->csum.
00060  *          NOTE: Even if device supports only some protocols, but
00061  *          is able to produce some skb->csum, it MUST use HW,
00062  *          not UNNECESSARY.
00063  *
00064  *      B. Checksumming on output.
00065  *
00066  *      NONE: skb is checksummed by protocol or csum is not required.

```

```

00067 *
00068 *      HW: device is required to csum packet as seen by hard_start_xmit
00069 *      from skb->h.raw to the end and to record the checksum
00070 *      at skb->h.raw+skb->csum.
00071 *
00072 *      Device must show its capabilities in dev->features, set
00073 *      at device setup time.
00074 *      NETIF_F_HW_CSUM - it is clever device, it is able to checksum
00075 *                      everything.
00076 *      NETIF_F_NO_CSUM - loopback or reliable single hop media.
00077 *      NETIF_F_IP_CSUM - device is dumb. It is able to csum only
00078 *                      TCP/UDP over IPv4. Sigh. Vendors like this
00079 *                      way by an unknown reason. Though, see comment above
00080 *                      about CHECKSUM_UNNECESSARY. 8)
00081 *
00082 *      Any questions? No questions, good.          --ANK
00083 */
00084
00085 #ifdef __i386__
00086 #define NET_CALLER(arg) (*((void **)&arg) - 1)
00087 #else
00088 #define NET_CALLER(arg) __builtin_return_address(0)
00089 #endif
00090
00091 #ifdef CONFIG_NETFILTER
00092 struct nf_conntrack {
00093     atomic_t use;
00094     void (*destroy)(struct nf_conntrack *);
00095 };
00096
00097 struct nf_ct_info {
00098     struct nf_conntrack *master;
00099 };
00100
00101 #ifdef CONFIG_BRIDGE_NETFILTER
00102 struct nf_bridge_info {
00103     atomic_t use;
00104     struct net_device *physindev;
00105     struct net_device *physoutdev;
00106     #if defined(CONFIG_VLAN_8021Q) || defined(CONFIG_VLAN_8021Q_MODULE)
00107     struct net_device *netoutdev;
00108     #endif
00109     unsigned int mask;
00110     unsigned long data[32 / sizeof(unsigned long)];
00111 };
00112 #endif
00113
00114 #endif
00115
00116 struct sk_buff_head {
00117     /* These two members must be first. */
00118     struct sk_buff *next;
00119     struct sk_buff *prev;
00120
00121     __u32 qlen;
00122     spinlock_t lock;
00123 };
00124
00125 struct sk_buff;
00126
00127 /* To allow 64K frame to be packed as single skb without frag_list */
00128 #define MAX_SKB_FRAGS (65536/PAGE_SIZE + 2)
00129
00130 typedef struct skb_frag_struct skb_frag_t;
00131
00132 struct skb_frag_struct {
00133     struct page *page;
00134     __u16 page_offset;
00135     __u16 size;
00136 };
00137
00138 /* This data is invariant across clones and lives at
00139  * the end of the header data, ie. at skb->end.
00140  */
00141 struct skb_shared_info {
00142     atomic_t dataref;

```

```

00143     unsigned int    nr_frags;
00144     unsigned short  tso_size;
00145     unsigned short  tso_segs;
00146     struct sk_buff  *frag_list;
00147     skb_frag_t      frags[MAX_SKB_FRAGS];
00148 };
00149
00191 struct sk_buff {
00192     /* These two members must be first. */
00193     struct sk_buff  *next;
00194     struct sk_buff  *prev;
00195
00196     struct sk_buff_head *list;
00197     struct sock        *sk;
00198     struct timeval     stamp;
00199     struct net_device  *dev;
00200     struct net_device  *real_dev;
00201
00202     union {
00203         struct tcphdr  *th;
00204         struct udphdr  *uh;
00205         struct icmphdr *icmph;
00206         struct igmpchr *igmpchr;
00207         struct iphdr   *iph;
00208         struct ipv6hdr *ipv6h;
00209         unsigned char  *raw;
00210     } h;
00211
00212     union {
00213         struct iphdr   *iph;
00214         struct ipv6hdr *ipv6h;
00215         struct arphdr  *arph;
00216         unsigned char  *raw;
00217     } nh;
00218
00219     union {
00220         struct ethhdr  *ethernet;
00221         unsigned char  *raw;
00222     } mac;
00223
00224     struct dst_entry  *dst;
00225     struct sec_path   *sp;
00226
00227     /*
00228      * This is the control buffer. It is free to use for every
00229      * layer. Please put your private variables there. If you
00230      * want to keep them across layers you have to do a skb_clone()
00231      * first. This is owned by whoever has the skb queued ATM.
00232      */
00233     char                cb[48];
00234
00235     unsigned int        len,
00236                       data_len,
00237                       mac_len,
00238                       csum;
00239     unsigned char       local_df,
00240                       cloned,
00241                       pkt_type,
00242                       ip_summed;
00243     __u32               priority;
00244     unsigned short      protocol,
00245                       security;
00246
00247     void                (*destructor)(struct sk_buff *skb);
00248 #ifdef CONFIG_NETFILTER
00249     unsigned long       nfmark;
00250     __u32               nfcache;
00251     struct nf_ct_info    *nfct;
00252 #ifdef CONFIG_NETFILTER_DEBUG
00253     unsigned int        nf_debug;
00254 #endif
00255 #ifdef CONFIG_BRIDGE_NETFILTER
00256     struct nf_bridge_info *nf_bridge;
00257 #endif
00258 #endif /* CONFIG_NETFILTER */
00259 #if defined(CONFIG_HIPPI)

```

```

00260     union {
00261         __u32         ifield;
00262     } private;
00263 #endif
00264 #ifdef CONFIG_NET_SCHED
00265     __u32         tc_index;          /* traffic control index */
00266 #endif
00267
00268     /* These elements must be at the end, see alloc_skb() for details. */
00269     unsigned int   truesize;
00270     atomic_t       users;
00271     unsigned char  *head,
00272                   *data,
00273                   *tail,
00274                   *end;
00275 };
00276
00277 #ifdef __KERNEL__
00278 /*
00279  *      Handling routines are only of interest to the kernel
00280  */
00281 #include <linux/slab.h>
00282
00283 #include <asm/system.h>
00284
00285 extern void        __kfree_skb(struct sk_buff *skb);
00286 extern struct sk_buff *alloc_skb(unsigned int size, int priority);
00287 extern void        kfree_skbmem(struct sk_buff *skb);
00288 extern struct sk_buff *skb_clone(struct sk_buff *skb, int priority);
00289 extern struct sk_buff *skb_copy(const struct sk_buff *skb, int priority);
00290 extern struct sk_buff *pskb_copy(struct sk_buff *skb, int gfp_mask);
00291 extern int         pskb_expand_head(struct sk_buff *skb,
00292                                     int nhead, int ntail, int gfp_mask);
00293 extern struct sk_buff *skb_realloc_headroom(struct sk_buff *skb,
00294                                             unsigned int headroom);
00295 extern struct sk_buff *skb_copy_expand(const struct sk_buff *skb,
00296                                       int newheadroom, int newtailroom,
00297                                       int priority);
00298 extern struct sk_buff *skb_pad(struct sk_buff *skb, int pad);
00299 #define dev_kfree_skb(a)      kfree_skb(a)
00300 extern void        skb_over_panic(struct sk_buff *skb, int len,
00301                                  void *here);
00302 extern void        skb_under_panic(struct sk_buff *skb, int len,
00303                                   void *here);
00304
00305 /* Internal */
00306 #define skb_shinfo(SKB)      ((struct skb_shared_info *)((SKB)->end))
00307
00314 static inline int skb_queue_empty(const struct sk_buff_head *list)
00315 {
00316     return list->next == (struct sk_buff *)list;
00317 }
00318
00326 static inline struct sk_buff *skb_get(struct sk_buff *skb)
00327 {
00328     atomic_inc(&skb->users);
00329     return skb;
00330 }
00331
00332 /*
00333  * If users == 1, we are the only owner and are can avoid redundant
00334  * atomic change.
00335  */
00336
00344 static inline void kfree_skb(struct sk_buff *skb)
00345 {
00346     if (atomic_read(&skb->users) == 1 || atomic_dec_and_test(&skb->users))
00347         __kfree_skb(skb);
00348 }
00349
00350 /* Use this if you didn't touch the skb state [for fast switching] */
00351 static inline void kfree_skb_fast(struct sk_buff *skb)
00352 {
00353     if (atomic_read(&skb->users) == 1 || atomic_dec_and_test(&skb->users))
00354         kfree_skbmem(skb);
00355 }

```

```

00356
00365 static inline int skb_cloned(const struct sk_buff *skb)
00366 {
00367     return skb->cloned && atomic_read(&skb_shinfo(skb)->dataref) != 1;
00368 }
00369
00377 static inline int skb_shared(const struct sk_buff *skb)
00378 {
00379     return atomic_read(&skb->users) != 1;
00380 }
00381
00395 static inline struct sk_buff *skb_share_check(struct sk_buff *skb, int pri)
00396 {
00397     might_sleep_if(pri & __GFP_WAIT);
00398     if (skb_shared(skb)) {
00399         struct sk_buff *nskb = skb_clone(skb, pri);
00400         kfree_skb(skb);
00401         skb = nskb;
00402     }
00403     return skb;
00404 }
00405
00406 /*
00407  * Copy shared buffers into a new sk_buff. We effectively do COW on
00408  * packets to handle cases where we have a local reader and forward
00409  * and a couple of other messy ones. The normal one is tcpdumping
00410  * a packet thats being forwarded.
00411  */
00412
00426 static inline struct sk_buff *skb_unshare(struct sk_buff *skb, int pri)
00427 {
00428     might_sleep_if(pri & __GFP_WAIT);
00429     if (skb_cloned(skb)) {
00430         struct sk_buff *nskb = skb_copy(skb, pri);
00431         kfree_skb(skb); /* Free our shared copy */
00432         skb = nskb;
00433     }
00434     return skb;
00435 }
00436
00450 static inline struct sk_buff *skb_peek(struct sk_buff_head *list_)
00451 {
00452     struct sk_buff *list = ((struct sk_buff *)list_)->next;
00453     if (list == (struct sk_buff *)list_)
00454         list = NULL;
00455     return list;
00456 }
00457
00471 static inline struct sk_buff *skb_peek_tail(struct sk_buff_head *list_)
00472 {
00473     struct sk_buff *list = ((struct sk_buff *)list_)->prev;
00474     if (list == (struct sk_buff *)list_)
00475         list = NULL;
00476     return list;
00477 }
00478
00485 static inline __u32 skb_queue_len(const struct sk_buff_head *list_)
00486 {
00487     return list_->qlen;
00488 }
00489
00490 static inline void skb_queue_head_init(struct sk_buff_head *list)
00491 {
00492     spin_lock_init(&list->lock);
00493     list->prev = list->next = (struct sk_buff *)list;
00494     list->qlen = 0;
00495 }
00496
00497 /*
00498  * Insert an sk_buff at the start of a list.
00499  *
00500  * The "__skb_xxxx()" functions are the non-atomic ones that
00501  * can only be called with interrupts disabled.
00502  */
00503
00514 extern void skb_queue_head(struct sk_buff_head *list, struct sk_buff *newsk);

```

```

00515 static inline void __skb_queue_head(struct sk_buff_head *list,
00516                                     struct sk_buff *newsk)
00517 {
00518     struct sk_buff *prev, *next;
00519
00520     newsk->list = list;
00521     list->qlen++;
00522     prev = (struct sk_buff *)list;
00523     next = prev->next;
00524     newsk->next = next;
00525     newsk->prev = prev;
00526     next->prev = prev->next = newsk;
00527 }
00528
00539 extern void skb_queue_tail(struct sk_buff_head *list, struct sk_buff *newsk);
00540 static inline void __skb_queue_tail(struct sk_buff_head *list,
00541                                     struct sk_buff *newsk)
00542 {
00543     struct sk_buff *prev, *next;
00544
00545     newsk->list = list;
00546     list->qlen++;
00547     next = (struct sk_buff *)list;
00548     prev = next->prev;
00549     newsk->next = next;
00550     newsk->prev = prev;
00551     next->prev = prev->next = newsk;
00552 }
00553
00554
00563 extern struct sk_buff *skb_dequeue(struct sk_buff_head *list);
00564 static inline struct sk_buff *__skb_dequeue(struct sk_buff_head *list)
00565 {
00566     struct sk_buff *next, *prev, *result;
00567
00568     prev = (struct sk_buff *) list;
00569     next = prev->next;
00570     result = NULL;
00571     if (next != prev) {
00572         result = next;
00573         next = next->next;
00574         list->qlen--;
00575         next->prev = prev;
00576         prev->next = next;
00577         result->next = result->prev = NULL;
00578         result->list = NULL;
00579     }
00580     return result;
00581 }
00582
00583
00584 /*
00585  *      Insert a packet on a list.
00586  */
00587 extern void      skb_insert(struct sk_buff *old, struct sk_buff *newsk);
00588 static inline void __skb_insert(struct sk_buff *newsk,
00589                                 struct sk_buff *prev, struct sk_buff *next,
00590                                 struct sk_buff_head *list)
00591 {
00592     newsk->next = next;
00593     newsk->prev = prev;
00594     next->prev = prev->next = newsk;
00595     newsk->list = list;
00596     list->qlen++;
00597 }
00598
00599 /*
00600  *      Place a packet after a given packet in a list.
00601  */
00602 extern void      skb_append(struct sk_buff *old, struct sk_buff *newsk);
00603 static inline void __skb_append(struct sk_buff *old, struct sk_buff *newsk)
00604 {
00605     __skb_insert(newsk, old, old->next, old->list);
00606 }
00607
00608 /*

```

```

00609  * remove sk_buff from list. _Must_ be called atomically, and with
00610  * the list known..
00611  */
00612 extern void      skb_unlink(struct sk_buff *skb);
00613 static inline void __skb_unlink(struct sk_buff *skb, struct sk_buff_head *list)
00614 {
00615     struct sk_buff *next, *prev;
00616
00617     list->qlen--;
00618     next      = skb->next;
00619     prev      = skb->prev;
00620     skb->next  = skb->prev = NULL;
00621     skb->list  = NULL;
00622     next->prev = prev;
00623     prev->next = next;
00624 }
00625
00626
00627 /* XXX: more streamlined implementation */
00628
00637 extern struct sk_buff *skb_dequeue_tail(struct sk_buff_head *list);
00638 static inline struct sk_buff *__skb_dequeue_tail(struct sk_buff_head *list)
00639 {
00640     struct sk_buff *skb = skb_peek_tail(list);
00641     if (skb)
00642         __skb_unlink(skb, list);
00643     return skb;
00644 }
00645
00646
00647 static inline int skb_is_nonlinear(const struct sk_buff *skb)
00648 {
00649     return skb->data_len;
00650 }
00651
00652 static inline unsigned int skb_headlen(const struct sk_buff *skb)
00653 {
00654     return skb->len - skb->data_len;
00655 }
00656
00657 static inline int skb_pagelen(const struct sk_buff *skb)
00658 {
00659     int i, len = 0;
00660
00661     for (i = (int)skb_shinfo(skb)->nr_frags - 1; i >= 0; i--)
00662         len += skb_shinfo(skb)->frags[i].size;
00663     return len + skb_headlen(skb);
00664 }
00665
00666 static inline void skb_fill_page_desc(struct sk_buff *skb, int i, struct page *page, int off, int size)
00667 {
00668     skb_frag_t *frag = &skb_shinfo(skb)->frags[i];
00669     frag->page = page;
00670     frag->page_offset = off;
00671     frag->size = size;
00672     skb_shinfo(skb)->nr_frags = i+1;
00673 }
00674
00675 #define SKB_PAGE_ASSERT(skb)      BUG_ON(skb_shinfo(skb)->nr_frags)
00676 #define SKB_FRAG_ASSERT(skb)     BUG_ON(skb_shinfo(skb)->frag_list)
00677 #define SKB_LINEAR_ASSERT(skb)   BUG_ON(skb_is_nonlinear(skb))
00678
00679 /*
00680  *      Add data to an sk_buff
00681  */
00682 static inline unsigned char *__skb_put(struct sk_buff *skb, unsigned int len)
00683 {
00684     unsigned char *tmp = skb->tail;
00685     SKB_LINEAR_ASSERT(skb);
00686     skb->tail += len;
00687     skb->len  += len;
00688     return tmp;
00689 }
00690
00700 static inline unsigned char *skb_put(struct sk_buff *skb, unsigned int len)
00701 {

```



```

00702     unsigned char *tmp = skb->tail;
00703     SKB_LINEAR_ASSERT(skb);
00704     skb->tail += len;
00705     skb->len += len;
00706     if (unlikely(skb->tail>skb->end))
00707         skb_over_panic(skb, len, current_text_addr());
00708     return tmp;
00709 }
00710
00711 static inline unsigned char *__skb_push(struct sk_buff *skb, unsigned int len)
00712 {
00713     skb->data -= len;
00714     skb->len += len;
00715     return skb->data;
00716 }
00717
00727 static inline unsigned char *skb_push(struct sk_buff *skb, unsigned int len)
00728 {
00729     skb->data -= len;
00730     skb->len += len;
00731     if (unlikely(skb->data<skb->head))
00732         skb_under_panic(skb, len, current_text_addr());
00733     return skb->data;
00734 }
00735
00736 static inline unsigned char *__skb_pull(struct sk_buff *skb, unsigned int len)
00737 {
00738     skb->len -= len;
00739     BUG_ON(skb->len < skb->data_len);
00740     return skb->data += len;
00741 }
00742
00753 static inline unsigned char *skb_pull(struct sk_buff *skb, unsigned int len)
00754 {
00755     return unlikely(len > skb->len) ? NULL : __skb_pull(skb, len);
00756 }
00757
00758 extern unsigned char *__pskb_pull_tail(struct sk_buff *skb, int delta);
00759
00760 static inline unsigned char *__pskb_pull(struct sk_buff *skb, unsigned int len)
00761 {
00762     if (len > skb_headlen(skb) &&
00763         !__pskb_pull_tail(skb, len-skb_headlen(skb)))
00764         return NULL;
00765     skb->len -= len;
00766     return skb->data += len;
00767 }
00768
00769 static inline unsigned char *pskb_pull(struct sk_buff *skb, unsigned int len)
00770 {
00771     return unlikely(len > skb->len) ? NULL : __pskb_pull(skb, len);
00772 }
00773
00774 static inline int pskb_may_pull(struct sk_buff *skb, unsigned int len)
00775 {
00776     if (likely(len <= skb_headlen(skb)))
00777         return 1;
00778     if (unlikely(len > skb->len))
00779         return 0;
00780     return __pskb_pull_tail(skb, len-skb_headlen(skb)) != NULL;
00781 }
00782
00789 static inline int skb_headroom(const struct sk_buff *skb)
00790 {
00791     return skb->data - skb->head;
00792 }
00793
00800 static inline int skb_tailroom(const struct sk_buff *skb)
00801 {
00802     return skb_is_nonlinear(skb) ? 0 : skb->end - skb->tail;
00803 }
00804
00813 static inline void skb_reserve(struct sk_buff *skb, unsigned int len)
00814 {
00815     skb->data += len;
00816     skb->tail += len;

```



```

00817 }
00818
00819 extern int __pskb_trim(struct sk_buff *skb, unsigned int len, int realloc);
00820
00821 static inline void __skb_trim(struct sk_buff *skb, unsigned int len)
00822 {
00823     if (!skb->data_len) {
00824         skb->len = len;
00825         skb->tail = skb->data + len;
00826     } else
00827         __pskb_trim(skb, len, 0);
00828 }
00829
00838 static inline void skb_trim(struct sk_buff *skb, unsigned int len)
00839 {
00840     if (skb->len > len)
00841         __skb_trim(skb, len);
00842 }
00843
00844
00845 static inline int __pskb_trim(struct sk_buff *skb, unsigned int len)
00846 {
00847     if (!skb->data_len) {
00848         skb->len = len;
00849         skb->tail = skb->data+len;
00850         return 0;
00851     }
00852     return __pskb_trim(skb, len, 1);
00853 }
00854
00855 static inline int pskb_trim(struct sk_buff *skb, unsigned int len)
00856 {
00857     return (len < skb->len) ? __pskb_trim(skb, len) : 0;
00858 }
00859
00868 static inline void skb_orphan(struct sk_buff *skb)
00869 {
00870     if (skb->destructor)
00871         skb->destructor(skb);
00872     skb->destructor = NULL;
00873     skb->sk = NULL;
00874 }
00875
00884 extern void skb_queue_purge(struct sk_buff_head *list);
00885 static inline void __skb_queue_purge(struct sk_buff_head *list)
00886 {
00887     struct sk_buff *skb;
00888     while ((skb = __skb_dequeue(list)) != NULL)
00889         kfree_skb(skb);
00890 }
00891
00904 static inline struct sk_buff *__dev_alloc_skb(unsigned int length,
00905                                                int gfp_mask)
00906 {
00907     struct sk_buff *skb = alloc_skb(length + 16, gfp_mask);
00908     if (likely(skb))
00909         skb_reserve(skb, 16);
00910     return skb;
00911 }
00912
00925 static inline struct sk_buff *dev_alloc_skb(unsigned int length)
00926 {
00927     return __dev_alloc_skb(length, GFP_ATOMIC);
00928 }
00929
00942 static inline int skb_cow(struct sk_buff *skb, unsigned int headroom)
00943 {
00944     int delta = (headroom > 16 ? headroom : 16) - skb_headroom(skb);
00945
00946     if (delta < 0)
00947         delta = 0;
00948
00949     if (delta || skb_cloned(skb))
00950         return pskb_expand_head(skb, (delta + 15) & ~15, 0, GFP_ATOMIC);
00951     return 0;
00952 }

```

```

00953
00966 static inline struct sk_buff *skb_padto(struct sk_buff *skb, unsigned int len)
00967 {
00968     unsigned int size = skb->len;
00969     if (likely(size >= len))
00970         return skb;
00971     return skb_pad(skb, len-size);
00972 }
00973
00982 extern int __skb_linearize(struct sk_buff *skb, int gfp);
00983 static inline int skb_linearize(struct sk_buff *skb, int gfp)
00984 {
00985     return __skb_linearize(skb, gfp);
00986 }
00987
00988 static inline void *kmap_skb_frag(const skb_frag_t *frag)
00989 {
00990     #ifdef CONFIG_HIGHMEM
00991         BUG_ON(in_irq());
00992         local_bh_disable();
00994     #endif
00995     return kmap_atomic(frag->page, KM_SKB_DATA_SOFTIRQ);
00996 }
00997
00998 static inline void kunmap_skb_frag(void *vaddr)
00999 {
01000     kunmap_atomic(vaddr, KM_SKB_DATA_SOFTIRQ);
01001     #ifdef CONFIG_HIGHMEM
01002         local_bh_enable();
01003     #endif
01004 }
01005
01006 #define skb_queue_walk(queue, skb) \
01007     for (skb = (queue)->next, prefetch(skb->next); \
01008          (skb != (struct sk_buff *) (queue)); \
01009          skb = skb->next, prefetch(skb->next))
01010
01011
01012 extern struct sk_buff *skb_recv_datagram(struct sock *sk, unsigned flags,
01013                                         int noblock, int *err);
01014 extern unsigned int datagram_poll(struct file *file, struct socket *sock,
01015                                  struct poll_table_struct *wait);
01016 extern int skb_copy_datagram(const struct sk_buff *from,
01017                              int offset, char *to, int size);
01018 extern int skb_copy_datagram_iovec(const struct sk_buff *from,
01019                                    int offset, struct iovec *to,
01020                                    int size);
01021 extern int skb_copy_and_csum_datagram(const struct sk_buff *skb,
01022                                       int offset, u8 *to, int len,
01023                                       unsigned int *csum);
01024 extern int skb_copy_and_csum_datagram_iovec(const
01025                                              struct sk_buff *skb,
01026                                              int hlen,
01027                                              struct iovec *iov);
01028 extern void skb_free_datagram(struct sock *sk, struct sk_buff *skb);
01029 extern unsigned int skb_checksum(const struct sk_buff *skb, int offset,
01030                                 int len, unsigned int csum);
01031 extern int skb_copy_bits(const struct sk_buff *skb, int offset,
01032                          void *to, int len);
01033 extern unsigned int skb_copy_and_csum_bits(const struct sk_buff *skb,
01034                                             int offset, u8 *to, int len,
01035                                             unsigned int csum);
01036 extern void skb_copy_and_csum_dev(const struct sk_buff *skb, u8 *to);
01037
01038 extern void skb_init(void);
01039 extern void skb_add_mtu(int mtu);
01040
01041 struct tux_req_struct;
01042
01043 #ifdef CONFIG_NETFILTER
01044 static inline void nf_conntrack_put(struct nf_ct_info *nfct)
01045 {
01046     if (nfct && atomic_dec_and_test(&nfct->master->use))
01047         nfct->master->destroy(nfct->master);
01048 }

```

```
01049 static inline void nf_conntrack_get(struct nf_ct_info *nfct)
01050 {
01051     if (nfct)
01052         atomic_inc(&nfct->master->use);
01053 }
01054
01055 #ifdef CONFIG_BRIDGE_NETFILTER
01056 static inline void nf_bridge_put(struct nf_bridge_info *nf_bridge)
01057 {
01058     if (nf_bridge && atomic_dec_and_test(&nf_bridge->use))
01059         kfree(nf_bridge);
01060 }
01061 static inline void nf_bridge_get(struct nf_bridge_info *nf_bridge)
01062 {
01063     if (nf_bridge)
01064         atomic_inc(&nf_bridge->use);
01065 }
01066 #endif
01067
01068 #endif
01069
01070 #endif /* __KERNEL__ */
01071 #endif /* _LINUX_SKBUFF_H */
```

---

Generated at Wed Sep 22 17:57:01 2004 for LINUX\_TCP\_STACK by  1.2.8.1 written by [Dimitri van Heesch](#), © 1997-

2001